

Build a Production SaaS Platform From Scratch

A rebuild-from-zero course in architecture, engineering practice,
and the decisions behind them

Contents

Preface	12
What this book is	12
Who this is for	12
What this book is not	12
The rhythm	13
How to use this book	13
 The platform at a glance	 13
 Lesson 0.1 — The mental model & the decision record	 17
1. What you are going to build	17
2. The four ideas everything else hangs on	17
3. The decision record — the habit that makes it stick	18
4. Exercise — write your ADR-000 and ADR-001	19
5. Checkpoint	19
 Lesson 0.2 — A reproducible machine	 20
1. Motivate — "works on my machine" is an architecture failure	20
2. Install the pinned toolchain	20
3. Infrastructure as one file — docker-compose	21
4. The <code>.env` contract — ADR-001 gets teeth</code>	22
5. Run it — verify everything	23
6. Architecture Decision	23
7. Checkpoint	23
 Lesson 1.1 — Solution, projects & supply chain	 27
1. Motivate — the layout is the first enforcement	27
2. Create the skeleton	27
3. One build policy for the whole solution	28
4. Central Package Management — one version table	29
5. Lock it — restore becomes reproducible	30
6. Run it	30
7. Architecture Decision	31

8. Checkpoint	31
Lesson 1.2 — First endpoint, first test	32
1. Motivate — why a "hello world" deserves a whole lesson	32
2. Red — the failing test comes first	32
3. Green — the minimum that passes	34
4. Refactor & harden — a scratchpad and a habit	34
5. Run it	35
6. Architecture Decision	35
7. Checkpoint	35
Lesson 1.3 — Database & the test container	36
1. Motivate — your tests are only as honest as their database	36
2. The context and its first entity	37
3. Migrations — schema changes as reviewable code	38
4. The fixture — a real database that can't leak between tests	39
5. Red → Green — prove the loop against real Postgres	40
6. Architecture Decision	41
7. Checkpoint	41
Lesson 1.4 — Configuration & the options pattern	42
1. Motivate — the 3 a.m. config bug	42
2. The config hierarchy — one mental model	42
3. Typed settings — validation at construction	43
4. Defaults are decisions	45
5. The gate — config that can't go undocumented	45
6. Run it — watch it fail correctly	46
7. Architecture Decision	46
8. Checkpoint	46
Lesson 1.5 — The error envelope	48
1. Motivate — the anatomy of a leaked stack trace	48
2. Green — the envelope	48
3. Harden — a shape is only a shape if it's the only one	49
4. Run it	50
5. Architecture Decision	50

6. Checkpoint	50
Lesson 1.6 — CI from commit one	51
1. Motivate — a gate you add at the end is a gate you never designed for	51
2. The workflow — build-test, from the real file	51
3. Two more jobs — the supply-chain gates go remote	53
4. The human gate — the PR template	54
5. Run it — push and watch	54
6. Architecture Decision	54
7. Checkpoint	55
Lesson 2.1 — Users & JWT access tokens	58
1. Motivate — the fork in the road	58
2. The `User` — smaller than you expect	59
3. What a JWT actually is — sixty seconds of anatomy	59
4. Green — the token service	60
5. Red first, though — the tests that drove it	61
6. Harden — the generator and hasher seams	61
7. Run it	62
8. Architecture Decision	63
9. Checkpoint	63
Lesson 2.2 — Rotating refresh tokens & sessions	64
1. Motivate — the session-length trilemma	64
2. The entity — a session record wearing a token costume	64
3. Issuing — the seams compose	65
4. The refresh endpoint — where rotation and detection live	65
5. The cookie — every attribute is a decision	67
6. Red — the tests that pin the properties	68
7. Run it	68
8. Architecture Decision	68
9. Checkpoint	69
Lesson 2.3 — Email I: the `IEmailSender` seam	70
1. Motivate — email is a vendor decision wearing a feature costume	70
2. The interface — small, but every word chosen	70

3. The implementation — and the sharp edges it rounds off	71
4. The containment test — a seam with a fence	73
5. Branded, localized — email HTML is a hostile country	73
6. Red → Green — Mailpit is your assertable inbox	74
7. Run it	75
8. Architecture Decision	75
9. Checkpoint	75

Lesson 2.4 — Passwordless: magic link + OTP 77

1. Motivate — passwords are a liability you can decline	77
2. One entity for both — `LoginToken`	78
3. Issue and redeem — the magic-link half	78
4. The OTP half — do the arithmetic first	79
5. What the caller learns — collapse the statuses	80
6. Red — the tests that drove it	81
7. Run it	81
8. Architecture Decision	82
9. Checkpoint	82

Lesson 2.5 — OAuth & account linking 83

1. Motivate — what OAuth actually outsources	83
2. The dance itself — buy, don't build	83
3. Normalize the claims — fail closed (R17)	84
4. The identity resolution — three cases, one guard	85
5. Explicit linking — the same guard from the other door	86
6. Red — the tests	86
7. Run it	87
8. Architecture Decision	87
9. Checkpoint	88

Lesson 2.6 — Tenancy I: the global query filter 89

1. Motivate — build the leak, then look at it	89
2. Red — a test that proves the leak	90
3. Green — make isolation structural	90
4. Harden — fail closed, then pin both properties	91
5. Run it	92

6. Architecture Decision	92
7. Checkpoint	93
Lesson 2.7 — Tenancy II: write-side stamping & `EnterTenant`	94
1. Motivate — the door you forgot you left open	94
2. Green — the interceptor	95
3. The trust boundary — who is `Guid.Empty`, really?	96
4. Red — the tests that pin all three behaviors	97
5. Run it	98
6. Architecture Decision	98
7. Checkpoint	98
Lesson 2.8 — Tenants & membership	99
1. Motivate — the model <i>*is*</i> the answer sheet	99
2. The entities — one boring, one load-bearing	99
3. Where tenants come from — the household of one	100
4. Transitions — where models go to die (or not)	101
5. The app-facing surface — and the rename seam	101
6. Red — the tests	102
7. Run it	102
8. Architecture Decision	102
9. Checkpoint	103
Lesson 2.9 — Invitations, dissolve & the contributor seam	104
1. Motivate — two lifecycle holes	104
2. Invitations — old parts, one new idea	104
3. The contributor seam — teardown that scales with features	106
4. Harden — the seam that can't be forgotten, provably	107
5. Red — the tests	107
6. Run it	108
7. Architecture Decision	108
8. Checkpoint	109
Lesson 3.1 — The repository seam	112
1. Motivate — two ways to give features data access, both wrong	112
2. The interface — safe by default, dangerous on purpose	112

3. The implementation — one registration covers every entity	114
4. Unit of work — where atomicity lives	114
5. Red — prove the seam keeps the promise	115
6. Run it	116
7. Architecture Decision	116
8. Checkpoint	116

Lesson 3.2 — Anatomy of a vertical slice (Notes) 118

1. Motivate — the shape is the product	118
2. The entity (step 1)	119
3. The endpoints — auth you can't forget (step 4a)	119
4. The handler — where the logic lives, and how little there is	120
5. The models — co-located, on purpose (step 4b)	121
6. The contributor — and the escape-hatch boundary (step 5)	121
7. Wiring and isolation — the rules that keep slices independent	122
8. Red — the slice's tests	123
9. Run it	123
10. Architecture Decision	123
11. Checkpoint	124

Lesson 3.3 — Injected clocks & the architecture tests 125

1. Motivate — the clock you can't test	125
2. `TimeProvider` — time as a seam	125
3. Green — ban the alternative	126
4. Zoom out — the architecture-test project	127
5. The philosophy — invariants that can't regress	128
6. Run it	128
7. Architecture Decision	129
8. Checkpoint	129

Lesson 3.4 — The web client & auth UI 130

1. Motivate — where does the UI live?	130
2. The client is an API client — the clean boundary, from the outside	131
3. The session, client-side — memory + silent refresh	131
4. The HTTP plumbing — two clients, one subtle DI cycle	132
5. The per-host seams — designed web-first, native for free	133

6. The pages — login and the callback	134
7. Single-origin, already	135
8. Run it — the first end-to-end human flow	135
9. Architecture Decision	135
10. Checkpoint	136

Lesson 3.5 — Localization 137

1. Motivate — the hard-coded string is a one-way door	137
2. Where strings live — resx + `LocalizedString`	137
3. The decision chain — which language, and why	138
4. Surviving a cold start — the web-first culture seam	138
5. UI vs email localization — the same resx, two resolution models	139
6. Red — proving it end to end	140
7. Run it	140
8. Architecture Decision	140
9. Checkpoint	141

Lesson 3.6 — The E2E harness 142

1. Motivate — the layer that catches what the others can't	142
2. The harness — a real browser against the real stack	142
3. Mailpit as an assertable inbox — the loop closes	143
4. The Page Object Model — journeys that read like intent	144
5. Red — the first journey	144
6. The economics — what earns an E2E test (the division of labor)	145
7. Run it	145
8. Architecture Decision	146
9. Checkpoint	146

Lesson 3.7 — Build your own slice (capstone) 147

1. The brief	147
2. Story first (the habit from `WAYS_OF_WORKING.md`)	147
3. The seven steps (your worksheet)	148
4. Test-drive it — Red before Green, every step	149
5. Your safety net — run the architecture tests	149
6. Definition of done	150
7. Checkpoint	150

Lesson 4.1 — The transactional outbox 153

1. Motivate — the write you can't make atomic	153
2. The message — a durable IOU	154
3. Enqueue — the effect rides the caller's commit	154
4. Handlers — typed, and necessarily idempotent	155
5. The processor — claiming work without stepping on yourself	155
6. Red — testing the core without the clock or the thread	157
7. Run it	158
8. Architecture Decision	158
9. Checkpoint	159

Lesson 4.2 — Email II: the outbox decorator 160

1. Motivate — the seam finally earns its keep	160
2. The decorator — same interface, new behavior	160
3. The handler — the real send, later	161
4. Idempotency, made concrete	162
5. Red — proving the swap	162
6. Run it	163
7. Architecture Decision	163
8. Checkpoint	164

Lesson 4.3 — The inbox 165

1. Motivate — the redelivery you can't wish away	165
2. The ledger — one row per delivery, ever	165
3. The claim — race-free by construction	166
4. The one rule that makes it correct — claim *inside* the work's transaction	167
5. Red — the tests that pin exactly-once-effect	168
6. Run it	168
7. Architecture Decision	168
8. Checkpoint	169

Lesson 4.4 — The scheduler 170

1. Motivate — the work nobody asks for	170
2. The job — an interface anyone can implement	170
3. The reference job — and a tenancy-flavored subtlety	171

4. The host — resilient, scoped, and testable	172
5. Red — deterministic scheduling without waiting	173
6. Run it	173
7. Architecture Decision	173
8. Checkpoint	174

Lesson 4.5 — Observability 175

1. Motivate — a production question you can't answer	175
2. Structured logging — logs you can query, scoped to who	175
3. Tracing — where the time actually went	176
4. Health — liveness vs readiness	177
5. Red — testing that telemetry enriches and health reports	178
6. Run it	178
7. Architecture Decision	179
8. Checkpoint	179

Lesson 4.6 — The append-only audit log 180

1. Motivate — the record that must not be editable	180
2. Recording — semantic events, staged on the caller's transaction	181
3. Explicit recording, not magic — a recorded decision	182
4. Immutability — enforced by the build, not the honor system	182
5. Red — proving append-only holds	183
6. Run it	183
7. Architecture Decision	183
8. Checkpoint	184

Lesson 5.1 — Billing abstraction & entitlements 187

1. Motivate — two ways billing goes catastrophically wrong	187
2. The provider seam — money mutations live elsewhere	188
3. The projection — a cache that fails closed	188
4. Plans are code, not data	189
5. Entitlements — the server-side gate, failing closed	189
6. The 402 gate — a distinct answer for "your plan doesn't include this"	190
7. The dev fake — and the rule that it must never ship	191
8. Red — the tests	191
9. Run it	192

10. Architecture Decision	192
11. Checkpoint	193

Lesson 5.2 — Stripe: checkout, webhook, portal 194

1. Motivate — three interactions, one hard one	194
2. Checkout & portal — redirect, grant nothing	194
3. The webhook handler — six primitives, one method	195
4. The endpoint — a system surface that logs its attackers	197
5. The provider implementation — Stripe behind the seam	198
6. Red — driving the whole flow with the fake	198
7. Run it	199
8. Architecture Decision	199
9. Checkpoint	199

Lesson 5.3 — Quotas, dunning & dissolve 201

1. Motivate — the last seat, sold twice	201
2. Two kinds of quota — seats and metered usage	201
3. Atomic consume — the check *is* the write	202
4. Dunning — nudging a lapsed tenant, once	203
5. Dissolve — stop charging a tenant that no longer exists	204
6. Red — the tests, concurrency first	205
7. Run it	206
8. Architecture Decision	206
9. Checkpoint	206

Preface

What this book is

This is a course disguised as a book. Over ten parts and fifty lessons, you will build — by typing every hand-written line yourself — a production-grade, multi-tenant SaaS platform: custom JWT authentication with passwordless sign-in, OAuth, and MFA; tenant isolation enforced by the type system *and* by the database; billing with Stripe; reliable asynchronous work via a transactional outbox; role-based access control; file storage; GDPR export and erasure; a public API; outbound webhooks; an admin back-office; full observability — and at the end you will deploy it to a real host behind a real CI/CD pipeline that you also built.

The code is real. Every excerpt in this book is drawn from a working reference platform (Perezosoft Platform), not invented for the page. A coverage gate — a script, in the same spirit as the tests you'll write — maps every file of that platform to the lesson that builds it, so nothing is skipped and nothing appears out of thin air.

But the code is not the point. The point is the two habits that separate an architect from a very fast typist:

1 Making invariants structural. Not "remember to filter by tenant" but "a query

cannot forget to filter by tenant." Not "please don't call this API in feature code" but "the build fails if you do." You will meet this move dozens of times — in query filters, interceptors, marker interfaces, architecture tests, CI gates, database policies — until reaching for it is a reflex.

1 Recording decisions. Every lesson ends with an *Architecture Decision* box: the

fork that was faced, the option chosen, the options rejected, and the price paid. You will write your own decision log from lesson 0.1 onward, before any code exists. Six months from now, "why is it like this?" will have an answer.

Who this is for

You are comfortable in C# — classes, interfaces, async/await, LINQ hold no fear — and you can find your way around a terminal and git. You have probably built applications, perhaps professionally, but you want the *architecture* to stop feeling like folklore: why layers, why seams, why tests first, why all these patterns with strange names, and — above all — when *not* to use them.

You do not need prior experience with multi-tenancy, EF Core internals, Stripe, Docker, OAuth, or GitHub Actions. Each is taught when the platform needs it, motivated by a concrete problem you will feel before you solve it.

What this book is not

Honesty up front: this is a course in **systems design and engineering practice**, not computational problem-solving. You will find no algorithm puzzles and almost no performance engineering here — no profilers, and caching is deliberately deferred (a recorded decision you'll read about). If you want to get better at data structures, this is the

wrong book; if you want to get better at *building software that survives contact with production and with other people*, it is the right one.

The rhythm

Every lesson follows the same beat, and the beat is the curriculum:

- 1 Motivate** — a concrete problem, ideally demonstrated by a failing or *leaking* test.
- 2 Red** — write the test that pins the behavior you want. It fails.
- 3 Green** — write the minimum code to pass it.
- 4 Refactor & harden** — extract the seam; add the guardrail that makes the mistake

impossible to repeat.

- 1 Run it** — commands and expected output; every lesson ends with the app working.
- 2 Architecture Decision** — the fork, the choice, the road not taken.
- 3 Checkpoint** — a git tag; what you should now be able to do.

Test-driven development is not a chapter in this book; it is the loop you will live roughly fifty times, exactly as the reference platform was actually built.

How to use this book


The companion repository. Alongside the book there is a companion repo built in teaching order, with one tagged checkpoint per lesson (`lesson-1.3`, `lesson-2.6`, ...). Type everything yourself — that is the course — but when you are stuck, `git diff` your work against the checkpoint, or check it out and keep moving. Falling behind the typing is recoverable; skipping the understanding is not.

Don't skip Part 0. Lesson 0.2 pins your toolchain and infrastructure so that the other forty-nine lessons behave identically on your machine. Most abandoned courses die at "works on my machine"; we kill that failure mode first.

When a lesson cites a rule (R1–R35) or an ADR, that is a pointer into the reference platform's own governance documents — the machine-enforced invariants and the decision log. You are not just building an app; you are building the *system that keeps the app honest*, and those citations show you where each piece of it came from.

Reading order is dependency order. Later lessons lean on earlier seams without re-explaining them. If Part 5 feels steep, the fix is usually two parts back, not a re-read of Part 5.

The platform at a glance



One sentence to keep for the whole book: **app data belongs to the tenant, not the user** — and by the end of Part 2 you will have made it impossible, at three separate layers, for any line of feature code to violate that sentence.

Part 0 — Orientation

Lesson 0.1 — The mental model & the decision record

Part 0 · Orientation. No code in this lesson — deliberately. You're about to type several hundred files over ~48 lessons. The two things that determine whether that makes you an architect or a very tired typist are (1) the mental model you hang each file on, and (2) the habit of recording *why* before *what*. Both are installed here.

Goal: you can draw the system from memory — its layers, its one non-negotiable invariant, and its unit of work — and you've written your project's first ADR.

Concepts & principles: multi-tenancy (tenant $\neq$ user); clean architecture's dependency rule; vertical slices vs. horizontal layers; Architecture Decision Records.

Maps to: ADR-C1/C2/C3, ADR-004 · docs/DECISIONS.md, docs/WAYS_OF_WORKING.md in the reference repo.

1. What you are going to build

A production-grade, multi-tenant SaaS **platform**: custom JWT auth (passwordless + OAuth + MFA), tenant isolation enforced by the type system, billing with Stripe, reliable async via a transactional outbox, RBAC, file storage, GDPR export/erasure, a public API, outbound webhooks, an admin back-office, observability — deployed to a real host at the end. Not a to-do app. The kind of chassis a real product bolts features onto.

You will build it **test-first, in vertical slices**, and every architectural rule will be enforced by a test you wrote — not by a comment hoping to be read.

2. The four ideas everything else hangs on

Idea 1 — Tenant $\neq$ user

The single most important sentence in this course: **app data belongs to the tenant, not the user**. A *tenant* is the paying unit — a household, a team, a company. Users are members of a tenant; they come and go, and the tenant's data stays. Get this backwards and you'll design tables, queries, and permissions around the wrong axis — and unwinding that later is a rewrite, not a refactor.

One consequence you'll meet in Lesson 2.6 and never stop meeting: every read and write of app data must be scoped to the current tenant, and that scoping must be **impossible to forget**, not merely encouraged.

Idea 2 — The dependency rule (clean architecture, minus the ceremony)

Three production layers, dependencies pointing strictly inward:

```
Api → Infrastructure → Core
(HTTP, auth, (EF Core, SMTP, (entities, abstractions –
  controllers, Stripe, S3, the depends on NOTHING)
  features) implementations)
```

Core holds the entities and the *seams* — small interfaces like `IMailSender`, `IFileStorage`, `IBillingProvider`. Infrastructure implements them with real technology (MailKit, S3, Stripe). Api orchestrates. The payoff is swappability where it matters (Mailpit in dev, Brevo in prod — same seam) and testability everywhere. The UI (Blazor WebAssembly) is *just another API client* — it never touches the database.

Idea 3 — Clean platform + vertical slices (the hybrid)

Pure layered architecture scatters one feature across four folders. Pure feature-folders duplicate the platform in every slice. This codebase does both, deliberately, and the split is the point:

- The **platform** — auth, tenancy, billing, email, persistence — is the durable chassis,

organized in clean layers. Built once, reused by everything.

- Each **app feature** lives in ONE folder (`src/Api/Features/<Feature>/`) containing its

endpoints, handler, DTOs, and data contributor. Features *reuse* the platform; they never reach into each other.

When you wonder "where does this file go?", the question to ask is: *chassis or feature?*

Idea 4 — Vertical slices as the unit of work

Every increment of work cuts through all layers — DB to UI — and leaves the app working. You'll never build "all the tables" then "all the endpoints." Each lesson in this course is itself a slice: it ends with a green build, passing tests, and a tagged checkpoint.

3. The decision record — the habit that makes it stick

Here's what separates a codebase you can *maintain* from one you can only *fear*: six months from now, the question is never "what does this code do?" (you can read it), it's "**why is it like this — and is it safe to change?**"

An **Architecture Decision Record** answers that in ~10 lines, written at the moment of decision:

```
## ADR-002 – Custom JWT + rotating refresh tokens (not ASP.NET Core Identity)

**Date:** 2026-06-XX • **Status:** Accepted

**Context:** We need auth for web + native clients with passwordless and OAuth.
Identity brings cookie-centric defaults, schema we don't control, and hides the
token lifecycle we need to reason about (MFA step-up, rotation, native flows).

**Decision:** Issue our own short-lived JWT access tokens + rotating refresh
tokens; store only token hashes.

**Consequences:** We own token security (rotation, reuse detection – must test
it ourselves). In exchange: full control of claims (tenant_id drives data
isolation), one auth model across web/native, no framework fight when MFA lands.
```

Note what an ADR is *not*: not documentation of how the code works (the code shows that), and never deleted when wrong — a reversed decision gets a **new dated ADR** superseding the old one. The history of your reasoning is itself an asset.

The reference repo runs on this: 20 app ADRs + amendments in `docs/DECISIONS.md`. Every lesson ahead ends with an **Architecture Decision box** — the fork faced, the option chosen, the roads not taken. By Part 4 you should be predicting them before reading them; that reflex *is* the architectural skill this course exists to build.

4. Exercise — write your ADR-000 and ADR-001

Create your project directory (empty — the solution comes in Lesson 1.1) and start the decision log:

```
mkdir my-saas && cd my-saas && git init
mkdir docs && $EDITOR docs/DECISIONS.md
```

Write two ADRs yourself — don't copy:

1 ADR-000 — Why I'm building this. Context (what you want to learn), decision (build the platform from scratch, test-first), consequences (weeks of effort; deep ownership of every line).

1 **ADR-001 — Local secrets live in `.env`, never in the repo.** You haven't built the loading mechanism yet (that's 0.2 and 1.4) — which is exactly the point: **the decision precedes the code**. Context: secrets in committed files leak (git history is forever). Decision: a gitignored `.env` is the single local source of truth; a committed `.env.example` documents every key; production uses real environment variables. Consequences: one more setup step per machine; zero secrets in history — enforced by a scanner you'll add in the next lesson.

5. Checkpoint

```
git add docs/DECISIONS.md && git commit -m "docs: ADR-000 and ADR-001 — the decision log
starts before the code"
git tag lesson-0.1
```

You should now be able to: state the tenant≠user invariant and its consequence for every future query; draw the three layers and their dependency direction; say which of the two organizational modes (chassis vs. feature folder) a given file belongs to; and write an ADR that captures context → decision → consequences in under a page.

Next: *Lesson 0.2 — A reproducible machine*, where your ADR-001 gets its mechanical teeth: `.env.example`, `docker-compose`, and a secret scanner.

Lesson 0.2 — A reproducible machine

Part 0 · Orientation. Most from-scratch courses die at "install PostgreSQL." This lesson exists so yours can't. At the end, your machine runs the exact infrastructure the app will use for the next 47 lessons — pinned, containerized, verifiable with three commands — and your ADR-001 (secrets never in the repo) is enforced by machinery, not memory.

Goal: `docker compose up -d` gives you a healthy Postgres 17 and a mail trap; `dotnet --version` prints the pinned SDK; a secret can't reach a commit.

Concepts & principles: pinned toolchains ("latest stable, never previews" — and never "whatever was on the machine"); infrastructure-as-code for dev; the `.env` contract (ADR-001); fail-loud health checks; defense-in-depth for secrets.

Maps to: ADR-001, ADR-C13 · Foundation theme R20–R22 (the config contract this enables) · repo: `docker-compose.yml`, `.env.example`, `.gitignore`, `.gitleaks.toml`.

Prerequisites: 0.1 (your repo exists and ADR-001 is written).

1. Motivate — "works on my machine" is an architecture failure

Two developers clone the same repo. One has Postgres 15 installed as a Windows service from a project two jobs ago; the other has nothing. One gets a subtle JSON-function bug that doesn't reproduce; the other gets a connection refused. Neither failure is *in the code* — and that's the point: **your dev environment is part of your architecture**, and undeclared dependencies are its silent killers.

The fix is the same one you'll apply to code all course long: make the implicit explicit, and make violations loud.

2. Install the pinned toolchain

Three tools go on the machine itself; everything else runs in containers:

Tool	Version	Check
.NET SDK	10.0.3xx (the line this course pins)	<code>dotnet --version</code>
Docker Desktop / Engine	current stable	<code>docker version</code>
EF Core tooling	current stable	<code>dotnet tool install --global dotnet-ef</code> (first used in 1.3)

Pin the SDK **in the repo**, not in a README nobody reads — create `global.json`:

```
{
  "sdk": { "version": "10.0.301", "rollForward": "latestFeature" }
}
```

Now a machine with the wrong SDK fails at `dotnet build` with an explicit message instead of failing three lessons later with an inscrutable one. This is the course's recurring move — *fail*

loud and early — applied for the first time, to the toolchain itself.

3. Infrastructure as one file — docker-compose

The app needs PostgreSQL and — less obviously — an **SMTP trap**. From Lesson 2.3 onward the platform sends real email (magic links, OTPs, invitations), and you need somewhere for it to land in dev that isn't a real inbox. [Mailpit](#) accepts SMTP on one port and shows every message in a web UI on another — and, critically for Lesson 3.6, exposes an **HTTP API your tests will query** to fish out OTP codes. Two infrastructure decisions, one committed file — create `docker-compose.yml`:

```
services:
  db:
    image: postgres:17
    restart: unless-stopped
    environment:
      POSTGRES_DB: ${DB_NAME:-dev_db}
      POSTGRES_USER: ${DB_USER:-dev}
      POSTGRES_PASSWORD: ${DB_PASSWORD:-devpassword}
    ports:
      - "${DB_PORT:-5432}:5432"
    volumes:
      - db_data:/var/lib/postgresql/data
    healthcheck:
      test: ["CMD-SHELL", "pg_isready -U ${DB_USER:-dev} -d ${DB_NAME:-dev_db}"]
      interval: 5s
      timeout: 5s
      retries: 5
      start_period: 10s

  mail:
    image: axllent/mailpit:latest
    restart: unless-stopped
    ports:
      - "${MAIL_SMTP_PORT:-1025}:1025" # SMTP — the app sends here
      - "${MAIL_UI_PORT:-8025}:8025" # Web UI + API — you and your tests read here
    healthcheck:
      test: ["CMD-SHELL", "wget -q --spider http://localhost:8025 || exit 1"]
      interval: 5s
      timeout: 5s
      retries: 5
      start_period: 5s

volumes:
  db_data:
```

Read it as a set of decisions, because each line is one:

- **postgres:17** — a pinned major version, mirroring `global.json`'s job for the SDK.

"Latest stable, never previews" also means *never unpinned*.

- **\${DB_PORT:-5432}** — every value is overridable via environment, with a working

default. A developer whose port 5432 is squatted by that old Windows service sets `DB_PORT=5433` in `.env` and moves on. (The reference repo defaults to 5433 for exactly that reason — port collisions are the #1 setup killer.)

- **healthcheck** — "container started" ≠ "Postgres accepting connections." Health

checks make readiness explicit, and anything that later depends_on these services waits for *healthy*, not merely *running*. You'll re-meet this exact idea as the API's own /health endpoint in Lesson 4.5.

- **volumes: db_data** — your data survives docker compose down. (Destroying it

becomes a deliberate act: `docker compose down -v`.)

4. The .env contract — ADR-001 gets teeth

Your ADR-001 said: secrets live in a gitignored .env; a committed .env.example documents every key. Build both halves now.

.gitignore (start; it grows with the solution in 1.1):

```
.env
bin/
obj/
storage/
```

.env.example — **committed**, and it is *documentation with a job*:

```
# Copy this file to .env and fill in your values. .env is gitignored — never commit it.
# Single source of local-dev config + secrets: read by docker-compose (containers) AND
# by the API via DotNetEnv (from Lesson 1.4). See docs/DECISIONS.md ADR-001.

# — Containers (docker-compose) —————
DB_NAME=dev_db
DB_USER=dev
DB_PASSWORD=devpassword
DB_PORT=5433

# Mailpit (SMTP trap)
MAIL_SMTP_PORT=1025
MAIL_UI_PORT=8025
```

Then `cp .env.example .env` and confirm `git status` does **not** list .env.

Notice the shape of the contract: *every* key the system reads appears here with an inline comment saying what it does, whether it's required, and what happens when unset. In Lesson 1.4 you'll write a test (DocAndConfigSyncTests) that **fails the build when a config key exists in code but not in this file** — turning "please document your config" from a review nag into a compile-adjacent guarantee (R20). The habit starts now.

Docker-compose reads .env automatically. So one file drives both the containers and — once DotNetEnv arrives in 1.4 — the app. Two consumers, one source of truth, zero drift.

The second layer: a secret scanner

.gitignore is advisory — it does nothing if someone hardcodes a connection string in a .cs file. Add [gitleaks](#) config, .gitleaks.toml:

```
# Secret scanning – the backstop behind ADR-001. Default rules + our allowlist.
[extend]
useDefault = true

[allowlist]
description = "Dev-only placeholder values that look like secrets but aren't"
paths = [''\.env\example'']
```

Run it locally (`gitleaks detect`) whenever you like; in Lesson 8.3 it becomes a CI gate. Layers, again: `.gitignore` prevents the accident, `gitleaks` catches what slips past, CI makes it policy. No single layer is trusted alone — the same defense-in-depth shape you'll apply to tenancy (filter + interceptor + arch test) from Part 2 on.

5. Run it — verify everything

```
docker compose up -d
docker compose ps # both services: status "healthy" (wait a few seconds)
dotnet --version # 10.0.3xx – matches global.json
docker compose exec db psql -U dev -d dev_db -c "select version();" # PostgreSQL 17.x
```

Open `http://localhost:8025` — Mailpit's empty inbox. In Lesson 2.4, your app's first magic-link email lands here; in 3.6, your Playwright tests read OTP codes out of it through its API.

6. Architecture Decision

ARCHITECTURE DECISION

The fork: where does dev infrastructure live? (a) Installed natively per machine, (b) a shared cloud dev database, (c) containers defined in the repo.

Chosen: containers (ADR-C13). The infrastructure's *versions and topology* are code, reviewed like code; a new machine is productive in minutes; two projects with different Postgres majors coexist.

Rejected: native installs — unpinned, colliding, undocumented by nature. Shared cloud dev DB — couples every developer's iteration loop to a network and to each other's data (and costs money before the first feature exists).

The trade: Docker becomes a hard dev dependency. Accepted knowingly — it's also the test-isolation strategy (Testcontainers, Lesson 1.3) and the deployment unit (8.2), so the dependency pays three times.

7. Checkpoint

```
git add global.json docker-compose.yml .gitignore .gitleaks.toml .env.example
git commit -m "chore: reproducible dev environment – pinned SDK, compose infra, env contract"
git tag lesson-0.2
```

You should now be able to: bring the full dev infrastructure up and down with one command; explain why the SDK and Postgres versions are pinned *in the repo* rather than in a README; state the `.env` / `.env.example` contract and the two layers that enforce it; and articulate why health checks matter before anything depends on these services.

Next: *Lesson 1.1 — Solution, projects & supply chain* — the .NET solution appears, with its dependencies pointing the right way and its packages locked.

Part 1 — The walking skeleton

Lesson 1.1 — Solution, projects & supply chain

Part 1 · The walking skeleton. Nothing runs yet — that's next lesson. Here we lay out the solution so that the architecture from Lesson 0.1 isn't a diagram you promised to respect but a set of project references the compiler enforces. We also make a decision most projects defer until it hurts: exactly which versions of everything we depend on, pinned so a build today and a build in six months produce the same bytes.

Goal: a solution with every project in place, dependencies physically pointing inward, one central version table, and a locked, reproducible restore.

Concepts & principles: the dependency rule made mechanical; Central Package Management; lockfiles and supply-chain hygiene; warnings-as-errors from day one.

Maps to: ADR-C3/C4/C5 · R25–R27 · repo: `Perezosoftware.slnx`, `Directory.Build.props`, `Directory.Packages.props`, `*/*.csproj`.

Prerequisites: 0.2 (pinned SDK; `dotnet --version` prints 10.0.3xx).

1. Motivate — the layout is the first enforcement

Every codebase claims to have an architecture. The honest question is: *what happens if someone violates it?* If the answer is "a reviewer might notice," you have a convention. If the answer is "it does not compile," you have a structure.

Project references are the cheapest structural enforcement .NET offers. When `Core` has no reference to `Infrastructure`, a domain entity *cannot* reach out to EF Core or SMTP — not "should not," cannot. The dependency rule from Lesson 0.1 stops being a poster on the wall the moment the reference graph encodes it. That's what we build in this lesson.

The second problem we solve is quieter and nastier. `dotnet add package` with no version pinned means your build depends on *when* you ran it. Two developers restore on different days and get different transitive graphs; CI builds an artifact subtly unlike the one you tested; and when a compromised package version hits NuGet — it has happened — nothing in your build even notices. Reproducibility isn't a nicety; it's the precondition for being able to say "this is what we shipped."

2. Create the skeleton

The solution uses the modern XML solution format (`.slnx` — human-readable, merge-friendly, supported by .NET 10 tooling):

```
dotnet new sln --format slnx -n Perezosoftware
dotnet new classlib -o src/Core -n Perezosoftware.Core
dotnet new classlib -o src/Infrastructure -n Perezosoftware.Infrastructure
dotnet new web -o src/Api -n Perezosoftware.Api
dotnet new blazorwasm -o src/Web -n Perezosoftware.Web
dotnet new razorclasslib -o src/Shared.Ui -n Perezosoftware.Shared.Ui
dotnet new xunit -o tests/Core.Tests -n Perezosoftware.Core.Tests
dotnet new xunit -o tests/Api.Tests -n Perezosoftware.Api.Tests
dotnet new nunit -o tests/E2E.Tests -n Perezosoftware.E2E.Tests
dotnet sln add src/Core src/Infrastructure src/Api src/Web src/Shared.Ui \
    tests/Core.Tests tests/Api.Tests tests/E2E.Tests
```

Delete the template's sample classes (Class1.cs, the Blazor counter page can stay for now — it goes when the real UI arrives in 3.4). Then wire the references — this is the architecture, so read it slowly:

```
dotnet add src/Infrastructure reference src/Core # Infra implements Core seams
dotnet add src/Api reference src/Core src/Infrastructure # Api composes both
dotnet add src/Web reference src/Shared.Ui # Web hosts the shared UI
dotnet add tests/Core.Tests reference src/Core
dotnet add tests/Api.Tests reference src/Api src/Infrastructure src/Core
```

What's deliberately **absent** is the point: Core references *nothing* of ours. Infrastructure cannot see Api. Web cannot see Core, Infrastructure, or Api at all — the UI will talk to the API over HTTP like any other client (Golden Rule: the clean API boundary), and the compiler now holds it to that. Try it: add `using Perezosoftware.Infrastructure;` to a file in Core — it does not build. That error message is the architecture working.

Each project file stays nearly empty. Here is the real `Perezosoftware.Core.csproj` in its entirety:

```
<Project Sdk="Microsoft.NET.Sdk">

  <PropertyGroup>
    <TargetFramework>net10.0</TargetFramework>
    <ImplicitUsings>enable</ImplicitUsings>
    <Nullable>enable</Nullable>
  </PropertyGroup>

</Project>
```

Nine lines. Everything else lives at the solution root, where it applies to *all* projects at once — which is our next move.

3. One build policy for the whole solution

MSBuild reads a file named `Directory.Build.props` from every ancestor directory of a project. Put one at the repo root and every project inherits it — a single place for build policy, impossible for a new project to forget. This is the reference platform's actual file:

```
<Project>

  <!--
    Solution-wide build hygiene. Nullable reference types on, and warnings are
    errors so the invariants the analyzers and the compiler check can't rot into
    ignored noise. This is half of "solid by construction" – the architecture
    tests are the other half, and CI runs both.
  -->
  <PropertyGroup>
    <Nullable>enable</Nullable>
    <TreatWarningsAsErrors>true</TreatWarningsAsErrors>
    <ImplicitUsings>enable</ImplicitUsings>
  <!--
    Supply-chain hardening. Restore writes a packages.lock.json next to each
    project pinning the full transitive graph; those files are committed and CI
    restores in locked mode so any drift fails the build. Paired with Central
    Package Management (Directory.Packages.props).
  -->
    <RestorePackagesWithLockFile>true</RestorePackagesWithLockFile>
  </PropertyGroup>

</Project>
```

Two of these settings deserve a hard look, because both are *decisions about failure*:

****TreatWarningsAsErrors**** — a warning is the compiler telling you something is probably wrong. On a codebase that tolerates warnings, the count creeps from 3 to 300, and the one that mattered drowns. Turning warnings into errors on day one costs nothing (there is no code yet!) and means the count stays at zero forever. Adopting this on an *existing* codebase is a month of cleanup — which is exactly why it must happen now, in the empty solution. Several rules you'll meet later (like nullability soundness) are only real because a violation refuses to compile.

****Nullable**** — nullable reference types turn "I forgot this could be null" from a production `NullReferenceException` into a compile-time squiggle — and with the previous setting, into a build failure. The domain modeling you'll do from Part 2 onward leans on this constantly: `string?` in an entity is a *documented decision* that the value is optional, checked by the compiler, rather than tribal knowledge.

4. Central Package Management — one version table

Without CPM, every `.csproj` carries its own `<PackageReference Include="X" Version="Y">` and the same package drifts across projects (the reference repo actually hit this: two test projects referencing different versions of the test SDK). With CPM, versions live in **one** committed table, `Directory.Packages.props`, and project files reference packages *without* versions. An excerpt of the real table:

```

<Project>
  <PropertyGroup>
    <ManagePackageVersionsCentrally>true</ManagePackageVersionsCentrally>
  </PropertyGroup>

  <ItemGroup>
    <!-- App / API -->
    <PackageVersion Include="DotNetEnv" Version="3.2.0" />
    <PackageVersion Include="Microsoft.AspNetCore.Authentication.JwtBearer"
Version="10.0.9" />

    <!-- Infrastructure -->
    <PackageVersion Include="MailKit" Version="4.17.0" />
    <PackageVersion Include="Npgsql.EntityFrameworkCore.PostgreSQL" Version="10.0.2" />

    <!-- Tests -->
    <PackageVersion Include="xunit" Version="2.9.3" />
    <PackageVersion Include="Testcontainers.PostgreSql" Version="4.12.0" />
    <!-- Consolidated: E2E referenced 17.14.0, the xUnit projects 17.14.1 → highest wins.
-->
    <PackageVersion Include="Microsoft.NET.Test.Sdk" Version="17.14.1" />
  </ItemGroup>
</Project>

```

Create yours with just the packages the walking skeleton needs (the test SDK, xunit); every later lesson that says "add package X" means: add the `<PackageVersion>` row here, add a version-less `<PackageReference>` in the project that uses it. One table to review when you upgrade; one diff line when a version changes; zero possibility of the same package at two versions (that's rule R27, and it's enforced by construction).

Notice that comment about 17.14.0 vs 17.14.1 — it's a fossil of the exact drift CPM exists to prevent, found *during* the consolidation and preserved in the reference repo as a comment. Real codebases carry these scars; good ones label them.

5. Lock it — restore becomes reproducible

`RestorePackagesWithLockFile` (already in our props file) makes restore write a `packages.lock.json` beside each project: the **full transitive closure** — every package your packages depend on, recursively, with exact versions and content hashes. Commit these files. Then the magic flag:

```

dotnet restore --use-lock-file # first restore: writes the lockfiles
git add **/packages.lock.json
dotnet restore --locked-mode # from now on (and in CI): verify, don't resolve

```

In locked mode, restore doesn't *resolve* versions — it *verifies* them. If anything in the graph differs from the lockfile (a floating version moved, a transitive dependency was swapped, a hash doesn't match), restore **fails loudly** instead of silently building something new. When lesson 1.6 stands up CI, `--locked-mode` is one of its first gates: a supply-chain change becomes a red build and a reviewable diff, not an invisible drift.

The mental model: your dependency graph is now *code*. It changes by commit, with review, or not at all.

6. Run it

```
dotnet build
# 0 Warning(s), 0 Error(s) – and it must stay that way: warnings ARE errors now
dotnet restore --locked-mode
# succeeds – graph matches the committed lockfiles
```

And prove a gate actually gates (pain now, so you trust it later): edit a `packages.lock.json` by hand — change any version digit — and run `dotnet restore --locked-mode` again. It fails with `NU1004`. Revert. That failure is what a supply-chain surprise will look like: loud, early, and in your terminal instead of in production.

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how do we manage dependency versions? (a) Per-project references with whatever versions were current at add-time; (b) a shared props file with version *variables*; (c) Central Package Management + committed lockfiles + locked-mode CI.

Chosen: (c) — CPM for one reviewable version table (R25), lockfiles for the transitive graph, locked restore so drift fails the build (R27). Also decided here: warnings-as-errors and nullable from the empty solution, because both are nearly free now and prohibitively expensive to retrofit.

Rejected: (a) is how version drift and silent supply-chain changes happen — it's the default, which is the problem. (b) centralizes the *numbers* but not the *graph*: transitive dependencies still float.

The trade: every dependency bump is now an explicit two-file diff (table + lock), and a new package is a two-step add. Accepted: that friction is the review point.

8. Checkpoint

```
git add -A && git commit -m "feat: solution skeleton – projects, inward refs, CPM, locked
restore"
git tag lesson-1.1
```

You should now be able to: explain why Core referencing nothing is the architecture's first enforcement; add a package the CPM way; say what a lockfile pins that a version number doesn't; and articulate why warnings-as-errors is a day-one decision, not a someday cleanup.

Next: Lesson 1.2 — *First endpoint, first test* — the API says hello, and the Red→Green loop you'll live fifty times gets its first turn.

Lesson 1.2 — First endpoint, first test

Part 1 · The walking skeleton. The solution compiles but does nothing. By the end of this lesson it answers HTTP — and, more importantly, a test *proves* it answers HTTP by booting the real application. This is the lesson where you live the Red→Green→Refactor loop for the first time, on the smallest possible stakes, so that when the stakes are tenant isolation or billing you already trust the rhythm.

Goal: GET /health returns 200, and a test in `Api.Tests` boots the actual app in-process and asserts it.

Concepts & principles: the walking skeleton; Red→Green→Refactor; in-process integration testing with `WebApplicationFactory`; testing the real composition, not a hand-assembled copy.

Maps to: ADR-C14 · repo: `src/Api/Program.cs`,
`tests/Api.Tests/Infrastructure/IntegrationTestFactory.cs`,
`tests/Api.Tests/Integration/HarnessSmokeTests.cs`, `src/Api/Perezosoft.Api.http`.

Prerequisites: 1.1 (solution builds, CPM in place).

1. Motivate — why a "hello world" deserves a whole lesson

A *walking skeleton* is the thinnest possible slice that exercises the whole path: request in, response out, test proving it. It looks trivial. It isn't, because it forces every piece of plumbing to exist and cooperate **before** any feature depends on that plumbing: the web host boots, routing routes, the test project can reach the app, CI (next lessons) can run it all. Every integration problem you flush out now is one that won't ambush you inside a complicated lesson later.

There's also a decision hiding in "a test proving it" — *what kind* of test? We could unit-test a handler function and call it a day. But the failures that hurt in web apps are rarely inside one function; they're in the *composition* — DI registrations missing, middleware ordered wrong, routes not mapped, config not read. So our default harness boots the **real** `Program` — the actual composition root, not a hand-assembled lookalike — in-process, and talks to it over real HTTP semantics. The reference platform's harness makes the reasoning explicit; from its actual doc comment:

Boots the REAL app (`Program`) via `WebApplicationFactory<T>` ... so tests exercise the full pipeline ..., not hand-assembled services. ... [this] de-risks routing/DI refactors by asserting routes stay stable end-to-end.

A test against a hand-wired `ServiceCollection` keeps passing while the real app is broken. A test against `Program` cannot.

2. Red — the failing test comes first

In `tests/Api.Tests`, add the package references (CPM style — versions come from the table you made in 1.1): `Microsoft.AspNetCore.Mvc.Testing`, and reference the `Api` project. Then write the test **before** the endpoint exists:

```
// tests/Api.Tests/HealthEndpointTests.cs
using System.Net;
using Microsoft.AspNetCore.Mvc.Testing;

namespace Perezosoft.Api.Tests;

public class HealthEndpointTests : IClassFixture<WebApplicationFactory<Program>>
{
    private readonly WebApplicationFactory<Program> _factory;

    public HealthEndpointTests(WebApplicationFactory<Program> factory) => _factory =
factory;

    [Fact]
    public async Task Health_endpoint_responds_200()
    {
        var client = _factory.CreateClient(); // in-process; no port, no network

        var response = await client.GetAsync("/health");

        Assert.Equal(HttpStatusCode.OK, response.StatusCode);
    }
}
```

Run it:

```
dotnet test tests/Api.Tests
# Failed! Health_endpoint_responds_200
# Assert.Equal() Failure: Expected OK, Actual NotFound
```

Red — the app boots (`WebApplicationFactory` found `Program` and started the real host) but there is no `/health` endpoint, so the request falls through to a 404. That's a *properly* failing test: it fails for exactly the reason it should, and the failure message describes the missing behavior.

One wrinkle to know about before it bites: our `Program.cs` is top-level statements (no explicit class), and the class the compiler generates for it is **internal**. On current SDKs the web template arranges for the test host to reach it anyway, so the test above compiles and runs — but if you ever hit error CS0246: The type or namespace name 'Program' could not be found, the fix is one line in `Perezosoft.Api.csproj`, and the reference platform carries it as standing hygiene:

```
<ItemGroup>
  <InternalsVisibleTo Include="Perezosoft.Api.Tests" />
</ItemGroup>
```

`InternalsVisibleTo` is a scalpel: it grants *one named assembly* access to internals, instead of making `Program` public for the whole world. You'll see this same "smallest-possible opening" instinct again and again — it's the security posture of the entire platform in miniature. Add it now even though the build didn't force you to; the next person's SDK might.

3. Green — the minimum that passes

src/Api/Program.cs, in full at this stage of its life:

```
var builder = WebApplication.CreateBuilder(args);

builder.Services.AddHealthChecks();

var app = builder.Build();

app.MapHealthChecks("/health");

app.Run();
```

```
dotnet test tests/Api.Tests
# Passed! - 1 test
```

That's the loop: a failing test that *specifies*, then the minimum code that *satisfies*. Resist the urge to add more — no controllers, no Swagger, no logging config. Every one of those will arrive in the lesson where a test demands it. (For the curious: this file grows to ~300 lines by Part 8 in the reference repo, and every block of it traces to a lesson.)

Why AddHealthChecks()/MapHealthChecks rather than a hand-rolled app.MapGet("/health", ...)? Because this endpoint has a *career* ahead of it: in Lesson 4.5 it splits into liveness (/health) and readiness (/health/ready) with a real database probe behind a tag filter, and in Lesson 8.3 the deploy pipeline's smoke gate calls it to decide whether a release goes out. Using the framework's health-check pipeline from day one means that growth is additive. The reference repo's final form, for a preview of where this line of code is headed:

```
app.MapHealthChecks("/health", new HealthCheckOptions { Predicate = _ => false });
app.MapHealthChecks("/health/ready", new HealthCheckOptions { Predicate = check =>
    check.Tags.Contains("ready") });
```

4. Refactor & harden — a scratchpad and a habit

No production refactor needed yet (there's barely any production). But two small work-habits start here.

****The .http scratchpad.**** Create src/Api/Perezosoftware.Api.http — using the https port from *your* src/Api/Properties/launchSettings.json (dotnet new web assigns a random one; the reference repo happens to sit on 7160):

```
@host = https://localhost:7160

GET {{host}}/health
```

Visual Studio and VS Code (REST Client) execute these files in-editor. Every lesson that adds an endpoint adds a request here — by Part 7 it's a living, executable catalog of the whole API surface, versioned with the code. It's the cheapest API documentation that can't go stale, because you *use* it to poke the app.

Run the real thing. Tests passing in-process is necessary, not sufficient — actually boot it once:

```
dotnet run --project src/Api
# listening on https://localhost:xxxx
curl -k https://localhost:xxxx/health
# Healthy
```

The response body `Healthy` comes from the health-check middleware's default writer.
In-process tests + a real boot = the two views you'll keep cross-checking all course.

5. Run it

```
dotnet build # 0 warnings, 0 errors – still, and always
dotnet test tests/Api.Tests
# 1 passed
```

6. Architecture Decision

ARCHITECTURE DECISION

The fork: what does the default test harness boot? (a) Nothing — unit-test handlers as plain functions; (b) a hand-assembled `ServiceCollection` with the pieces each test needs; (c) the real `Program` via `WebApplicationFactory`, in-process.

Chosen: (c) as the default for anything HTTP-shaped (ADR-C14's integration layer). The composition root is where web apps actually break — missing registrations, middleware order, unmapped routes — and only booting the real one tests it. In-process means no ports, no network flakes, parallel-safe.

Rejected: (a) alone leaves the wiring untested (we still unit-test pure logic — `Core.Tests` exists for exactly that). (b) is the seductive trap: it *looks* like an integration test but silently diverges from the real app — a test passing against a copy proves nothing about the original.

The trade: booting the real app makes some things awkward on purpose — config the app reads at startup must genuinely exist (you'll feel this in 1.4, and the reference harness has a documented seam for it: the JWT secret must be an *environment variable* because `Program` reads it before the factory's config hooks run). That awkwardness is information: it's your startup contract made visible.

7. Checkpoint

```
git add -A && git commit -m "feat: walking skeleton – health endpoint + real-app test harness"
git tag lesson-1.2
```

You should now be able to: explain what a walking skeleton de-risks; write the failing test first and feel why the order matters; say why `WebApplicationFactory<Program>` beats a hand-assembled service collection; and explain what `InternalsVisibleTo` grants and why it's preferable to making things public.

Next: *Lesson 1.3 — Database & the test container* — Postgres joins the skeleton, and the tests get a real database that appears and disappears on demand.

Lesson 1.3 — Database & the test container

Part 1 · The walking skeleton. The skeleton gets a spine: PostgreSQL via EF Core, a first migration, and — the real prize — a test fixture that gives every test run a genuine Postgres that appears in seconds and vanishes without trace. The single most consequential testing decision of the course happens here: **we do not use the EF in-memory provider. Ever.** By the end you'll know exactly why.

Goal: AppDbContext exists with one throwaway entity, a migration creates its schema, and a test writes to and reads from a *real* Postgres spun up by the test run itself.

Concepts & principles: EF Core migrations discipline; Testcontainers; why in-memory database fakes lie; model-derived test resets (isolation that can't be forgotten).

Maps to: ADR-C6/C7 · R11 · repo: `src/Infrastructure/Persistence/AppDbContext.cs`, `src/Api/AppDbContextFactory.cs`, `tests/Api.Tests/Infrastructure/PostgresFixture.cs`, `tests/Api.Tests/Infrastructure/PostgresTestBase.cs`, `tests/Api.Tests/MigrationsTests.cs`.

Prerequisites: 1.2 (harness runs); 0.2 (Docker running — Testcontainers needs the daemon, and you already have it for compose).

1. Motivate — your tests are only as honest as their database

Here's the trap nearly every .NET team falls into once. EF Core ships an in-memory provider; it's right there; tests using it need no Docker and run fast. So the test suite grows on it. Then one day a query that passed a thousand test runs fails in production, because the in-memory provider isn't a database — it's a `List<T>` wearing a trench coat. The reference platform's fixture states the consequence as a matter of record, in its actual doc comment:

Real Postgres (not the EF in-memory provider) is required because several services branch on `Database.IsRelational()` and use relational-only constructs — `ExecuteUpdateAsync`, savepoints — that the in-memory provider silently skips.

"Silently skips" is the killer phrase. And the list is longer than that comment: the in-memory provider has no transactions worth the name, no unique constraints, no `SKIP LOCKED` (Lesson 4.1 depends on it), case-insensitive string comparisons where Postgres is case-sensitive — and, fatally for this course, its handling of **global query filters** (Lesson 2.6, the keystone of tenant isolation) diverges from a real relational provider. A tenant-isolation test that passes in-memory proves *nothing*.

The classical escape was a shared "test database" on somebody's server — slow to reset, polluted by yesterday's runs, fought over by CI agents. Testcontainers dissolves the dilemma: the test run itself starts a pristine `postgres:17` in a Docker container, uses it, and destroys it. Real engine, zero residue.

2. The context and its first entity

In Infrastructure, add the EF packages (CPM rows + version-less references — `Npgsql.EntityFrameworkCore.PostgreSQL`; `Microsoft.EntityFrameworkCore.Design` goes in the Api project for tooling). Then the context — humble at this stage:

```
// src/Infrastructure/Persistence/AppDbContext.cs
using Microsoft.EntityFrameworkCore;

namespace Perezosoft.Infrastructure.Persistence;

public class AppDbContext : DbContext
{
    public AppDbContext(DbContextOptions<AppDbContext> options) : base(options) { }

    protected override void OnModelCreating(ModelBuilder builder)
    {
        base.OnModelCreating(builder);
        // Per-entity mapping lives in one IEntityTypeConfiguration<TEntity> class each,
        // discovered from this assembly – no central registry to forget.
        builder.ApplyConfigurationsFromAssembly(typeof(AppDbContext).Assembly);
    }
}
```

That `ApplyConfigurationsFromAssembly` line is a pattern you met in 0.2 (health checks) and will meet a dozen more times: **discovery over registration**. Entity configurations will be *found* because they exist, not because someone remembered to list them. (In 2.6, tenant filters attach the same way — by reflection over the model — and in this lesson's fixture, the reset list will too. One idea, three appearances already.)

For a first entity we need something disposable — the real platform's first entities (User, Tenant) arrive with Part 2 and deserve their own lessons. Use a placeholder `Probe` entity with an `Id` and a `Note` — its only job is to prove the pipeline, and you'll delete it in 2.1 without ceremony. Give it an `IEntityTypeConfiguration<Probe>` so the discovery pattern gets exercised.

Now scroll your build output before moving on — there's a landmine here worth defusing in daylight. Depending on package versions, adding the EF stack can produce:

```
warning MSB3277: Found conflicts between different versions of
"Microsoft.EntityFrameworkCore" that could not be resolved.
```

...and the build still says **succeeded**. Didn't we make warnings errors in 1.1? We did — for *compiler* warnings. MSB3277 is an **MSBuild** warning, a different pipeline that `TreatWarningsAsErrors` does not cover. Your zero-warnings bar has a blind spot, and this is the classic thing that hides in it: the `Npgsql` provider pulling one EF Core version while the `Design/test` packages pull another. The fix is to pin the shared assembly explicitly — add `Microsoft.EntityFrameworkCore.Relational` at the same version as your other EF packages to the CPM table and reference it from `Api.Tests` (the reference repo carries exactly this pin, with a comment saying why). The transferable lesson: **"the build is green" and "the build is clean" are different claims** — read the output when you add a dependency, because version-unification warnings are how runtime `MethodNotFound` surprises introduce themselves politely first.

3. Migrations — schema changes as reviewable code

The EF tooling is a separate global tool, not part of the SDK — install it once (0.2's toolchain table now has a third row):

```
dotnet tool install --global dotnet-ef
dotnet ef migrations add InitialCreate --project src/Infrastructure --startup-project
src/Api
```

...and it fails, or worse, half-works, because EF's design-time tooling wants to *boot your app* to obtain a configured `DbContext` — and your app (from 1.4 onward) will refuse to boot without validated config. The reference platform closes this with a **design-time factory**, and its doc comment explains the whole story:

Design-time factory so EF tooling ... can build the context WITHOUT booting the API host — routing through the host runs startup config validation (e.g. `Jwt:Secret`) that isn't present at design time. The connection string comes from `ConnectionStrings__DefaultConnection`, the SAME single source the running app uses ... There is no hardcoded fallback.

```
// src/Api/AppDbContextFactory.cs (abridged from the real file)
public class AppDbContextFactory : IDesignTimeDbContextFactory<AppDbContext>
{
    public AppDbContext CreateDbContext(string[] args)
    {
        try { DotNetEnv.Env.TraversePath().Load(); } catch { /* no .env present */ }

        var connectionString =
            Environment.GetEnvironmentVariable("ConnectionStrings__DefaultConnection")
            ?? throw new InvalidOperationException(
                "ConnectionStrings__DefaultConnection is not set. Add it to .env ...");

        var options = new DbContextOptionsBuilder<AppDbContext>()
            .UseNpgsql(connectionString).Options;
        return new AppDbContext(options);
    }
}
```

Note what it *refuses* to do: no hardcoded `Host=localhost`; ... fallback. The connection string lives in exactly one place — your `.env` (ADR-001, given teeth in 0.2) — and if it's missing, the error says precisely what to fix. A fallback would "work" and thereby hide misconfiguration until it mattered.

Which means two files gain a line right now, honoring 0.2's contract that every key the system reads is documented. In `.env` (so the tooling works on your machine) and in `.env.example` (so it works on the next machine):

```
# REQUIRED — Postgres connection for the API + EF tooling. Mirror the DB_* values above
# (note the port: 5433 if you kept 0.2's collision-avoiding default).
ConnectionStrings__DefaultConnection="Host=localhost;Port=5433;Database=dev_db;Username=dev
;Password=devpassword"
```

(When 1.4's doc-sync gate arrives, an undocumented key like this would fail the build — we're just early.)

Run the migration command again; a `Migrations/` folder appears in `Infrastructure`. Read the generated file once, top to bottom — you'll rarely need to again, but you should know what the tool writes on your behalf: an `Up()` building your tables, a `Down()` reversing it, and a model snapshot the *next* migration diffs against. Migrations are why "the schema" is never a rumor: it's the sum of committed, reviewed, replayable scripts.

4. The fixture — a real database that can't leak between tests

Now the centerpiece. Add `Testcontainers.PostgreSql` to `Api.Tests` and build the fixture (this is the reference platform's real code, trimmed of the tenancy parts that arrive in Part 2):

```
// tests/Api.Tests/Infrastructure/PostgresFixture.cs
public sealed class PostgresFixture : IAsyncLifetime
{
    private readonly PostgreSQLContainer _container = new
    PostgreSQLBuilder("postgres:17").Build();

    public string ConnectionString => _container.GetConnectionString();

    public async Task InitializeAsync()
    {
        await _container.StartAsync();
        await using var db = CreateContext();
        // Build the schema from the EF model (fast; no migration history needed for
tests).
        await db.Database.EnsureCreatedAsync();
    }

    public Task DisposeAsync() => _container.DisposeAsync().AsTask();

    public AppDbContext CreateContext()
    {
        var options = new DbContextOptionsBuilder<AppDbContext>()
            .UseNpgsql(ConnectionString).Options;
        return new AppDbContext(options);
    }

    /// <summary>Truncates every table so each test starts from a clean slate. The table
    /// list is DERIVED from the EF model, so a new entity is reset automatically – no
    /// hand-maintained list to forget.</summary>
    public async Task ResetAsync()
    {
        await using var db = CreateContext();
        var tables = db.Model.GetEntityTypes()
            .Select(e => e.GetTableName())
            .Where(name => name is not null)
            .Distinct()
            .Select(name => $"\"{name}\"");
        var sql = "TRUNCATE TABLE " + string.Join(", ", tables) + " RESTART IDENTITY
CASCADE;";
        await db.Database.ExecuteSqlRawAsync(sql);
    }
}

[CollectionDefinition(PostgresCollection.Name)]
public sealed class PostgresCollection : ICollectionFixture<PostgresFixture>
{
    public const string Name = "postgres";
}
```

Three design decisions in forty lines, each load-bearing:

One container per run, not per test. Container startup costs seconds; a suite of hundreds of tests can't pay it each time. The xUnit *collection fixture* shares one container across every test class marked `[Collection(PostgresCollection.Name)]`, and xUnit runs a collection's members serially — so sharing is safe by construction.

Reset between tests, and make the reset un-forgettable. Shared container means shared state, so every test must start clean. The rookie version is a hand-maintained list of tables to truncate — which goes stale the day someone adds an entity and forgets the list (this exact bug is why rule R11 exists; the reference repo found it in an audit). The fix is the discovery pattern again: `**derive the table list from db.Model.GetEntityTypes()**. New entity → automatically in the reset. And rather than trusting each test to call ResetAsync(), a tiny base class makes isolation structural:`

```
// tests/Api.Tests/Infrastructure/PostgresTestBase.cs — real file, in full
public abstract class PostgresTestBase(PostgresFixture fixture) : IAsyncLifetime
{
    protected PostgresFixture Fixture { get; } = fixture;

    public Task InitializeAsync() => Fixture.ResetAsync(); // before EVERY test
    public Task DisposeAsync() => Task.CompletedTask;
}
```

Inherit it and there is no reset call to forget. Same move as the tenant filter to come: correct by default, not correct by discipline.

`**EnsureCreatedAsync` here, migrations elsewhere.** The fixture builds schema straight from the model — fast, no history table. But that means the *migrations themselves* could rot untested. So one more test exists solely to apply every migration, in order, against a scratch database (MigrationsTests in the reference repo — and the integration harness from 1.2 will boot the app with its real startup-migration path in Part 2). Schema-from-model for speed; migrations proven separately; nothing unverified.

5. Red → Green — prove the loop against real Postgres

```
// tests/Api.Tests/ProbePersistenceTests.cs
[Collection(PostgresCollection.Name)]
public class ProbePersistenceTests(PostgresFixture fixture) : PostgresTestBase(fixture)
{
    [Fact]
    public async Task Probe_roundtrips_through_real_postgres()
    {
        await using (var db = Fixture.CreateContext())
        {
            db.Add(new Probe { Note = "hello from a container" });
            await db.SaveChangesAsync();
        }
        await using (var db2 = Fixture.CreateContext()) // fresh context: no cache tricks
        {
            var probe = await db2.Set<Probe>().SingleAsync();
            Assert.Equal("hello from a container", probe.Note);
        }
    }
}
```



```
dotnet test tests/Api.Tests
# Passed! - 2 tests (health + probe) ... first run pulls postgres:17, be patient
```

Watch `docker ps` during the run: a postgres container blinks into existence, then is gone. No cleanup script, nothing to remember, nothing left behind.

6. Architecture Decision

ARCHITECTURE DECISION

The fork: what database do tests run against? (a) EF in-memory provider; (b) SQLite in-memory ("at least it's relational"); (c) a shared long-lived test server; (d) Testcontainers — throwaway real Postgres per run.

Chosen: (d), with schema-from-model for speed and a separate migrations-apply test. Tests must exercise the engine production runs, because this platform *leans* on Postgres-specific behavior: global query filters as the tenancy wall, `SKIP LOCKED` job claiming, savepoints, and eventually row-level security. (ADR-C7/C13; R11 for the derived reset.)

Rejected: (a) silently skips the exact constructs we depend on — a green suite proving nothing. (b) is a different engine with different SQL, different concurrency, no RLS: it converts "provider lies" into "dialect lies". (c) reintroduces shared mutable state across runs and machines — the thing test isolation exists to kill.

The trade: Docker becomes a hard requirement for running the suite, and the first test run pays an image pull. Accepted consciously in 0.2 — the same daemon already runs your dev infrastructure, and CI (1.6) has it natively.

7. Checkpoint

```
git add -A && git commit -m "feat: EF Core + Postgres — first migration, Testcontainers
fixture, model-derived reset"
git tag lesson-1.3
```

You should now be able to: list three concrete ways the in-memory provider lies; explain the container-per-run / reset-per-test economy and why the reset list is derived, not written; say what a design-time factory is for and why it must not have a connection fallback; and describe how migrations stay verified even though the fixture doesn't use them.

Next: *Lesson 1.4 — Configuration & the options pattern* — the app learns to read config, validate it at startup, and fail loudly at boot rather than mysteriously at 3 a.m.

Lesson 1.4 — Configuration & the options pattern

Part 1 · The walking skeleton. The app can answer HTTP and talk to Postgres; now it learns to be *configured* — and, more to the point, to refuse to start when it isn't. Configuration is where "works on my machine" goes to hide: a missing key that defaults to something plausible, a secret pasted into a JSON file, a setting the code reads that nobody documented. This lesson closes all three holes, and closes them with machinery, not vigilance.

Goal: the API loads `.env` in dev, exposes typed, validated settings objects, fails at *startup* (not first-use) on missing/invalid config, and a test fails the build when code reads a config key that isn't documented.

Concepts & principles: the config hierarchy; secrets never in the repo (ADR-001 mechanized); typed settings with construction-time validation; fail-fast startup; the doc-sync gate.

Maps to: ADR-001 · R20–R22 · repo: `src/Api/Configuration/SettingsProvider.cs`, `SettingsRegistration.cs`, `src/Api/Program.cs` (the `DotNetEnv` line), `tests/Api.Tests/DocAndConfigSyncTests.cs`, `.env.example`, `appsettings.json`.

Prerequisites: 1.3; the `.env` contract from 0.2.

1. Motivate — the 3 a.m. config bug

Picture the failure mode we're designing against. A service reads `config["Jwt:Secret"]` deep inside a request handler. On the machine where it's missing, nothing happens at startup — the app boots, health checks pass, the deploy is declared good. Hours later the first user tries to sign in and gets a 500, and the stack trace points at token-signing code, three layers away from the actual problem: a key nobody set. Now multiply by every config value in the system.

The root defect isn't the missing key — machines will always occasionally be misconfigured. The defect is **when** and **how** the system tells you. Config errors should surface at *boot*, in a message that names the key and says where to set it. That transformation — from runtime mystery to startup message — is what this lesson builds.

And beneath it, a second rule from Lesson 0.1's ADR-001: secrets never live in committed files. Which raises the practical question: where *does* the app get them, and how do committed files and secret files cooperate?

2. The config hierarchy — one mental model

ASP.NET Core builds `IConfiguration` from layered sources, later layers overriding earlier ones. Our arrangement (dev):

```

appsettings.json committed – non-secret defaults, shape documentation
appsettings.Development.json committed – dev-only tweaks (verbose logging, etc.)
environment variables NOT committed – secrets & machine-specifics ...
    ▲
    └─ in dev, populated from .env by DotNetEnv at startup

```

The elegant part is the last line. Production sets real environment variables; dev keeps them in the gitignored `.env` from 0.2, and a single line at the very top of `Program.cs` — before anything reads config — loads them. From the real file:

```

// Local dev: load secrets/config from the repo-root .env (the single local source of
// truth – see docs/DECISIONS.md). TraversePath walks up to find it regardless of the
// working dir; the try/catch makes it a no-op when there's no .env (e.g. production,
// which uses real environment variables).
try { DotNetEnv.Env.TraversePath().Load(); } catch { /* no .env present */ }

var builder = WebApplication.CreateBuilder(args);

```

One mechanism, two environments: to the app, dev and prod look identical — env vars either way. Nested keys use the `Section__Sub` double-underscore convention (`Jwt__Secret` in `.env` becomes `Jwt:Secret` to `IConfiguration`), which is the standard env-var transport for hierarchical config.

`appsettings.json` keeps the **non-secret** shape: section names, defaults worth committing, values a reviewer should see. It also serves as *documentation of what exists* — which the gate in §5 exploits.

3. Typed settings — validation at construction

Reading `config["Jwt:Secret"]` at call sites scatters string literals through the code, defers errors to first use, and re-parses on every read. The platform's pattern: one small class per config section, reading and validating **in its constructor**, registered as a singleton. The real `JwtSettings`:

```

// src/Api/Configuration/SettingsProvider.cs
public class JwtSettings : IJwtSettings
{
    public string SecretKey { get; }
    public string Issuer { get; }
    public int ExpiryMinutes { get; }

    public JwtSettings(IConfiguration config)
    {
        SecretKey = config["Jwt:Secret"]
            ?? throw new InvalidOperationException(
                "Jwt:Secret not configured (set Jwt__Secret in .env for dev)");
        Issuer = config["Jwt:Issuer"] ?? "Perezosoftware";
        ExpiryMinutes = config.GetValue("Jwt:ExpiryMinutes", 60);

        const int MinSecretKeyLength = 32;
        if (SecretKey.Length < MinSecretKeyLength)
            throw new InvalidOperationException(
                $"Jwt:Secret must be at least {MinSecretKeyLength} characters");
    }
}

```

Read the error messages as UX, because they are: each names the key **and the fix** ("set Jwt__Secret in .env for dev"). Whoever hits this — possibly you, in six months, on a new laptop — gets handed the solution, not a scavenger hunt. And notice the *semantic* validation: not just "present" but "at least 32 characters," because an HMAC key shorter than that is a security bug that would otherwise hide until token validation mysteriously failed.

Registration happens once, in one place, and the singletons are constructed **eagerly** so a bad value kills the boot:

```
// src/Api/Configuration/SettingsRegistration.cs (abridged – real file)
public static IJwtSettings AddAppSettings(this IServiceCollection services, IConfiguration
config)
{
    // Build JwtSettings once; the same instance backs both the DI singleton and the
    // JWT-bearer TokenValidationParameters – a single source of truth.
    var jwtSettings = new JwtSettings(config);

    services.AddSingleton<IJwtSettings>(jwtSettings);
    services.AddSingleton<IRefreshTokenSettings>(new RefreshTokenSettings(config));
    services.AddSingleton<IApplicationSettings>(new ApplicationSettings(config));
    // ...
    return jwtSettings;
}
```

Two subtleties worth pausing on. The settings are constructed with *new*, *right here*, not lazily by the container — so startup **is** the validation pass; a missing key can't lurk until first resolve. And *JwtSettings* is built into a local and *returned*, because *Program.cs* needs the same values to configure JWT-bearer validation — one instance, two consumers, zero chance of drift between "the settings DI hands out" and "the settings auth actually validates with." (This is the platform's blessed shape per rule R22; .NET's *IOptions* + *ValidateOnStart* is a fine alternative — the *decision* is eager validation and one pattern everywhere, not which library performs it.)

Consumers then depend on *IJwtSettings* — an interface, so tests hand in a stub without touching *IConfiguration* at all. Config-reading is quarantined to these constructors the way *MailKit* is quarantined to *Infrastructure/Email*: one place, one pattern.

And your 1.2 test just broke — good. *WebApplicationFactory* boots the *real* *Program*, which now demands *Jwt:Secret* at startup; the health test dies with the very *InvalidOperationException* you wrote. This is fail-fast config working exactly as designed — the test boots the app, the app validates its config, no exceptions for test mode. The seam is the one the reference harness documents: *Program* reads config at *CreateBuilder* time, *before* the factory's config hooks run, so the value must exist as an **environment variable** that early. A static constructor on the test class is the cleanest place:

```
public class HealthEndpointTests : IClassFixture<WebApplicationFactory<Program>>
{
    static HealthEndpointTests() =>
        Environment.SetEnvironmentVariable("Jwt__Secret",
            "integration-test-jwt-secret-key-0123456789"); // ≥32 chars; value irrelevant
    // ...
}
```

The dummy only needs to *satisfy the validation* (length), because nothing in this test signs or verifies a token. When the integration harness grows up in Part 2, this seam moves into its

constructor — same idea, one shared home.

4. Defaults are decisions

Look again: Issuer defaults to "Perezosoft", ExpiryMinutes to 60 — but SecretKey throws. Why the asymmetry? Because **a default is a decision that absence is safe**. A missing issuer name has a sensible universal answer. A missing *secret* does not — any default would be catastrophic, so absence must be fatal. When you write a settings class, sort every value into one of these bins deliberately: safe default, or refuse to boot. (Part 7 adds a third bin for optional *features*: absent config means the feature is **off** — R21's config-gated default-off. Same principle: absence maps to the safe state, never the surprising one.)

5. The gate — config that can't go undocumented

Now the structural guarantee. Nothing so far stops a developer from reading `config["Fancy:NewKnob"]` in code and telling no one — until someone deploys without it. The reference platform closes this with a test that cross-references **code against documentation**, from the real `DocAndConfigSyncTests`:

```
[Fact]
public void ConfigKeys_ReadInCode_AreDocumented()
{
    // R20: every Section:Key literal read via IConfiguration is documented — either in
    // an appsettings*.json (as a nested path) or in .env.example (as Section__Key).
    // Catches config drift where code reads a key nobody declared.
    var documented = DocumentedConfigPaths(); // parsed from appsettings*.json +
    .env.example

    var readKeys = new HashSet<string>(StringComparer.Ordinal);
    var access = new Regex(
        @"(?:GetSection|GetValue<[^>*>|GetValue|configuration|config|Configuration)\s*[\\(\\
    ]\s*"([A-Za-z][A-Za-z0-9]*(?:[A-Za-z0-9]+)*)""");
    foreach (var f in SourceFiles(Path.Combine(RepoRoot(), "src")))
        foreach (Match m in access.Matches(File.ReadAllText(f)))
            readKeys.Add(m.Groups[1].Value);

    var undocumented = readKeys.Where(k => !documented.Any(d =>
        d.Equals(k) || d.StartsWith(k + ":") || k.StartsWith(d + ":")))
        .OrderBy(k => k).ToList();

    Assert.True(undocumented.Count == 0,
        $"Config keys read in code but not in appsettings*.json or .env.example:
    {string.Join(", ", undocumented)}");
}
```

It's a source scan — grep with a build verdict. Read a key in code without adding it to `.env.example` or an `appsettings*.json`, and the suite goes red with the key's name in the failure message. (One adaptation when you write `RepoRoot()`: anchor the upward directory walk on a file your repo actually has at its root — `global.json` works perfectly and 0.2 guaranteed it exists.) The documentation *cannot* drift from the code, because drift fails CI. This is the third face of one idea you now hold: discovery over registration (1.3), isolation by base class (1.3), documentation by gate (here). None of them ask a human to remember anything.

Add the corresponding entries to `.env.example` as you go — from 0.2 it's the contract: every key, a comment saying what it does, required-or-default status. `Jwt__Secret=` goes in now

(empty value, comment "REQUIRED — ≥32 chars").

6. Run it — watch it fail correctly

The best test of fail-fast config is failing it on purpose:

```
# comment out Jwt__Secret in your .env, then:
dotnet run --project src/Api
# Unhandled exception: System.InvalidOperationException:
# Jwt:Secret not configured (set Jwt__Secret in .env for dev)
```

Boot refused, key named, fix stated. Restore the key; boots clean. Then run the suite — including your new doc-sync test — and prove the gate: read a fake key (config["Nope:Missing"]) anywhere in src/, watch the test name it, delete the line.

```
dotnet test tests/Api.Tests # all green, including ConfigKeys_ReadInCode_AreDocumented
```

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how does configuration reach the code? (a) config["key"] at call sites; (b) .NET's options pattern (IOptions<T> + BindConfiguration + ValidateOnStart); (c) hand-rolled typed settings classes validating in their constructors, registered eagerly as singletons.

Chosen: (c) as this platform's single blessed pattern (R22), with ADR-001's .env as the dev secret source and a doc-sync test (R20) making documentation mandatory. Constructor validation gives fail-at-boot with zero framework ceremony, interfaces make consumers trivially testable, and *one pattern everywhere* means a new reader learns it once.

Rejected: (a) scatters literals, defers failure to first use, and is exactly what the R20 gate exists to catch. (b) is genuinely fine — validateOnStart achieves the same fail-fast — but brings IOptions<T> indirection at every consumer and *two* ways to do one thing once any hand-rolled class exists. Uniformity won. The pattern choice matters less than the properties: eager, validated, typed, documented.

The trade: hand-rolled means we own the validation code, and each new section is a small class rather than one Bind call. Accepted for a platform this size; revisit if settings classes multiply into the dozens.

8. Checkpoint

```
git add -A && git commit -m "feat: typed validated config — .env loading, fail-fast
settings, doc-sync gate"
git tag lesson-1.4
```

You should now be able to: draw the config layering and say where secrets live in dev vs prod; explain why validation belongs in the constructor and registration must be eager; sort a config value into "safe default" vs "refuse to boot" and defend the sort; and describe how the doc-sync gate turns documentation from a request into an invariant.

Next: *Lesson 1.5 — The error envelope* — deciding, once, what every error in the system looks like from the outside.

Lesson 1.5 — The error envelope

Part 1 · The walking skeleton. The shortest lesson in the course, on purpose — the artifact is five lines. But it's a five-line *treaty*: a decision about what every error in the system looks like from the outside, made once, before there are any errors to regret. Small decisions made early are cheap; the same decisions retrofitted across forty endpoints are a migration project.

Goal: one shared error shape for the whole API; a written rule that exception internals never cross the boundary; and clarity on what an error *code* is for versus an error *message*.

Concepts & principles: the error envelope; machine-readable vs human-readable error channels; information leakage through error text; deciding API contracts before the API grows.

Maps to: R16, R18 · repo: `src/Api/Services/ErrorResponse.cs`, its uses across every controller.

Prerequisites: 1.2 (an API exists to have errors).

1. Motivate — the anatomy of a leaked stack trace

Two failure modes, one root cause.

Failure one: the inconsistent zoo. Without a decision, each endpoint invents its own error shape: `{ "error": "..."} here, { "message": "..."} there, { "errors": [ ... ] }` from whoever wrote validation, a bare string from the intern's endpoint, HTML from the framework when nobody handled it at all. Every client — your own Blazor app first of all — grows a thicket of "well, if the body has *this* shape..." parsing. The API's error surface becomes something you *reverse-engineer* rather than read.

Failure two: the confession. Somewhere, a handler does the natural thing:

```
catch (Exception ex)
{
    return StatusCode(500, new { error = ex.Message }); // ← the natural thing. Don't.
}
```

Now consider what `ex.Message` can contain, on a bad day: "28P01: password authentication failed for user \"app_rt\" — a Postgres error naming an internal role. "Connection refused (10.0.3.17:5432)" — internal topology. "The token 'eyJhbGciOi...' is malformed" — a credential fragment. An attacker probing your API doesn't need a vulnerability report; your exceptions *are* one, streamed on demand. This is why rule R16 exists and is checked in review: **client-facing error bodies never contain raw exception text; detail stays in server logs**, where you have access control and the attacker doesn't.

Both failures have the same fix: decide the shape once, make the safe thing the easy thing.

2. Green — the envelope

The reference platform's entire `ErrorResponse.cs`:


```
namespace Perezosoft.Api.Services;

/// <summary>The standard API error envelope. Serializes to
/// <c>{ "error": ..., "message": ... }</c> (camelCase) — the shape every error
/// response uses.</summary>
public record ErrorResponse(string Error, string Message);
```

That's the artifact. Its power is entirely in the *convention it anchors*. Usage, from the real `AuthController`:

```
return Unauthorized(new ErrorResponse("no_refresh_token", "Refresh token not found"));
return Unauthorized(new ErrorResponse("invalid_refresh_token", "Refresh token is invalid or
expired"));
return StatusCode(500, new ErrorResponse("refresh_failed", "Failed to refresh token"));
```

Two fields, two audiences, and the discipline is in keeping them apart:

****error is for machines.**** A stable, snake_case code — `invalid_refresh_token`, `seat_limit_reached` — that clients branch on. It's *API contract*: renaming one is a breaking change, exactly like renaming a JSON property. When the Blazor client (3.4) decides whether to bounce a user to login or show "try again," it switches on this field, never on prose.

****message is for humans.**** A safe, self-contained sentence for logs and debugging. It describes the *category* of failure in words you chose — never words an exception chose. Note what "Failed to refresh token" doesn't say: no exception type, no stack frame, no connection string. The interesting detail went to the server log at the catch site, correlated by the request's trace ID (which Lesson 4.5 makes automatic).

Why not `ProblemDetails` (RFC 7807), the standards-blessed answer? It's a reasonable fork — see the decision box. The short version: the envelope predates it here, is smaller, and the *rule* (one shape, no leaks) matters far more than which shape.

3. Harden — a shape is only a shape if it's the only one

A convention with exceptions is a suggestion. Two enforcement layers make this one hold:

Status codes carry the first bit of meaning. The envelope rides on top of correct HTTP semantics, and this platform is unusually disciplined about three of them — you'll build each: **401** "we don't know who you are" (2.1), **403** "we know you, your *role* may not do this" (6.1, `RequirePermission`), **402** "we know you, your *plan* doesn't include this" (5.1, `RequireEntitlement`). Three different problems, three different client reactions (sign in / hide the button / show the upgrade page), distinguishable before reading a byte of body. The envelope then says *which* permission, *which* limit.

A scan for defectors. Rule R18: all error responses use the shared helper — no anonymous `{ error = ... }` objects. It's enforced the same way as 1.4's config gate: a source scan that fails when an anonymous error shape appears in a controller. You now have three of these grep-with-a-verdict tests (MailKit containment, config sync, error shape) — cheap, blunt, effective. When you write your own (3.3 formalizes the habit into the architecture-test project), keep the pattern: scan for the *forbidden* form, fail with the file name and the fix.

One honest caveat: model-binding failures (malformed JSON, missing required fields) produce ASP.NET's built-in `ValidationProblemDetails` before your code runs. The platform

leaves that be — framework-generated 400s are consistent *with each other*, and bending the framework's validation pipeline into the envelope isn't worth the ceremony. The envelope rule governs *your* error paths. Record the exception to the rule as part of the rule; an undocumented exception is how conventions rot.

4. Run it

```
curl -k https://localhost:7160/api/nope
# 404 — and once real endpoints exist, error bodies all read:
# { "error": "not_found", "message": "..."} }
```

The real payoff is invisible today. It arrives every time this course adds an endpoint and there is exactly one way an error can look — and it arrives hardest in 3.4, when the client's error handling is a single small function instead of a parser museum.

5. Architecture Decision

ARCHITECTURE DECISION

The fork: what do API errors look like? (a) Whatever each endpoint returns; (b) RFC 7807 ProblemDetails everywhere; (c) a minimal shared envelope (`{ error, message }`) + correct status codes, with framework validation left native.

Chosen: (c). The binding decisions are *one shape* and *no exception text* (R16, R18); a two-field record delivers both with zero ceremony, and the machine-readable `error` code gives clients a stable contract the status code alone can't.

Rejected: (a) isn't a decision, it's an absence of one — and it's what naturally happens if this lesson doesn't exist. (b) is genuinely respectable: standard, extensible, tooling-friendly. It lost on weight — type URLs and extension members nobody would populate — and on the platform's preference for the smallest thing that holds the rule. Migrating later is mechanical precisely because every error already flows through one type.

The trade: a custom shape means API consumers read *your* docs, not an RFC. Accepted for a platform whose first-party clients dominate; the public API surface (7.3) documents the envelope in its OpenAPI schema.

6. Checkpoint

```
git add -A && git commit -m "feat: shared error envelope – one shape, no exception leakage"
git tag lesson-1.5
```

You should now be able to: name the two audiences of an error and which field serves each; explain why `ex.Message` in a response body is a security defect, not a style issue; distinguish 401/402/403 and say who reacts to each; and defend deciding contract shapes *before* the surface grows.

Next: Lesson 1.6 — *CI from commit one* — everything this part built (locked restore, warnings-as-errors, the gates) starts running on every push, forever.

Lesson 1.6 — CI from commit one

Part 1 · The walking skeleton. The final lesson of the skeleton, and the one that makes everything before it *permanent*. So far, the gates you've built — locked restore, warnings-as-errors, the doc-sync scan — run when you remember to run them. This lesson puts them on a machine that never forgets: a pipeline that executes on every push, on a clean runner, with no local state to flatter you. The pipeline is born small and — like the architecture-test suite it partners with — **grows a job per part** for the rest of the course.

Goal: a GitHub Actions workflow that restores in locked mode, builds with warnings-as-errors, runs the tests (Testcontainers included), scans for secrets and copyleft licenses — red on any violation, on every push and pull request.

Concepts & principles: CI as a growing gate; clean-room verification ("if it only works on your machine, it doesn't work"); supply-chain gates in the pipeline; the PR checklist as process.

Maps to: R26, R27 (+ every earlier rule, now enforced remotely) · repo:

`.github/workflows/ci.yml` (jobs build-test, secret-scan, license-scan),
`.github/forbidden-licenses.json`, `.github/pull_request_template.md`.

Prerequisites: 1.1–1.5 (there must be something worth gating); a GitHub repo to push to.

1. Motivate — a gate you add at the end is a gate you never designed for

Here's what happens to teams that "add CI later." Later arrives; someone writes the workflow; it fails — of course it fails, nothing was ever built on a clean machine. The tests assume a database that isn't there, the build has 87 warnings that were always going to be errors *someday*, a connection string is hardcoded somewhere, and one package resolves differently than it did in March. Fixing all of it is now a project, so the workflow gets watered down — warnings allowed, the flaky tests skipped — and the team learns that CI is a formality. The gate exists, but the gate gates nothing.

Now run the tape the other way. The pipeline exists from commit one, when the entire system is a health endpoint and a probe entity. It's green in five minutes of work, because there's nothing to be red *about*. From then on, every lesson that adds a capability adds its check while the diff is small and the author is present. The pipeline never has a catching-up crisis, because it never fell behind. That's the whole argument, and it's the same argument as warnings-as-errors in 1.1: **adopt the standard when meeting it is free.**

There's a second, subtler payoff. Your machine has your `.env`, your Docker images, your half-remembered global tools. A green build there is a claim contaminated by local state. CI is a *clean room*: fresh runner, no `.env`, nothing but the repo and the workflow. When it's green, the claim "anyone can build this from a checkout" has actually been tested — which is, notice, the same claim Lesson 0.2 made about dev machines. CI is 0.2's reproducibility promise, verified continuously.

2. The workflow — build-test, from the real file

Create `.github/workflows/ci.yml`. This is the reference platform's actual first job, and every step is a decision you already made — now enforced remotely:

```
name: CI

on:
  push:
    branches: [main, develop]
  pull_request:
    branches: [main, develop]

jobs:
  build-test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v5

      - uses: actions/setup-dotnet@v5
        with:
          dotnet-version: "10.0.x"

      # Supply-chain: restore in LOCKED mode against the committed packages.lock.json
      # files. --locked-mode fails if the resolved graph differs from the lockfile, so a
      # dependency change can't slip in without updating the (reviewed) lockfile.
      - name: Restore (locked mode)
        run: |
          dotnet restore src/Api/Perezosoftware.Api.csproj --locked-mode
          dotnet restore tests/Core.Tests/Perezosoftware.Core.Tests.csproj --locked-mode
          dotnet restore tests/Api.Tests/Perezosoftware.Api.Tests.csproj --locked-mode

      # Warnings-as-errors comes from Directory.Build.props – the same build policy
      # locally and here, one definition. --no-restore: reuse the locked restore above
      # (don't let the build re-restore loosely).
      - name: Build (warnings-as-errors)
        run: dotnet build src/Api/Perezosoftware.Api.csproj -c Release --no-restore

      # Api.Tests spins up Postgres via Testcontainers – Docker is available on
      # ubuntu-latest, so the same tests run here as locally, unchanged.
      - name: Test (Core + Api)
        run: |
          dotnet test tests/Core.Tests/Perezosoftware.Core.Tests.csproj -c Release --no-restore
          dotnet test tests/Api.Tests/Perezosoftware.Api.Tests.csproj -c Release --no-restore

      # Fail the build if the EF model has drifted from the latest migration/snapshot.
      - name: No pending EF model changes
        env:
          ConnectionStrings__DefaultConnection:
            "Host=localhost;Port=5432;Database=ef_designtime;Username=ci;Password=ci"
        run: |
          dotnet tool install --global dotnet-ef
          dotnet ef migrations has-pending-model-changes \
            --project src/Infrastructure/Perezosoftware.Infrastructure.csproj \
            --startup-project src/Api/Perezosoftware.Api.csproj
```

Walk the details, because each is a small lesson:

- ****--locked-mode then --no-restore.**** The restore step verifies the graph against

the lockfiles (1.1's gate, now unskippable); the build and test steps then *reuse* that restore. Without `--no-restore`, `dotnet build` would quietly re-restore in normal mode — and your

carefully locked graph would be decoration.

- **Testcontainers just works.** The GitHub runner has Docker, so the suite you built in

1.3 — real Postgres and all — runs identically here. This is the moment the Testcontainers decision pays its second dividend: no "CI database" to provision, no service-container YAML to keep in sync with the fixture.

- **The EF drift check** closes a gap you couldn't feel locally: change an entity,

forget `dotnet ef migrations add`, and everything *builds* — the model and the migrations have simply parted ways, to be discovered at the next deploy. `has-pending-model-changes` compares model to snapshot and fails loudly. Note the placeholder connection string: the check opens no DB connection, but the design-time factory (1.3) *requires the key to exist* — no fallback, remember. CI documents that contract by satisfying it explicitly.

3. Two more jobs — the supply-chain gates go remote

Secret scan. 0.2 installed gitleaks as a local backstop; policy is when it runs whether you remember or not. From the real workflow:

```
secret-scan:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v5
      with:
        fetch-depth: 0 # full history so the scan covers every commit, not just the tip

    - name: gitleaks
      env:
        GITLEAKS_VERSION: "8.21.2"
      run: |
        curl -sSfL "https://github.com/gitleaks/gitleaks/releases/download/v${GITLEAKS_VERSION}/gitleaks_${GITLEAKS_VERSION}_linux_x64.tar.gz" \
          | tar -xz gitleaks
        ./gitleaks detect --no-banner --redact --config .gitleaks.toml
```

Two deliberate choices worth stealing: `fetch-depth: 0` scans **all history**, because a secret committed and then "deleted" is still a leaked secret — git remembers; and `--redact`, because a scanner that prints the secrets it finds into a CI log has just created a second leak. (It runs the pinned OSS binary directly rather than a marketplace action — one less third party in the trust chain of the job whose *purpose* is trust.)

License scan. Rule R26: no copyleft in the dependency graph. Not because copyleft is bad — because for *this* platform (a commercial SaaS template) a GPL dependency imposes obligations the product can't meet, and a transitive one arrives silently with some innocent-looking package. The job inventories every license, direct **and transitive**, and fails on a prohibited list (`.github/forbidden-licenses.json`: GPL/AGPL/LGPL/SSPL):

```

license-scan:
  runs-on: ubuntu-latest
  steps:
    - uses: actions/checkout@v5
    - uses: actions/setup-dotnet@v5
      with:
        dotnet-version: "10.0.x"
    - name: dotnet-project-licenses (fail on copyleft)
      run: |
        dotnet restore src/Api/Perezosoftware.Api.csproj
        dotnet tool install --global dotnet-project-licenses
        export PATH="$PATH:$HOME/.dotnet/tools"
        dotnet-project-licenses --input src/Api/Perezosoftware.Api.csproj \
          --include-transitive --use-project-assets-json \
          --forbidden-license-types .github/forbidden-licenses.json

```

Note the shape: a **prohibited-list, not an allow-list** — MIT/Apache/BSD and friends never trip a false positive, so the gate stays quiet until it has something real to say. A gate that cries wolf gets disabled; this one is designed to be believed. It's also a *separate job*, deliberately: a license failure and a build failure are different conversations, and the red X should say which one you're having.

4. The human gate — the PR template

One more file, `.github/pull_request_template.md`, auto-loaded by GitHub into every pull request: the platform's checklist — tests written first, tenant-scoping verified, docs updated, no direct UI→DB access. It gates nothing mechanically; it's the *review script* — the conversation every change must survive, written down. The division of labor is worth stating: **machines check what machines can check** (this lesson, the arch tests), **the checklist scripts what humans must check** (rules like R16's "no exception text" — review-enforced until a scan exists). Process is a gate too; it just runs on people.

5. Run it — push and watch

```

git add .github && git commit -m "ci: build-test + secret-scan + license-scan on every
push"
git push -u origin main
# → GitHub → Actions: three jobs, all green

```

Then prove the gates gate (one minute, permanent trust): bump any package version in `Directory.Packages.props` *without* updating the lockfile, push a branch, and watch Restore (locked mode) fail with NU1004. Revert. That red X is 1.1's promise, kept by a machine that doesn't know you and doesn't take your word for anything.

Where this file goes from here: e2e journeys join in 3.6; the Docker image build in 8.2; deploy-to-staging with a post-deploy smoke, then gated production, in 8.3; native platform legs in the appendix. In the reference repo this file ends at ~800 lines and fourteen jobs — every one of which you'll have written in the lesson that gave it a reason to exist.

6. Architecture Decision

ARCHITECTURE DECISION

The fork: when does CI arrive, and how much does it check? (a) Later, once there's "something real to test"; (b) now, minimal — build only, expand someday; (c) now, with every gate the codebase already claims to honor: locked restore, warnings-as-errors, tests, EF drift, secrets, licenses.

Chosen: (c). A claimed-but-unchecked standard decays silently; every gate in this job list is one the previous five lessons *already committed to*, so enforcing it remotely costs minutes now and prevents the catch-up crisis later. Growth by accretion — a job per capability, added with the capability — keeps every addition small and owned.

Rejected: (a) is how teams end up with the watered-down formality pipeline — descoped at exactly the moment it's needed most. (b) defers the cheap-now checks to when they're expensive, which is (a) with better intentions.

The trade: every push now takes a few minutes to go green, and a red X blocks merging even when "it's fine, it's just the license tool." That friction is the product. Accepted — and never bypassed with a skip flag; a gate you can wave through is a gate in name only.

7. Checkpoint

```
git tag lesson-1.6
```

Part 1 complete. You have a walking skeleton with a spine, a nervous system, and an immune system: an API answering HTTP, real-database tests, validated config, one error shape — and a clean-room pipeline that re-verifies all of it on every push, forever.

You should now be able to: argue for CI-from-day-one against "we'll add it when there's something to test"; explain `--locked-mode + --no-restore` as a pair; say why the secret scan reads full history and redacts; and describe the prohibited-list design and why quiet gates are trustworthy gates.

Next: *Part 2 — Identity & tenancy.* The skeleton gets people: users, tokens, and the tenant model everything else stands on.

Part 2 — Identity & tenancy

Lesson 2.1 — Users & JWT access tokens

Part 2 · Identity & tenancy. The platform meets its first real domain concept — a user — and its most consequential early decision: **we build authentication ourselves** instead of adopting ASP.NET Core Identity. That choice ripples through the next twenty lessons, so this lesson earns it properly: what a JWT actually is, what belongs in one, and why "store hashes, never tokens" is the reflex to install before any token exists.

Goal: a User entity persists; the API can issue a signed JWT access token carrying identity claims; a protected endpoint returns 401 without one and 200 with one.

Concepts & principles: JWT anatomy (header/payload/signature); claims as the identity contract; custom auth vs framework Identity (ADR-002); hash-only credential storage; the token generator/haser seams.

Maps to: ADR-002 · repo: `src/Core/Entities/User.cs`, `src/Api/Services/JwtTokenService.cs`, `JwtClaims.cs`, `TokenGenerator.cs`, `TokenHasher.cs`, `src/Api/Configuration/JwtValidation.cs`, `src/Api/Controllers/AuthControllerBase.cs`, `AuthController.cs` (begins), `tests/Api.Tests/TokenServiceTests.cs`, `UserServiceTests.cs`.

Prerequisites: Part 1 complete (config validation especially — `Jwt:Secret` from 1.4 is about to become load-bearing). Delete the `Probe` entity from 1.3 now; its job is done.

1. Motivate — the fork in the road

Type `dotnet new webapp --auth Individual` and you get ASP.NET Core Identity: user tables, password hashing, cookie login, two-factor scaffolding — batteries included. So the burden of proof is on *not* using it. Here's the case the reference platform recorded in ADR-002, and it's worth internalizing because it's really a case about **owning your core domain**:

This platform's authentication is not a checkbox; it's a *product surface*. Passwordless magic links (2.4), OAuth with account-linking rules (2.5), a `tenant_id` claim that drives all data isolation (2.6), MFA step-up on multiple sign-in paths (6.5), native clients with a different token transport (appendix) — every one of these needs precise control of *when tokens are issued, what they carry, and how sessions extend*. Identity can be bent to each of these, but you end up fighting a framework whose defaults are cookie-centric and whose schema you don't own — and the fights compound. Meanwhile the part of Identity that's genuinely hard to hand-roll — password hashing — **we don't need**, because this platform never stores passwords at all: sign-in is passwordless or via OAuth, full stop.

Strip out passwords and "authentication" reduces to two well-understood mechanics: **mint a signed claim set** (this lesson) and **manage its renewal** (2.2). Both sit on primitives the BCL provides. That's a domain worth owning.

The honest counter-case: owning auth means owning its mistakes — rotation, reuse detection, replay windows are now *your* tests (and this course writes them). If your product's auth is a checkbox — internal tool, password login is fine, nothing custom — take Identity and

spend your innovation budget elsewhere. Architecture is choosing what to own; this platform chooses auth, deliberately.

2. The `User` — smaller than you expect

```
// src/Core/Entities/User.cs – the real entity
public class User
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public required string Email { get; set; }

    /// <summary>Display name from the OAuth provider, refreshed each sign-in.</summary>
    public string? DisplayName { get; set; }

    /// <summary>True only when the provider asserts a verified email claim.</summary>
    public bool EmailVerified { get; set; }

    /// <summary>Preferred UI language (e.g. "en", "es"), or null to fall back to the
    /// browser/OS culture. A per-user preference that follows the user across
    devices.</summary>
    public string? Locale { get; set; }

    // Tenant membership is the source of truth for which tenant a user belongs to;
    // resolve it via TenantMembership (one tenant per user). See ITenantRepository.

    public DateTimeOffset CreatedAt { get; set; }
    public DateTimeOffset UpdatedAt { get; set; }

    /// <summary>OAuth identities linked to this account (one per provider).</summary>
    public ICollection<UserLogin> Logins { get; set; } = new List<UserLogin>();
}
```

Read it for what's **missing**. No password hash — decided above. No `TenantId` — the comment points at the deliberate indirection: membership is a separate entity (2.8), so "which tenant" has one source of truth and a user's account survives leaving a household. No roles — membership owns those too. Every nullable is an ADR in miniature: `DisplayName?` because providers may not send one; `Locale?` because null means "follow the device." And `Guid.CreateVersion7()` — UUIDv7 ids are time-ordered, so Postgres b-tree inserts append instead of splattering randomly (a real index-locality win over v4), while staying globally unique and unguessable enough for URLs. Add the `IEntityTypeConfiguration<User>` (unique index on `Email` — one account per address is a *rule*, so it's a *constraint*), and a migration.

3. What a JWT actually is — sixty seconds of anatomy

A JWT is three base64url segments: `header.payload.signature`. The payload is just JSON claims — **readable by anyone who holds the token** (paste one into any decoder). The signature is an HMAC-SHA256 over the first two segments with your `Jwt:Secret`: it makes the token *tamper-evident*, not *secret*. Two consequences drive everything else:

1 Never put secrets in claims. Claims are identity statements, not storage.

2 Validation is pure computation. Any API instance can verify a token with just the

signing key — no session table, no auth-server round-trip. That's why JWTs scale horizontally for free, and also their weakness: *you can't revoke pure math*. A stolen access token is valid until it expires. The mitigation is the two-token design: access tokens are **short-lived** (60

minutes here) and renewal runs through a *revocable* refresh token — which is exactly Lesson 2.2.

4. Green — the token service

The real issuer, trimmed to its spine:

```
// src/Api/Services/JwtTokenService.cs
public class JwtTokenService(IJwtSettings settings, TimeProvider clock,
                             ILogger<JwtTokenService> logger) : IJwtTokenService
{
    public string IssueAccessToken(Guid userId, string email, string provider,
                                   string? displayName = null, string? tenantName = null,
                                   string? locale = null, Guid? tenantId = null)
    {
        var key = new SymmetricSecurityKey(Encoding.UTF8.GetBytes(settings.SecretKey));
        var credentials = new SigningCredentials(key, SecurityAlgorithms.HmacSha256);

        var claims = new List<Claim>
        {
            new(ClaimTypes.NameIdentifier, userId.ToString()),
            new(ClaimTypes.Email, email),
            new(JwtClaims.Provider, provider),
            // jti is an opaque uniqueness token, not a DB key — random Guid, not UUIDv7.
            new(JwtRegisteredClaimNames.Jti, Guid.NewGuid().ToString()),
        };
        claims.Add(new Claim(ClaimTypes.Name,
                             string.IsNullOrEmpty(displayName) ? email : displayName));
        if (tenantId is { } tid)
            claims.Add(new Claim(JwtClaims.TenantId, tid.ToString()));
        // ... tenant_name, locale: same optional pattern

        var token = new JwtSecurityToken(
            issuer: settings.Issuer,
            audience: settings.Issuer,
            claims: claims,
            expires: clock.GetUtcNow().UtcDateTime.AddMinutes(settings.ExpiryMinutes),
            signingCredentials: credentials);

        return new JwtSecurityTokenHandler().WriteToken(token);
    }
}
```

Details that are decisions:

- ****IJwtSettings and TimeProvider injected**** — 1.4's validated config and the

injectable clock (formally taught in 3.3; note it's *already* here — expiry math must be testable, so ambient `DateTime.UtcNow` never appears).

- **The claim set is the contract.** `JwtClaims` centralizes the custom names as

constants — and its doc comment notes they *mirror the client's* `AppClaims`: these strings cross the wire, so they're API surface, named once. The `tenant_id` claim is optional here and earns its keep in 2.6, where it becomes the input to all data isolation. The `jti` comment repays reading: a *uniqueness* nonce wants randomness (v4), a *database key* wants ordering (v7) — the platform uses each where it belongs.

- **Validation parameters come from one place**

(JwtValidation/settings.CreateParameters()):

the JWT-bearer middleware and the service's own `ValidateToken` share them, so "what makes a token valid" can't fork into two subtly different answers.

Wire up JWT-bearer in `Program.cs` (`AddAuthentication().AddJwtBearer(...)` using the same `IJwtSettings` instance `AddAppSettings` returned — the single-source trick from 1.4), add `app.UseAuthentication()/UseAuthorization()`, and a first protected endpoint: `GET /api/auth/me` on the nascent `AuthController` returning the caller's claims.

5. Red first, though — the tests that drove it

TDD order restored for the record — these exist before the service does (`TokenServiceTests` in the reference repo):

```
[Fact]
public void Issued_token_roundtrips_and_carries_identity()
{
    var svc = CreateService(); // stub IJwtSettings + FakeTimeProvider
    var token = svc.IssueAccessToken(userId, "a@b.test", provider: "test");

    var principal = svc.ValidateToken(token);

    Assert.NotNull(principal);
    Assert.Equal("a@b.test", principal!.FindFirstValue(ClaimTypes.Email));
}

[Fact]
public void Expired_token_is_rejected()
{
    var clock = new FakeTimeProvider(start);
    var svc = CreateService(clock);
    var token = svc.IssueAccessToken(userId, "a@b.test", "test");

    clock.Advance(TimeSpan.FromMinutes(61)); // past ExpiryMinutes=60

    Assert.Null(svc.ValidateToken(token));
}

[Fact]
public void Tampered_token_is_rejected()
{
    var svc = CreateService();
    var token = svc.IssueAccessToken(userId, "a@b.test", "test");
    var tampered = token[..^4] + "AAAA"; // corrupt the signature tail

    Assert.Null(svc.ValidateToken(tampered));
}
```

Expiry, tamper, roundtrip — the three properties that *are* a token service. The `FakeTimeProvider` makes the expiry test deterministic and instant; that's the injected clock paying rent already. At the HTTP level, the 1.2 harness gets its first auth-shaped assertions: `/api/auth/me` → 401 bare, 200 with a minted token. (In the reference repo those live in the integration harness you'll grow in 2.6's wire tests.)

6. Harden — the generator and hasher seams

Beyond JWTs, the platform is about to need many *opaque* tokens — refresh tokens (2.2), magic-link tokens (2.4), invitations (2.9). Two five-line services, built now, set the pattern for all of them:

```
public class TokenGenerator : ITokenGenerator
{
    public string GenerateToken()
        => Convert.ToBase64String(RandomNumberGenerator.GetBytes(64)); // CSPRNG, 512 bits
}

public class TokenHasher : ITokenHasher
{
    public string HashToken(string token)
        => Convert.ToBase64String(SHA256.HashData(Encoding.UTF8.GetBytes(token)));

    public bool Verify(string rawToken, string storedHash)
    {
        var computed = SHA256.HashData(Encoding.UTF8.GetBytes(rawToken ?? string.Empty));
        byte[] stored;
        try { stored = Convert.FromBase64String(storedHash ?? string.Empty); }
        catch (FormatException) { return false; }
        // FixedTimeEquals is constant-time and returns false (not throws) on a length
        // mismatch, so a malformed stored hash is simply rejected.
        return CryptographicOperations.FixedTimeEquals(computed, stored);
    }
}
```

Three security reflexes, in fifteen lines:

- ****RandomNumberGenerator, never Random.**** Random is predictable by design; a

guessable token is no token. CSPRNG or nothing, for anything security-bearing.

- **The database stores hashes, never tokens.** The rule that names the lesson: when a refresh token or magic link lands in a table, it lands as SHA-256. A leaked database then yields nothing replayable — the attacker holds hashes of credentials, not credentials. (Why plain SHA-256, when passwords demand bcrypt/argon2? Because these are 512-bit *random* values — brute-forcing the preimage is hopeless. Slow hashes exist to protect *low-entropy* secrets. Using the right hash for the entropy at hand is itself the lesson.)

- ****FixedTimeEquals.**** Naïve == returns early on the first differing byte; timing

the difference leaks how much of a guess was right. Constant-time comparison closes the side channel — and note the *fail-closed* posture of `Verify`: malformed stored hash → `false`, not an exception.

7. Run it

```
dotnet test tests/Api.Tests # token roundtrip/expiry/tamper + 401/200 wire tests
dotnet run --project src/Api
# then, with a token minted via a test helper or the temporary dev endpoint:
curl -k https://localhost:7160/api/auth/me -H "Authorization: Bearer eyJ..."
```

Paste the token into a JWT decoder and *look* at it: your claims, in plain sight. Internalize that — everything in a JWT is public; only its integrity is protected.

8. Architecture Decision

ARCHITECTURE DECISION

The fork: adopt ASP.NET Core Identity, or issue our own JWTs + refresh tokens?

Chosen: custom (ADR-002). The platform is passwordless-or-OAuth — Identity's hardest-won feature (password storage) is dead weight, while the things this product needs precise control over (claim contents driving tenancy, multiple sign-in paths converging on one issue-point, MFA step-up, native token transport) are exactly where a framework's opinions cost most. Token mechanics without passwords are small, well-understood, and — decisive point — *fully testable by us* (expiry, tamper, rotation, reuse all have first-class tests).

Rejected: Identity — for *this* product. The trade is real: we own every auth bug, and several later lessons exist to pay that debt honestly. Also rejected: an external IdP (Auth0/Entra) — a fine choice in many orgs, but it outsources the exact seams this course exists to teach, and its per-user pricing fights a template meant to run free.

The trade: ~1,500 lines of auth code across Part 2 that Identity would have given us "free," in exchange for auth that bends to the product instead of the reverse — and zero framework fights in lessons 2.4, 2.5, 6.5, and the native appendix.

9. Checkpoint

```
git add -A && git commit -m "feat: users + JWT access tokens – custom auth foundation (ADR-002)"
git tag lesson-2.1
```

You should now be able to: decode a JWT by eye and say which parts are secret (none) and which are protected (integrity); justify custom-auth-here without defending custom-auth-everywhere; explain hash-only storage and why *fast* SHA-256 is correct for *high-entropy* tokens; and name the three properties every token service must prove by test.

Next: Lesson 2.2 — *Rotating refresh tokens & sessions* — how a 60-minute token becomes a 30-day session without becoming a 30-day theft window.

Lesson 2.2 — Rotating refresh tokens & sessions

Part 2 · Identity & tenancy. Lesson 2.1 left a deliberate tension: access tokens expire in 60 minutes (because they're unrevocable), but nobody will use a product that logs them out hourly. The classical resolution is the **refresh token** — and the interesting engineering isn't the happy path, it's the two attacker-facing properties: rotation (a refresh token works *once*) and **reuse detection** (a replayed old token doesn't just fail — it burns the thief's entire access).

Goal: POST /api/auth/refresh exchanges a valid refresh token for a fresh access token + a *new* refresh token; the used token is dead; replaying a dead token revokes every session the user has.

Concepts & principles: the two-token session model; rotation; reuse-as-theft-signal; the HttpOnly cookie as token transport (and each cookie attribute as a decision); server-side revocability.

Maps to: ADR-002 · repo: src/Core/Entities/RefreshToken.cs, src/Api/Services/RefreshTokenService.cs, SessionService.cs, CookieService.cs, src/Api/Controllers/AuthController.cs (/refresh, /logout), tests/Api.Tests/SessionServiceTests.cs, CookieServiceTests.cs.

Prerequisites: 2.1 (JWTs, TokenGenerator/TokenHasher).

1. Motivate — the session-length trilemma

Pick two: long sessions, revocability, stateless validation. A pure-JWT design that stretches expiry to 30 days gets long sessions and statelessness — and a stolen token that works for a month with no kill switch. Short-lived JWTs alone get revocability by expiry — and hourly logouts. The two-token model takes all three by *splitting the job*:

- The **access token** (2.1) stays short-lived and stateless — validated by signature on

every request, no database touch, horizontally free.

- The **refresh token** is long-lived (30 days) but *stateful*: an opaque random value

whose hash lives in a table, checked against that table exactly once per hour-ish, when the client exchanges it. Server-side state means server-side control: revoke the row, kill the session.

The pattern's cost is that the refresh token is now the crown jewel — steal one and you own the session for a month. So the whole design bends toward protecting *it*: it rides in an HttpOnly cookie JavaScript can't read, it's stored only as a hash, and — the crucial move — it's **single-use**.

2. The entity — a session record wearing a token costume


```
// src/Core/Entities/RefreshToken.cs – the real entity
public class RefreshToken
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public Guid UserId { get; set; }
    public required string TokenHash { get; set; } // never the token itself
    public DateTimeOffset IssuedAt { get; set; }
    public DateTimeOffset ExpiresAt { get; set; }
    public bool IsRevoked { get; set; }
    public required string IssuedFromIp { get; set; }

    /// <summary>OAuth provider used at sign-in for this session; carried through
    /// rotation so refreshed JWTs keep an accurate provider claim.</summary>
    public required string Provider { get; set; }
}
```

Notice what this table is: the platform's session store. Each row is one device's session; `RevokeAllForUser` is "sign me out everywhere"; `IssuedFromIp` is forensics. And notice `IsRevoked` — rotated tokens are **flagged, not deleted**. That's not housekeeping laziness; it's the entire reuse-detection mechanism, as you're about to see. (A scheduled cleanup job eventually prunes long-dead rows — Lesson 4.4, and now you know why that job exists.)

3. Issuing — the seams compose

```
// src/Api/Services/RefreshTokenService.cs (core of the real file)
public async Task<IssuedRefreshToken> IssueRefreshTokenAsync(
    Guid userId, string ipAddress, string provider, CancellationToken ct = default)
{
    var rawToken = tokenGenerator.GenerateToken(); // 512-bit CSPRNG (2.1)
    var tokenHash = tokenHasher.HashToken(rawToken); // SHA-256 (2.1)
    var now = clock.GetUtcNow();

    var refreshToken = new RefreshToken
    {
        Id = Guid.CreateVersion7(),
        UserId = userId,
        TokenHash = tokenHash,
        IssuedAt = now,
        ExpiresAt = now.AddDays(settings.ExpiryDays),
        IsRevoked = false,
        IssuedFromIp = ipAddress,
        Provider = provider,
    };

    var created = await repository.CreateAsync(refreshToken, ct);
    return new IssuedRefreshToken(rawToken, created); // raw → client; hash → DB
}
```

The return type says the security model in one line: `IssuedRefreshToken(RawToken, Token)` — the raw value goes to the client and is *never seen again by the server as plaintext*; the entity holds only the hash. Its doc comment states the payoff: "a database leak cannot be used to forge cookies." The generator and hasher you built in 2.1 as ten throwaway lines are now load-bearing — that's what a seam buys.

4. The refresh endpoint — where rotation and detection live

The service classifies; the controller responds. First the classifier — a presented token is exactly one of four things:

```
public enum RefreshTokenStatus
{
    Valid, // active, unexpired – safe to rotate
    Expired, // known token past its expiry – reject (no theft signal)
    Unknown, // no such hash – reject as a plain bad token
    Reuse, // a previously-rotated token is being REPLAYED – theft signal
}

public async Task<RefreshTokenInspection> InspectRefreshTokenAsync(string rawToken, ...)
{
    var token = await repository.GetByHashAsync(tokenHasher.HashToken(rawToken), ct);
    if (token is null) return new(RefreshTokenStatus.Unknown, null);
    // A revoked row whose hash is being presented = a rotated-out token replayed → theft.
    if (token.IsRevoked) return new(RefreshTokenStatus.Reuse, token);
    if (token.ExpiresAt <= clock.GetUtcNow()) return new(RefreshTokenStatus.Expired,
    token);
    return new(RefreshTokenStatus.Valid, token);
}
```

And the endpoint (abridged from the real `AuthController.Refresh`):

```
var inspection = await refreshTokenService.InspectRefreshTokenAsync(rawToken, ct);
if (inspection.Status == RefreshTokenStatus.Reuse)
{
    // Replay of a rotated-out token assume theft: revoke every session for the user.
    // Client still gets the generic error below, so the reuse signal isn't leaked.
    await refreshTokenService.RevokeAllUserTokensAsync(inspection.Token!.UserId, ct);
    logger.LogWarning("Refresh-token reuse detected for user {UserId}; revoked all
    sessions", ...);
}
if (inspection.Status != RefreshTokenStatus.Valid)
    return Unauthorized(new ErrorResponse("invalid_refresh_token", "Refresh token is
    invalid or expired"));

// Rotate: revoke the used token, then issue a fresh session.
await refreshTokenService.RevokeRefreshTokenAsync(validToken.Id, ct);
var session = await sessionService.IssueAsync(user, validToken.Provider, ClientIp, native,
ct);
cookieService.SetRefreshTokenCookie(Response, session.RefreshToken, Request);
return Ok(session.Response);
```

Walk the threat model, because every line answers an attack:

Why rotate at all? If refresh tokens were reusable, a stolen one is a 30-day skeleton key and you'd never know. Rotation makes every token single-use — so after a theft, *two parties* hold single-use tokens: the attacker and the legitimate user. Whichever redeems first wins... and the *second* redemption trips `Reuse`.

****Why is Reuse a five-alarm signal?**** Because in a correct client, a rotated-out token is *gone* — replaced the moment it was used. The only party that presents a dead token's hash is someone who copied it earlier: a thief (or, rarely, a client bug — which also deserves investigation). The response is deliberately maximal: revoke **all** the user's sessions. Attacker *and* user get logged out everywhere; the user re-authenticates and is safe; the attacker holds nothing. One stolen token cost the attacker their entire foothold.

****Why the same generic 401 for Reuse and Unknown?**** Look at the comment: "the reuse signal isn't leaked." If reuse returned a distinguishable response, an attacker could *probe*: present a captured token and learn from the error whether it's been used yet — i.e., whether it's still worth redeeming, or whether their theft has been noticed. Same principle as 1.5's error envelope: what an error *doesn't say* is part of its design. The signal goes where it belongs — the audit log and a `LogWarning` — not to the caller.

****Why Expired ≠ Reuse?**** An expired-but-never-rotated token is a user who left a laptop closed for a month — normal life, not theft. Treating it as an attack would mass- revoke sessions on every stale device. Classifying failure modes *separately, then responding differently* is the deep skill this endpoint teaches.

5. The cookie — every attribute is a decision

The refresh token travels in a cookie, and the real `CookieService` sets each attribute for an articulable reason:

```
response.Cookies.Append(RefreshTokenCookieName, token, new CookieOptions
{
    HttpOnly = true, // JS can never read it → XSS can't exfiltrate it
    Secure = request.IsHttps, // HTTPS-only transport (see the dev nuance below)
    SameSite = SameSiteMode.Lax, // sent on same-site + top-level nav; CSRF-resistant
    Path = "/api/auth", // sent ONLY to auth endpoints – nothing else needs it
    // MaxAge/Expires: matches the token's ExpiresAt
});
```

- ****HttpOnly**** is the reason the refresh token lives in a cookie at all rather than

`localStorage`: a script-injection (XSS) can call your API as the user *while the page is open*, but it cannot *steal the durable credential* — the browser withholds it from JavaScript entirely.

- ****Path=/api/auth**** is least-privilege for cookies: the token accompanies exactly the

three endpoints that need it (refresh, logout, sign-in) and never rides along on every API call, shrinking both attack surface and bytes-per-request. (The real file carries battle scars here — a long comment about expiring a legacy `Path=/` cookie that would otherwise *shadow* the scoped one. Cookies have sharp edges; the comments are the platform remembering a real 3-hour debug.)

- ****SameSite=Lax + Secure = request.IsHttps**** encode a dev/prod reality: in dev the

client (:7008) and API (:7160) are different *ports* — same-site, so `Lax` works — and `Secure` follows the actual scheme because a hardcoded `Secure=true` silently drops the cookie on any non-HTTPS leg. The comment in the real file names the failure: "silently breaks refresh." In production this whole class of problem is *deleted* rather than configured — single-origin hosting (8.1) puts client and API on one origin. That decision is being planted here, two parts early, and this cookie is why.

Notice also what the endpoint does *not* require: `[Authorize]`. Refresh runs on the cookie alone — by design, because its whole purpose is to work *after* the access token expired. (Native clients, which have no cookie jar, send the token in the body under an `X-Native-Client` header — same endpoint, different transport; the appendix picks this up.)

6. Red — the tests that pin the properties

From `SessionServiceTests`/the integration suite — the three that matter, in Gherkin first (the story format from `WAYS_OF_WORKING.md`):

```
Scenario: refresh rotates the token
  Given a signed-in user with refresh token R1
  When the client exchanges R1
  Then it receives a new access token and refresh token R2
  And exchanging R1 again fails

Scenario: replaying a rotated token burns all sessions
  Given tokens R1 (rotated out) and R2 (current) for the same user
  When an attacker presents R1
  Then the response is the same generic 401 as any bad token
  And R2 no longer works either

Scenario: expired token is rejected without the theft response
  Given a token past ExpiresAt that was never rotated
  When the client presents it
  Then the response is 401
  And the user's other sessions remain valid
```

The second scenario is the one teams skip — and it's the *entire point* of rotation. Write it with two clients: rotate via client A, replay the old token via client B, then assert client A's *new* token is also dead. `FakeTimeProvider` drives the third.

7. Run it

```
dotnet test tests/Api.Tests
# then live, via the .http scratchpad: sign in (temporary dev seed), then
# POST {{host}}/api/auth/refresh → 200, Set-Cookie with a NEW value
# POST again with the OLD cookie → 401 invalid_refresh_token (and a warning in the API log)
```

Watching the reuse warning appear in the log while the client sees only a generic 401 — that's the design, observable.

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how do sessions extend? (a) Long-lived JWTs (30-day expiry, no server state); (b) server sessions for everything (session id cookie, every request hits the store); (c) short JWT + rotating single-use refresh token with reuse detection.

Chosen: (c) — the standard modern trade (ADR-002). Per-request validation stays stateless (JWT signature); the stateful check runs once per access-token lifetime, at the refresh exchange, which is also the *choke point* where revocation, rotation, and theft detection all naturally live. Reuse-response = revoke-all: maximal, cheap, and user-recoverable (they just sign in again).

Rejected: (a) — unrevocable month-long credentials; a breach response of "rotate the signing key and log out *every user on the platform*" is not a plan. (b) — a DB hit on every API call buys revocation-latency we don't need at the cost of the JWT's free horizontal scaling; reasonable elsewhere, wrong here.

The trade: more moving parts than either pure option — four token states, cookie path subtleties, a cleanup job (4.4), and the client must handle mid-flight 401s by refreshing and retrying (3.4 builds exactly that). Each part is testable, and each got its test before its code.

9. Checkpoint

```
git add -A && git commit -m "feat: rotating refresh tokens - sessions, reuse detection,
HttpOnly cookie transport"
git tag lesson-2.2
```

You should now be able to: draw the two-token model and say which token is stateless and why; explain rotation as the mechanism that *converts theft into a detectable event*; justify revoke-all-on-reuse and the deliberately indistinguishable 401; and defend every attribute on the refresh cookie, including the path scoping.

Next: *Lesson 2.3 — Email I: the `IMailSender` seam* — before magic links can be mailed, the platform needs exactly one way to send email.

Lesson 2.3 — Email I: the `ISender` seam

Part 2 · Identity & tenancy. Before the next lesson can mail a magic link, the platform needs to send email — and *how* you introduce that capability is a masterclass in the seam pattern. One interface in Core; MailKit quarantined behind it so hard that a test fails if it leaks; Mailpit as the dev transport so your inbox stays a test fixture; and branded, localized templates with the one logo-embedding technique real mail clients actually render. Email is also this book's best *long-game* example: the five-line interface you write today gets decorated into crash-safety in Lesson 4.2 without touching a single call site.

Goal: `ISender` exists in Core; an SMTP implementation (MailKit) lives in Infrastructure/Email/ and nowhere else — enforced by a test; a branded test email lands in Mailpit's inbox.

Concepts & principles: the seam (swap-point interface); dependency quarantine with a containment test; SMTP dev/prod duality (Mailpit ↔ Brevo, same code); email HTML as a hostile rendering target; CID inline images; server-side localization by explicit culture.

Maps to: auth rules ("ISender is the only way to send email") · repo:
`src/Core/Abstractions/ISender.cs`, `src/Infrastructure/Email/SmtpSender.cs`,
`SmtpSettings.cs`, `BrandedEmail.cs`, `EmailStrings.resx`, `Assets/logo.png`,
`tests/Api.Tests/DocAndConfigSyncTests.cs` (MailKit_StaysBehindInfrastructureEmail),
`tests/Api.Tests/SmtpSettingsTests.cs`.

Prerequisites: 0.2 (Mailpit is already running in your compose stack); 1.4 (options pattern).

1. Motivate — email is a vendor decision wearing a feature costume

"Send an email" sounds like a feature. It's actually a *vendor relationship*: dev wants a fake inbox, staging wants a free SMTP relay (Brevo), production might want SES or Postmark next year, and tests want to assert "an email was sent" without any network at all. If new `SmtpClient()` calls are scattered through your services, every one of those is a migration. If instead there is exactly **one** interface and callers know nothing else, every one of those is a config change or one new class.

That's the seam pattern, and email is its cleanest first outing because the swap pressure is real and immediate — you'll swap transports *today* (Mailpit locally, Brevo when staging arrives in 8.2) without either "side" of the seam noticing.

2. The interface — small, but every word chosen

```
// src/Core/Abstractions/EmailSender.cs – the real file
public interface IEmailSender
{
    /// <summary>
    /// Sends an HTML email. <paramref name="inlineImages"/> are embedded in the message
    /// (multipart/related) and referenced from the HTML via <c>cid:{ContentId}</c> – the
    /// only logo-embedding approach mainstream clients (Gmail/Outlook) render reliably.
    /// <para>Throws <see cref="EmailSendException"/> if delivery to the SMTP server fails;
    /// a hung server is bounded by a per-send timeout rather than stalling the request
    /// indefinitely.</para>
    /// </summary>
    Task SendAsync(string to, string subject, string htmlBody,
        IReadOnlyList<EmailInlineImage>? inlineImages = null, CancellationToken
        cancellationToken = default);
}

public sealed record EmailInlineImage(string ContentId, string FileName, byte[] Content,
    string MediaType);

public sealed class EmailSendException(string message, Exception innerException)
    : Exception(message, innerException);
```

Study the *placement* first: this lives in **Core**, which references nothing — so the interface mentions no MailKit type, no MimeMessage, nothing SMTP-flavored. It speaks the domain's language ("send this HTML to this address, these images inline") and lets Infrastructure translate. Note also the custom EmailSendException: callers catch a *domain* exception, not MailKit.Net.Smtp.SmtpCommandException — because catching a vendor's exception type is a vendor reference in disguise, and it would quietly weld call sites to the implementation the interface exists to hide.

3. The implementation — and the sharp edges it rounds off

```
// src/Infrastructure/Email/SmtpEmailSender.cs (the real file, lightly trimmed)
public class SmtpEmailSender(IOptions<SmtpSettings> options, ILogger<SmtpEmailSender>
logger) : IEmailSender
{
    private readonly SmtpSettings _settings = options.Value;

    public async Task SendAsync(string to, string subject, string htmlBody,
        IReadOnlyList<EmailInlineImage>? inlineImages = null, CancellationToken ct =
default)
    {
        var message = new MimeMessage();
        message.From.Add(new MailboxAddress(_settings.FromName, _settings.FromAddress));
        message.To.Add(MailboxAddress.Parse(to));
        message.Subject = subject;

        var builder = new BodyBuilder { HtmlBody = htmlBody };
        if (inlineImages is not null)
            foreach (var image in inlineImages)
            {
                var resource = builder.LinkedResources.Add(
                    image.FileName, image.Content, ContentType.Parse(image.MediaType));
                resource.ContentId = image.ContentId; // referenced from HTML as
cid:{ContentId}
            }
        message.Body = builder.ToMessageBody();

        using var client = new SmtpClient
        {
            // Bound a hung server: without this, a stuck Connect/Send stalls the request
            Timeout = _settings.TimeoutSeconds * 1000,
            CheckCertificateRevocation = _settings.CheckCertificateRevocation,
        };
        try
        {
            await client.ConnectAsync(_settings.Host, _settings.Port,
SecureSocketOptions.Auto, ct);
            if (!string.IsNullOrEmpty(_settings.Username))
                await client.AuthenticateAsync(_settings.Username, _settings.Password ??
string.Empty, ct);
            await client.SendAsync(message, ct);
            await client.DisconnectAsync(true, ct);
        }
        catch (OperationCanceledException)
        {
            throw; // genuine cancellation – let it propagate, don't mask as a send failure
        }
        catch (Exception ex)
        {
            logger.LogError(ex, "SMTP send to {Host}:{Port} failed", _settings.Host,
_settings.Port);
            throw new EmailSendException($"Failed to send email via
{_settings.Host}:{_settings.Port}.", ex);
        }
    }
}
```

Three details that separate production code from sample code:

- **The timeout.** SMTP servers hang; without a bound, the *user's sign-in request*

hangs with them. The `Timeout` line converts "infinite stall" into "bounded failure" — and its value is `config(SmtpSettings.TimeoutSeconds)`, because 1.4 taught us where numbers like this live. (`CheckCertificateRevocation` is likewise a knob — added when a corporate network's blocked CRL endpoint made every send fail; a real scar with a real test, `SmtpSettingsTests`.)

- ****The `OperationCanceledException` pass-through.**** A cancelled request isn't a failed

email; catching it into `EmailSendException` would lie to callers and pollute logs with phantom SMTP errors. Distinguishing *cancellation* from *failure* is a habit worth stealing for every I/O wrapper you ever write.

- **Auth only when configured.** Mailpit takes no credentials; Brevo requires them. One

if makes the same code serve both — the dev/prod duality lives in config, not code.

Add `SmtpSettings` (host/port default to Mailpit — `localhost:1025` — so a fresh checkout sends to the trap with *zero* config), the `.env.example` block for a real provider (commented out, per the contract), and register in DI: `services.AddScoped<IEmailSender, SmtEmailSender>()`.

4. The containment test — a seam with a fence

An abstraction "everyone should use" lasts until the first hurried Tuesday. The platform makes the quarantine mechanical (from the real `DocAndConfigSyncTests`):

```
[Fact]
public void MailKit_StayBehindInfrastructureEmail()
{
    // The IEmailSender abstraction is the only way to send email; the concrete
    // MailKit/MimeKit dependency must never leak outside src/Infrastructure/Email/.
    var offenders = SourceFiles(Path.Combine(RepoRoot(), "src"))
        .Where(f => !f.Contains(Path.Combine("Infrastructure", "Email")))
        .Where(f => File.ReadAllText(f) is var t && (t.Contains("using MailKit")
            || t.Contains("using MimeKit") || t.Contains("MailKit.") ||
            t.Contains("MimeKit.")))
        .Select(Path.GetFileName).ToList();

    Assert.True(offenders.Count == 0,
        $"MailKit/MimeKit must stay behind IEmailSender: {string.Join(", ", offenders)}");
}
```

You've now built this move three times (config sync, error shapes, and this) — but note what's different here: this is the first time a *dependency* is fenced rather than a practice. The rule "only Infrastructure/Email knows MailKit exists" stops being a convention the moment this test exists; it's a property of the codebase. When 4.2 wraps this sender in an outbox decorator, and when a future maintainer swaps MailKit for something else entirely, the blast radius is one folder — *provably*.

5. Branded, localized — email HTML is a hostile country

A magic-link email is often the *first* thing your product ever shows a user — it should look like the product. `BrandedEmail` builds the HTML, and its doc comment is a compact field guide to why email templates can't just reuse your web CSS:

Email HTML is its own world — table-based layout, inline styles, web-safe fonts only — so this does NOT reuse the app's CSS. The logo is embedded via CID (multipart/related), the one approach Gmail and Outlook render reliably (they block data-URI images).

Unpack the two traps it names. **Layout:** mail clients render HTML with engines that are years behind browsers and aggressively sanitized — no external stylesheets, patchy flexbox, `<style>` blocks stripped by some clients. Table layout + inline styles isn't retro taste; it's the *only* dialect with universal support. **The logo:** you can't hotlink it (blocked-by-default remote images) and you can't data-URI it (Gmail/Outlook block those too). The working answer is **CID embedding**: the PNG travels *inside* the message as a multipart/related part with a Content-ID, and the HTML references `cid:perezosoftware-logo`. That's why IEmailSender has the `inlineImages` parameter — the seam's shape was dictated by a Gmail rendering rule, which is exactly the kind of knowledge that belongs *behind* a seam.

And localization — note the design, because it differs from UI localization in one important way:

```
private static readonly ResourceManager Rm =
    new("Perezosoftware.Infrastructure.Email.EmailStrings", typeof(BrandedEmail).Assembly);

public static CultureInfo ResolveCulture(string? locale)
{
    if (string.IsNullOrEmpty(locale)) return DefaultCulture; // "en"
    try { return CultureInfo.GetCultureInfo(locale.Trim()); }
    catch (CultureNotFoundException) { return DefaultCulture; } // fail soft, not 500
}
```

Emails are rendered **server-side, in an explicit culture** — the requester's UI language for an OTP, the *inviter's* saved locale for an invitation — so the code keys ResourceManager by a passed-in culture and never trusts the ambient thread culture (which on a server belongs to nobody in particular). EmailStrings.resx + EmailStrings.es.resx hold the strings; Lesson 3.5 adds the UI's resx pipeline and they'll rhyme. Also file away the *rebrand trap* the platform's docs call out: this inline logo and palette are brand touchpoints that live nowhere near the UI folder — lesson 9.1's checklist exists because everyone forgets them.

6. Red → Green — Mailpit is your assertable inbox

The test story for email is two-layered. Unit level: a FakeEmailSender recording (to, subject, body) — services in later lessons assert against it. Integration level: actually send, and *read the result back out of Mailpit* — it exposes a REST API over its inbox (GET `/api/v1/messages`), which is precisely why 0.2 chose it:

```
[Fact]
public async Task Branded_email_arrives_with_inline_logo()
{
    var sender = CreateRealSender(); // pointed at localhost:1025
    var body = BrandedEmail.MagicLink("https://example.test/x",
    BrandedEmail.ResolveCulture("es"));

    await sender.SendAsync("probe@test.local", body.Subject, body.Html, body.InlineImages);

    var msg = await Mailpit.LatestMessageTo("probe@test.local"); // Mailpit's REST API
    Assert.Contains("cid:", body.Html);
    Assert.Equal(body.Subject, msg.Subject); // localized subject
}
```

In 3.6 this same trick graduates to E2E: Playwright journeys will fish OTP codes out of Mailpit to complete real sign-ins, browser-to-inbox-to-browser.

7. Run it

```
docker compose up -d # Mailpit already in your stack from 0.2
dotnet test tests/Api.Tests # containment + settings + send tests
# then send one for real (temporary dev endpoint or test), and open:
# http://localhost:8025 → your branded email, logo rendered, in Spanish if you asked
```

Seeing the actual rendered email in Mailpit's UI is worth thirty assertions: fonts, layout, logo — this is what your user's first impression looks like.

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how does the platform send email? (a) SMTP client calls inline where needed; (b) a vendor SDK (SendGrid/SES) used directly; (c) one Core seam (`IEmailSender`), an SMTP implementation behind it, containment enforced by test.

Chosen: (c), with **SMTP as the protocol** rather than any vendor's API — SMTP is the one interface every provider (Mailpit, Brevo, SES, Postmark) speaks, so even the *implementation* is vendor-portable; switching providers is a config change (8.2 does exactly this for staging). The seam keeps callers domain-pure and gives 4.2 its decoration point.

Rejected: (a) welds a vendor and its failure modes into every call site — the anti-seam. (b) buys nicer delivery analytics at the price of a per-vendor rewrite and mocks that mimic a vendor SDK; if those analytics are ever needed, they arrive as *another implementation of the same seam*.

The trade: SMTP is lowest-common-denominator — no template hosting, no delivery webhooks (bounces land in a provider dashboard, not your code). Accepted; a platform can graduate individual needs later without touching a single caller.

9. Checkpoint

```
git add -A && git commit -m "feat: email seam - IEmailSender, MailKit quarantined, branded  
localized templates"  
git tag lesson-2.3
```

You should now be able to: place an interface on the correct side of the dependency rule and say why the exception type matters as much as the method signature; explain CID embedding and why email HTML can't share the app's CSS; state why emails localize by explicit culture; and write a containment test for any dependency you ever decide to quarantine.

Next: *Lesson 2.4 — Passwordless: magic link + OTP* — the seam gets its first real customer, and the platform gets its first sign-in.

Lesson 2.4 — Passwordless: magic link + OTP

Part 2 · Identity & tenancy. The platform gets its first sign-in — and it's not a password form. Two passwordless credentials, one entity: the **magic link** (a long random token in an emailed URL) and the **OTP** (a short numeric code you type). The magic link is easy. The OTP is where this lesson earns its place, because a 6-digit code has a million possibilities and an attacker has a keyboard — the interesting engineering is entirely in the *brute-force arithmetic* and the lockout design that answers it, including one genuinely subtle bug ("resend-proof lockout") the platform's audit caught.

Goal: a user with nothing but an email address can sign in twice over — click a mailed link, or type a mailed code — and both credentials are single-use, time-limited, stored only as hashes, and brute-force-proof.

Concepts & principles: possession-based auth (email as the factor); high- vs low-entropy credentials and their different defenses; cumulative lockout windows; enumeration-resistant issuance *and* error mapping; account creation at redemption.

Maps to: ADR-C15, auth rules (LoginToken + PasswordlessService, lifetimes in config) · repo: src/Core/Entities/LoginToken.cs, src/Api/Services/PasswordlessService.cs, AuthController issue/redeem endpoints, tests/Api.Tests/PasswordlessServiceTests.cs, OtpErrorMappingTests.cs, tests/Core.Tests/LoginTokenTests.cs.

Prerequisites: 2.1 (generator/hasher, JWT issue), 2.2 (sessions to mint on success), 2.3 (email to deliver the credentials).

1. Motivate — passwords are a liability you can decline

Every password your system stores is a promise: to hash it with an expensive KDF, to survive credential-stuffing from every other site's breach, to run "forgot password" flows that are themselves attack surface, to nag about rotation. Decline the deal and the liability evaporates: **prove control of the inbox instead.** Anyone who can read ana@example.test's email is — for this product's threat model — Ana; email is already the recovery root for password systems anyway, so passwordless just removes the middleman.

Two ergonomic forms of the same proof:

	Magic link	OTP
Credential	512-bit random token in a URL	6 random digits
UX	one click, same device	type a code — survives cross-device (read on phone, type on desktop)
Entropy	$\sim 10^{77}$ possibilities	10^6 possibilities
Defense	its own entropy	<i>everything else in this lesson</i>

That entropy row is the lesson's spine. Nobody brute-forces the link. The code, though — a million possibilities, five tries a code, and unlimited codes on request? Do the math before writing the code; we'll do it in §4.

2. One entity for both — LoginToken

```
// src/Core/Entities/LoginToken.cs – the real entity
public class LoginToken
{
    public Guid Id { get; set; } = Guid.CreateVersion7();

    /// <summary>Normalized (lower-cased) email the credential was issued to.</summary>
    public required string Email { get; set; }

    /// <summary>SHA-256 hash of the raw token/code.</summary>
    public required string CodeHash { get; set; }

    /// <summary><see cref="LoginTokenPurpose"/> – magic link or OTP.</summary>
    public required string Purpose { get; set; }

    public DateTimeOffset ExpiresAt { get; set; }

    /// <summary>Set when the credential is redeemed (or locked out); null while
    usable.</summary>
    public DateTimeOffset? ConsumedAt { get; set; }

    /// <summary>Failed verification attempts – used to lock out OTP
    brute-forcing.</summary>
    public int AttemptCount { get; set; }

    public DateTimeOffset CreatedAt { get; set; }

    // Derived, never stored. The *At(now) overloads are the deterministic core (testable
    // with an explicit clock); the parameterless properties are ambient-time conveniences.
    public bool IsConsumed => ConsumedAt.HasValue;
    public bool IsExpiredAt(DateTimeOffset now) => now >= ExpiresAt;
    public bool IsValidAt(DateTimeOffset now) => !IsConsumed && !IsExpiredAt(now);
}
```

Familiar decisions compounding: hash-only storage (2.1), UUIDv7 keys, lifetimes from config (1.4 — Auth:MagicLink:TokenLifespanMinutes, Auth:Otp:*). Two new ones:

- ****IsValid is computed, never stored.**** There is no IsValid column to forget to

update — validity *derives* from ConsumedAt + ExpiresAt at the moment of asking (Golden Rule: derived values are computed, not stored as stale flags). And note the shape: IsValidAt(now) is the deterministic core the unit tests drive with an explicit clock; the ambient-time property is a convenience that delegates. Logic first, sugar second.

- ****ConsumedAt is a timestamp, not a boolean.**** Same storage cost, strictly more

information — *when* it was redeemed is forensics the platform gets for free.

3. Issue and redeem — the magic-link half

```
// src/Api/Services/PasswordlessService.cs (real code)
public async Task<string> IssueMagicLinkTokenAsync(string email, CancellationToken ct =
default)
{
    email = Normalize(email); // trim + lower — one canonical form
    await repository.InvalidateActiveAsync(email, LoginTokenPurpose.MagicLink, ct);

    var raw = tokenGenerator.GenerateToken(); // 512-bit CSPRNG (2.1's seam)
    await repository.AddAsync(NewToken(email, LoginTokenPurpose.MagicLink, raw,
settings.MagicLinkLifespanMinutes), ct);

    return raw; // → BrandedEmail → IEmailSender (2.3)
}

public async Task<User?> RedeemMagicLinkAsync(string email, string token, CancellationToken
ct = default)
{
    email = Normalize(email);
    var hash = tokenHasher.HashToken(token);
    var record = await repository.GetActiveByHashAsync(email, LoginTokenPurpose.MagicLink,
hash, ct);
    if (record is null) return null;

    record.ConsumedAt = clock.UtcNow(); // single-use: consumed atomically with success
    await repository.UpdateAsync(record, ct);

    return await userService.GetOrCreateByEmailAsync(email, ct); // ← account born HERE
}
```

Three quiet decisions:

- **Issue invalidates predecessors.** One active credential per email per purpose —

"request, request, request" leaves one live token, not a growing pile of valid ones.

- ****The account is created at *redemption*, never issuance.**** The service's doc comment

states why: "a typo'd or probed email leaves no trace." Anyone can type any address into a login box; if issuance created accounts, your Users table would fill with strangers' addresses and typos — and an attacker could *enumerate* which emails are registered by watching for "account already exists" behavior. Sign-in and sign-up collapse into one flow, and the collapse is a security feature.

- **Redemption looks up by hash** — the raw token exists only in the email and the redeeming request, never in a query log or a table.

The controller endpoints glue this to 2.3's email and, on success, to 2.2's session issuance: redeem → SessionService.IssueAsync → access token + refresh cookie. The routes are worth a glance for a subtlety: POST /api/auth/magic-link/send issues, but verify is GET /api/auth/magic-link/verify — a **GET**, because it's reached by *clicking a link in an email*, not by a scripted request. (OTP, typed by hand, is symmetric POSTs: otp/send and otp/verify.) All sign-in paths in this platform converge on that same session-mint point — remember that convergence; MFA (6.5) will exploit it.

4. The OTP half — do the arithmetic first

A 6-digit code: 10^6 possibilities. Guessing has a 1-in-a-million shot per try. Sounds fine — until you multiply. Five attempts per code and free code-reissue gives an attacker who

requests a fresh code every lockout: 5 tries, resend, 5 more tries against a *new* code with a *reset* counter... The per-code counter never trips, and over hours a patient script accumulates thousands of attempts. Worse, each resend gives *this* code its own 1-in-200,000 ($5/10^6$) chance. The naive design is soft.

The platform's answer, from the real `RedeemOtpAsync` — read the comment first, it's the audit finding that shaped the code:

```
// Cumulative, resend-proof lockout: count failed attempts across EVERY code issued to
// this email in the window, not just the current one. Issuing a fresh code
// (AttemptCount=0) can't hand the attacker another budget, and the lock holds even
// against the correct code until the window elapses (CONF-5).
var windowStart = clock.UtcNow().AddMinutes(-settings.OtpLockoutWindowMinutes);
var failuresInWindow = await repository.CountFailedAttemptsSinceAsync(
    email, LoginTokenPurpose.Otp, windowStart, ct);
if (failuresInWindow >= settings.OtpMaxAttempts)
    return new OtpResult(OtpStatus.TooManyAttempts, null);

var record = await repository.GetLatestActiveAsync(email, LoginTokenPurpose.Otp, ct);
if (record is null)
    return new OtpResult(OtpStatus.Expired, null); // none active → expired or never issued

// Timing-safe comparison — the OTP code is low-entropy, so don't leak it via
// early-exit string equality.
if (tokenHasher.Verify(code, record.CodeHash)) { /* consume + get-or-create user */ }

record.AttemptCount++; // wrong code — count it
```

The design in three moves:

- ****The budget belongs to the *email*, not the code.**** Failures are counted across every

code in the window (`CountFailedAttemptsSinceAsync`), so "resend" no longer resets anything. Five failures in fifteen minutes = locked, *even for the correct code*, until the window drains. Redo the math: 5 attempts per 15 minutes = 480/day against a code that *itself expires in 10 minutes*. The attack is dead.

- **The lockout is state the attacker can't launder.** Note what the audit comment

implies about the bug it fixed: any per-code counter is attacker-resettable *by design* (they can always request a fresh code). When you design a lockout, ask: *what action resets my counter, and can the attacker perform it?*

- **Timing-safe verify, even for six digits.** Low-entropy secrets are exactly where a

timing oracle helps most — leaking "first digit right" collapses 10^6 to 10^5 fast. 2.1's `FixedTimeEquals` hasher was built for this moment.

And the numeric code generator — one line worth pausing on:

```
digits[i] = (char)('0' + RandomNumberGenerator.GetInt32(0, 10));
```

`GetInt32(0, 10)` per digit — CSPRNG *and* modulo-bias-free, versus the classic `randomBytes % 10` bug that skews digit distribution. Small crypto details are still crypto details.

5. What the caller learns — collapse the statuses

Internally there are four OTP outcomes. Externally, there are fewer — deliberately. From the real `OtpErrors`:

```
/// "No active code" (Expired) and "wrong code" (Invalid) collapse to the SAME
/// invalid_code so a caller can't probe whether an address has an outstanding OTP;
/// the full status stays server-side.
public static string ClientCode(OtpStatus status) => status switch
{
    OtpStatus.Invalid or OtpStatus.Expired => "invalid_code",
    OtpStatus.TooManyAttempts => "too_many_attempts",
    ...
};
```

Same principle as 2.2's indistinguishable 401: the error surface is *part of the attack surface*.

"Does this address have a code outstanding?" is information (it reveals someone is mid-sign-in). Model statuses precisely inside; collapse them at the boundary (1.5's error codes); keep the distinction in logs. There's a dedicated test — `OtpErrorMappingTests` — because a *mapping* is exactly the kind of code a refactor breaks silently.

One more boundary defense wired in the controller: issuance is **rate-limited per IP** (`Auth:RateLimit:PasswordlessPermitLimit`, default 5/min) — because "email anyone a code, free, forever" is otherwise a spam cannon pointed at your SMTP reputation.

6. Red — the tests that drove it

The Gherkin scenarios (per `WAYS_OF_WORKING.md`, written before the service):

```
Scenario: magic link signs in and is single-use
  Given "ana@example.test" requested a magic link
  When she redeems it
  Then she is signed in and an account exists for her email
  And redeeming the same link again fails

Scenario: OTP lockout survives a resend
  Given an attacker has failed 5 OTP attempts for "ana@example.test"
  When a fresh code is issued and the attacker tries again — even the CORRECT code
  Then verification returns too_many_attempts until the window elapses

Scenario: expiry is enforced
  Given a magic link issued 16 minutes ago (lifespan 15)
  When it is redeemed
  Then sign-in fails
```

`PasswordlessServiceTests` drives all of these with `FakeTimeProvider` (advance past expiry; drain the lockout window) and the fake email sender from 2.3 capturing raw codes. The resend-proof scenario is the one that *distinguishes this design* — if you write one test today, write that one. In 3.6, the E2E suite re-proves the happy path end-to-end by fishing the real OTP out of Mailpit with a browser driving.

7. Run it

```
dotnet test tests/Api.Tests --filter Passwordless
dotnet run --project src/Api
# .http scratchpad: POST /api/auth/otp/send { "email": "you@test.local" }
# → open http://localhost:8025, read the branded code from Mailpit
# → POST /api/auth/otp/verify { "email": "...", "code": "123456" }
# → 200: access token + refresh cookie. You just signed in with no password in the system.
```

8. Architecture Decision

ARCHITECTURE DECISION

The fork: primary sign-in mechanism? (a) Passwords (+email reset); (b) OAuth only ("Sign in with Google, period"); (c) passwordless email credentials + OAuth as an accelerator (next lesson).

Chosen: (c) (ADR-C15). Email possession is the root of trust either way — password reset already reduces to it — so passwordless removes an entire liability class (storage, stuffing, reset flows) rather than defending it. Magic link and OTP share one entity and one service because they are one concept at two entropy levels; the low-entropy one is wrapped in cumulative lockout, rate limiting, timing-safe compare, and enumeration-resistant errors.

Rejected: (a) — maximum liability for zero differentiation; this platform would spend its auth budget defending a mechanism users increasingly don't want. (b) — excludes anyone unwilling to link a big-tech identity and makes *your* account system hostage to *their* account system; fine for internal tools, wrong default for a SaaS template.

The trade: sign-in now depends on email deliverability (SMTP down = login down — which is why 2.3 bounded that timeout, and why 4.2 will make sends crash-safe), and inbox compromise = account compromise. That last is true of password systems too — passwordless is just honest about it. MFA (6.5) exists for the users who need more.

9. Checkpoint

```
git add -A && git commit -m "feat: passwordless sign-in – magic link + OTP, cumulative
resend-proof lockout"
git tag lesson-2.4
```

You should now be able to: run the brute-force arithmetic for any short credential and size a lockout to beat it; explain why the lockout budget must attach to the email, not the code; justify account-creation-at-redemption as enumeration defense; and articulate which OTP statuses collapse at the boundary and why the *mapping itself* deserves a test.

Next: Lesson 2.5 — OAuth & account linking — "sign in with Google" in one line per provider, and the takeover-shaped edge case hiding inside account linking.

Lesson 2.5 — OAuth & account linking

Part 2 · Identity & tenancy. "Sign in with Google" looks like a convenience feature. It's actually a *federation decision* — you're accepting identity assertions from a third party — and the dangerous part isn't the OAuth dance (the framework does that); it's the moment two identities claim the same email address. This lesson builds the one-line-per-provider registration the platform advertises, and then spends its real energy where the real risk is: the **account-linking takeover guard** and the fail-closed claim normalization behind it.

Goal: Google and Microsoft sign-in work end-to-end; adding a future provider is one constant + one `.AddXxx()` block; a same-email sign-in from a *new* provider links to the existing account — but **only** when the email is provably verified.

Concepts & principles: OAuth as identity federation (what you delegate, what you keep); provider identity vs email identity; the unverified-email takeover; fail-closed normalization (R17); trust decisions as explicit, testable objects.

Maps to: ADR-002 ("new provider = one line", takeover guard) · R17 · repo:
`src/Api/Services/AuthProviders.cs`, `ClaimsExtractor.cs`, `ProviderEmailTrust.cs`,
`UserService.cs` (`GetOrCreateUserAsync`), `src/Core/Entities/UserLogin.cs`,
`LinkTokenService.cs`, `src/Infrastructure/ServiceCollectionExtensions.cs` (provider
 registration), tests: `ClaimsExtractorTests`, `ProviderEmailTrustTests`, `AuthProvidersTests`.

Prerequisites: 2.1–2.4 (all sign-in paths converge on the same session mint).

1. Motivate — what OAuth actually outsources

When a user clicks "Sign in with Google," you delegate *authentication* — Google proves who they are, with all its infrastructure (password vaults, device signals, its own MFA). What you do **not** outsource is *identity*: the decision "which account in *my* system is this person" remains entirely yours. That's the part frameworks can't do for you, and it's the part with the security cliff.

Here's the cliff, concretely. Ana signed up months ago via magic link as `ana@example.test`. Today someone signs in via OAuth-Provider-X, and X asserts their email is `ana@example.test`. If you match by email and link this new credential to Ana's account — congratulations, you may have just handed Ana's account to an attacker, because **some providers let users type any email into their profile, unverified**. The attacker didn't hack Ana; they hacked *your merge rule*.

Everything in this lesson exists to let the convenient thing (same email = same account, no duplicate accounts) happen *only* when it's safe.

2. The dance itself — buy, don't build

The OAuth redirect flow (authorize → callback → code exchange → claims) is standardized and hostile to hand-rolling; ASP.NET Core's authentication handlers do it properly. The platform's entire per-provider cost, from the real `ServiceCollectionExtensions`:

```
if (!string.IsNullOrEmpty(configuration["Authentication:Google:ClientId"]))
    auth.AddGoogle(google =>
    {
        google.ClientId = configuration["Authentication:Google:ClientId"]!;
        google.ClientSecret = configuration["Authentication:Google:ClientSecret"]!;
        // callback path, scopes...
    });
```

Note the guard clause: **a provider registers only when its ClientId is configured**. That's R21's config-gating instinct applied to auth — an undeployed provider isn't a broken button, it's an absent scheme. `AuthProviders.EnabledAsync` closes the loop by asking the runtime which schemes actually registered, so the login UI renders buttons for exactly the providers this deployment configured — no dead, 500-ing "Sign in with Microsoft" on an instance that never set it up. And `AuthProviders` itself:

```
public static class AuthProviders
{
    public const string Google = "google";
    public const string Microsoft = "microsoft";
    public static readonly IReadOnlyList<string> Supported = [Google, Microsoft];

    /// <summary>The authentication scheme to challenge for a provider, or null if
    unsupported.</summary>
    public static string? SchemeFor(string provider) => provider.ToLowerInvariant() switch
    {
        Google => GoogleDefaults.AuthenticationScheme,
        Microsoft => MicrosoftAccountDefaults.AuthenticationScheme,
        _ => null,
    };
}
```

One place for provider names instead of string literals scattered through controllers — "new OAuth provider = one line" is really "one constant, one switch arm, one `.AddXxx()`," and nothing else in the platform changes. That's the promise; keeping it depends on the next two sections *not* being provider-specific.

The callback flow: the provider's handler materializes the external identity as a short-lived **cookie principal on a dedicated "External" scheme** — a carrier, not a session. The callback endpoint reads it, runs the identity resolution below, mints the platform's own session (2.2), and discards the carrier. External assertion in; *our* tokens out; the two never blur.

3. Normalize the claims — fail closed (R17)

Providers emit claims differently: `email_verified` as "true", as true, or — the dangerous case — not at all. Normalization is one tiny class, and its posture is the whole point. From the real `ClaimsExtractor`:

```
public bool IsEmailVerified(ClaimsPrincipal principal)
{
    // Fail closed: verified ONLY when the provider explicitly asserts
    // email_verified="true". Absent / empty / any other value not verified. An
    // absent claim must not be trusted – not every provider asserts it, and the
    // platform advertises "new OAuth provider = one line", so a silent
    // absent-verified default would hand any such provider the account-takeover
    // merge path.
    var claim = principal.FindFirst("email_verified")?.Value;
    return string.Equals(claim, "true", StringComparison.OrdinalIgnoreCase);
}
```

Read the comment's chain of reasoning — it connects a *marketing promise* to a *security posture*. Because adding a provider is deliberately trivial, the system must assume some future provider will be added carelessly. If absence-of-claim defaulted to "verified," that careless one-liner would silently open the takeover path. Fail-closed normalization means the careless default is the *safe* default. This is rule R17, and it has its own unit tests (ClaimsExtractorTests): absent claim → false, "True" → true, "1" → false, garbage → false. Boring tests; load-bearing boundary.

4. The identity resolution — three cases, one guard

The heart of the lesson, from the real `UserService.GetOrCreateUserAsync`. A provider identity arrives as (provider, providerUserId, email, emailVerified) — note that ****providerUserId is the real identity****; email is metadata that can change or lie:

```
// 1. Known provider identity → existing account
var existingUser = await repository.GetByLoginAsync(provider, providerUserId, ct);
if (existingUser != null)
    return existingUser; // the (provider, id) pair is ground truth

// 2. Same email from a new provider → same account; link the identity.
// Guard: an UNVERIFIED email claim must never attach a new credential to an
// existing account (takeover vector) – refuse outright.
var userByEmail = await repository.GetByEmailAsync(email, ct);
if (userByEmail != null && !emailVerified)
{
    logger.LogWarning("Refused unverified-email merge for {Email} via {Provider}", email,
        provider);
    throw new UnverifiedEmailConflictException(email);
}
if (userByEmail != null)
{
    await repository.AddLoginAsync(new UserLogin { UserId = userByEmail.Id,
        Provider = provider, ProviderUserId = providerUserId }, ct);
    return userByEmail; // safe merge: verified email, linked
}

// 3. Brand new user – create the tenant ("household of one"), the user, and an
// owner membership linking them, atomically.
```

The three cases in threat-model terms:

- **Case 1 — the returning user.** (provider, providerUserId) matched a `UserLogin`

row: this exact Google account signed in before. No email logic runs at all — even if the user changed their email at Google, they get *their* account. The `UserLogin` entity (one account,

many provider identities) is what makes "Ana signs in with Google on Monday and Microsoft on Tuesday" one person instead of two accounts.

- **Case 2 — the merge, guarded.** New provider identity, known email. Convenience says

link; the guard says *only if the provider vouches for the email*. Unverified → refuse loudly (a typed exception the callback turns into a safe error page — the legitimate owner's account is untouched, and the refusal is logged as the security-relevant event it is).

- **Case 3 — the newcomer.** No matches: create user + their tenant + owner membership

in one transaction. (The "household of one" is a Part 2.8 concept arriving early — every user has a tenant from birth, which is what lets the `tenant_id` claim of 2.6 be non-optional in practice.)

One nuance sits on top: `ProviderEmailTrust`. Microsoft's v2 endpoint *omits* `email_verified` even for accounts whose email Microsoft itself verified (personal accounts). Strictly fail-closed would force those users through case-2 refusal forever. The platform's answer is a tiny, explicit trust object:

```
public bool TrustsEmailWithoutClaim(string provider) => provider.ToLowerInvariant() switch
{
    AuthProviders.Google => true, // Google always asserts the claim anyway
    AuthProviders.Microsoft => IsMicrosoftConsumersTenant(), // personal accounts only
    _ => false, // unknown/future providers stay fail-closed
};
```

Study the *shape* more than the rules: the exception to fail-closed is not an `if` buried in the merge — it's a named, injectable, unit-tested policy object (`ProviderEmailTrustTests`) whose default arm is still `false`, and whose Microsoft arm narrows to exactly the sub-case with a documented justification (consumers-tenant emails are Microsoft-verified) that *reverts automatically* if the deployment flips the configured tenant to work/school accounts. When you must carve an exception into a security rule, this is what the carving should look like: explicit, narrow, self-reverting, tested.

5. Explicit linking — the same guard from the other door

Settings will eventually offer "Connect your Google account" for an already-signed-in user (`LinkTokenService` carries the intent through the OAuth round-trip as — of course — a hashed, single-use, time-limited token; the 2.4 pattern reused). Linking *while authenticated* proves control of the platform account, so the email-match question doesn't arise — but the *other* direction still bites: the OAuth identity being attached must not already belong to another user. Same lesson, different door: every path that can attach a credential to an account needs the takeover question asked. The reference repo keeps both paths' checks in `UserService` so there's one answer.

6. Red — the tests

OAuth's redirect dance is miserable to test through a real provider — so the platform tests the *decisions* as units and the *flow* with a stand-in. The integration harness (1.2/2.6) swaps the External scheme's handler for one that authenticates from test headers — the callback then runs *its real code* against controlled identities:

```

Scenario: same verified email links, doesn't duplicate
  Given an account exists for "ana@example.test" (created via magic link)
  When a Google identity asserting email_verified=true for that email signs in
  Then no new account is created
  And the Google identity is linked to Ana's account

Scenario: unverified email must not merge (the takeover test)
  Given an account exists for "ana@example.test"
  When a provider identity with that email but NO verified claim signs in
  Then sign-in is refused with a safe error
  And Ana's account has no new login attached

Scenario: returning provider identity ignores email drift
  Given a user who previously signed in with Google id "g-123"
  When "g-123" signs in again asserting a different email
  Then they get their original account

```

The middle one is this lesson's reason to exist. If your suite has one OAuth test, make it that one.

7. Run it

Real-provider setup is config, not code: create OAuth apps in Google/Microsoft consoles, put ClientId/Secret in `.env` (the `.env.example` block from 1.4 documents the keys), run, click the button, land signed in. Then verify the *linking* behavior: sign in via OTP as you@. . ., sign out, sign in via Google with the same address — one account, two rows in `UserLogins`.

```
dotnet test tests/Api.Tests --filter "ClaimsExtractor|ProviderEmailTrust|AuthProviders"
```

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how do provider identities map to accounts? (a) Separate account per provider ("you signed up with Google" forever); (b) auto-merge on email match, always; (c) merge on email match **only when verified** (explicit claim or a narrow, tested trust policy), with (provider, providerUserId) as the primary identity.

Chosen: (c) (ADR-002 + R17). Duplicate accounts are a genuine product failure (split data, "where are my things"), so merging matters; the guard makes it safe, and fail-closed normalization keeps it safe against the *next* provider someone adds in one careless line. `ProviderEmailTrust` handles the real-world wrinkle (Microsoft consumers) as an explicit policy object rather than a buried special case.

Rejected: (a) — punts the hard problem onto users, who end up with two half accounts and a support ticket. (b) — is the account-takeover vector, full stop; it's what several real-world breaches were made of.

The trade: legitimate users on providers that don't assert verification (outside the trust list) can't auto-merge — they land in a refusal instead of a silent link. Correct trade: the failure mode is a support case, not a takeover. The error page tells them to sign in with their original method and link explicitly from Settings.

9. Checkpoint

```
git add -A && git commit -m "feat: OAuth sign-in – one-line providers, fail-closed claims, takeover-guarded linking"
git tag lesson-2.5
```

You should now be able to: state precisely what OAuth delegates and what it can never delegate; explain why `(provider, providerUserId)` is identity and email is testimony; reproduce the takeover scenario from memory and name the two lines that prevent it; and justify when a security exception deserves to be a policy object instead of an `if`.

Next: *Lesson 2.6 — Tenancy I: the global query filter.* The `tenant_id` claim has been riding along in every token since 2.1. Time for it to earn its keep — the keystone lesson of the entire course.

Lesson 2.6 — Tenancy I: the global query filter

Part 2 · Identity & tenancy. The keystone of the whole platform. Everything you build after this — billing, files, webhooks, every feature slice — leans on the guarantee installed here: *a query cannot accidentally read another tenant's data*. We earn that guarantee the hard way: build the leak first, prove it with a test, then close it structurally. The `tenant_id` claim that has ridden every JWT since 2.1 finally does its job today.

Goal: every entity marked `ITenantScoped` is filtered to the current tenant automatically, so a handler that "forgets" to scope its query still cannot read another tenant's rows — and an unauthenticated caller sees *nothing*, not *everything*.

Concepts & principles: structural vs conventional enforcement; EF Core global query filters; marker interfaces driving platform behavior; fail-closed defaults; invariant tests (pinning a platform guarantee, not a method's output).

Maps to: ADR-003 · Golden Rule 1 · R2 · repo: `src/Core/Entities/ITenantScoped.cs`, `src/Core/Abstractions/ICurrentTenant.cs`, `src/Api/Services/HttpCurrentTenant.cs`, `src/Infrastructure/Persistence/AppDbContext.cs`, tests: `TenantScopeFilterTests.cs`, `TenantInvariantTests.cs`, `HttpCurrentTenantTests.cs`, `HarnessSmokeTests.cs` (the wire-level proof).

Prerequisites: 2.1 (the `tenant_id` claim), 1.3 (the Postgres fixture — this lesson is *why* tests run against a real relational engine).

1. Motivate — build the leak, then look at it

Every SaaS bug that makes the news is some version of "*customer A saw customer B's data*." The naive defense is discipline: "always add `.Where(x => x.TenantId == me)` to every query." That fails the first time a tired developer forgets — and you cannot grep your way back to confidence across a growing codebase, because absence of a `.Where` is invisible in review.

Let's feel it. Suppose a `Widget` entity and the obvious handler:

```
// first attempt – looks multi-tenant, is not
public class Widget
{
    public Guid Id { get; set; }
    public Guid TenantId { get; set; } // we store it...
    public string Name { get; set; } = "";
}

public Task<List<Widget>> ListAsync() =>
    _db.Set<Widget>().ToListAsync(); // ...but nothing filters on it
```

The `TenantId` column exists, so the schema *looks* multi-tenant. Nothing enforces it. A stored column is data; isolation is a *behavior*, and no behavior has been built.

2. Red — a test that proves the leak

Seed two tenants, a widget in each, act as tenant A, and assert you see only A's rows. Against the naive handler this fails — and that failure is the whole point:

```
// tests/Api.Tests/TenantScopeFilterTests.cs (shape of the real test)
[Collection(PostgresCollection.Name)]
public class TenantScopeFilterTests(PostgresFixture fixture) : PostgresTestBase(fixture)
{
    [Fact]
    public async Task Listing_never_returns_another_tenants_rows()
    {
        var (tenantA, tenantB) = (Guid.CreateVersion7(), Guid.CreateVersion7());
        await SeedWidgetAsync(tenantA, "A's widget");
        await SeedWidgetAsync(tenantB, "B's widget");

        await using var db = Fixture.CreateContext(tenantId: tenantA); // act as A
        var visible = await db.Set<Widget>().ToListAsync(); // "forgetful" query

        Assert.Single(visible);
        Assert.Equal("A's widget", visible[0].Name); // naive: returns BOTH
    }
}
```

Note what this test is *not*: it isn't testing a handler's logic. It's pinning an **invariant of the platform** — "a tenant-scoped read is isolated *no matter how carelessly it's written*." That's why the fix must live in the platform, not in the handler. (And note it runs on the real Postgres fixture: the EF in-memory provider's handling of global query filters diverges — a green isolation test on a fake database proves nothing. Lesson 1.3's decision was made *for this moment*.)

3. Green — make isolation structural

Three pieces. First, the marker — and the real file's doc comment deserves reading in full, because it teaches the *boundary* of the rule, not just the rule:

```
// src/Core/Entities/ITenantScoped.cs – the real file
/// <summary>
/// Marks an entity as belonging to a single tenant. Every entity implementing this is
/// automatically filtered to the current tenant by a global EF query filter (see
/// AppDbContext.OnModelCreating), so feature/domain queries are tenant-scoped by
/// default – you cannot forget to scope a read.
/// <para>
/// Platform entities that are intentionally cross-tenant or used before a tenant is
/// resolved (User, TenantMembership, auth tokens) deliberately do NOT implement this.
/// The rare cross-tenant/pre-auth lookup opts out with IgnoreQueryFilters().
/// </para>
/// </summary>
public interface ITenantScoped
{
    Guid TenantId { get; }
}
```

Second, *who is the current tenant?* A request-scoped abstraction the DbContext can depend on — and read its null contract carefully, it's the fail-closed hinge:

```
// src/Core/Abstractions/ICurrentTenant.cs – the real file
/// <summary>
/// The tenant the current request acts within, resolved from the authenticated
/// principal (the JWT tenant_id claim). This is the single, request-scoped tenancy
/// entry point: it drives the global tenant query filter in the DbContext, and
/// feature slices inject it to read their tenant id.
/// <para>TenantId is null when the caller is unauthenticated or has no tenant; the
/// global filter then matches no tenant-scoped rows (fail closed).</para>
/// </summary>
public interface ICurrentTenant
{
    Guid? TenantId { get; }
}
```

The HTTP implementation (HttpCurrentTenant) reads the tenant_id claim off the validated JWT — the claim you've been issuing since 2.1. Tests swap in the trivial settable fake from 1.3's fixture.

Third — the move that closes the leak for every entity at once, from the real AppDbContext:

```
// src/Infrastructure/Persistence/AppDbContext.cs (the real code)
/// <summary>Tenant the global query filter scopes to. Guid.Empty when there is no
/// current tenant – it matches no real (UUIDv7) row, so unauthenticated/tenant-less
/// callers see no tenant-scoped data (fail closed).</summary>
public Guid CurrentTenantId => _currentTenant.TenantId ?? Guid.Empty;

protected override void OnModelCreating(ModelBuilder builder)
{
    base.OnModelCreating(builder);
    builder.ApplyConfigurationsFromAssembly(typeof(AppDbContext).Assembly);

    // Tenant isolation as a structural guarantee: every ITenantScoped entity is
    // filtered to CurrentTenantId by default, so feature/domain queries can't forget
    // to scope. Discovered from the model – not hand-listed.
    foreach (var entityType in builder.Model.GetEntityTypes())
        if (typeof(ITenantScoped).IsAssignableFrom(entityType.ClrType))
            ApplyTenantFilterMethod.MakeGenericMethod(entityType.ClrType).Invoke(this,
[builder]);
}

private void ApplyTenantFilter<TEntity>(ModelBuilder builder) where TEntity : class,
ITenantScoped
    => builder.Entity<TEntity>().HasQueryFilter(e => e.TenantId == CurrentTenantId);
```

Mark Widget : ITenantScoped, re-run the leak test: **green**. The unchanged, still-careless ToListAsync() now emits WHERE tenant_id = @currentTenant — the handler never mentioned the tenant; the platform did. One subtlety in that lambda repays study: the filter references CurrentTenantId — a *property of the context* — so EF captures it as a parameter evaluated **per query**, not a constant baked at startup. One model, every request correctly scoped.

This is the discovery pattern's fourth appearance (configurations 1.3, fixture reset 1.3, config gate 1.4, now this) and its highest-stakes one: a new entity is protected the moment it implements the marker. No registration step exists to forget.

4. Harden — fail closed, then pin both properties

Look hard at one line: `_currentTenant.TenantId ?? Guid.Empty`. Why not throw? Why not skip filtering? Because tenant ids are UUIDv7 — `**Guid.Empty matches no real row**`. The failure mode of "no tenant" is therefore *see nothing, never see everything*. A misconfigured caller gets an empty screen, not a data breach. Pin it:

```
[Fact]
public async Task With_no_current_tenant_scoped_reads_return_nothing()
{
    await SeedWidgetAsync(Guid.CreateVersion7(), "someone's widget");
    await using var db = Fixture.CreateContext(tenantId: null); // nobody
    Assert.Empty(await db.Set<Widget>().ToListAsync()); // fail closed
}
```

Then pin the *coverage* of the rule itself. The filter protects `ITenantScoped` entities — but who checks that entities which *should* be marked, are? Rule R2's invariant test (TenantInvariantTests in the reference repo) walks the EF model and asserts every entity is either global-filtered or on an explicit, commented allowlist of deliberately cross-tenant types (User, TenantMembership, auth tokens, outbox plumbing). Forgetting the marker on a new entity fails the build with the entity's name — the allowlist makes the *exceptions* as reviewable as the rule.

Finally, the wire-level proof. Unit-level isolation could still be undone by a wrong registration or a middleware gap, so the integration harness (1.2) asserts the whole chain — from the real HarnessSmokeTests:

```
// Two owners, two tenants. Each calls the identical route with their own token; the
// JWT tenant_id claim drives the global filter, so neither can see the other's data.
var bodyA = await
_factory.CreateClientFor(a).GetFromJsonAsync<HouseholdDto>("/api/household");
var bodyB = await
_factory.CreateClientFor(b).GetFromJsonAsync<HouseholdDto>("/api/household");
Assert.DoesNotContain(bodyA.Members, m => m.UserId == b.UserId);
Assert.DoesNotContain(bodyB.Members, m => m.UserId == a.UserId);
```

Token → claim → `HttpContext.Tenant` → filter → SQL: proven at the wire, per commit.

5. Run it

```
dotnet test tests/Api.Tests --filter "TenantScope|TenantInvariant|HttpContext.Tenant"
# isolation, fail-closed, coverage invariant – the three-legged stool
```

Optional but worth it once: log the SQL (`LogTo(Console.WriteLine)`) and see the injected WHERE clause on a query whose C# has no where at all.

6. Architecture Decision

ARCHITECTURE DECISION

The fork: how is tenant isolation enforced on reads? (a) Discipline — every query adds `.Where(TenantId == ...)`, reviewers catch misses; (b) repository-layer scoping — data access goes through repositories that add the predicate; (c) a global EF query filter keyed on a marker interface, discovered from the model, fail-closed on no tenant.

Chosen: (c) (ADR-003). Isolation enforced by the type system + one central filter can't be forgotten per-query; discovery means new entities are protected by construction; ?? `Guid.Empty` turns misconfiguration into an empty result instead of a breach. (b) arrives *as well* in 3.1 — the repository seam adds ergonomics and a sanctioned escape hatch — but the guarantee must not *depend* on everyone using the repository; defense in depth starts by not trusting your own conventions.

Rejected: (a) — it's not a mechanism, it's a hope; the news is full of its failures. Also considered: database-enforced isolation (row-level security) — *not* rejected, deferred: it arrives in 8.4 as the second wall beneath this one, once the platform is worth two walls. And schema-per-tenant / database-per-tenant: real isolation, but operationally heavy (migrations × tenants) and wrong for this platform's "thousands of small tenants" shape.

The trade: query filters are an EF feature — raw SQL escapes them (a rule for later: raw SQL is platform code, never slice code), and the filter must be *provably attached* to every entity that needs it, which is exactly what the R2 invariant test exists to hold. Write-side isolation is a separate problem — deliberately: it's the next lesson.

7. Checkpoint

```
git add -A && git commit -m "feat: tenancy reads - ITenantScoped + global query filter, fail
closed (ADR-003)"
git tag lesson-2.6
```

You should now be able to: explain why a stored `TenantId` column is not multi-tenancy; add a global query filter keyed on a marker interface and say why it's discovered, not listed; defend ?? `Guid.Empty` as a security decision; distinguish an invariant test from a behavior test; and name which entities deliberately live *outside* the filter and why.

Next: Lesson 2.7 — *Tenancy II: write-side*. The filter scopes what you can *read*. Nothing yet stops a buggy handler from *writing* a row stamped with someone else's tenant — or from updating cross-tenant with `ExecuteUpdate`. Isolation needs both doors.

Lesson 2.7 — Tenancy II: write-side stamping & EnterTenant

Part 2 · Identity & tenancy. Lesson 2.6 locked the read door; this lesson locks the write door — because a filter on *reads* does nothing to stop a buggy handler from *inserting* a row stamped with someone else's tenant, or updating a row it loaded through the escape hatch. Isolation is only structural when **both directions** are. And once writes are guarded, a new question appears: how do legitimate tenant-less callers — a Stripe webhook, a future admin console — write into a tenant *without* bypassing all the guarantees? The answer is the platform's most elegant small primitive: `EnterTenant`.

Goal: a tenant-scoped entity written under the wrong tenant never reaches the database; new rows are stamped automatically; and system-authenticated operations enter a tenant scope that gives them the *same* isolation an authenticated request gets.

Concepts & principles: write-side stamping; EF `SaveChangesInterceptor`; the system-context trust level; scoping vs authorizing; the sanctioned escape hatch and its CI ban (R5); ambient context with disposable restore.

Maps to: ADR-003 (amendment) · R5, R28 · repo:

`src/Infrastructure/Persistence/TenantStampingInterceptor.cs`,
`src/Core/Abstractions/ITenantContext.cs`, `HttpCurrentTenant` (`EnterTenant` impl), tests:
`TenantStampingInterceptorTests.cs`, `EnterTenantScopingTests.cs`, `ArchitectureTests.cs`
(the `IgnoreQueryFilters` ban).

Prerequisites: 2.6 (the filter, `ICurrentTenant`, fail-closed `Guid.Empty`).

1. Motivate — the door you forgot you left open

The 2.6 test suite proves tenant A cannot *read* tenant B's rows. Write this handler, though, and watch the guarantee evaporate:

```
public async Task TransferAsync(Guid widgetId, Guid toTenantId) // hypothetical bug
{
    var widget = await _db.Set<Widget>().SingleAsync(w => w.Id == widgetId);
    widget.TenantId = toTenantId; // ← quietly re-homes A's data into B
    await _db.SaveChangesAsync(); // ← nothing objects
}
```

Or subtler: a handler builds a new `Widget { TenantId = someIdFromTheRequest }` — trusting client input for the one field that must never come from the client. Or subtler still (this one is the v2 audit's ADV-1 finding): platform code legitimately loads a row *cross-tenant* through the escape hatch, then a later code path innocently `SaveChanges`es it — an update landing under the wrong tenant with no one ever having "written" a tenant id at all.

Reads have one shape (a query) so one filter covers them. Writes have three shapes (insert, update, delete) and arrive through the change tracker — so the write-side guard lives where the change tracker converges: **`**SaveChanges**`**.

2. Green — the interceptor

EF Core's `SaveChangesInterceptor` sees every pending change just before it becomes SQL. The real `TenantStampingInterceptor`, core logic in full:

```
// src/Infrastructure/Persistence/TenantStampingInterceptor.cs
private static void Stamp(AppDbContext? db)
{
    if (db is null) return;

    var currentTenantId = db.CurrentTenantId;
    if (currentTenantId == Guid.Empty) return; // system/seed/cross-tenant context – not
enforced

    foreach (var entry in db.ChangeTracker.Entries<ITenantScoped>())
    {
        if (entry.State is not (EntityState.Added or EntityState.Modified or
EntityState.Deleted))
            continue;

        var tenantProperty = entry.Property(nameof(ITenantScoped.TenantId));
        var tenantId = (Guid)tenantProperty.CurrentValue!;

        // A new row may leave TenantId unset → stamp the current tenant.
        if (entry.State == EntityState.Added && tenantId == Guid.Empty)
        {
            tenantProperty.CurrentValue = currentTenantId;
            continue;
        }

        // An explicitly-tenanted insert, or any update/delete, must belong to the current
tenant.
        if (tenantId != currentTenantId)
            throw new InvalidOperationException(
                $"Refusing to {entry.State.ToString().ToLowerInvariant()} a
{entry.Entity.GetType().Name} " +
                $"for tenant {tenantId} while the current tenant is {currentTenantId}.
...");
    }
}
```

Two behaviors, both worth naming:

- **Stamping is a gift.** A handler creating an entity simply doesn't set `TenantId` —

the platform stamps it. The *correct* code is the code with *less* in it: feature authors literally cannot get this field wrong because they never touch it. (Notice the interplay with 2.6's fail-closed choice: `Guid.Empty` means "unset," which is exactly why real tenant ids being UUIDv7 matters twice.)

- **Validation is a wall.** Any insert/update/delete whose `TenantId` differs from the

current tenant throws — *before* SQL is generated. The exception message is written for the developer who hits it: what was refused, both tenant ids, and what to do instead.

Wrong-tenant writes are now a *test failure during development*, not a data corruption in production.

And one registration subtlety from the real `AppDbContext` — the interceptor is attached in `OnConfiguring`, *inside the context*, not (only) in DI:

```
protected override void OnConfiguring(DbContextOptionsBuilder optionsBuilder)
{
    base.OnConfiguring(optionsBuilder);
    optionsBuilder.AddInterceptors(TenantStampingInterceptor.Instance, ...);
}
```

Why there? Because tests (and tools) construct contexts directly, bypassing DI — and a guard that only exists on the DI path is a guard your test double forgot. Attaching in `OnConfiguring` means **every** `AppDbContext`, however constructed, enforces the rule. The interceptor is deliberately stateless (it reads the tenant from the context being saved), so one shared `Instance` serves them all.

3. The trust boundary — who is `Guid.Empty`, really?

Look again at the early return: *no current tenant* → *no enforcement*. Isn't that a hole? It's a **trust level**, and drawing it consciously is this lesson's second idea.

A context with no current tenant is the *system* context: migrations, seeders, the outbox dispatcher (Part 4), tenant-dissolution teardown (2.9). System code is exactly the code that legitimately touches many tenants — and it's held to a different standard: it must be *platform* code, reviewed as such, and the read-side twin of this trust level (the `IgnoreQueryFilters()` / `QueryAllTenants()` escape hatch) is **banned from feature slices by an architecture test** (R5 — the scan you'll formalize in 3.3 fails the build if the string appears under `src/Api/Features/`). Feature code gets the guarded world with no exit; platform code gets the trusted world with review. Two tiers, machine-separated.

But there's a third kind of caller, and it's the interesting one: code that has **no JWT but a trusted tenant id**. A Stripe webhook arrives bearing a cryptographic signature (5.2) and *names* a tenant; the future admin console (7.5) impersonates into one. The lazy answer is "run them as system and be careful." The platform's answer is better — from the real `ITenantContext`:

```
// src/Core/Abstractions/ITenantContext.cs
/// Enters a tenant context for system/integration operations that have no JWT but a
/// TRUSTED tenant id — e.g. a Stripe-signature-verified webhook. While the returned
/// scope is active, ICurrentTenant.TenantId resolves to the entered tenant, so the
/// write-stamping interceptor and the global read filter scope to it. Such operations
/// therefore get the SAME structural tenant isolation an authenticated request gets,
/// instead of bypassing it with the cross-tenant escape hatch.
///
/// This primitive only SCOPES; it does not AUTHORIZE. The caller MUST have already
/// authenticated the tenant id by other means (a verified webhook signature, an admin
/// authz check).
public interface ITenantContext
{
    IDisposable EnterTenant(Guid tenantId);
}
```

Usage reads like what it is — a scoped assumption:


```
// inside the (signature-verified) billing webhook handler, Part 5:
using (tenantContext.EnterTenant(tenantIdFromVerifiedEvent))
{
    // in here, reads are filtered to this tenant and writes are stamped/validated
    // exactly as if a user of this tenant had made an authenticated request
    subscription.Status = newStatus;
    await db.SaveChangesAsync();
} // dispose restores the previous tenant (nesting supported)
```

The doc comment's sharpest sentence deserves a highlight: **"only scopes; does not authorize."** `EnterTenant` never decides *whether* you may act as tenant X — the webhook's signature check or the admin's authz check did that. It decides that, having been authorized, you operate under the *full* isolation regime rather than the trusted bypass. The alternative — webhook code running system-level with hand-rolled `WHERE` clauses — is precisely the conventional-discipline world 2.6 was built to abolish. This primitive exists so the platform never has to choose between "no JWT" and "no guardrails."

Implementation-wise, `HttpContext.Current.Tenant` doubles as the `ITenantContext`: `EnterTenant` swaps the ambient tenant and returns an `IDisposable` that restores the previous value (nesting supported) — the same `enter/restore` shape you saw in the test double back in 1.3, which is no accident: the fixture's `TestCurrentTenant` implements both interfaces so services under test see one coherent ambient tenant.

4. Red — the tests that pin all three behaviors

From the real `TenantStampingInterceptorTests` / `EnterTenantScopingTests`, as Gherkin:

```
Scenario: new rows are stamped
  Given the current tenant is A
  When a Widget with unset TenantId is added and saved
  Then the persisted row's TenantId is A

Scenario: wrong-tenant writes are refused – all three states
  Given the current tenant is A and a row belonging to B (loaded via the escape hatch)
  When the row is modified / deleted / a new row is added with TenantId = B
  Then SaveChanges throws before any SQL runs
  And the message names both tenants

Scenario: EnterTenant gives a webhook the same isolation as a request
  Given no current tenant (system context)
  When the handler enters tenant A, writes, and disposes the scope
  Then the write is stamped/validated as A
  And after dispose the context is system again (and nested scopes restore correctly)
```

The middle scenario is ADV-1's regression test — the audit finding that widened the interceptor from "validate inserts" to "validate every state." Audits feed tests; tests hold the line.

One honest boundary the interceptor's own comment draws: ****bulk operations (ExecuteUpdate/ExecuteDelete) bypass the change tracker**** — no `SaveChanges`, no interceptor. The global *read* filter still constrains what they touch (they compose onto the filtered query), but the write-side wall has this known gap at the edges, held by review and by 8.4's database-level backstop. Structural enforcement doesn't mean pretending the structure covers everything; it means *knowing exactly where it ends*.

5. Run it

```
dotnet test tests/Api.Tests --filter "TenantStamping|EnterTenant"
```

Then feel the wall once: in any handler, set a `TenantId` to a random `Guid` and `SaveChanges`. The exception you get — naming the entity, both tenants, and the rule — is the platform teaching your future teammates in the only classroom that scales.

6. Architecture Decision

ARCHITECTURE DECISION

The fork: how are *writes* kept tenant-correct? (a) Trust handlers to set `TenantId` (review catches mistakes); (b) require every write to pass through a repository method that takes the tenant explicitly; (c) a `SaveChanges` interceptor that stamps unset inserts and refuses cross-tenant writes, plus `EnterTenant` for trusted tenant-less callers.

Chosen: (c) (ADR-003 amendment). The change tracker is the one choke point every tracked write crosses, so enforcement there is complete-by-construction for tracked writes; stamping makes correct code *smaller* than incorrect code; and `EnterTenant` extends the same regime to signature-authenticated work instead of granting it the bypass. Registered in `OnConfiguring` so no construction path escapes.

Rejected: (a) — the same hope 2.6 rejected for reads, with worse blast radius. (b) — helps ergonomics (and 3.1 builds it!) but can't be the *guarantee*: any code holding the `DbContext` can sidestep a repository, and the platform refuses to found isolation on "everyone used the blessed API."

The trade: an ambient-context primitive (`EnterTenant`) is ambient state — the classic critique applies (invisible in signatures). Contained by its shape: scoped, disposable, nest-restoring, and used only at authenticated entry points, never as a general parameter-passing dodge. And the tracked-write wall knowingly excludes bulk ops — a gap documented in code, patrolled by review, and closed for good by the RLS backstop in 8.4.

7. Checkpoint

```
git add -A && git commit -m "feat: tenancy writes – stamping interceptor + EnterTenant
primitive (ADR-003)"
git tag lesson-2.7
```

You should now be able to: name the three write shapes and where they converge; explain why stamping-plus-refusal beats "set it yourself"; articulate the system trust level and the difference between *scoping* and *authorizing*; and state precisely where the tracked-write guarantee ends and what covers the remainder.

Next: *Lesson 2.8 — Tenants & membership* — the entities behind the claim: what a tenant actually *is*, and the membership model that decides everything from "who's in my household" to "what happens when the owner leaves."

Lesson 2.8 — Tenants & membership

Part 2 · Identity & tenancy. Two lessons of machinery have enforced "the current tenant" without ever defining what a tenant *is*. Time to model it — and the modeling lesson here is about restraint and invariants: the Tenant is four properties; all the interesting decisions live in **membership**, the link entity that answers who belongs where, in what role, and — the questions that actually bite — what happens on remove, transfer, and leave. Lifecycle transitions are where domain models earn or lose their integrity.

Goal: Tenant and TenantMembership exist with their invariants (one tenant per user, exactly one owner per tenant) enforced; the app-facing `/api/household` surface supports `get/rename/remove-member/transfer/leave`, each transition preserving every invariant.

Concepts & principles: link entities as decision records; invariants as constraints + result enums; the "membership, not `User.TenantId`" indirection; result types over exceptions for domain outcomes; lifecycle design ("nobody is ever tenant-less").

Maps to: ADR-003, ADR-C1/C2 · repo: `src/Core/Entities/Tenant.cs`, `TenantMembership.cs`, `src/Api/Services/TenantService.cs`, `TenantModels.cs`, `src/Api/Controllers/HouseholdController.cs`, `src/Core/Repositories/ITenantRepository.cs`, tests: `UserServiceTests` (tenant-of-one), `integration HarnessSmokeTests (/api/household)`, later `MemberRoleManagementTests`.

Prerequisites: 2.6/2.7 (the machinery this model feeds), 2.1 (users).

1. Motivate — the model is the answer sheet

Product questions this lesson must answer, each of which has burned a real SaaS:

- Can a user belong to two tenants? (Your data model answers even if you never decided.)
- Ana leaves the household — does the shopping list she created leave with her?
- The owner deletes their account — is the household now headless? Orphaned? Gone?
- Two admins demote each other simultaneously — who wins?

Answer these *in entities and constraints* and the platform behaves consistently forever; answer them ad hoc per endpoint and every feature reinvents the rules until two of them disagree. So: model first, endpoints second.

2. The entities — one boring, one load-bearing

```
// src/Core/Entities/Tenant.cs – the real file, in full
public class Tenant
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public required string Name { get; set; }
    public DateTimeOffset CreatedAt { get; set; }
    public DateTimeOffset UpdatedAt { get; set; }
}
```

Four properties. That's not underdesign — it's the discipline of putting state where it belongs: billing state will live in a Subscription projection (Part 5), member count is *derived* (never a stored counter that drifts — Golden Rule 4), settings arrive when a feature needs them. A tenant is an identity to hang ownership off; resist making it a junk drawer.

```
// src/Core/Entities/TenantMembership.cs — the real file
/// <summary>
/// Links a user to their tenant with a role. A user belongs to exactly one tenant at a
/// time (unique on UserId); a tenant has exactly one owner. Identity (logins, inbox)
/// stays user-scoped; app data is scoped to the tenant. Membership — not User.TenantId —
/// is the source of truth for tenant resolution.
/// </summary>
public class TenantMembership
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public Guid TenantId { get; set; }
    public Guid UserId { get; set; }
    public string Role { get; set; } = TenantRoles.Member; // owner | admin | member
    public DateTimeOffset JoinedAt { get; set; } = DateTimeOffset.UtcNow;
}
```

The doc comment is a compressed ADR — unpack its three commitments:

"Exactly one tenant at a time (unique on UserId)." The biggest product decision in Part 2, and it's enforced as a **unique index**, not a code path. Multi-tenant-per-user (Slack-style workspace switching) is a legitimate other product — with tenant pickers, per-tenant tokens, "which tenant is this request for" on every call. This platform's domain (a household product) wants one home per user, and buys enormous simplification with it: the JWT carries *the* tenant_id, not a selected one. Know which product you're building; put the answer in the schema.

*****"Membership — not User.TenantId — is the source of truth."***** A TenantId column on User would be simpler — and would weld identity to belonging. As a separate entity, membership can carry *belonging-specific* state (Role, JoinedAt), be deleted without touching the account, and later support history. The user's account — logins, locale, notifications — survives every household transition. Identity is user-scoped; app data is tenant-scoped; membership is the hinge (ADR-C1/C2, in code).

"Exactly one owner." Not a preference — an invariant every transition below must conserve. TenantRoles pins the vocabulary (owner > admin > member, ordered — the matrix that gives the ordering teeth is RBAC's, Lesson 6.1; today only *owner* gates anything).

Neither entity implements ITenantScoped — they're on 2.6's documented allowlist. Membership is the thing tenant resolution derives from; filtering it by the thing it defines would be circular. First-hand proof the allowlist isn't a formality.

3. Where tenants come from — the household of one

No "create tenant" endpoint exists. Look back at 2.5's case 3 and 2.4's redemption: both funnel into one private CreateUserWithTenantAsync, which writes **user + tenant + owner membership atomically** — three rows, one IUnitOfWork transaction, all-or-nothing (a crash mid-way must not birth a user with no household). Every user is thus born an owner of their own tenant-of-one; there is no moment — not one request — in which an authenticated user

has no tenant. That non-existence of a "tenant-less user" state kills an entire class of null-handling and onboarding UX, and it's why 2.6's claim `tenant_id` can be relied on rather than defended against.

This "nobody is ever tenant-less" invariant now becomes the design key for every transition below.

4. Transitions — where models go to die (or not)

The service surface, from the real `TenantService` — notice it speaks in **result enums**, not exceptions:

```
public enum RemoveMemberResult { Removed, NotAMember, CannotRemoveOwner }
public enum TransferResult { Transferred, TargetNotMember, ConcurrentModification }
public enum LeaveOutcome { Left, MustTransferFirst, ConfirmationRequired, Dissolved }
```

A member being non-removable-because-owner isn't *exceptional* — it's a legal outcome the UI must explain. Exceptions are for broken invariants; enums are for domain answers. (The controller maps each to 1.5's envelope: `CannotRemoveOwner` → { "error": "cannot_remove_owner", ... } — machine-readable all the way out.) Now the transitions, each read against the invariants:

Remove a member. From the real doc comment: *"lands them in a fresh tenant-of-one (owner) so they are never tenant-less. Their contributed data stays with the tenant."* Two decisions in one sentence. Re-homing conserves the no-tenant-less invariant even under an *involuntary* transition. And the data question — the shopping list Ana created — is answered by tenancy itself: app data belongs to the *tenant* (it's `ITenantScoped`, stamped with the tenant's id, 2.7). Ana's authorship was participation, not ownership; she leaves, the list stays. The model made this decision back in 2.6 — this lesson just makes it visible. And `CannotRemoveOwner`: removal would leave the tenant headless; ownership *moves*, never vanishes.

Transfer ownership. Demote-owner + promote-target — two rows that must flip *atomically* (a crash between them = zero or two owners; a transaction makes it all-or-nothing). And note `ConcurrentModification`: transfer is guarded against the racing-admins problem with optimistic concurrency — if the membership rows changed under you, the transfer reports it instead of silently double-applying. Invariants under *concurrency*, not just under happy sequences.

Leave. The richest one — walk `LeaveOutcome`'s four arms as a decision tree: a non-owner leaves (re-homed, data stays); an owner with members can't leave until they transfer (`MustTransferFirst` — the invariant, stated as a 409); a *sole* owner leaving means destroying the tenant, which is irreversible and therefore demands an explicit confirmation round-trip (`ConfirmationRequired` → client re-sends with `confirm_dissolve: true` → `Dissolved`). That confirm-then-wipe dance previews 2.9, where "wipe" gets its real machinery.

5. The app-facing surface — and the rename seam

The controller binds it all to HTTP — GET `/api/household`, PUT `rename`, DELETE `members/{id}`, POST `transfer-ownership`, POST `leave` — with one naming oddity worth explaining rather than hiding: the route says **household**, the code says **tenant**. Deliberate bilingualism. *Tenant* is platform vocabulary — every seam, filter, and rule from 2.6/2.7 speaks it, and it stays stable

across products built from this template. *Household* is this app's word for it — user-facing, changeable per product (a team app renames it "workspace" by touching the controller layer and UI strings only; the rebrand checklist in 9.1 includes exactly this). Keep the dictionary small and one-directional: domain term inside, product term at the boundary.

Owner-only actions (rename, remove, transfer) are enforced at the controller via the membership role — full RBAC arrives in 6.1; today the check is "is caller the owner," which is the only distinction the model needs yet. The wire tests you met in 2.6's harness (`/api/household` returning *your* household, cross-tenant invisibility) now read as what they were all along: this lesson's integration suite.

6. Red — the tests

```
Scenario: removing a member re-homes them
  Given a tenant with owner Ana and member Ben, where Ben created a Widget
  When Ana removes Ben
  Then Ben is the owner of a fresh tenant-of-one
  And the Widget still belongs to Ana's tenant

Scenario: ownership transfer is atomic and conserves single-owner
  Given owner Ana and member Ben
  When Ana transfers ownership to Ben
  Then Ben is owner and Ana is admin – and at no observable point were there zero or two owners

Scenario: sole owner leaving requires explicit confirmation
  Given Ana alone in her tenant
  When she leaves without confirmation
  Then the outcome is ConfirmationRequired and nothing changed
  When she confirms
  Then the tenant and its data are gone and Ana is re-homed
```

Each scenario is an invariant wearing a story. The membership-lifecycle E2E journey (built with the UI in Part 3/6) replays them through a real browser.

7. Run it

```
dotnet test tests/Api.Tests --filter "Tenant|Household"
# .http: GET /api/household → your tenant-of-one, role "owner", member list of you
```

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how do users relate to tenants? (a) `User.TenantId` — a column; (b) many-to-many memberships (multi-workspace); (c) a membership link entity with a **unique-per-user** constraint — one tenant at a time, roles on the link.

Chosen: (c) (ADR-003). The link entity keeps identity and belonging separable (accounts survive transitions; roles/joined-at live where they belong) while the unique index keeps the *product's* answer — one home — enforced by the schema. Transitions are modeled as result enums with re-homing so no path strands a user tenant-less; ownership is conserved by construction (transfer-only, atomic).

Rejected: (a) — welds identity to belonging; every lifecycle transition turns into account surgery, and role-on-membership has nowhere to live. (b) — a different product; adopting its complexity (tenant pickers, per-request tenant selection, $N \times M$ permission surfaces) without its requirement is pure cost. The link-entity shape deliberately keeps (b) *reachable* — drop the unique index and add a picker — which is the correct way to keep an option open: pay for it only if exercised.

The trade: "one tenant per user" is a real product constraint (no guest memberships across households), and re-homing means removals *create* tenants — tenant count $\neq$ paying customers, which Part 5's billing must (and does) treat knowingly.

9. Checkpoint

```
git add -A && git commit -m "feat: tenants + membership – invariants, transitions,
household surface (ADR-003)"
git tag lesson-2.8
```

You should now be able to: defend membership-as-entity against the simpler column; name both schema-enforced invariants and every transition that must conserve them; explain re-homing and whose data stays where (and why the model already decided); and argue result-enums-over-exceptions for domain outcomes.

Next: Lesson 2.9 — *Invitations, dissolve & the contributor seam*. How anyone ever joins someone *else's* household — and the decentralized-teardown seam that GDPR will thank us for in Part 7.

Lesson 2.9 — Invitations, dissolve & the contributor seam

Part 2 · Identity & tenancy — finale. Two jobs close out the chassis. First, invitations: the only way a tenant ever gains a second member, built entirely from parts you already own (hashed single-use tokens, branded email, membership transitions). Second — the deeper one — *dissolve*: when a tenant dies, every feature's data must die with it, including the data of **features that don't exist yet**. That impossible-sounding requirement has a beautiful answer, the `ITenantDataContributor` seam — the platform's purest example of the Open/Closed principle, and the hook GDPR will hang the entire export/erasure epic on in Part 7.

Goal: an owner can invite by email; whoever proves control of the invite token joins (with correct re-homing of *their* old tenant); a dissolving tenant's data is wiped feature-by-feature through a discovered seam — with a test that fails any future feature that forgets to register.

Concepts & principles: composing prior seams into a feature; identity-at-redemption (the invite token authenticates the *invitation*, the login authenticates the *person*); Open/Closed via DI fan-out; teardown inside one transaction; guarding solo-owner data abandonment.

Maps to: ADR-003/ADR-011 (the seam GDPR extends) · repo:

```
src/Core/Entities/TenantInvitation.cs, src/Api/Services/TenantInvitationService.cs,
src/Api/Controllers/HouseholdInvitationsController.cs,
src/Core/Abstractions/ITenantDataContributor.cs,
src/Api/Services/TenantDissolutionService.cs, tests: TenantInvitationTests (Core),
WipeDataTests.cs.
```

Prerequisites: 2.8 (membership + transitions), 2.4 (the token pattern), 2.3 (email).

1. Motivate — two lifecycle holes

Hole one: every tenant so far is a household of one. There is no path by which Ana's household ever contains Ben — and "Ben types Ana's household id" is obviously not it. Joining must be *invitation-shaped*: initiated by an owner, addressed to a person, consensual on both sides, expiring, revocable.

Hole two, quieter and worse: 2.8's `Dissolved` arm claims the tenant's data gets destroyed. *Which* data? Today, a `Note` here, a `Widget` there. Next month, files, notification prefs, webhook subscriptions, billing counters — written by features that don't exist yet. A central `WipeTenant()` method listing every table is doomed to incompleteness: each new feature must remember to edit it, and the one that forgets leaves orphaned rows — data that legally *must* die (wait for GDPR) quietly surviving. The platform needs teardown that **new features join automatically**.

2. Invitations — old parts, one new idea

The entity is 2.4's pattern wearing a new purpose — hashed single-use token, expiry, computed validity — plus a status lifecycle:

```
// src/Core/Entities/TenantInvitation.cs (the real shape)
public class TenantInvitation : ITenantScoped // ← note the marker!
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public Guid TenantId { get; set; }
    public required string InvitedEmail { get; set; } // normalized lower-case
    public Guid InvitedByUserId { get; set; }
    public string Status { get; set; } = InvitationStatuses.Pending; //
    pending|accepted|revoked|expired
    public required string TokenHash { get; set; } // raw token revealed once, at creation
    public DateTimeOffset CreatedAt { get; set; }
    public DateTimeOffset ExpiresAt { get; set; }

    public bool IsValidAt(DateTimeOffset now) => Status == InvitationStatuses.Pending &&
    !IsExpiredAt(now);
}
```

`ITenantScoped` — the first *feature-owned* entity to wear the marker since the machinery went in. Invitations belong to the inviting tenant: the owner lists *their* pending invites, 2.6's filter scopes them for free, 2.7 stamps them on insert. The chassis just works; this is what building *on* it feels like.

The flow composes seams end to end: owner posts an email address → service invalidates prior pending invites for it, mints token via `ITokenGenerator`, stores the hash, lifespan from config (`Auth:Invitation:LifespanDays, 1.4`) → `BrandedEmail` renders the invite **in the inviter's saved locale** (2.3's explicit-culture design paying off — the recipient sees the language of the household inviting them) → `ISender` delivers.

The one genuinely new idea is at redemption, and the entity's doc comment states it precisely: "*redeemed by whoever is signed in and presents the single-use token — **identity is the login, not the bare email.***" Unpack that, because it's a subtle authorization design:

- The invite link lands on `/join`. The redeemer must **sign in first** (any method —

OTP, Google, whatever *they* have). The token proves they hold the invitation; the login proves who they are. Two proofs, two different questions.

- Consequence: the invite can be *forwarded*. Ana invites `ben@work.test`; Ben redeems

while signed in as `ben@personal.test` — and joins as his real, existing account rather than growing a duplicate. The addressed email is *delivery routing*, not identity. (Compare 2.4, where email possession *is* the identity claim — same token machinery, opposite trust conclusion, both correct for their question.)

- Joining triggers 2.8's transitions: Ben's old tenant-of-one is handled by the same

leave semantics — and here the contributor seam makes its entrance early. What if Ben's solo tenant *has data*? Silently discarding it on join would be the data-abandonment bug. The platform checks `HasDataAsync` across all contributors and requires explicit confirmation — "joining will delete your current household's data" — before dissolving it. The seam serves its first customer before dissolve even ships.

Revocation (owner cancels a pending invite) and expiry round out the lifecycle; `TenantInvitationTests` in `Core.Tests` drive the validity state machine with explicit clocks — pure domain logic, no database.

3. The contributor seam — teardown that scales with features

Now hole two. The platform's answer, in full:

```
// src/Core/Abstractions/ITenantDataContributor.cs – the real file
/// <summary>
/// A feature's contribution to tenant-data presence and teardown. Each vertical slice
/// that owns tenant-scoped tables registers one of these in DI; the platform queries
/// all contributors instead of a hard-coded method – so adding a feature never means
/// editing a central "has data" / "wipe data" method.
/// </summary>
public interface ITenantDataContributor
{
    /// <summary>True if this contributor holds any data for the tenant.</summary>
    Task<bool> HasDataAsync(Guid tenantId, CancellationToken ct = default);

    /// <summary>Removes this contributor's data for the tenant. Called within the
    /// dissolve transaction, so it shares the caller's unit of work – do not commit
    here.</summary>
    Task WipeAsync(Guid tenantId, CancellationToken ct = default);

    /// <summary>The section name this contributor's data appears under in a tenant
    /// export (GDPR-1), e.g. "notes". Stable and unique across contributors.</summary>
    string ExportKey { get; }

    /// <summary>A JSON-serializable snapshot for a data export. Null when nothing to
    /// include. Identifiers and content only – never secrets.</summary>
    Task<object?> ExportAsync(Guid tenantId, CancellationToken ct = default);
}
```

The inversion is the entire idea. The platform doesn't know what features exist; it asks DI for *all* `ITenantDataContributors` and fans out. A new feature ships its own contributor (five minutes: "do I have rows for this tenant / delete them / here's my export"), registers it, and is *automatically* part of dissolve, of the join-time data check, and — Part 7 — of GDPR export.

The platform is closed to modification, open to extension: Open/Closed not as a poster slogan but as a DI query.

Notice the interface already carries `ExportKey/ExportAsync`. Historically the seam was born two-method and GDPR *widened* it — and because widening an interface breaks every implementor at compile time, each existing feature was forced, by the build, to answer "what do I export?" That's the seam working even while changing: the compiler runs the migration checklist. (The course teaches the final four-member shape from the start; you get the anecdote instead of the breakage.)

The orchestrator stays platform-small:

```
// src/Api/Services/TenantDissolutionService.cs (the real shape)
/// Fan out over every ITenantDataContributor so each feature wipes its own domain
/// data first, then run the platform's core teardown (memberships, invitations, the
/// tenant row). Deliberately does NOT open its own transaction – callers invoke it
/// inside their existing dissolve transaction, so the whole thing commits atomically.
public async Task DissolveAsync(Guid tenantId, CancellationToken ct = default)
{
    foreach (var contributor in contributors) // injected:
        await contributor.WipeAsync(tenantId, ct);
    await tenantRepository.WipeDataAsync(tenantId, ct); // core: memberships, invites,
    tenant row
}
```

Two disciplines in the comments deserve promotion to principles. **Order:** features first, core last — contributors may need the tenant's rows (FKs) to still exist while they wipe.

Transactionality: the service *enlists* in the caller's transaction rather than owning one — dissolve is one atomic event alongside the re-home and audit steps that surround it (2.8's leave flow); a crash mid-wipe rolls back to "still dissolved-nothing," never "half-dissolved." A teardown that can partially complete is a data-integrity bug with a delay timer. And who *calls* dissolve? Only system-level flows — the sole-owner leave and (later) account erasure — running on the system context from 2.7; the contributors' `WipeAsync` implementations are exactly the sanctioned cross-tenant-write code that trust level exists for.

4. Harden — the seam that can't be forgotten, provably

The seam's Achilles heel: a feature that *doesn't* register a contributor silently opts out of teardown. Convention again — unless a test closes it. The reference repo's `WipeDataTests` does exactly that, with 1.3's discovery trick pointed at a new target: walk the EF model for `ITenantScoped` entity types, dissolve a seeded tenant, then assert **zero rows remain in any tenant-scoped table**:

```
[Fact]
public async Task Dissolve_leaves_no_tenant_scoped_rows_behind()
{
    var tenantId = await SeedEveryFeaturesSampleDataAsync(); // incl. YOUR new feature

    await dissolution.DissolveAsync(tenantId);

    foreach (var entityType in db.Model.GetEntityTypes()
        .Where(t => typeof(ITenantScoped).IsAssignableFrom(t.ClrType)))
        Assert.Equal(0, await CountRowsForTenantAsync(entityType, tenantId));
}
```

A feature that adds a tenant-scoped table without wiring teardown turns this red with the *table's name in the failure*. Model-derived, so it covers tables that don't exist yet — the same move as the fixture reset (R11) and the filter-coverage invariant (R2), now guarding deletion. Three guards, one pattern: **derive the checklist from the model; never hand-maintain it.**

5. Red — the tests

```

Scenario: invitation joins the redeemer's real account
  Given Ana invited ben@work.test
  And Ben signs in as ben@personal.test and presents the token
  Then Ben's account joins Ana's household as member
  And Ben's empty solo tenant is dissolved

Scenario: joining with data requires explicit consent to lose it
  Given Ben's solo household contains a Note
  When Ben redeems Ana's invite without confirmation
  Then the join is refused with has_data
  And nothing changed

Scenario: dissolve is total and atomic
  Given a tenant with data in every feature
  When the sole owner leaves with confirmation
  Then no tenant-scoped rows remain for that tenant
  And a wipe failure in any contributor rolls the entire dissolve back

Scenario: invitations are single-use and revocable
  Given a pending invitation
  When it is redeemed / revoked / expires
  Then subsequent redemption attempts fail with the same generic error

```

The third scenario's rollback half is worth the extra test plumbing (a contributor rigged to throw): atomicity claims are exactly the claims that rot silently.

6. Run it

```

dotnet test tests/Api.Tests --filter "Invitation|WipeData"
# live: invite yourself at a second address; open Mailpit; redeem on /join while
# signed in as another test account; GET /api/household → two members.

```

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how does per-feature data participate in tenant teardown (and later, export)? (a) A central `WipeTenant()` that each feature edits; (b) database cascade deletes from the tenant row; (c) a DI-discovered contributor seam + a model-derived completeness test.

Chosen: (c) (ADR-003 lineage; ADR-011 widens it). Fan-out inverts the dependency: the platform stops needing to know its features. The completeness test converts the seam's one weakness (forgetting to register) into a named build failure. The same seam answers three lifecycle questions — has-data (join guard), wipe (dissolve), export (GDPR) — so a feature declares its data relationship *once*.

Rejected: (a) — the forget-to-edit design; its failure mode is silent orphaned data with legal implications. (b) — cascades handle *rows* but not the rest (files on disk in Part 6, provider-side subscriptions in Part 5 — dissolve is not only SQL), offer no has-data or export answer, and hide teardown semantics in migration DDL where no test names them.

The trade: every feature carries ~30 lines of contributor boilerplate, and dissolve is $O(\text{features})$ sequential queries — perfectly acceptable for an event as rare as tenant death, and the boilerplate is the checklist (3.2's slice template includes it pre-written).

8. Checkpoint

```
git add -A && git commit -m "feat: invitations + dissolve – contributor seam, atomic  
teardown, completeness gate"  
git tag lesson-2.9
```

Part 2 complete. The chassis stands: identity (tokens, passwordless, OAuth), tenancy (read filter, write stamping, EnterTenant), the tenant model with its invariants, and lifecycle machinery that future features join by construction. Every lesson from here *rides* this.

You should now be able to: separate "who holds the invitation" from "who is joining" and say why forwarding is safe; implement the contributor seam and its completeness test from memory; explain wipe ordering and transaction enlistment; and recognize Open/Closed when it's a DI query rather than a slogan.

Next: *Part 3 — The slice pattern & the UI.* The chassis gets its assembly manual: the repository seam, the anatomy of a vertical slice, and the first pixels.

Part 3 — The slice pattern & the UI

Lesson 3.1 — The repository seam

Part 3 · The slice pattern & the UI. Part 2 built the chassis; Part 3 is its *assembly manual* — the repeatable way features bolt on. This first lesson supplies the data-access seam every future slice uses, and it carries a quiet responsibility: it must give features an ergonomic way to read and write their own tables **without letting them undermine the tenant isolation** you fought for in 2.6/2.7. One generic interface, two query methods — one safe by default, one deliberately dangerous and greppable — and a commit model that makes atomicity a choice rather than an accident.

Goal: a generic `IRepository<T>` lets a slice read (tenant-filtered) and write its own entity with no bespoke repository code; an explicit, named escape hatch exists for the rare cross-tenant path; and a unit-of-work seam groups multi-write operations into a transaction.

Concepts & principles: the generic repository; open-generic DI registration; the safe-default / explicit-escape-hatch pattern; unit of work and the shared-context commit model; ergonomics that *preserve* guarantees rather than bypass them.

Maps to: ADR-003/ADR-004 · R5 (the escape-hatch ban, formalized 3.3) · repo:

`src/Core/Repositories/IRepository.cs, IUnitOfWork.cs,`
`src/Infrastructure/Repositories/EfRepository.cs, EfUnitOfWork.cs,`
`tests/Api.Tests/RepositoryScopingTests.cs.`

Prerequisites: 2.6 (the query filter this seam must preserve), 2.7 (the write interceptor), 2.9 (the contributor wipe that first needed cross-tenant access).

1. Motivate — two ways to give features data access, both wrong

Wrong way one: a bespoke repository per entity. `INoteRepository`, `IWidgetRepository`, each with `GetById`, `List`, `Add`, `Delete` — ninety percent identical, hand-written every slice, and every one an opportunity to write a query that *forgets the tenant filter* or subtly diverges from its siblings. Multiply by thirty features and you have a maintenance tax paid in boilerplate and isolation risk.

****Wrong way two: hand features the `DbContext` directly.**** Maximum flexibility, maximum danger. A feature holding `AppDbContext` can call `IgnoreQueryFilters()`, can `SaveChanges` on entities it shouldn't, can reach into other features' `DbSets`. The chassis's guarantees become "please don't."

The seam threads the needle: **one generic repository** that eliminates the boilerplate *and* channels every feature read through the tenant-filtered path by default — while making the cross-tenant escape a separate, named, greppable method that a code reviewer (and an arch test) can see from orbit.

2. The interface — safe by default, dangerous on purpose


```
// src/Core/Repositories/IRepository.cs – the real interface
/// <summary>
/// A thin generic repository for FEATURE/domain entities, so a vertical slice can read
/// and write its own data without authoring a bespoke repository pair. Reads via Query()
/// are automatically tenant-scoped for ITenantScoped entities (the AppDbContext global
/// query filter), so a slice can't forget to scope.
/// </summary>
public interface IRepository<TEntity> where TEntity : class
{
    /// <summary>Queryable over the set – tenant-filtered for ITenantScoped entities.
    /// This is the normal feature read path; it cannot leak across tenants.</summary>
    IQueryable<TEntity> Query();

    /// <summary>Cross-tenant queryable: the global tenant filter is OFF, so this sees
    /// EVERY tenant's rows. The audited, deliberately-named escape hatch ... Feature
    /// slices use Query(); reach for this only when crossing tenants is the explicit
    /// intent, and always re-constrain with a Where(x => x.TenantId == ...).
    /// Greppable on purpose.</summary>
    IQueryable<TEntity> QueryAllTenants();

    Task AddAsync(TEntity entity, CancellationToken ct = default);
    void Update(TEntity entity);
    void Remove(TEntity entity);

    Task<int> SaveChangesAsync(CancellationToken ct = default);
}
```

The whole design philosophy is in the contrast between the two query methods:

- **Query()** is the paved road. It returns `db.Set<TEntity>()` — which, for anything

`ITenantScoped`, already carries 2.6's global filter. A feature writes `Query().Where(n => n.Pinned).ToListAsync(ct)` and gets *only its tenant's* pinned notes, having written nothing about tenants. The guarantee rides underneath the ergonomics. `IQueryable` (not `Task<List<T>>`) is deliberate: the slice composes its own typed LINQ — filtering, projection, paging — and the repository never becomes a bag of `GetByFooBar` methods.

- **QueryAllTenants()** is the clearly-marked exit. Its whole job is to turn the filter

off. Read the doc comment's tone — "EVERY tenant's rows," "deliberately-named," "Greppable on purpose." This is the sanctioned version of the escape hatch you first met in 2.9, when a dissolve contributor had to wipe a tenant *other* than the current one. The enforcement is more precise than "banned in features," and the precision is the lesson (three arch-test rules you'll write in 3.3): **raw IgnoreQueryFilters() is banned everywhere under src/Api/Features/** — so the *only* way a slice can cross tenants is this named, greppable method; and **QueryAllTenants()** itself is allowed only inside a slice's `*DataContributor.cs` (dissolve and export are inherently cross-tenant), never on the request path — an endpoint or handler that reaches for it fails the build. So a feature's *read path* has no exit at all; its *teardown hook* has exactly one, and it announces itself in a file named for the purpose. The escape hatch and its bans are halves of one decision.

Forward glance (8.4). The real `EfRepository.QueryAllTenants()` also chains `.TagWith(RlsTags.CrossTenant)`. That tag matters only once the Postgres row-level security backstop exists (Lesson 8.4) — it tells the *database* "this cross-tenant read is sanctioned." Until then your implementation is just `IgnoreQueryFilters()`; the tag is a seam we'll thread through later. It's here now so you recognize it when it returns.

3. The implementation — one registration covers every entity

```
// src/Infrastructure/Repositories/EfRepository.cs — the real file
public class EfRepository<TEntity>(AppDbContext db) : IRepository<TEntity> where TEntity :
class
{
    public IQueryable<TEntity> Query() => db.Set<TEntity>();

    public IQueryable<TEntity> QueryAllTenants() =>
        db.Set<TEntity>().IgnoreQueryFilters(); // + .TagWith(RlsTags.CrossTenant) in 8.4

    public async Task AddAsync(TEntity entity, CancellationToken ct = default) =>
        await db.Set<TEntity>().AddAsync(entity, ct);

    public void Update(TEntity entity) => db.Set<TEntity>().Update(entity);
    public void Remove(TEntity entity) => db.Set<TEntity>().Remove(entity);

    public Task<int> SaveChangesAsync(CancellationToken ct = default) =>
        db.SaveChangesAsync(ct);
}
```

Thirty lines, and every entity type in the platform is served — because of one **open-generic registration**:

```
services.AddScoped(typeof(IRepository<>), typeof(EfRepository<>));
```

That single line means `IRepository<Note>`, `IRepository<Widget>`, and every future `IRepository<Whatever>` resolve with no per-entity wiring. Ask the container for `IRepository<Note>` and it constructs `EfRepository<Note>` on demand. This is the boilerplate-killer: adding a feature entity adds *zero* data-access code.

Notice what `EfRepository` injects: `AppDbContext` — the **request-scoped** one, the same instance the tenant filter (2.6) and stamping interceptor (2.7) are attached to. The repository is a thin pass-through, not a second layer of caching or logic. Its value is entirely in *what it refuses to expose*: a feature holding `IRepository<Note>` cannot reach `IRepository<OtherFeaturesThing>`'s data by accident, can't touch platform tables, and defaults to the filtered read path.

4. Unit of work — where atomicity lives

`SaveChangesAsync` on the repository flushes the shared context. But what about an operation that must write *several* things all-or-nothing? That's a different concern — transaction boundaries — and it gets its own seam, whose doc comment is the clearest statement of the platform's commit model:

```
// src/Core/Repositories/IUnitOfWork.cs
public interface IUnitOfWork
{
    Task<ITransactionScope> BeginTransactionAsync(CancellationToken ct = default);
}
public interface ITransactionScope : IAsyncDisposable
{
    Task CommitAsync(CancellationToken ct = default);
}
```

The subtlety the doc comment nails: ***every repository shares the same request-scoped DbContext, so any SaveChangesAsync flushes all pending tracked changes, not just the caller's.*** That sounds alarming until you see how the two layers compose:

- **A single self-contained write** — one entity, no cross-entity invariant — just calls

SaveChangesAsync and rides EF's implicit per-call transaction. No ceremony.

- **A grouped operation** — several writes that must commit together — wraps them:

await using var tx = await uow.BeginTransactionAsync(ct); ... ; await tx.CommitAsync(ct);. Inside that scope, every intermediate SaveChangesAsync *enlists in the one open transaction*, and disposing the scope without committing rolls the whole group back.

You already saw this contract consumed before you built it: 2.8's atomic user+tenant+membership creation, and 2.9's dissolve that "enlists in the caller's transaction rather than opening its own." Both rely on exactly this model — one shared context, transactions layered on top when grouping is needed. The await using + commit-or-rollback shape is the same disposable-scope discipline as 2.7's EnterTenant; the platform reuses a small number of good shapes deliberately.

(The EfUnitOfWork implementation wraps EF's Database.BeginTransactionAsync; on a non-relational provider it degrades to a no-op so the same service code runs everywhere — though this course always tests on real Postgres, where the transaction is real.)

5. Red — prove the seam keeps the promise

The seam's entire reason for existing is that Query() *cannot leak*. So the test that matters re-proves 2.6's isolation, but *through the repository* a feature will actually use (RepositoryScopingTests in the reference repo):

```
[Fact]
public async Task Query_is_tenant_scoped_but_QueryAllTenants_sees_everything()
{
    await SeedNoteAsync(tenantA, "A's note");
    await SeedNoteAsync(tenantB, "B's note");

    var repoAsA = RepositoryActingAs(tenantA); // scoped context, tenant A

    Assert.Single(await repoAsA.Query().ToListAsync()); // only A's
    Assert.Equal(2, await repoAsA.QueryAllTenants().CountAsync()); // the escape hatch sees both
}
```

Two assertions, two halves of the contract: the paved road is isolated; the marked exit genuinely exits. (The *bans* on taking that exit are enforced separately, by the arch tests in 3.3 — raw `IgnoreQueryFilters()` nowhere in features, `QueryAllTenants()` only in `*DataContributor.cs`. This test proves the mechanism works; those prove features can't misuse it.)

6. Run it

```
dotnet test tests/Api.Tests --filter Repository
```

No user-facing surface yet — this is infrastructure the next lesson consumes. The proof that it's right is the isolation test above plus the fact that 3.2's Notes slice will contain *no data-access code of its own*.

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how do feature slices access data? (a) A bespoke repository per entity; (b) the raw `DbContext` injected into handlers; (c) one generic `IRepository<T>` with a tenant-filtered `Query()` and a named, greppable `QueryAllTenants()` escape hatch, plus a separate `IUnitOfWork` for grouped transactions.

Chosen: (c) (ADR-004). It deletes per-entity boilerplate via open-generic registration, channels every feature read through 2.6's filter *by default*, and makes the dangerous cross-tenant path a distinct method that review and R5's arch test can police. Atomicity is factored into its own seam so the common single-write case stays ceremony-free while grouped writes get real transactions.

Rejected: (a) — boilerplate tax plus N chances to write an unfiltered query; the generic repo makes "forgot to scope" unreachable from the paved road. (b) — hands features the keys to the whole database and turns every chassis guarantee into a convention; the seam exists precisely to *not* do this.

The trade: a generic `IQueryable`-returning repository is less prescriptive than named query methods — a slice *can* compose an inefficient query. Accepted: the alternative (a method per access pattern) is the boilerplate we're eliminating, and slice code is reviewed. Platform/auth entities keep their dedicated repositories (they have bespoke needs and predate this seam); `IRepository<T>` is for feature tables.

8. Checkpoint

```
git add -A && git commit -m "feat: generic repository seam – tenant-safe Query, named
escape hatch, unit of work"
git tag lesson-3.1
```

You should now be able to: explain why a generic repository beats both per-entity repositories and raw `DbContext` access; describe how `Query()` inherits tenant isolation for free and why `QueryAllTenants()` is deliberately conspicuous; state the shared-context commit model and when a slice needs an explicit transaction; and register an open-generic service.

Next: *Lesson 3.2 — Anatomy of a vertical slice (Notes)* — the reference slice that consumes this seam, and the seven-step shape every feature you ever build will follow.

Lesson 3.2 — Anatomy of a vertical slice (Notes)

Part 3 · The slice pattern & the UI. This is the lesson the whole platform was built to make possible. Every feature you add from here — and every feature a *real product* built on this template adds — follows one shape: a self-contained folder that bolts onto the chassis and reuses everything Part 2 gave you. We learn that shape from Notes, the reference slice, which exists to be copied and then deleted. Five small files, one DbSet, one registration — no changes to platform *behavior*. Master this and you can add a feature in an afternoon; the rest of the course is mostly this shape at higher stakes.

Goal: a complete CRUD feature — list/create/delete notes — reachable at `/api/notes`, tenant-isolated, wired into dissolve and export, driven by tests, and touching platform code only by *addition* — a DbSet on `AppDbContext` and a few registration lines in `Program.cs` — never changing platform *behavior*.

Concepts & principles: the vertical-slice feature folder; minimal-API feature groups vs controllers; auth-by-construction (`MapTenantFeatureGroup`); the thin handler; co-located DTOs; the four-member contributor; slice isolation rules; the request-path vs contributor escape-hatch boundary.

Maps to: ADR-004 · R6–R10, R34/R35 · `docs/WAYS_OF_WORKING.md` (the add-a-slice checklist) · repo: `src/Api/Features/Notes/*`, `src/Core/Entities/Note.cs`, `src/Api/Endpoints/FeatureEndpointExtensions.cs`, `src/Api/Program.cs` (the `MapNotes()` + DI lines), tests: `tests/Api.Tests/NotesSliceTests.cs`, `FeatureAuthorizationTests.cs`.

Prerequisites: 3.1 (the repository seam), all of Part 2 (the chassis this reuses).

1. Motivate — the shape is the product

The platform's value proposition, stated plainly: *adding a feature should be boring*. No decisions about auth, no decisions about tenant scoping, no editing a central router or a central teardown method, no fear of breaking another feature. If adding a feature is exciting, your architecture is charging you interest.

`docs/WAYS_OF_WORKING.md` distills the shape into a seven-step checklist — memorize it, because it *is* the rest of this course:

- 1 Entity** → `src/Core/Entities/<Entity>.cs`, implementing `ITenantScoped`.
- 2 DbSet + config** → add to `AppDbContext`, add an `IEntityTypeConfiguration`.
- 3 Migration** → `dotnet ef migrations add Add<Entity>`.
- 4 DI wiring** → register the handler + map the group in `Program.cs`.
- 5 Contributor** → register the `ITenantDataContributor` (all four members).
- 6 Fixture reset** → (automatic — the reset list is model-derived, 1.3).
- 7 UI** → nav entry + component in the `Shared.Ui` RCL + resx strings (Part 3.4/3.5).

Notes is a worked example of steps 1–5. Let's build it, reading each file as an instance of a principle you already know.

2. The entity (step 1)

```
// src/Core/Entities/Note.cs
public class Note : ITenantScoped
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public Guid TenantId { get; set; } // ← the marker's one requirement
    public required string Title { get; set; }
    public string Content { get; set; } = "";
    public DateTimeOffset CreatedAt { get; set; }
    public DateTimeOffset UpdatedAt { get; set; }
}
```

: `ITenantScoped` is the entire cost of joining the tenancy regime. The moment this compiles, 2.6's global filter scopes every read of `Note`, 2.7's interceptor stamps and guards every write, the R2 invariant test counts it as covered, and 2.9's wipe- completeness gate expects it to be torn down. One interface, and the chassis adopts your entity. Step 2 (a `DbSet<Note>` on `AppDbContext` + a `NoteConfiguration` setting the `Title` max length) and step 3 (dotnet ef migrations add `AddNotes`) are mechanical — the discovery patterns from 1.3 pick the configuration up automatically.

3. The endpoints — auth you can't forget (step 4a)

```
// src/Api/Features/Notes/NotesEndpoints.cs
public static class NotesEndpoints
{
    public static IEndpointRouteBuilder MapNotes(this IEndpointRouteBuilder app)
    {
        var group = app.MapTenantFeatureGroup("/api/notes"); // ← the one true way in

        group.MapGet("/", async (NotesHandler handler, CancellationToken ct) =>
            Results.Ok(await handler.ListAsync(ct)));

        group.MapPost("/", async (CreateNoteRequest request, NotesHandler handler,
            CancellationToken ct) =>
        {
            var created = await handler.CreateAsync(request, ct);
            return created is null
                ? Results.BadRequest(new { error = "invalid_request", message = "A note title is required" })
                : Results.Created($"{api/notes/{created.Id}", created);
        });

        group.MapDelete("/{id:guid}", async (Guid id, NotesHandler handler,
            CancellationToken ct) =>
            await handler.DeleteAsync(id, ct) ? Results.NoContent() : Results.NotFound());

        return app;
    }
}
```

The load-bearing line is the first one. `MapTenantFeatureGroup` — recall its real definition from 3.1's neighbor file:

```
public static RouteGroupBuilder MapTenantFeatureGroup(this IEndpointRouteBuilder app,
string prefix) =>
    app.MapGroup(prefix).RequireAuthorization(AuthPolicies.TenantApi);
```

It is the *only* sanctioned way a feature registers routes, and the reason is R6: a slice must **not** write `app.MapGroup("/api/notes").RequireAuthorization(...)` by hand, because a hand-rolled group can forget the `.RequireAuthorization`, and a feature endpoint with forgotten auth is an unauthenticated data leak. The helper makes "authenticated, tenant-scoped" the *only* thing you can express — an arch test (3.3) fails any feature group not created through it. This is 2.6's philosophy applied to routing: don't ask developers to remember the guarantee; make the guaranteed path the only path.

Two more notes on style: features are **minimal-API groups**, while the *platform* stays controllers (auth, household, billing — the durable chassis). And the endpoints are *thin* — parse, delegate to the handler, map the result to an HTTP status. No logic lives here.

4. The handler — where the logic lives, and how little there is

```
// src/Api/Features/Notes/NotesHandler.cs
public class NotesHandler(IRepository<Note> notes, ICurrentTenant tenant, TimeProvider
clock)
{
    public async Task<IReadOnlyList<NoteResponse>> ListAsync(CancellationTok
en ct)
    {
        // Query() is tenant-scoped by the global filter – no explicit Where(TenantId)
        needed.
        var rows = await notes.Query().OrderByDescending(n => n.CreatedAt).To
ListAsync(ct);
        return rows.Select(NoteResponse.From).ToList();
    }

    public async Task<NoteResponse?> CreateAsync(CreateNoteRequest request,
CancellationTok
en ct)
    {
        if (tenant.TenantId is not { } tenantId) return null; // no tenant on the token
        if (string.IsNullOrEmpty(request.Title)) return null; // endpoint maps null →
400

        var now = clock.GetUtcNow();
        var note = new Note { Id = Guid.CreateVersion7(), TenantId = tenantId,
            Title = request.Title.Trim(), Content = request.Content?.Trim() ?? "",
            CreatedAt = now, UpdatedAt = now };
        await notes.AddAsync(note, ct);
        await notes.SaveChangesAsync(ct);
        return NoteResponse.From(note);
    }

    public async Task<bool> DeleteAsync(Guid id, CancellationTok
en ct)
    {
        // Tenant-scoped lookup: a note from another tenant simply isn't found here.
        var note = await notes.Query().FirstOrDefaultAsync(n => n.Id == id, ct);
        if (note is null) return false;
        notes.Remove(note);
        await notes.SaveChangesAsync(ct);
        return true;
    }
}
```


Its own doc comment says it best: "*Note how little it needs.*" Three injected dependencies, all from the chassis: `IRepository<Note>` (3.1's seam), `ICurrentTenant` (2.6), and `TimeProvider` (the injected clock, formalized next lesson). Read the tenancy into it — this is where all of Part 2 pays out:

- `**ListAsync` writes no `Where(TenantId ==)`. `**Query()` is already filtered (2.6). A

developer *cannot* accidentally list another tenant's notes; the safe query is the natural one.

- `**DeleteAsync` gets tenant-scoped isolation for free. `**Query().FirstOrDefault(id)`

looks up by id *within the current tenant* — so a caller trying to delete another tenant's note by guessing its id gets `null` → 404, indistinguishable from "no such note." Cross-tenant access isn't forbidden with an error; it's *invisible*. That's the filter turning an authorization question into a not-found.

- `**CreateAsync` sets `TenantId` from the token's tenant**, and even if it didn't, 2.7's

interceptor would stamp it and refuse a foreign value. Belt and suspenders, by construction.

The handler contains exactly the logic that is *this feature's* — nothing about auth, tenancy plumbing, or persistence mechanics. That's the whole point of a chassis.

5. The models — co-located, on purpose (step 4b)

```
// src/Api/Features/Notes/NotesModels.cs
public record CreateNoteRequest(string? Title, string? Content);

public record NoteResponse(Guid Id, string Title, string Content, DateTimeOffset CreatedAt)
{
    public static NoteResponse From(Note n) => new(n.Id, n.Title, n.Content, n.CreatedAt);
}
```

The DTOs live *in the slice*, not in a shared `Models` project. This is deliberate (R7/R10): a shared DTO project becomes a coupling magnet — two features start sharing a DTO "to avoid duplication," and now they can't change independently or be deleted independently. A slice owns its request/response shapes; if two slices need a similar shape, they each have their own. The tiny duplication buys total independence — which is what makes the delete-me property real.

Note too that `NoteResponse` is a *projection*, not the entity: the API exposes `Id/Title/Content/CreatedAt`, deliberately omitting `TenantId` and `UpdatedAt`. The wire contract is chosen, never "just serialize the entity" — the entity is free to grow fields the API doesn't expose.

6. The contributor — and the escape-hatch boundary (step 5)

```
// src/Api/Features/Notes/NotesDataContributor.cs
public class NotesDataContributor(IRepository<Note> notes) : ITenantDataContributor
{
    public Task<bool> HasDataAsync(Guid tenantId, CancellationToken ct = default) =>
        // QueryAllTenants: dissolve runs for a tenant OTHER than the current one, so this
        // crosses tenants by design – the audited escape hatch, re-constrained to the
        target.
        notes.QueryAllTenants().AnyAsync(n => n.TenantId == tenantId, ct);

    public async Task WipeAsync(Guid tenantId, CancellationToken ct = default) =>
        await notes.QueryAllTenants().Where(n => n.TenantId ==
tenantId).ExecuteDeleteAsync(ct);

    public string ExportKey => "notes";

    public async Task<object?> ExportAsync(Guid tenantId, CancellationToken ct = default)
=>
        await notes.QueryAllTenants().Where(n => n.TenantId == tenantId)
            .OrderBy(n => n.CreatedAt)
            .Select(n => new { n.Id, n.Title, n.Content, n.CreatedAt, n.UpdatedAt })
            .ToListAsync(ct);
}
```

Registering this (one line in `Program.cs`) is *all* it takes for Notes to participate in dissolve (2.9) and GDPR export (7.1) — no central method to edit. That's Open/Closed again, from the feature's side.

Now the subtle, precise thing this file teaches — the one I want you to carry into every slice. ****The contributor uses `QueryAllTenants()`, and that is allowed here and *only* here.**** Dissolve and export run for a tenant *other* than the ambient one (the system context has no current tenant), so they genuinely must cross tenants — hence the escape hatch, immediately re-constrained with `Where(n => n.TenantId == tenantId)`. But the arch test (3.3) enforces a three-part boundary that's worth stating exactly:

- Raw `IgnoreQueryFilters()` is banned **everywhere** in `Features/**` — even the contributor must go through the named `QueryAllTenants()`, never the raw EF call.
- `QueryAllTenants()` is permitted ****only in `*DataContributor.cs`**** — an **endpoint or handler** that calls it fails the build (v2 audit ADV-2). The request path has no exit.
- So `NotesHandler` (request path) uses `Query()`; `NotesDataContributor` (teardown) uses `QueryAllTenants()`. The boundary runs between the two files, and a machine patrols it.

That granularity — cross-tenant access allowed for teardown, forbidden for serving requests — is what mature tenant isolation looks like: not "never cross tenants" (some operations must) but "cross them only in named places a reviewer and a test can find."

7. Wiring and isolation — the rules that keep slices independent

Registration, the whole of step 4/5, in `Program.cs`:

```
builder.Services.AddScoped<NotesHandler>();
builder.Services.AddScoped<ITenantDataContributor, NotesDataContributor>();
// ... after app is built:
app.MapNotes();
```

And the isolation rules that make a slice a *unit* (R7–R9, each an arch test in 3.3):

- **A slice may not reference another slice's namespace** (R7) — Features/Notes can't

reach into Features/Billing. Slices compose only through the platform, never through each other, so any one can be deleted without breaking the rest.

- ****Only Program.cs imports Features.**** — the composition root wires slices; nothing

else knows they exist.

- ****Platform tests must not depend on the Notes sample**** — the reference slice is

deletable *including from the test suite's assumptions*. (There's literally an arch test listing `Set<Note>`, `IRepository<Note>`, etc. as banned patterns outside the Notes tests themselves.)

Which is the last thing to internalize: Notes is marked **DELETE-ME** in every file. It is a *worked example*, not a feature — you copy its shape for your first real slice and then delete it (Lesson 9.1's checklist includes the deletion, and the arch tests above guarantee nothing secretly depended on it). The reference implementation ships the teacher and a promise that the teacher can leave.

8. Red — the slice's tests

```
Scenario: notes are tenant-isolated end to end
  Given tenant A created a note and tenant B created a note
  When A lists /api/notes
  Then A sees only A's note
  And A deleting B's note id returns 404 (not 403 – it's invisible, not forbidden)

Scenario: a note requires a title
  When A posts a note with a blank title
  Then the response is 400 invalid_request

Scenario: the feature group requires auth
  When an unauthenticated caller hits /api/notes
  Then the response is 401
```

NotesSliceTests drives the behavior; FeatureAuthorizationTests proves the group is guarded (the auth-by-construction claim, tested). The isolation scenario is the one that matters — and note the assertion is **404, not 403**: cross-tenant access is invisible, per §4.

9. Run it

```
dotnet test tests/Api.Tests --filter "Notes|FeatureAuthorization"
dotnet run --project src/Api
# .http: POST /api/notes {"title":"first"} → 201 ; GET /api/notes → your note
```

10. Architecture Decision

ARCHITECTURE DECISION

The fork: how are app features structured? (a) Controllers alongside the platform's; (b) a service/repository layer cake shared across features; (c) self-contained vertical-slice folders (endpoints + handler + DTOs + contributor), minimal-API groups, DTOs co-located, one DI registration.

Chosen: (c) (ADR-004). A slice is independently understandable, independently testable, and independently *deletable*; it reuses the chassis and edits no platform code;

MapTenantFeatureGroup makes auth unforgettable; the contributor plugs into lifecycle without a central method. The isolation arch tests (R7–R9) keep the independence real rather than aspirational.

Rejected: (a) — controllers work, but the platform reserves them for its durable surface and gives features the lighter minimal-API shape with the auth guarantee baked into the group helper; mixing both freely would blur "platform vs feature." (b) — shared layers are where coupling accretes: the shared DTO, the shared service "just this once," until slices can't move independently. Co-location's small duplication is the price of independence, and it's worth it.

The trade: minimal-API groups are less ceremonious than controllers (no attribute routing, model binding is more manual) and DTO co-location duplicates similar shapes across slices. Both accepted deliberately: the first keeps features light, the second keeps them independent. The platform's own controllers remain for the cases that want the heavier machinery.

11. Checkpoint

```
git add -A && git commit -m "feat: Notes reference slice – the vertical-slice shape every
feature copies"
git tag lesson-3.2
```

You should now be able to: recite the seven-step add-a-slice checklist; explain why MapTenantFeatureGroup exists and what R6 prevents; describe how a handler stays tiny by reusing Query()/ICurrentTenant; state the request-path vs contributor escape-hatch boundary precisely; and justify co-located DTOs and the delete-me convention.

Next: Lesson 3.3 — *Injected clocks & the architecture tests* — the TimeProvider you've been quietly using gets its rationale, and the arch-test project that has been enforcing all these rules gets built for real.

Lesson 3.3 — Injected clocks & the architecture tests

Part 3 · The slice pattern & the UI. Two threads that have been running quietly since Part 1 finally get their own lesson. First, the **injected clock**: every service you've written took a `TimeProvider` instead of calling `DateTime.UtcNow`, and it's time to say why — and to ban the alternative with a test. Second, the **architecture- test project**: throughout Parts 1–2 I kept promising "an arch test enforces this in 3.3." This is 3.3. We build the file that has been the invisible enforcer of every structural rule — and establish the pattern by which it *grows a test per rule* for the rest of the course.

Goal: all server code takes an injected clock; a build-failing test bans ambient `DateTime.UtcNow` in server projects; and the `ArchitectureTests` project exists, holding the invariants Parts 1–3 established, ready to accrete more.

Concepts & principles: deterministic time as a dependency; `TimeProvider` / `FakeTimeProvider`; architecture tests (source scans + model reflection); invariants as executable, build-failing assertions; audit findings hardened into permanent guards.

Maps to: R15 (+ R2, R6–R10, R13, R24, R34/R35 — the whole arch suite) · repo: `tests/Api.Tests/ArchitectureTests.cs`, and every service that injects `TimeProvider`.

Prerequisites: the Part-1 grep-gates (config sync 1.4, MailKit containment 2.3, error shape 1.5) — this lesson generalizes them; and 2.6/3.2 whose rules it enforces.

1. Motivate — the clock you can't test

Cast back to Lesson 2.2. You needed to prove a refresh token is rejected *after* it expires. How? With `DateTime.UtcNow` buried inside the service, your options are: sleep for the token's lifetime in the test (absurd — a 30-day token?), or monkey-patch a static (you can't, cleanly, in C#). Time-dependent logic built on ambient `UtcNow` is **structurally untestable** — the one input that determines the outcome is a global you can't control from the test.

That's the acute pain. The chronic one is subtler: ambient `UtcNow` scattered across a codebase means *many* clocks. A token minted reading one `UtcNow`, validated against another, logged with a third — and in a distributed system, on different machines with different skew. "What time is it" should be *one* answer, injected, so every collaborator in a request shares it. The fix is to make time a dependency like any other.

2. `TimeProvider` — time as a seam

.NET's `TimeProvider` (built into the modern BCL) is the abstraction. Inject it; call `GetUtcNow()`. You've been doing this since 2.1 without comment — here it is in the `NotesHandler` you just built:

```
public class NotesHandler(IRepository<Note> notes, ICurrentTenant tenant, TimeProvider
clock)
{
    public async Task<NoteResponse?> CreateAsync(CreateNoteRequest request,
CancellationToken ct)
    {
        var now = clock.GetUtcNow(); // ← the injected clock, not DateTimeOffset.UtcNow
        var note = new Note { /* ... */ CreatedAt = now, UpdatedAt = now };
        // ...
    }
}
```

In production you register the real one — `services.AddSingleton(TimeProvider.System)` — and `GetUtcNow()` returns the wall clock. In tests you inject `FakeTimeProvider` (from `Microsoft.Extensions.TimeProvider.Testing`, the package your test project already references), and time becomes a value you *control*:

```
var clock = new FakeTimeProvider(DateTimeOffset.Parse("2026-01-01T00:00:00Z"));
var svc = new PasswordlessService(/* ... */, clock);
var code = await svc.IssueOtpAsync("a@b.test", ct); // issued "now"

clock.Advance(TimeSpan.FromMinutes(11)); // jump past the 10-min lifespan

var result = await svc.RedeemOtpAsync("a@b.test", code, ct);
Assert.Equal(OtpStatus.Expired, result.Status); // deterministic, instant
```

Every expiry test you met in Part 2 — refresh-token expiry (2.2), OTP lifespan and logout window (2.4), magic-link expiry — works *because* of this. `Advance` makes "wait eleven minutes" happen in microseconds, with no flake. Deterministic time isn't a testing nicety bolted on; it's what makes an entire category of logic testable at all.

3. Green — ban the alternative

A convention followed everywhere is one `DateTime.UtcNow` away from a hole. So the platform makes it a build failure, from the real `ArchitectureTests`:

```
[Fact]
public void ServerServices_UseInjectedClock_NotAmbientUtcNow()
{
    // R15: server code takes TimeProvider, so a cookie/URL/token lifetime can't drift
    // from the injected clock. Entity default-value helpers and the WASM client (no
    // injected clock) live in Core/Shared.Ui and are outside this scan; migrations
    // excluded.
    var dirs = new[] { Path.Combine(RepoRoot(), "src", "Api"),
                      Path.Combine(RepoRoot(), "src", "Infrastructure") };

    var offenders = dirs.SelectMany(d => SourceFiles(d))
        .Where(f =>
            !f.Contains($"{Path.DirectorySeparatorChar}Migrations{Path.DirectorySeparatorChar}"))
        .Where(f => File.ReadAllText(f) is var t &&
            (t.Contains("DateTime.UtcNow") || t.Contains("DateTimeOffset.UtcNow")))
        .Select(Path.GetFileName)
        .ToList();

    Assert.True(offenders.Count == 0,
        $"Server services must use an injected TimeProvider, not ambient UtcNow: {string.Join(", ", offenders)}");
}
```

Write `DateTimeOffset.UtcNow` in any `Api` or `Infrastructure` file and the build goes red with the file's name. But read the comment's *scope*, because the exclusions are where the real thinking is — a blunt "ban `UtcNow` everywhere" would be wrong:

- **Core entity defaults are exempt.** `Note.CreatedAt` isn't set by a default helper, but

some entities use `DateTimeOffset.UtcNow` as a property initializer. Core has no request, no DI, no injected clock — so a *default* there is acceptable; what matters is that the **service** overwrites it with `clock.GetUtcNow()` on the write path (as `NotesHandler` does). The rule targets the code that *makes decisions* from time.

- **The WASM client is exempt.** The Blazor client (`Shared.Ui`) runs in the browser with

no server clock to inject; "now" there is display, not a security decision.

- **Migrations are exempt.** Generated code, timestamped by tooling.

That list of exclusions is the difference between a rule people respect and one they route around. A guard that fires on legitimate cases gets suppressed; a guard scoped to exactly the risk (server-side *decisions* from time) stays trusted. You saw this instinct in 1.6's prohibited-list license gate — same discipline, applied to time.

4. Zoom out — the architecture-test project

That clock test is one of a family. `ArchitectureTests.cs` is where the platform's structural guarantees stop being prose in `CLAUDE.md` and become assertions the build runs. It has been referenced in nearly every lesson — "an arch test enforces this in 3.3" — and now you build it. Two mechanisms, both of which you've already met in miniature:

Source scans — `grep` with a verdict. The exact shape of 1.4's config-sync gate, 2.3's MailKit containment, 1.5's error-shape check. Read files, look for a forbidden (or required) pattern, fail with the offending file name. The tenant-filter-bypass test (the three-part `IgnoreQueryFilters / QueryAllTenants / RlsTags` rule from 3.1–3.2), the "only `Program.cs` imports `Features.*`"

isolation rule (R8), "features use MapTenantFeatureGroup not raw MapGroup" (R6), "each file declares a type matching its name" (R24) — all source scans.

Model reflection — ask the EF model about itself. This is the more powerful kind, because it covers entities *that don't exist yet*. The keystone example:

```
[Fact]
public void EveryTenantScopedEntity_HasAGlobalQueryFilter()
{
    using var ctx = new AppDbContext(modelOnlyOptions, new TestCurrentTenant());

    var missing = ctx.Model.GetEntityTypes()
        .Where(e => typeof(ITenantScoped).IsAssignableFrom(e.ClrType))
        .Where(e => e.GetDeclaredQueryFilters().Count == 0)
        .Select(e => e.ClrType.Name)
        .ToList();

    Assert.True(missing.Count == 0,
        $"ITenantScoped entities missing the global tenant query filter: {string.Join(", ", missing)}");
}
```

This is 2.6's R2 invariant, and it's *self-extending*: add an `ITenantScoped` entity that somehow escapes the filter, and it names itself in the failure — no one has to remember to add a test for the new entity, because the test iterates the model. The same technique guards distinct table mappings (R35), unique export keys across contributors (R13), unique route prefixes (R35), and — arriving with their epics — that every user-keyed entity is wired into GDPR erasure (7.1) and every controller derives from a tenant/admin base (6.1).

5. The philosophy — invariants that can't regress

Here's the throughline of the entire course, stated once. Every structural guarantee in this platform exists at three levels: **prose** (the ADR / `CLAUDE.md` — the *why*), **a mechanism** (the query filter, the interceptor, the seam — the *how*), and **a test** (the arch test — the *proof it stays true*). The prose can be ignored; the mechanism can be bypassed; only the test, wired into 1.6's CI, makes the guarantee *permanent*.

Look at where these tests came from. Several cite v2 audit ADV-1, ADV-2, TR-1 in their comments — they are **audit findings turned into guards**. Someone once found that a rotated-out row could be updated cross-tenant (2.7's ADV-1), or that a handler could call `QueryAllTenants` (3.2's ADV-2); the fix was code, but the *durable* fix was a test so the hole can never silently reopen. That is what "the architecture-test project grows a test per rule" means in practice: every time the platform learns something — from an audit, from a bug, from a new epic's invariant — the lesson is recorded as an executable assertion. Build the file now with the Part 1–3 rules; every part after this adds its own alongside the feature that motivates it.

6. Run it

```
dotnet test tests/Api.Tests --filter Architecture
# the whole structural suite, green
```


Then feel two of them bite (pain now, trust forever): write `DateTimeOffset.UtcNow` in a handler → the clock test names your file; add a raw `app.MapGroup("/api/x")` in a `Features/` file → the R6 test names it. Delete both; green returns. These failures are the platform refusing to let a guarantee rot.

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how are structural rules enforced? (a) Documentation + code review; (b) a dedicated architecture-testing library (NetArchTest, ArchUnitNET); (c) plain xUnit tests using source scans + EF-model reflection.

Chosen: (c). No new dependency (the scans are `File.ReadAllText` + a `Contains`; the model checks are `ctx.Model.GetEntityTypes()`), the model-reflection tests are *self-extending* (they cover future entities for free), and every rule reads as a plain test with a failure message that names the offender and the fix. They run in the same `dotnet test` as everything else, gated by CI.

Rejected: (a) — review is human and forgets; the whole point is guarantees that *can't* be forgotten. (b) — capable libraries, but they add a dependency and a second DSL to learn, and for the mostly-textual and EF-model-shaped rules here, a `Contains` scan is more transparent than a fluent `Types().That().ResideInNamespace()` chain. A larger codebase with deep layering rules might reasonably choose (b).

The trade: source scans are textual — they match strings, so they can be fooled (a `QueryAllTenants` in a comment trips the scan; an obfuscated call might slip it) and they're coarser than a real type analysis. Accepted: these guard against *accidents and drift*, not a determined adversary editing around them, and the model-reflection tests (which use real type info) cover the highest-stakes rules precisely.

8. Checkpoint

```
git add -A && git commit -m "test(arch): injected-clock ban + the architecture-test project
(invariants as tests)"
git tag lesson-3.3
```

You should now be able to: explain why ambient `UtcNow` is structurally untestable and how `TimeProvider/FakeTimeProvider` fix it; justify the clock ban's *exclusions* as carefully as the ban; distinguish source-scan from model-reflection arch tests and say which is self-extending; and articulate the prose→mechanism→test trichotomy that makes a guarantee permanent.

Next: Lesson 3.4 — *The web client & auth UI* — the first pixels. Blazor WebAssembly, the shared RCL, and the client that consumes everything Part 2 built over HTTP.

Lesson 3.4 — The web client & auth UI

Part 3 · The slice pattern & the UI. The first pixels. Everything so far has been API — testable, headless, invisible. Now we build the Blazor WebAssembly client that consumes it, and three architectural decisions dominate the lesson: **UI lives in a shared class library** (not the web app), so non-web clients get it for free; the **access token lives in memory only**, with sessions surviving reloads via the silent refresh that completes 2.2's rotation story; and the two things that differ between web and native are hidden behind **seams designed now, web-first**, so the MAUI shells in the appendix are cheap. The client is just another API consumer — held to the same clean-boundary rule as any other.

Goal: a Blazor WASM app where an unauthenticated visitor lands on a login page, signs in (magic link / OTP / OAuth) against the Part-2 API, stays signed in across reloads, and sees their notes — with every component living in the shared RCL.

Concepts & principles: Blazor WASM + the shared RCL (why UI-in-a-library); the client as an API client (clean boundary); in-memory access token + silent refresh; per-host seams (*ISessionStore*, *IOAuthInitiator*) designed web-first; delegating handlers and a DI-cycle-avoiding client split; single-origin readiness.

Maps to: ADR-C4/C5, Golden Rule 3 · repo: `src/Web/Program.cs`, `src/Web/Http/CookieHandler.cs/AuthHeaderHandler.cs/CookieSessionStore.cs`, `src/Shared.Ui/Auth/AuthService.cs/ISessionStore.cs/IOAuthInitiator.cs/AppClaims.cs`, `src/Shared.Ui/App.razor`, `Layout/MainLayout.razor`, `Pages/Login.razor`, `Pages/AuthCallback.razor`, arch test `WebApp_HasNoInlineBlazorComponents`.

Prerequisites: Part 2 (the whole auth API), 3.2 (the Notes API the UI will show).

1. Motivate — where does the UI live?

The obvious place for a Blazor page is the Blazor web app. The platform puts it somewhere else — a **Razor Class Library** (*Shared.Ui*) — and the web app becomes a nearly-empty host. Why fight the obvious?

Because of a promise made in Lesson 0.1: this platform ships (eventually) desktop and mobile shells (MAUI), and *web-first* means "build the feature on web, then extend the native shells cheaply." That's only possible if the native shells can *reuse the actual UI*. If your login page, your notes list, your settings screen live in the web app, the MAUI apps must reimplement all of it. If they live in a class library, every host — web and native — references the same components and gets the same UI for free. So the rule (Golden Rule 3): **Blazor components live in the shared RCL, never inline in the web app**. The web app carries only bootstrap markup.

This is enforced, of course — you built the mechanism in 3.3:

```
[Fact]
public void WebApp_HasNoInlineBlazorComponents()
{
    var allowed = new HashSet<string> { "_Imports.razor", "App.razor", "Routes.razor" };
    var inline = SourceFiles(Path.Combine(RepoRoot(), "src", "Web"), "*.razor")
        .Select(Path.GetFileName).Where(name => !allowed.Contains(name!)).ToList();
    Assert.True(inline.Count == 0,
        $"Blazor components belong in Shared.Ui, not the web app: {...}");
}
```

Put a .razor component in the web app and the build fails, naming it. The web-first promise is kept by a test, not a habit.

2. The client is an API client — the clean boundary, from the outside

The web app talks to the API the same way any third party would: HTTP + JSON, bearer token, no shortcuts. It has **no reference to Infrastructure, no DbContext, no shared entity types** — it can't; the project references forbid it (1.1). This is the clean API boundary (Golden Rule 2) seen from the client's side: the UI cannot reach around the API to the database even if a developer wanted to, because the types simply aren't there. When the Notes list renders, it does so from a GET /api/notes and a DTO it deserializes locally — never from a Note entity.

That constraint is a feature. It means the API's contract is the *only* contract, the same one the public API (7.3) and the native shells honor. There is no privileged client.

3. The session, client-side — memory + silent refresh

Here's the decision that surprises people: **the access token is never persisted**. Not in localStorage, not in sessionStorage, not anywhere on disk. From AuthService's own summary:

The access token is held in memory only — never persisted. Sessions survive reloads/restarts via the refresh token: on startup we silently exchange it for a fresh access token.

Why refuse to store the access token? Because localStorage is readable by any JavaScript on the page — one XSS and your bearer token is exfiltrated. Keeping it in a private field means a reload *loses* it — which sounds like a bug until you remember 2.2: the durable credential is the **HttpOnly refresh cookie**, which JavaScript *cannot* read and the browser *does* persist. So the session model is:

```
// src/Shared.Ui/Auth/AuthService.cs (trimmed to the web path)
private string? _accessToken; // memory only
private Task<bool>? _refreshInFlight;

public bool IsAuthenticated => !string.IsNullOrEmpty(_accessToken) &&
!IsTokenExpired(_accessToken);

/// Resolves the session on app startup: with no token in memory, silently exchange the
/// HttpOnly refresh cookie for an access token. Called ONCE from MainLayout.
public async Task InitializeAsync()
{
    if (IsAuthenticated) return;
    await TryRefreshAsync();
}

public Task<bool> TryRefreshAsync() => _refreshInFlight ??= RunRefreshAsync();

private async Task<bool> RunRefreshAsync()
{
    // Web: the CookieHandler attaches the HttpOnly refresh cookie; no body.
    var response = await httpClient.PostAsync("/api/auth/refresh", null);
    if (!response.IsSuccessStatusCode) { await ClearSessionAsync(); return false; }
    var payload = await response.Content.ReadFromJsonAsync<TokenResponse>();
    _accessToken = payload!.AccessToken; // fresh token, in memory
    return true;
}
```

Trace a page reload: memory is wiped, `_accessToken` is null, but the refresh cookie survived. `MainLayout` calls `InitializeAsync` before rendering the body; it silently POST `/api/auth/refresh`; the cookie rides along (via the handler in §4); the API rotates the token (2.2) and returns a fresh access token; the app is signed in again — invisibly. The user never sees a login screen they didn't ask for. This is the *client half* of the rotation dance you built server-side in 2.2; now you can see why the refresh token was worth all that care.

One subtlety worth its own sentence — `TryRefreshAsync` coalesces concurrent calls into one in-flight `Task`:

```
public Task<bool> TryRefreshAsync() => _refreshInFlight ??= RunRefreshAsync();
```

Because refresh **rotates** (2.2), two overlapping refreshes would race: the first burns the cookie, the second presents the now-revoked token and trips reuse detection — logging the user out of their own session. Sharing one `Task` makes concurrent callers await the same refresh. It's a direct consequence of a server-side decision rippling into client design — exactly the kind of cross-layer reasoning this course exists to build.

4. The HTTP plumbing — two clients, one subtle DI cycle

The client needs two things attached to its requests: the **refresh cookie** (for `/refresh`) and the **bearer token** (for everything authenticated). Delegating handlers do this. But there's a trap, and the platform's comment documents the escape:

```
// src/Web/Program.cs
// "Api" — general-purpose client used by components: credentials + Bearer token.
builder.Services.AddHttpClient("Api", c => c.BaseAddress = new Uri(apiBase))
    .AddHttpMessageHandler<CookieHandler>() // sends the HttpOnly cookie
    .AddHttpMessageHandler<AuthHeaderHandler>(); // attaches the in-memory JWT

// "ApiAuth" — used ONLY by AuthService for refresh/logout. Credentials only (no Bearer
// handler) to avoid a DI cycle (AuthHeaderHandler depends on AuthService).
builder.Services.AddHttpClient("ApiAuth", c => c.BaseAddress = new Uri(apiBase))
    .AddHttpMessageHandler<CookieHandler>();
```

The cycle: `AuthHeaderHandler` needs `AuthService` (to read the token), but `AuthService` needs an `HttpClient` (to call `/refresh`). If `AuthService`'s client carried the bearer handler, resolving it would loop. The fix is two clients: components use "Api" (cookie + bearer); `AuthService` uses "ApiAuth" (cookie only — it doesn't need a bearer to refresh, the cookie is its credential). Recognizing and breaking a DI cycle is a rite of passage; here it's a two-line decision with a comment so the next person doesn't "simplify" it back into a loop.

`CookieHandler` itself is tiny but non-obvious — WASM must *opt in* to sending cookies:

```
public class CookieHandler : DelegatingHandler
{
    protected override Task<HttpResponseMessage> SendAsync(HttpRequestMessage request,
        CancellationToken ct)
    {
        request.SetBrowserRequestCredentials(BrowserRequestCredentials.Include);
        return base.SendAsync(request, ct);
    }
}
```

Without `Include`, the browser withholds cookies on the fetch and every refresh silently fails — the kind of one-liner that costs an afternoon if you don't know it exists.

(A registration detail with a real reason: `AuthService` is a **singleton**. `IHttpClientFactory` resolves handlers in their own DI scope, so a *scoped* `AuthService` would hand the bearer handler a different instance with an empty token — and the header would never attach. One token store, shared by app and handler, requires singleton. The comment in `Program.cs` spells this out; it's a subtle Blazor-DI fact worth filing away.)

5. The per-host seams — designed web-first, native for free

Two things genuinely differ between the web browser and a native MAUI shell, and the platform isolates *exactly those two* behind interfaces — introduced now, on web, so the appendix's native shells only implement the interface rather than re-architect the app.

Where the refresh token lives — `ISessionStore`:

```
// src/Shared.Ui/Auth/ISessionStore.cs
/// Abstracts where the refresh token lives, the one thing that differs between hosts.
public interface ISessionStore
{
    bool UsesBodyTransport { get; } // web: false (cookie) · native: true (body)
    Task<string?> GetRefreshTokenAsync(); // web: always null · native: from secure store
    Task SaveRefreshTokenAsync(string token); // web: no-op · native: persist rotated token
    Task ClearAsync(); // web: no-op · native: clear
}
```

On web, the browser owns the HttpOnly cookie, so the store is a **no-op** and `UsesBodyTransport` is false — the refresh call carries no body and relies on the cookie (`CookieSessionStore` in the web project). On native (appendix), there's no shared cookie jar, so the token lives in the OS secure store and travels in the request body. Look back at `RunRefreshAsync`: the branch on `sessionStore.UsesBodyTransport` is the *entire* web-vs-native difference in the auth flow, and it's one `if`. That's the payoff of finding the true seam.

How interactive OAuth starts — `IOAuthInitiator`: on web, "Sign in with Google" is a *full-page navigation* to the API's challenge URL (the browser handles the whole OAuth round-trip and returns to the app), so the web host **never registers or resolves** this service. On native (appendix), there's no full-page nav, so an implementation opens the system browser and captures the redirect. Designing the seam now, while building web, is what "web-first" operationally means — not "ignore native," but "shape the abstractions so native slots in."

A warnings-as-errors wrinkle worth internalizing (it bites the moment you try this): define the **interface** now — that's the seam — but do *not* yet add a constructor parameter for it to `AuthService`. The full reference takes `IOAuthInitiator? oauth = null` because its native methods *use* it; at this web-only stage those methods don't exist, so an injected-but-unread parameter trips **CS9113 ("parameter is unread") → build failure** under 1.1's `TreatWarningsAsErrors`. The parameter returns in the appendix, alongside the native code that reads it. The lesson: a seam is an *interface you define early*, not necessarily a *dependency you inject early* — inject it when something consumes it, or the compiler (rightly) objects to the dead wire.

6. The pages — login and the callback

Two components carry the auth UX (both in the RCL, both routable):

- `**Login.razor**` renders what the *deployment* supports: an email box for magic-link /

OTP, and OAuth buttons for exactly the providers `GET /api/auth/providers` reports (2.5's `EnabledAsync` — so an unconfigured provider shows no dead, 500-ing button; the UI fails closed to fewer options). Magic link and OTP post to the Part-2 endpoints (2.4); OAuth buttons do the full-page navigation of §5.

- `**AuthCallback.razor**` is where an OAuth (or magic-link) round-trip lands back in the

SPA. Because the browser did the redirecting, the app "boots fresh" here — so the callback simply runs the same `TryRefreshAsync` (the API set the refresh cookie during the callback), obtains an access token, and routes onward. The whole external-flow return reuses the silent-refresh machinery; there's no separate "accept the OAuth token" path on web.

One detail on the login page repays a callback to 2.4. When a sign-in *fails*, the client must choose what to say — and it must not undo the server's enumeration safety. A tiny shared helper decides:

```
// src/Shared.Ui/Auth/AuthErrorCopy.cs — maps a server error code to the copy key to show
public static string OtpErrorKey(string? errorCode) =>
    errorCode == "too_many_attempts" ? "Login_ErrTooManyAttempts" :
    "Login_ErrCodeIncorrect";
```

Recall 2.4: the server deliberately collapses "wrong code" and "no active code" into one generic `invalid_code` (enumeration safety — CONF-6), but reports lockout as a *distinct* `too_many_attempts` (CONF-5, so the user learns retrying won't help until the window elapses). The client mirrors exactly that shape: lockout gets its own message, everything else stays on the deliberately-generic line — it does **not** invent a more specific "wrong code" message the server refused to give, which would leak the very distinction the server hid. The helper is shared across *every* sign-in surface (web now, native OTP in the appendix) so they can't drift into disagreeing about what an error means. (The two copy keys land as localized strings in 3.5.)

`MainLayout` is the single auth entry point: it awaits `InitializeAsync` before rendering `@Body`, so every page below it can assume the session is resolved — no page triggers its own refresh, which (with the coalescing in §3) keeps rotation sane.

7. Single-origin, already

One line in `Program.cs` quietly prepares for deployment (Part 8):

```
var apiBase = builder.Configuration["ApiBaseUrl"] is { Length: > 0 } configured
    ? configured : builder.HostEnvironment.BaseAddress; // fall back to our own origin
```

In dev, `ApiBaseUrl` points the client at the API on a different port. In production (8.1), the API *serves this very bundle*, so the client and API share one origin and the fallback kicks in — same-origin, which deletes an entire class of cross-site-cookie pain (the reason 2.2's cookie fussed over `SameSite`). The client is written origin-agnostic so the deployment topology is a config value, not a code change.

8. Run it — the first end-to-end human flow

No new unit tests here (UI behavior is proven by the E2E harness, next lesson) — but this is the first lesson you *run and click*:

```
docker compose up -d
dotnet run --project src/Api # terminal 1
dotnet run --project src/Web # terminal 2 (dev: different port)
# browse to the Web dev URL → Login → request an OTP → read it from Mailpit (:8025)
# → you're in, you see your (empty) notes → reload the page → STILL in (silent refresh)
```

That reload-and-stay-signed-in moment is the whole session design working in front of you: memory was wiped, the cookie survived, the refresh was invisible.

9. Architecture Decision

ARCHITECTURE DECISION

The fork (two forks, really): (rendering) Blazor WebAssembly, Blazor Server, or a separate JS SPA framework? (token storage) access token in `localStorage`, or in memory with silent refresh?

Chosen: Blazor WASM + a shared RCL, with the access token in memory (ADR-C4/C5, Golden Rule 3). WASM keeps the UI a *pure API client* (no server-side render state, honors the clean boundary, and — decisively — the same C# components run in the MAUI shells). In-memory tokens + the `HttpOnly` refresh cookie put the durable credential where XSS can't reach it, at the cost of a silent refresh on startup — a cost the rotation design already required.

Rejected: *Blazor Server* — a stateful circuit per user, a persistent `WebSocket`, and a server that renders UI: it breaks "the UI is just a client," complicates scaling, and can't be a MAUI shell. A *JS SPA* (React/etc.) — fine technically, but a second language and toolchain for a .NET team, and no code-sharing with native. **localStorage tokens** — the standard footgun: convenient, and one XSS from total credential theft.

The trade: WASM ships a .NET runtime to the browser (a larger first load than a lean JS app) and needs the silent-refresh choreography (§3) instead of a token just sitting in storage. Both accepted: the download is one-time and cached, and the choreography buys XSS-resistant sessions plus one codebase across every client.

10. Checkpoint

```
git add -A && git commit -m "feat: web client + auth UI - RCL components, in-memory token +
silent refresh, host seams"
git tag lesson-3.4
```

You should now be able to: explain why UI lives in the RCL and how the arch test keeps it there; describe the in-memory-token + silent-refresh model and connect it to 2.2's rotation; identify and break the `AuthHeaderHandler`→`AuthService` DI cycle; and name the two genuine web-vs-native differences and the seams (`ISessionStore`, `IOAuthInitiator`) that isolate them.

Next: *Lesson 3.5 — Localization* — the UI (and the emails from 2.3) learn to speak more than one language, the right way.

Lesson 3.5 — Localization

Part 3 · The slice pattern & the UI. A short but genuinely architectural lesson. "Support Spanish" sounds like a translation task; it's actually three separate decisions — where strings live, how the app *decides* which language to show, and how that decision survives a cold start on each host. Do it now, while the UI is small, and every future component is born translatable. Retrofit it later and you're hunting hard-coded English through fifty files. And there's a satisfying payoff: the emails you built in 2.3 already localized themselves; this lesson makes the *UI* rhyme with them.

Goal: the UI renders in English or Spanish from resource files; a language switcher persists the choice per host; a signed-in user's saved locale follows them across devices; and adding a language is adding a `.resx`.

Concepts & principles: externalized strings (`resx` + `IStringLocalizer`); the culture decision chain (stored → user claim → default); the web-first culture-persistence seam; UI vs email localization (ambient culture vs explicit culture); fail-soft culture resolution.

Maps to: ADR-C(i18n) · docs/LOCALIZATION.md · repo:

`src/Shared.Ui/Resources/AppStrings.cs` + `AppStrings.resx/AppStrings.es.resx`,
`src/Shared.Ui/Components/LanguageSwitcher.razor`, `ICulturePersistence.cs`,
`LocalStorageCulturePersistence.cs`, `src/Web/Program.cs` (culture bootstrap),
`src/Core/Entities/User.cs` (Locale), tests: `tests/E2E.Tests/I18nTests.cs`.

Prerequisites: 3.4 (the client + `MainLayout`), 2.1 (`User.Locale`), 2.3 (localized emails).

1. Motivate — the hard-coded string is a one-way door

Write `<h1>Sign in</h1>` in a component and you've made a decision you didn't know you were making: this app is English-only, and undoing it means finding every literal across the codebase. Externalize from the first component and the marginal cost of a second language is a translation file — no code changes. The choice is *when* you pay: a little now (look strings up by key), or a lot later (a codebase-wide string hunt). i18n is cheap only if it's early, which is why it's a Part-3 lesson and not a "someday" ticket.

2. Where strings live — `resx` + `IStringLocalizer`

.NET's localization is resource files keyed by string, resolved through a typed marker. The marker is a nearly-empty class whose only job is to *name* a resource set:

```
// src/Shared.Ui/Resources/AppStrings.cs – the whole file
/// <summary>Marker type for the app's shared UI strings. Inject
/// IStringLocalizer<AppStrings> and look up keys (e.g. Localizer["Login_Title"]).
/// Translations live in co-located AppStrings.resx (English) +
AppStrings.{culture}.resx.</summary>
public sealed class AppStrings;
```

Beside it: `AppStrings.resx` (neutral/English) and `AppStrings.es.resx` (Spanish), each a key→string map — `Login_Title` → "Sign in" / "Iniciar sesión". Components inject the localizer

and look strings up:

```
@inject IStringLocalizer<AppStrings> L
<h1>@L["Login_Title"]</h1>
<button>@L["Login_SendOtp"]</button>
```

Registered once (`builder.Services.AddLocalization()` in `Program.cs`, 3.4), and — because the components live in the RCL (3.4) — the *same* strings serve web and the native shells. Adding French is adding `AppStrings.fr.resx` with the same keys; no component changes. (The reference platform ships EN + ES live, with FR/DE/PT scaffolded — [docs/LOCALIZATION.md](#).)

A discipline worth adopting now: **keys are structured, not sentences**. `Login_Title`, not "Sign in", as the key — so a copy change ("Sign in" → "Log in") is a resx edit, not a key rename that ripples through components. The key names the *slot*; the resx fills it.

3. The decision chain — which language, and why

"Show Spanish" requires the app to *decide* Spanish, and the platform resolves it through a deliberate chain, most-specific-wins:

1 A signed-in user's saved locale. `User.Locale` (2.1) rides in the JWT as the

`locale` claim; `MainLayout` reconciles the running culture to it after sign-in. So a user who chose Spanish sees Spanish **on every device** — the preference is server-side, following the account, not the browser. (This is the one per-*user* preference the platform allows, per Golden Rule 4's "only preferences are per-user" — localization is its canonical example.) Setting it needs a small endpoint: when a signed-in user picks a language, the switcher PUTs it to a `.../locale` endpoint (validated against the supported set — an unknown code is a 400 `unsupported_locale`, never silently accepted) that updates `User.Locale`, so the *next* token carries the new claim. That endpoint is the first member of the account surface that GDPR (7.1) later extends with export and erasure.

1 A stored host choice, for anyone not signed in (a visitor on the login page). This

is the `LanguageSwitcher`'s job, and where the host seam comes in (§4).

1 Default (English) when nothing else is known.

The important architectural point: locale is a *user* attribute that happens to have a *host* fallback — not a browser setting the server is unaware of. That's what makes "same language across devices" work, and it's why the claim exists.

4. Surviving a cold start — the web-first culture seam

Here's the subtle part. A Blazor WASM app must apply the culture **before it renders**, or the first paint flashes English then re-renders. But *where the chosen culture is stored* differs by host — and, exactly like 3.4's `ISessionStore`, that difference is the only one, so it gets a seam:

```
// src/Shared.Ui/ICulturePersistence.cs – the whole interface
/// <summary>Persists the user's chosen UI culture so the host can re-apply it on the next
/// cold start. Each host owns a store its startup path can actually read: web writes
/// localStorage["app_culture"] (read pre-render by Web/Program.cs); MAUI writes OS
/// Preferences (read in MauiProgram – WebView localStorage doesn't exist yet when the
/// native process boots). Callers still set the in-process culture themselves; this is
/// only the durable half.</summary>
public interface ICulturePersistence
{
    Task PersistAsync(string cultureCode);
}
```

The comment is the whole lesson in miniature — read *why* the seam exists rather than just *that* it does: on native (MAUI), the WebView's localStorage **doesn't exist yet** when the process boots, so the culture must live in OS Preferences, which the native startup path *can* read before the WebView spins up. Web can use localStorage because its startup (Program.cs) runs in the browser where localStorage is available. Same seam pattern as 3.4: the shared LanguageSwitcher component calls PersistAsync; each host's implementation puts the value where *its* boot path will find it. The web implementation and the bootstrap:

```
// src/Web/Program.cs – apply the saved culture BEFORE the app renders
var stored = await js.InvokeAsync<string?>("localStorage.getItem",
LocalStorageCulturePersistence.StorageKey);
try
{
    var culture = string.IsNullOrEmpty(stored) ? "en" : stored;
    CultureInfo.DefaultThreadCurrentCulture = CultureInfo.DefaultThreadCurrentUICulture =
new CultureInfo(culture);
}
catch (CultureNotFoundException) { /* corrupt stored value – keep the default culture */ }
```

Note the **fail-soft** catch: a garbage stored value (hand-edited localStorage, version skew) falls back to English rather than crashing the boot. Same instinct as 2.5's ResolveCulture for emails — an unknown culture is a shrug, never a 500. Designing this seam now, on web, is once again what makes the native shell (appendix) a matter of *one small implementation* rather than a re-architecture.

5. UI vs email localization — the same resx, two resolution models

You've now localized strings twice — emails in 2.3, UI here — and the *difference* is instructive. Both use .resx, but they resolve culture differently, for a good reason:

- ****The UI uses IStringLocalizer and the *ambient* thread culture.**** There's one user

looking at one browser; "the current culture" is unambiguous, so the localizer reads it implicitly. @L["Login_Title"] just works.

- ****Emails use ResourceManager keyed by an *explicit* culture**** (2.3's BrandedEmail).

A server rendering an email has no meaningful "ambient" culture — it's processing many users' requests, and the recipient's language (the requester's for an OTP, the *inviter's* for an invite) is a parameter, not an ambient fact. So emails pass the culture in.

Same translation files, two resolution strategies — because "what language is this?" has an obvious answer in a browser and a *chosen* answer on a server. Recognizing when ambient context is legitimate versus when it's a bug is a recurring judgment; here you see both correct answers side by side.

6. Red — proving it end to end

Localization is a UI behavior, so its test is an E2E one (the harness arrives fully in 3.6; `I18nTests` is its first real journey):

```
Scenario: switching language re-renders the UI and persists
  Given a visitor on the login page in English ("Sign in")
  When they choose Español in the language switcher
  Then the page shows "Iniciar sesión"
  And after a full reload it is STILL in Spanish
```

The reload assertion is the one that proves the *seam* works, not just the switch — exactly the shape of 3.4's "reload and stay signed in." A unit-ish test also guards the `resx` themselves: `assert AppStrings.es.resx` has a value for every key in the neutral file, so a missing translation is caught before a user sees an English word in a Spanish page. (Keys present in one culture but not another is the classic `i18n` rot; a test that diffs the key sets stops it.)

7. Run it

```
dotnet run --project src/Api & dotnet run --project src/Web
# Login page → language switcher → Español → UI flips → reload → still Español
# Sign in as a user who saved "es" on another browser → lands in Spanish (the claim)
```

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how is language chosen and stored? (a) Browser Accept-Language only; (b) a per-host setting (cookie/localStorage), server-unaware; (c) a per-user locale (server-persisted, in the JWT) with a host-stored fallback for anonymous visitors, all resolving `.resx` via `IStringLocalizer`.

Chosen: (c). Language is a user preference that should follow the account across devices (hence the claim + `User.Locale`), with a host store so a not-yet-signed-in visitor still gets their choice. `IStringLocalizer` + `resx` is the platform-native path, and putting components in the RCL means web and native share one string set.

Rejected: (a) — Accept-Language is a decent *default* but can't honor an explicit user choice and doesn't cross devices predictably. (b) — host-only settings mean a user re-picks their language on every new device and the server can't localize *emails* to them (it doesn't know their language) — which is exactly why 2.3 needed a server-side locale to begin with.

The trade: a per-user locale adds a `User` column, a claim, and a reconcile step in `MainLayout` — more moving parts than a cookie. Accepted: it's the only model where a user's language is truly *theirs*, and it's what lets emails and UI agree.

9. Checkpoint

```
git add -A && git commit -m "feat: localization - resx + IStringLocalizer, per-user locale, web-first culture seam"
git tag lesson-3.5
```

You should now be able to: externalize a string and explain why the key is structured, not a sentence; describe the culture decision chain and why locale is per-user with a host fallback; explain the `ICulturePersistence` seam and the native cold-start reason behind it; and contrast ambient-culture (UI) with explicit-culture (email) resolution.

Next: *Lesson 3.6 — The E2E harness* — Playwright, the Page Object Model, and reading real OTPs out of Mailpit to drive sign-in through a real browser.

Lesson 3.6 — The E2E harness

Part 3 · The slice pattern & the UI. The third leg of the test pyramid. You have unit tests (Core logic) and integration tests (the real app + real Postgres, via `WebApplicationFactory`). Now the top: **end-to-end tests that drive a real browser against the running stack**, proving the journeys a *user* actually takes. The star technique closes a loop three lessons in the making — the E2E suite reads the OTP code out of Mailpit, exactly as you did by hand in 2.4, so a full passwordless sign-in runs browser-to-inbox-to-browser with no human. And we settle, explicitly, the question of *which* tests belong here versus in the faster layers below.

Goal: a Playwright/NUit harness that signs a user in through the real UI + API + database, using Mailpit as an assertable inbox; a Page Object Model that keeps journeys readable; and a clear rule for what earns an E2E test.

Concepts & principles: the test pyramid and its economics; E2E with Playwright; Page Object Model; stable `data-testid` selectors; the mail-trap-as-fixture trick; choosing the right test layer (the E2E-vs-integration division of labor).

Maps to: ADR-C14 · `docs/WAYS_OF_WORKING.md` (testing strategy) · repo: `tests/E2E.Tests/E2ETestBase.cs`, `Mailpit.cs`, `Pages/BasePage.cs/LoginPage.cs`, `AuthFlowTests.cs`, `playwright.runsettings`.

Prerequisites: 3.4 (the running UI), 2.4 (OTP sign-in), 0.2 (Mailpit in the stack).

1. Motivate — the layer that catches what the others can't

Your integration tests boot the real app and prove `/api/notes` is tenant-isolated. They do *not* prove that a human can click "Sign in," receive a code, type it, and land on their notes — because they never render a pixel or run a line of the Blazor client. The wiring *between* the browser and the API — routing, the silent-refresh choreography (3.4), a `data-testid` that got renamed, a CORS misstep — lives in a gap no server-side test sees. E2E tests live in that gap: they drive a real browser through the real UI against the real API and database, asserting the *journey*, not the unit.

They are also, honestly, the most expensive tests you own — slow (a browser boots, pages load), heavier to run (the whole stack must be up), and the most prone to flake (timing, async rendering). That cost is exactly why the *last* section of this lesson is about using them sparingly. First, build one.

2. The harness — a real browser against the real stack

Playwright drives Chromium/Firefox/WebKit from .NET; the NUnit integration gives each test a fresh browser context. The base class, from the real `E2ETestBase`:

```
public abstract class E2ETestBase : PageTest
{
    protected static string BaseUrl =>
        TestContext.Parameters.Get("PLAYWRIGHT_BASE_URL")
        ?? Environment.GetEnvironmentVariable("PLAYWRIGHT_BASE_URL")
        ?? "https://localhost:7008"; // the Web app's dev https profile

    // Dev runs on a self-signed cert, so ignore HTTPS errors for the test browser.
    public override BrowserNewContextOptions ContextOptions() => new()
    {
        BaseURL = BaseUrl,
        IgnoreHTTPSErrors = true,
    };

    protected static string UniqueEmail(string role) =>
        $"e2e-{role}-{Guid.NewGuid():N}@example.com";
}
```

Two decisions already: the base URL is *overridable* (a runsettings parameter, then an env var, then the dev default) so the same suite runs against localhost and against CI's single-origin deployment (8.3) unchanged; and every test mints a **unique email**, so tests don't collide on account state — the E2E equivalent of 1.3's per-test database reset, achieved by never sharing an identity. Unlike the integration harness, E2E tests run against a stack *you* start (docker compose up, then the API, then the Web app) — they're testing the deployed reality, not an in-process factory.

3. Mailpit as an assertable inbox — the loop closes

This is the technique the whole passwordless design was quietly building toward. In 2.4 you signed in by reading an OTP out of Mailpit's web UI *by hand*. The E2E suite does the same thing *programmatically*, via Mailpit's REST API — from the real Mailpit helper:

```
/// Minimal client for the dev Mailpit REST API so E2E tests can read the OTP code / magic
/// link the app "sends" — the same trick a manual tester uses.
public static async Task<string> WaitForOtpAsync(string toEmail, TimeSpan timeout)
{
    var deadline = DateTime.UtcNow + timeout;
    while (DateTime.UtcNow < deadline)
    {
        var code = await TryGetLatestCodeAsync(toEmail); // GET /api/v1/messages, find,
        regex the 6 digits
        if (code is not null) return code;
        await Task.Delay(500);
    }
    throw new TimeoutException($"No OTP email for {toEmail} within
{timeout.TotalSeconds:0}s.");
}
```

Look back at 0.2, where we chose Mailpit specifically because it "exposes an HTTP API your tests will query." *This is that*. The email seam (2.3) sends to Mailpit; the E2E test polls Mailpit's API for the message; a regex lifts the code; the browser types it. A real sign-in, no human, no real email provider — the mail trap is a **test fixture with an API**. One production-grade detail in the real code worth noting: it filters on the *subject* ("verification code") before regexing six digits, because an invitation email (delivered late by the outbox, Part 4) to the same address could otherwise satisfy a bare `\d{6}` — a real flake, prevented by matching the message you

actually want.

4. The Page Object Model — journeys that read like intent

Raw Playwright calls (`page.Locator("#email").FillAsync(...)`) scattered through tests are unreadable and brittle — every selector duplicated, every UI tweak a shotgun edit. The **Page Object Model** wraps each screen in a class exposing intent-named locators and actions. From the real `LoginPage`:

```
public class LoginPage(IPage page) : BasePage(page)
{
    public override string Path => "/login";

    public ILocator Email => Page.GetByTestId("login-email");
    public ILocator SendOtp => Page.GetByTestId("login-send-otp");
    public ILocator OtpCode => Page.GetByTestId("login-otp-code");
    public ILocator VerifyOtp => Page.GetByTestId("login-verify-otp");

    /// Requests an OTP, reads it from Mailpit, enters it, and submits.
    public async Task SignInWithOtpAsync(string email)
    {
        await Email.FillAsync(email);
        await SendOtp.ClickAsync();
        var code = await Mailpit.WaitForOtpAsync(email, TimeSpan.FromSeconds(20));
        await OtpCode.FillAsync(code);
        await VerifyOtp.ClickAsync();
    }
}
```

Two disciplines make this durable:

- ****Selectors are data-testid, never CSS classes or visible text.****
`GetByTestId("login-email")`

binds to a stable hook the component author placed *for testing*, not to a class name a restyle will change or to the label text (which 3.5 just made *translatable* — a test keyed on "Sign in" would break the moment the UI switches to Spanish!). This is the concrete reason data-testids matter here more than in most apps: your visible text is localized, so it is *by design* not a stable selector.

- **Tests speak in page methods, not clicks.** `E2ETestBase.SignInAsync` composes

`login.GotoAsync() + SignInWithOtpAsync` + a wait for the app shell (`GetByTestId("sign-out")`). A test reads "sign in as this user, then..." — the *what*, with the *how* encapsulated. When the login flow changes, one page object changes, not twenty tests.

5. Red — the first journey

`AuthFlowTests` is the first E2E file — one file per epic, mirroring `docs/stories/`:


```

Scenario: a new user signs in with a one-time code (happy path)
  Given a visitor on the login page
  When they request an OTP for a fresh email and enter the code from their inbox
  Then they land on the app shell, signed in

Scenario: an expired or wrong code is rejected
  When they enter a wrong code
  Then they see the login error and remain signed out

```

That happy path exercises *everything* at once: the Blazor client, the OTP endpoints (2.4), the email seam (2.3) → Mailpit, session issuance (2.2), the silent-refresh shell (3.4). One test, the entire vertical proven to work *together* — which no lower-layer test can claim. This is what E2E is *for*.

6. The economics — what earns an E2E test (the division of labor)

Now the rule this lesson exists to establish, because it governs every part after. E2E tests are slow, heavy, and flake-prone, so they are **rationed**: reserve them for critical user *journeys*, and push everything else down to the faster, sturdier layers. `docs/WAYS_OF_WORKING.md` draws the line, and it's worth internalizing:

Belongs in E2E (Playwright)	Belongs lower (xUnit + Testcontainers)
The happy path of a critical journey (sign-in, invite-and-join, upgrade, export)	Domain logic and derived rules (Core.Tests)
A key <i>user-visible</i> unhappy path (wrong OTP shows an error)	Tenant-isolation invariants, permission/entitlement boundaries
Cross-layer wiring only a browser exercises (silent refresh, i18n reload)	Concurrency, redelivery, quota atomicity — anything needing precise control

The principle: **an xUnit test that can prove it will always beat an E2E test that can**. Tenant isolation *could* be shown through a browser — but you proved it far faster and more reliably as an integration test in 2.6, seeding two tenants directly and asserting at the wire. Reserve the browser for what *only* the browser can see. Concretely, this is why the reference platform's billing quota logic lives in `QuotaServiceTests` (xUnit, precise concurrency control) while only the *402-becomes-a-visible-upgrade-prompt* moment is an E2E journey (5.3). Same feature, two layers, each doing what it's best at.

A corollary you'll feel in later parts: when an epic was built integration-first and lacks an E2E journey, that's a deliberate ration, not an oversight — the invariants are proven below, and a browser test would add cost without adding coverage. Write the E2E for the *journey*; trust the lower layers for the *logic*.

7. Run it

First-time Playwright setup installs the browser binaries; then the suite needs the full stack up (per `tests/E2E.Tests/README.md`):

```
dotnet build tests/E2E.Tests
pwsh tests/E2E.Tests/bin/Debug/net10.0/playwright.ps1 install # once
docker compose up -d
dotnet run --project src/Api & dotnet run --project src/Web &
dotnet test tests/E2E.Tests
# AuthFlowTests: a real Chromium signs in via an OTP read from Mailpit — green
```

One gotcha the reference `.env` note calls out: OTP E2E tests need email to land in *Mailpit*, so the API must use the dev SMTP default — if your local `.env` points SMTP at a real provider (Brevo), the tests can't read the code. Override back to Mailpit for E2E runs.

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how are user journeys tested? (a) Manual QA only; (b) component tests (bUnit) rendering Blazor components in isolation; (c) full E2E with Playwright against the running stack, Page Object Model, Mailpit as the inbox — reserved for critical journeys.

Chosen: (c), *rationed* (ADR-C14). Only a real browser against the real stack proves the cross-layer journeys that break in production; the POM + `data-testid` keep the suite readable and resilient to restyles and localization; Mailpit makes passwordless flows testable without a mail provider. The rationing (§6) keeps the slow layer small.

Rejected: (a) — manual QA doesn't scale or gate CI (though a *manual QA plan* still exists for release readiness — 8.3; automation and manual QA are complements, not substitutes). (b) — bUnit tests a component's render logic in isolation, which is useful but proves nothing about the component talking to the real API through the real client pipeline; it's a different (lower) layer, not a substitute for the journey.

The trade: E2E is the slowest, flakiest layer and needs the whole stack running — so it's excluded from the fast `build-test` CI job (1.6) and gated separately, and it's rationed hard. Accepted deliberately: a handful of journey tests catch the integration failures nothing else can, and the discipline of §6 keeps the cost bounded.

9. Checkpoint

```
git add -A && git commit -m "test(e2e): Playwright harness — POM, Mailpit-as-inbox, first
sign-in journey"
git tag lesson-3.6
```

You should now be able to: place a test on the right pyramid layer and defend the choice; build a Page Object with `data-testid` selectors and say why localized text can't be a selector; explain the Mailpit-as-fixture trick and trace the passwordless journey through every layer it touches; and state the rule for what earns an E2E test.

Next: *Lesson 3.7 — Build your own slice (capstone)* — no reference code this time. You apply the seven-step checklist and the full Red→Green→refactor loop, unaided, to a feature that's yours.

Lesson 3.7 — Build your own slice (capstone)

Part 3 · The slice pattern & the UI — capstone. Every other lesson handed you the code. This one doesn't. You've read the chassis, reproduced the reference slice, and built the UI and test harness around it — now you build a *new* feature end to end, unaided, and discover that the platform makes it almost boring. That "almost boring" is the entire point of Parts 1–3: if adding a feature is calm, the architecture is doing its job. This is where the knowledge becomes a skill.

Goal: ship one complete, tested, tenant-isolated, localized feature slice — API to UI to E2E — adding to platform code only a `DbSet` and registrations (never changing platform behavior), and have every architecture test pass *because you followed the pattern*, not because you studied the tests.

Concepts & principles: synthesis. The seven-step checklist (3.2), TDD (1.2/3.3), the repository seam (3.1), tenancy (2.6/2.7), the contributor seam (2.9), RCL UI (3.4), localization (3.5), an E2E journey (3.6), and the arch tests as your safety net (3.3).

Maps to: everything in Parts 1–3 · `docs/WAYS_OF_WORKING.md` (slice + story format).

Prerequisites: all of Parts 1–3. Do not read ahead to Part 4 first — the point is to prove *this* foundation is solid in your hands.

1. The brief

Build a **Bookmarks** feature: a tenant can save, list, and delete links.

- A bookmark has a `Title`, a `Url`, and a `CreatedAt`.
- Anyone in the tenant can list the tenant's bookmarks (newest first), add one, delete one.
- Expose a **derived** value the API computes and never stores: the link's `Host` (e.g.

`github.com` from `https://github.com/x/y`). This exercises Golden Rule 4 — derived, not stored — and gives you a genuine unit of domain logic to TDD.

- Validation: a blank title or a non-absolute URL is a `400`, using the shared error

envelope (1.5).

If Bookmarks bores you, pick your own single-entity feature — Reminders, Tags, Saved Searches — the shape is identical. Keep it to **one entity**: the goal is to run the loop cleanly, not to design a subsystem. (Anything needing *new platform machinery* — files, background work, a second entity with relationships — is out of scope; those capabilities arrive in Parts 4+. A capstone uses the chassis you have.)

2. Story first (the habit from `WAYS_OF_WORKING.md`)

Before code, write `docs/stories/bookmarks.md` — intent + Gherkin acceptance criteria. This is the spec your tests will mirror:

```

### BOOK-1 – save and manage bookmarks

**As a** tenant member **I want** to save links **So that** my household keeps shared
references.

Scenario: save and list a bookmark
  Given I am signed in
  When I save a bookmark titled "Docs" for "https://learn.microsoft.com/dotnet"
  Then it appears at the top of my bookmark list
  And its displayed host is "learn.microsoft.com"

Scenario: bookmarks are tenant-isolated
  Given tenant A and tenant B each saved a bookmark
  When A lists bookmarks
  Then A sees only A's bookmark
  And A deleting B's bookmark id returns 404

Scenario: a bookmark needs a title and an absolute URL
  When I save a bookmark with a blank title (or "not-a-url")
  Then the response is 400 invalid_request

```

Writing the unhappy paths *now* — isolation, validation — is what stops you from "forgetting" to test them once the happy path works.

3. The seven steps (your worksheet)

Work the checklist from 3.2. For each step, the lesson that taught it is in parentheses — revisit it if you stall, but try from memory first:

1 Entity — `src/Core/Entities/Bookmark.cs, : ITenantScoped, TenantId +`

`Title/Url/CreatedAt`. Compute `Host` as a property from `Url` (or in the response DTO) — never a stored column. (2.6, 3.2, Golden Rule 4)

1 DbSet + config — `DbSet<Bookmark> ON ApplicationDbContext; a BookmarkConfiguration`

with a sensible `Title/Url` max length. (1.3)

1 Migration — `dotnet ef migrations add AddBookmarks`. (1.3)

2 Slice folder — `src/Api/Features/Bookmarks/: BookmarksEndpoints` (via

`MapTenantFeatureGroup` — *not* `raw MapGroup`), `BookmarksHandler` (inject `IRepository<Bookmark>, ICurrentTenant, TimeProvider; Query()` for reads), `BookmarksModels` (co-located DTOs; the response carries `Host`). Register the handler and `app.MapBookmarks()` in `Program.cs`. (3.1, 3.2)

1 Contributor — `BookmarksDataContributor : ITenantDataContributor` with all four

members; `ExportKey => "bookmarks"`; use `QueryAllTenants()` here (and *only* here). Register it. (2.9, 3.2)

1 Fixture reset — nothing to do; the reset list is model-derived (1.3). Confirm your entity is truncated by running the suite.

1 UI + i18n — a `Bookmarks.razor` page in **Shared.Ui** (not the web app), nav entry, strings via `IStringLocalizer<AppStrings>` with EN + ES resx keys. (3.4, 3.5)

4. Test-drive it — Red before Green, every step

- **Unit (Core.Tests):** the Host derivation. Feed it `https://a.b.com/x?y=1` → `a.b.com`

(that's `Uri.Host` — the *whole* host, matching \$1's `github.com` definition; extracting a registrable domain like `b.com` needs a public-suffix list and is out of scope). Feed it junk → your defined behavior (empty? the raw string?). This is pure logic — no database — so it's a fast unit test, and writing it first *forces you to define* the edge cases. (1.2, 3.3)

- **Integration (Api.Tests):** the slice against real Postgres — the three Gherkin

scenarios. The isolation one is non-negotiable: seed two tenants, act as A, assert A sees one bookmark and gets 404 (not 403) deleting B's. Use the harness patterns from 2.6/3.2. (1.3, 2.6, 3.2)

- **E2E (E2E.Tests):** *one* journey — sign in, add a bookmark, see it listed with its

host. Reuse `E2ETestBase.SignInAsync` and the POM style; add `data-testid` hooks to your page. Don't E2E the validation or isolation — those are proven faster above (§6 of 3.6 is the rule you're now applying for real). (3.6)

5. Your safety net — run the architecture tests

Here's the reward for everything Part 1–3 taught. You didn't study the arch tests to game them; you followed the pattern, so they should pass. Run them:

```
dotnet test tests/Api.Tests --filter Architecture
```

What they're checking about *your* slice, and what a failure would be telling you:

- `EveryTenantScopedEntity_HasAGlobalQueryFilter / EveryEntityWithATenantId_IsScopedOrAllowlisted` —

green means your `Bookmark` correctly wears `ITenantScoped`. Red means you added a `TenantId` without the marker — an isolation hole. (2.6)

- `FeatureSlices_DoNotBypassTheTenantFilter` — green means your handler used `Query()` and

only your `*DataContributor.cs` used `QueryAllTenants()`. Red means the request path reached for the escape hatch. (3.1, 3.2)

- `FeatureFiles_RegisterRoutesViaMapTenantFeatureGroup` — green means you didn't hand-roll

a `MapGroup`. Red means your endpoints could have forgotten auth. (3.2)

- `WebApp_HasNoInlineBlazorComponents` — green means your page is in the RCL. (3.4)
- `TenantDataContributors_HaveUniqueExportKeys / WipeDataTests` — green means your

contributor is registered with a unique key and your table tears down on dissolve. (2.9)

- `TenantScopedEntities_MapToDistinctTables / RouteGroupPrefixes_AreUnique` —

green means bookmarks didn't collide with an existing table or route. (3.3)

A red arch test here is not an obstacle — it's the platform catching a real mistake at build time, with the offending file named, exactly as designed. That's the whole thesis of the course made personal: the guardrails you built now guard *you*.

6. Definition of done

Check every box honestly (it's the PR template from 1.6):

- ☐ Story written first, with unhappy paths, before any code.
- ☐ Tests written before production code (Red→Green), at the right layers (\$4).
- ☐ `dotnet build` — 0 warnings (1.1's bar).
- ☐ `dotnet test` — all green, including the full architecture suite (\$5).
- ☐ Tenant isolation proven (the two-tenant integration test; 404, not 403).
- ☐ The derived `Host` is computed, never stored (Golden Rule 4).
- ☐ UI in `Shared.Ui`, EN + ES strings present, no hard-coded English.
- ☐ One E2E journey; validation/isolation left to the faster layers.
- ☐ Touched platform code only by addition — the `DbSet` on `AppDbContext` and the

`Program.cs` registrations — and changed no platform *behavior*.

- ☐ Any real decision you made is recorded as a short ADR (0.1).

If all ten hold, you didn't just copy *Notes* — you *internalized* the platform. That last box matters: you made your own small decisions here (how to handle a malformed URL, what the empty-list UI says). Record them. The habit is the point.

7. Checkpoint

```
git add -A && git commit -m "feat(bookmarks): my first slice – end to end on the chassis"
git tag lesson-3.7
```

You should now be able to: build a complete feature slice from an empty folder without referring to *Notes*; sequence the work test-first across the three layers; and read a failing architecture test as feedback rather than an obstacle. If this felt calm — that calm is what a good architecture *feels like*, and it's what the rest of the course protects.

Part 3 complete. You have the assembly manual and the proof you can use it. From here the course returns to the chassis itself — the platform capabilities that make this a *product*, not just an app.

Next: *Part 4 — Reliability & operations.* The transactional outbox, idempotency, the scheduler, observability, and the append-only audit log — how the platform behaves when things fail, retry, and must be explained after the fact.

Part 4 — Reliability & operations

Lesson 4.1 — The transactional outbox

Part 4 · Reliability & operations. The chassis can now authenticate, isolate tenants, and grow features. Part 4 makes it *reliable* — it teaches the platform how to behave when things fail, retry, and must be explained afterward. We open with the problem behind almost every "the charge went through but no email arrived" bug: the **dual-write problem**. You cannot atomically commit a database row *and* perform an external side effect — one of them will eventually be lost. The transactional outbox is the standard, load-bearing answer, and it's the foundation the next five lessons (email, webhooks, notifications, billing) all stand on.

Goal: a side effect enqueued *in the same transaction* as the business change that caused it, delivered later by a background dispatcher that claims work safely, retries with backoff, and dead-letters what it can't deliver.

Concepts & principles: the dual-write problem; the transactional outbox pattern; at-least-once delivery + mandatory idempotency; FOR UPDATE SKIP LOCKED work claiming; exponential backoff + dead-lettering; a testable core split from a background service.

Maps to: ADR-007 · repo: `src/Core/Abstractions/IOutbox.cs`, `IOutboxHandler.cs`, `src/Core/Entities/OutboxMessage.cs`, `src/Infrastructure/Outbox/EfOutbox.cs`, `OutboxProcessor.cs`, `OutboxDispatcher.cs`, `OutboxOptions.cs`, tests: `tests/Api.Tests/Outbox/OutboxProcessorTests.cs`.

Prerequisites: 1.3 (real Postgres — SKIP LOCKED is why the in-memory provider was banned), 3.1 (the unit of work whose transaction the outbox rides).

1. Motivate — the write you can't make atomic

A handler does the most natural thing in the world:

```
await repository.SaveChangesAsync(ct); // 1. persist the order
await emailSender.SendAsync(...); // 2. email the receipt
```

Now enumerate the failure modes. The process crashes between line 1 and line 2 → order saved, no receipt, *forever* (nothing will ever retry it). Line 2 throws → do you roll back the order that already committed? You can't; it's committed. You reorder them — email first, then save — and now a crash leaves a receipt for an order that doesn't exist. There is **no ordering that makes a database commit and an external call atomic**, because they're two different systems with two different notions of "done." This is the dual-write problem, and it's not a bug you can carefully code around — it's a property of distributed systems.

The insight that dissolves it: you *can* make two writes to the **same** database atomic — that's what a transaction is. So don't try to perform the side effect inline. Instead, in the same transaction as the business change, **write a durable record that the side effect is owed**. Commit both together or neither. Then a separate process reads those records and performs the effects. The atomic unit becomes "order + intent-to-email," which a single transaction *can* guarantee; the actual email happens later, driven by a record that cannot be lost.

2. The message — a durable IOU

```
// src/Core/Entities/OutboxMessage.cs (the real entity)
/// A durable record of a side effect to perform out-of-band (email, a webhook call, ...).
/// It is written in the SAME transaction as the business change that produced it, so the
/// effect is atomic with the data – no "saved the row but lost the email" (ADR-007).
public class OutboxMessage
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public required string Type { get; set; } // handler discriminator, e.g. "email"
    public required string Payload { get; set; } // opaque JSON for the matching handler
    public Guid? TenantId { get; set; } // context for the handler, NOT a scoping key
    public string Status { get; set; } = OutboxStatus.Pending; // pending | sent | dead
    public int AttemptCount { get; set; }
    public DateTimeOffset CreatedAt { get; set; }
    public DateTimeOffset NextAttemptAt { get; set; } // earliest dispatch; advanced on
    retry (backoff)
    public DateTimeOffset? ProcessedAt { get; set; }
    public string? LastError { get; set; } // last failure detail; never a secret
}
```

Two decisions worth pausing on, both callbacks to earlier lessons:

- **It is deliberately *not* `ITenantScoped`.** Read the entity's own comment: the outbox

is *platform infrastructure* and may carry system (tenant-less) effects, so it sits **outside** the global tenant filter (2.6). And notice how it stays clear of the R2 invariant test (2.6/3.3) *without* an allowlist entry: that test flags an entity only when its `TenantId` is a **non-nullable `Guid`** (a real scoping key someone forgot to scope). Here `TenantId` is `Guid?` — *nullable context passed to the handler*, not a scoping key — so the test doesn't consider it a tenant-scoped table at all, and no allowlist entry is needed. (The allowlist is reserved for the genuinely awkward cases, like `TenantMembership`, that *do* carry a non-nullable `TenantId` yet mustn't be filtered.)

Choosing `Guid?` here is thus both semantically right and what keeps the guardrail quiet — recognizing that a platform table is legitimately cross-tenant is the same judgment you exercised in 2.6.

- **`LastError` says "never a secret."** The same 1.5 discipline: failure detail is for

operators, and a payload or exception can carry sensitive data — so what lands in a stored, queryable column is bounded and scrubbed.

The `Type/Payload` shape is deliberately generic: the outbox doesn't know what an "email" *is*. It's a typed queue; the meaning lives in handlers (§4). That genericity is why the *same* outbox serves email (4.2), webhooks (7.4), and notifications (7.2).

3. Enqueue — the effect rides the caller's commit

```
// src/Infrastructure/Outbox/EfOutbox.cs (the whole enqueue side)
public async Task EnqueueAsync(string type, string payload, Guid? tenantId = null,
Cancellation token ct = default)
{
    var now = clock.UtcNow();
    await db.Set<OutboxMessage>().AddAsync(new OutboxMessage
    {
        Type = type, Payload = payload, TenantId = tenantId,
        Status = OutboxStatus.Pending, CreatedAt = now, NextAttemptAt = now,
    }, ct);
    // Intentionally no SaveChanges – persistence rides the caller's commit.
}
```

The most important line is the comment: ****no SaveChanges.**** `EnqueueAsync` *stages* the message on the shared request-scoped `AppDbContext` (3.1) and stops. It becomes durable when the caller's transaction commits — the *same* transaction as the business change. So a handler that does `repository.Add(order); outbox.EnqueueAsync("email", ...); uow.Commit()` gets order-and-intent atomically, for free, because they share one unit of work (3.1's commit model, now paying off precisely as designed). The `IOutbox` interface documents the flip side too: a caller with *no* ambient transaction (the email decorator in 4.2) must persist explicitly — the outbox doesn't presume to own the commit.

4. Handlers — typed, and necessarily idempotent

Delivery is somebody's job, split by message type:

```
// src/Core/Abstractions/IOutboxHandler.cs
/// MUST be idempotent: delivery is at-least-once, so the same message may be handed to
/// HandleAsync more than once (e.g. after a crash between send and commit). Throw to
/// signal failure – the dispatcher retries with backoff and dead-letters after the cap.
public interface IOutboxHandler
{
    string Type { get; } // the type this handler claims
    Task HandleAsync(OutboxMessage message, Cancellation token ct = default);
}
```

Sit with the phrase **at-least-once**. The outbox guarantees an effect happens *at least* once — never zero times (that's the whole point) — but it *cannot* guarantee exactly once, because the handler could succeed at sending the email and then crash *before* recording that success, so the message stays pending and is retried. Therefore **every handler must be idempotent**: handling the same message twice must equal handling it once. This isn't a nicety; it's a contract the pattern forces on you. (It's also why 4.3 builds the *inbox* — the mirror-image dedup for effects you can't make naturally idempotent, like "charge this card.") One handler claims each `Type`; the dispatcher routes by it.

5. The processor — claiming work without stepping on yourself

Here's the crux, and it's a genuinely beautiful use of Postgres. Multiple dispatcher instances (or just concurrent polls) must not grab the *same* message and deliver it twice-over. The naive `SELECT ... WHERE pending LIMIT 1` then `UPDATE` has a race: two pollers read the same row before either claims it. The real `OutboxProcessor`:

```
// src/Infrastructure/Outbox/OutboxProcessor.cs (core)
public async Task<bool> ProcessNextAsync(CancellationTokentoken ct = default)
{
    if (!db.Database.IsRelational())
        throw new InvalidOperationException("The outbox requires a relational provider (FOR UPDATE SKIP LOCKED).");

    var now = clock.UtcNow();
    await using var tx = await db.Database.BeginTransactionAsync(ct);

    // Claim the oldest due+pending row, skipping any another poller already holds.
    var message = (await db.Set<OutboxMessage>()
        .FromSql($"
            SELECT * FROM "OutboxMessages"
            WHERE "Status" = {OutboxStatus.Pending} AND "NextAttemptAt" <= {now}
            ORDER BY "CreatedAt"
            LIMIT 1
            FOR UPDATE SKIP LOCKED
            """)
        .ToListAsync(ct)).FirstOrDefault();

    if (message is null) { await tx.RollbackAsync(ct); return false; }

    try
    {
        if (!_handlers.TryGetValue(message.Type, out var handler))
            throw new InvalidOperationException($"No IOutboxHandler registered for outbox
type '{message.Type}'.");
        await handler.HandleAsync(message, ct);
        message.Status = OutboxStatus.Sent;
        message.ProcessedAt = now;
        message.LastError = null;
    }
    catch (Exception ex)
    {
        message.AttemptCount++;
        message.LastError = Truncate(ex.Message, 1000);
        if (message.AttemptCount >= options.MaxAttempts)
            message.Status = OutboxStatus.DeadLettered; // give up, keep the record
        else
            message.NextAttemptAt = now + Backoff(message.AttemptCount); // exponential
        backoff
    }

    await db.SaveChangesAsync(ct);
    await tx.CommitAsync(ct);
    return true;
}
```

Everything here is load-bearing:

- ****FOR UPDATE SKIP LOCKED**** is the whole concurrency story in three words. FOR UPDATE

row-locks the claimed message inside the transaction; SKIP LOCKED tells Postgres "if a row is already locked by another poller, *skip it*, don't wait." So N pollers each grab a *different* due row and never block on each other — safe horizontal scaling with zero coordination code. This is a Postgres feature the EF in-memory provider cannot emulate, which is precisely why 1.3 banned it and why this method throws on a non-relational provider rather than pretending.

- ****The FromSql + ToListAsync().FirstOrDefault() shuffle**** has a real reason (the

comment explains it): the raw SQL already `LIMIT 1s`, but if you let EF apply `.FirstOrDefault()` server-side it layers its own order-less `First` operator on top and logs a noisy false-positive warning every poll. Materializing the (at most one) row, then taking `First` in memory, keeps the log clean. A small scar, labeled.

- **Backoff and dead-lettering** encode a failure policy: a transient failure (SMTP

hiccup) retries on an exponential schedule (`NextAttemptAt` pushed further out each time); after `MaxAttempts` the message is dead-lettered — *kept*, not deleted, so an operator can inspect `LastError` and decide. Failure is a state, not a disappearance.

- **The whole thing is one transaction.** Claim, handle, record-outcome, commit — so

either the message is durably marked `sent/retrying` or nothing happened and it stays claimable. No outcome is ever lost, which is the property the whole pattern exists for.

Note the deliberate split: `OutboxProcessor` is the *testable core* (no timing, no threads), separate from the `OutboxDispatcher BackgroundService` that merely loops it:

```
// src/Infrastructure/Outbox/OutboxDispatcher.cs (the loop)
using var scope = scopeFactory.CreateScope(); // fresh scope per pass (scoped DbContext)
var processor = scope.ServiceProvider.GetRequiredService<OutboxProcessor>();
var processed = await processor.ProcessDueAsync(stoppingToken);
if (processed == 0) await Task.Delay(options.PollInterval, stoppingToken); // idle only
when empty
```

A build detail you'll hit the moment you write this: `BackgroundService` and `IServiceScopeFactory` live in the ASP.NET Core shared framework, not the base runtime — so `Infrastructure` (a plain class library) can't see them yet. Add a framework reference to `Perezosoftware.Infrastructure.csproj`, exactly as the reference project does:

```
<ItemGroup>
  <FrameworkReference Include="Microsoft.AspNetCore.App" />
</ItemGroup>
```

That's a deliberate, minimal opening — a `FrameworkReference` pulls in the shared framework's *hosting* abstractions without turning the library into a web app. (And per 1.1's locked-restore discipline: adding it changes the transitive graph, so regenerate `packages.lock.json` — `dotnet restore --use-lock-file` — and commit it, or CI's `--locked-mode` will fail.)

Its comment states the operational stance: one in-process poller is the baseline, the `SKIP LOCKED` claim keeps it correct *if* ever run multi-instance, and a real broker/scheduler is a documented swap-in — the seam is the `IOutbox/IOutboxHandler` pair, so replacing the *transport* never touches a caller. (A wrinkle the free-tier deployment must respect — an instance that sleeps pauses the poller — is a recorded trade-off you'll meet in Part 8, not a surprise.)

6. Red — testing the core without the clock or the thread

Because the processor is split out, its tests are fast and deterministic — real Postgres (1.3) for `SKIP LOCKED`, `FakeTimeProvider` (3.3) for backoff:

```

Scenario: a due message is delivered and marked sent
  Given a pending outbox message due now
  When the processor runs one pass
  Then the matching handler received it and the row is "sent"

Scenario: a failing handler retries with backoff, then dead-letters
  Given a handler that always throws and MaxAttempts = 3
  When the processor runs, advancing the clock past each NextAttemptAt
  Then attempts increment, NextAttemptAt pushes out exponentially, and after the 3rd the
  row is "dead"

Scenario: two concurrent claims never take the same row
  Given two pending messages and two processors racing
  Then each processor claims a different row (SKIP LOCKED), none delivered twice

```

The third scenario is the one that *needs* real Postgres and would silently pass-but-lie on any faked database — the concrete payoff of 1.3's decision, cashed two parts later.

OutboxProcessorTests drives all three; the backoff test is pure FakeTimeProvider.Advance, no waiting.

****A FakeTimeProvider trap worth one sentence of warning.**** A fresh FakeTimeProvider starts at the **year 2000**, not "now." So if you enqueue a message with the *real* system clock (NextAttemptAt = today) but run the processor with a fake clock at 2000, the message looks *far-future-dated* and is never "due" — the test hangs on nothing. The fix is the discipline 3.3 was really about: use the **same injected clock** to seed the message and to run the processor. Seed NextAttemptAt at the fake clock's instant, then Advance from there. Mixing a real clock into a fake-clock test is the single most common way these deterministic-time tests go confusingly wrong.

7. Run it

```

dotnet test tests/Api.Tests --filter Outbox
# then live: enqueue in a handler, watch the OutboxMessages table go pending → sent
# on the next dispatcher pass (and a forced failure climb attempts → dead)

```

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how are side effects performed reliably? (a) Inline, in the request (Save then Send); (b) fire-and-forget on a background Task/queue in memory; (c) a transactional outbox — persist intent in the business transaction, deliver asynchronously with at-least-once + idempotent handlers.

Chosen: (c) (ADR-007). It's the only option that makes the effect *atomic* with the data (dual-write solved) and *durable* across crashes; `SKIP LOCKED` gives safe concurrency on the database we already run, with no broker to operate; the processor/dispatcher split keeps the logic unit-testable. The genericity (`Type + Payload`) means every later async feature reuses it.

Rejected: (a) — the dual-write bug, unfixable by ordering; a crash loses effects silently. (b) — an in-memory queue evaporates on restart (same data-loss, just faster) and still isn't atomic with the commit. A dedicated broker (Rabbit/Kafka/SQS) is a fine *scale-up* — and the seam permits it — but adding an operational dependency to a Postgres-only, run-cheap platform is cost the baseline doesn't need to pay yet.

The trade: effects are now *eventually* consistent (a poll interval of latency) and every handler carries the idempotency burden (§4) — plus a background poller to run and, on a sleeping free-tier instance, to notice is paused (Part 8). Accepted: eventual-and- guaranteed beats immediate-and-lossy for side effects, every time.

9. Checkpoint

```
git add -A && git commit -m "feat: transactional outbox – atomic side effects, SKIP LOCKED
claiming, backoff + dead-letter"
git tag lesson-4.1
```

You should now be able to: state the dual-write problem and why no ordering fixes it; explain how staging-in-the-caller's-transaction makes an effect atomic with its cause; say what `FOR UPDATE SKIP LOCKED` buys and why it needs real Postgres; and justify why at-least-once delivery makes handler idempotency non-optional.

Next: Lesson 4.2 — *Email II: the outbox decorator* — the `ISender` seam from 2.3 gets wrapped so that sending an email becomes crash-safe, without any caller noticing.

Lesson 4.2 — Email II: the outbox decorator

Part 4 · Reliability & operations. Back in 2.3 I made a promise: the five-line `ISender` seam would one day gain crash-safety "without touching a single call site." This is the lesson that keeps it. We wrap the email seam in a **decorator** that enqueues onto the outbox (4.1) instead of sending inline — and passwordless sign-in (2.4), invitations (2.9), and every future caller become reliable without changing one line. It's the shortest lesson in Part 4 and the clearest demonstration in the whole course of why seams are worth the discipline.

Goal: `ISender.SendAsync` enqueues the email as an outbox message and returns immediately; the real SMTP send happens later on the dispatcher; no caller changes, and a request no longer blocks on — or fails because of — the mail server.

Concepts & principles: the decorator pattern; transparent behavior-swap via a shared interface; keyed DI registration to avoid a self-wrapping loop; the idempotency trade-off made concrete.

Maps to: ADR-007 · repo: `src/Infrastructure/Email/OutboxEmailSender.cs`, `EmailOutboxHandler.cs`, the DI registration in `ServiceCollectionExtensions`, tests: `tests/Api.Tests/Outbox/OutboxEmailTests.cs`.

Prerequisites: 2.3 (the `ISender` seam + `SmtplibSender`), 4.1 (the outbox).

1. Motivate — the seam finally earns its keep

Recall what sending email looked like after 2.3: every caller injected `ISender` and called `SendAsync`, which talked to SMTP *inline, during the request*. That has two problems 4.1 taught us to see. First, it's a dual-write: the magic-link token is saved to the database *and* the email is sent, two systems, no atomicity — crash between them and the user has a token row but no email (or an email for a token that rolled back). Second, it couples the *request's* fate to the *mail server's*: SMTP is slow, so the user waits; SMTP is down, so sign-in fails. 2.3 bounded that with a timeout, but a bounded failure is still a failure.

The outbox (4.1) fixes both — *if* email goes through it. The question is how to route every existing caller through the outbox without a find-and-replace across the codebase. The answer is the reason `ISender` was an interface in the first place.

2. The decorator — same interface, new behavior

A **decorator** implements the same interface it wraps, adding behavior while remaining substitutable. `OutboxEmailSender` is an `ISender`, but instead of sending, it enqueues:


```
// src/Infrastructure/Email/OutboxEmailSender.cs (the whole thing)
/// IEmailSender that ENQUEUES the email onto the transactional outbox instead of sending
/// it inline; the real SMTP send happens later on the dispatcher via EmailOutboxHandler →
/// SmtplibEmailSender (ADR-007). This is the app-facing IEmailSender, so existing call sites
/// (passwordless, invitations) are unchanged – but a request no longer blocks on SMTP or
/// fails if the mail server is momentarily down.
public sealed class OutboxEmailSender(IOutbox outbox, AppDbContext db) : IEmailSender
{
    public const string MessageType = "email";

    public async Task SendAsync(string to, string subject, string htmlBody,
        IReadOnlyList<EmailInlineImage>? inlineImages = null, CancellationToken ct =
        default)
    {
        var payload = JsonSerializer.Serialize(new EmailOutboxPayload(to, subject,
            htmlBody, inlineImages));
        await outbox.EnqueueAsync(MessageType, payload, cancellationTokens: ct);

        // The email call sites aren't wrapped in an explicit unit of work, so flush now to
        // make the
        // enqueue durable before the request returns. SaveChanges on the shared context
        // also commits
        // any pending business change alongside the message – exactly the atomicity we
        // want.
        await db.SaveChangesAsync(ct);
    }
}

public sealed record EmailOutboxPayload(
    string To, string Subject, string HtmlBody, IReadOnlyList<EmailInlineImage>?
    InlineImages);
```

Look at what this *doesn't* require: PasswordlessService (2.4) still injects `IEmailSender` and calls `SendAsync`; it has no idea it's now talking to a decorator that enqueues rather than sends. The behavior of *every* email in the platform changed from "inline, fragile" to "queued, durable" by adding one class and changing one DI line. That is the entire return on the investment of making email a seam in 2.3 — a promise made three parts ago, paid in full here.

The `SaveChanges` deserves its comment (§1's atomicity, made concrete): the passwordless and invitation call sites don't run inside an explicit unit of work, so the decorator flushes to make the enqueue durable before returning. And because it's the *shared* context (3.1), that flush commits any pending business change — the `LoginToken`, the `TenantInvitation` — *in the same transaction* as the outbox row. Token-and-email-intent, atomic. The dual-write is gone, and the caller never knew there was one.

3. The handler — the real send, later

The dispatcher (4.1) eventually hands the message to the type-"email" handler, which does the actual SMTP send using the *undecorated* sender:

```
// src/Infrastructure/Email/EmailOutboxHandler.cs (the whole thing)
public sealed class EmailOutboxHandler(IEmailSender smtpSender) : IOutboxHandler
{
    public string Type => OutboxEmailSender.MessageType; // "email"

    public async Task HandleAsync(OutboxMessage message, CancellationToken ct = default)
    {
        var payload = JsonSerializer.Deserialize<EmailOutboxPayload>(message.Payload)
            ?? throw new InvalidOperationException($"Outbox message {message.Id} has an
unreadable email payload.");
        await smtpSender.SendAsync(payload.To, payload.Subject, payload.HtmlBody,
payload.InlineImages, ct);
    }
}
```

Now — a subtle DI trap, and its fix. The handler needs *an* `IEmailSender`, but if it resolved the app-facing one it would get `OutboxEmailSender` and enqueue the email *again* — an infinite loop of a message that re-queues itself forever. So the registration distinguishes the two by a **key**:

```
// the app-facing IEmailSender is the decorator; the real one is keyed "smtp"
services.AddKeyedScoped<IEmailSender, SmtplibEmailSender>("smtp");
services.AddScoped<IEmailSender, OutboxEmailSender>(); // what callers get
// EmailOutboxHandler injects the KEYED "smtp" sender, not the decorator
```

Callers ask for `IEmailSender` and get the decorator (enqueue); the handler asks for the keyed "smtp" one and gets the real `SmtplibEmailSender` (send). This is the same "break-the-cycle-with-two-registrations" move you met in 3.4 (the two `HttpClient`s) — a recurring shape whenever a decorator and the thing it decorates share an interface. Recognize it once, and you'll reach for it automatically.

4. Idempotency, made concrete

4.1 insisted every outbox handler be idempotent because delivery is at-least-once. Here's what that means for *this* handler, from its own doc comment:

At-least-once delivery means a duplicate dispatch can re-send a transactional email — accepted here (a second magic link / invitation copy is harmless); the dispatcher only retries on failure.

This is the honest, non-magical version of idempotency. `EmailOutboxHandler` doesn't do anything clever to *prevent* a double-send; it makes the *judgment* that a duplicate transactional email is harmless — worst case, a user gets two copies of the same magic link (both single-use; the first consumed wins, per 2.4). So "idempotent enough" is achieved by *choosing effects whose repetition doesn't matter*, not by deduplication machinery. Contrast that with an effect where repetition *does* matter — "charge this card twice" is not harmless — which is exactly the case the **inbox** (4.3) exists to handle. Email sits on the easy side of the idempotency line; knowing which side an effect is on is the design skill.

5. Red — proving the swap

```

Scenario: sending email enqueues rather than sends inline
  When a caller invokes IEmailSender.SendAsync
  Then no SMTP send has happened yet
  And one pending outbox message of type "email" exists carrying the payload

Scenario: the dispatcher performs the real send
  Given a pending "email" outbox message
  When the dispatcher runs a pass
  Then the underlying SmtplibEmailSender received the exact to/subject/body/images
  And the message is marked sent

```

OutboxEmailTests drives both — using a fake `SmtplibEmailSender` (recording sends) as the keyed sender, so the test asserts the round-trip payload without real SMTP. Note the first scenario proves the *decoupling* (the request did its work with zero SMTP contact); the second proves the *delivery*. Together they show the behavior swapped end to end while `SendAsync`'s signature — and every caller — stayed put.

6. Run it

```

dotnet test tests/Api.Tests --filter Outbox
dotnet run --project src/Api
# request a magic link → the request returns instantly (no SMTP wait); the OutboxMessages
# table shows a pending "email"; the next dispatcher pass sends it → Mailpit (:8025) shows
it

```

Stop your SMTP (or Mailpit) and request a link: the *request still succeeds* and the message waits as pending, retrying with backoff until the mail server returns. That's the reliability win made visible — sign-in no longer depends on email being up *right now*.

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how do existing email callers get outbox reliability? (a) Change every call site to enqueue instead of send; (b) put an "if async" branch inside `SmtplibEmailSender`; (c) a decorator implementing `IEmailSender` that enqueues, swapped in via DI, with the real sender kept as a keyed dependency for the handler.

Chosen: (c) (ADR-007). The decorator is *transparent* — callers are untouched, the change is one class plus one registration, and the seam's substitutability (2.3) is what makes it possible. Keyed registration cleanly separates "what callers get" (enqueue) from "what the handler uses" (send), avoiding the self-wrapping loop.

Rejected: (a) — a codebase-wide edit that must be repeated for every future email caller, and easy to forget (a new caller would silently go back to fragile inline sends). (b) — a mode flag inside `SmtplibEmailSender` conflates two responsibilities (talk to SMTP / decide whether to defer) in one class, and every caller would still need to know which mode to ask for. The decorator keeps each class doing one thing.

The trade: email is now *eventually* delivered (a poll interval late) and a transactional email can, rarely, arrive twice — both accepted in §4. And there are now two `IEmailSender` registrations to understand; the keying + comments carry that weight.

8. Checkpoint

```
git add -A && git commit -m "feat: email outbox decorator – crash-safe sends, zero  
call-site changes"  
git tag lesson-4.2
```

You should now be able to: define the decorator pattern and explain how it swaps behavior without touching callers; set up keyed DI to keep a decorator from wrapping itself; articulate "idempotent enough" as *choosing repetition-safe effects* rather than building dedup; and point to this lesson as the concrete payoff of 2.3's seam.

Next: *Lesson 4.3 — The inbox* — the mirror image: idempotency for inbound effects you *can't* just re-run safely, like a billing webhook that must not double-charge.

Lesson 4.3 — The inbox

Part 4 · Reliability & operations. The outbox (4.1) handles effects *we* send; the inbox handles effects *we receive*. They're mirror images of one at-least-once problem. 4.2 got away with "idempotent enough" because a duplicate email is harmless — but the platform is about to receive webhooks where a duplicate is *not* harmless: a billing provider that redelivers "subscription activated" must not activate twice, and one that retries "payment succeeded" must not be read as two payments. When an inbound effect isn't naturally repetition-safe, you need a dedup gate. That's the inbox, and its correctness hinges on one subtle transaction rule most implementations get wrong.

Goal: an inbound delivery is processed *exactly once* even when the sender delivers it many times — claimed by a (source, key) ledger, race-free under concurrency, and rolled back with its work if the work fails so the inevitable redelivery re-processes.

Concepts & principles: inbound at-least-once & redelivery; idempotency via a dedup ledger; the claim-and-work-commit-together rule (and the classic autocommit bug); Postgres `INSERT ... ON CONFLICT DO NOTHING`; the outbox-vs-inbox contrast (`SKIP LOCKED` vs `ON CONFLICT`).

Maps to: ADR-007 · repo: `src/Core/Abstractions/IInbox.cs`,
`src/Core/Entities/InboxMessage.cs`, `src/Infrastructure/Inbox/EfInbox.cs`, tests:
`tests/Api.Tests/Inbox/InboxTests.cs`.

Prerequisites: 4.1 (at-least-once, the shared-transaction commit model), 3.1 (unit of work).

1. Motivate — the redelivery you can't wish away

External systems that call you — payment providers, webhook senders — deliver **at-least-once**, for the same reason your outbox does: they can't tell "my call succeeded" from "my call succeeded but the response got lost," so when in doubt they resend. Stripe will deliver the same event twice; that's not a bug in Stripe, it's the honest behavior of any reliable sender. So your webhook endpoint *will* receive duplicates.

For some effects that's fine (4.2's email — send the receipt twice, shrug). For others it's a disaster: "activate the subscription" applied twice, "payment of \$50 received" counted as \$100, "grant admin" run again after it was revoked. These effects are *not* naturally idempotent, and — critically — you often *can't make them so* by choosing a safer effect the way 4.2 did, because the effect is dictated by the business event, not by you. So you need the other tool: a gate that recognizes "I've already handled this exact delivery" and refuses to do the work a second time. That gate is the inbox.

2. The ledger — one row per delivery, ever

```
// src/Core/Entities/InboxMessage.cs (the real entity)
/// A dedup ledger entry for an inbound at-least-once delivery (e.g. a Stripe webhook).
/// Recording a Source + IdempotencyKey exactly once is what makes inbound processing
/// idempotent: the first delivery claims a row and does the work in the same transaction;
/// a redelivery of the same key finds the row already present and is a no-op (ADR-007).
public class InboxMessage
{
    public Guid Id { get; set; } = Guid.CreateVersion7();
    public required string Source { get; set; } // namespaces the key, e.g. "stripe"
    public required string IdempotencyKey { get; set; } // the provider's delivery id, e.g.
    evt_...
    public DateTimeOffset ReceivedAt { get; set; }
}
```

The design is almost entirely in one constraint: a ****unique index on (Source, IdempotencyKey)****. Two decisions inside that:

- ****Source namespaces the key.**** Stripe's evt_123 and some other provider's evt_123

must not collide, so the identity is the *pair*, not the key alone. Cheap now, essential the day a second integration arrives.

- ****IdempotencyKey is the *provider's* id, not one you invent.**** Stripe stamps each event

evt_... and reuses it across redeliveries of the same event — so "same event" is *their* identity, which is exactly what you must dedup on. Inventing your own key (hashing the body, say) is fragile: a re-serialized-but-identical redelivery could hash differently and slip through. Use the sender's stable id.

Like the outbox, InboxMessage is deliberately ****not ITenantScoped**** — its own comment: an inbound event is reconciled to its tenant *by the handler* (e.g. via the Stripe customer id), so the ledger is platform infrastructure on 2.6's allowlist. (That reconciliation is where 5.2's EnterTenant comes in — the webhook has no JWT, so it enters the tenant it resolves. Two Part-4 primitives converging in Part 5.)

3. The claim — race-free by construction

```
// src/Infrastructure/Inbox/EfInbox.cs (the whole implementation)
public async Task<bool> TryClaimAsync(string source, string idempotencyKey,
    CancellationToken ct = default)
{
    var rows = await db.Database.ExecuteSqlAsync($"
        INSERT INTO "InboxMessages" ("Id", "Source", "IdempotencyKey", "ReceivedAt")
        VALUES ({Guid.CreateVersion7()}, {source}, {idempotencyKey}, {clock.GetUtcNow()})
        ON CONFLICT ("Source", "IdempotencyKey") DO NOTHING
        """, ct);

    return rows == 1; // 1 = first delivery (claimed, do the work); 0 = duplicate (no-op)
}
```

The naïve approach — SELECT to check if the key exists, then INSERT if not — has a race: two concurrent redeliveries both SELECT (neither sees a row), both INSERT, and you've either double-processed or hit a duplicate-key crash. INSERT ... ON CONFLICT DO NOTHING makes claim-or-skip a **single atomic statement**: the insert affects one row on the first delivery and zero on a duplicate, and concurrent claims of the same key *serialize on the unique index* — exactly one wins, no crash, no double. The return is just "did I insert a row" → "is this the first

time." Race-freedom comes from the database constraint, not from application locking.

This is a nice contrast to hold against 4.1. Both the outbox and the inbox solve concurrency with a Postgres feature rather than app-level locks, but *different* features for *different* shapes:

	Outbox (4.1)	Inbox (4.3)
Job	pick the next item off a queue	recognize a duplicate delivery
Postgres idiom	FOR UPDATE SKIP LOCKED	INSERT ... ON CONFLICT DO NOTHING
"Winner"	one poller claims each <i>pending</i> row	one insert claims each <i>unique key</i>

Reaching for the right database primitive — a lock-and-skip for queue draining, a unique-constraint upsert for dedup — is the transferable skill; both beat hand-rolled locking.

4. The one rule that makes it correct — claim inside the work's transaction

Here is the subtlety, and it's the whole lesson. It is not enough to claim the key and then do the work. Read `IInbox`'s doc comment, which is really a warning:

Use inside the caller's transaction. The claim and the work it guards must commit together: wrap `TryClaimAsync` + the processing in an `IUnitOfWork` transaction and commit only after the work succeeds. If the work throws and the transaction rolls back, the claim rolls back too, so the inevitable redelivery re-processes — that's the correct exactly-once-*effect* behaviour. (Claiming in autocommit and then doing the work would mark a delivery handled even if the work fails — don't.)

Walk the failure the wrong way exposes. Suppose you claim in its own committed statement, *then* start the work, and the work throws (a bug, a transient DB error). The claim is already committed — so the ledger says "handled" — but the work never happened. The sender redelivers (at-least-once, remember), your inbox sees the key is "already claimed," and **no-ops** — permanently dropping an event that was never processed. You've converted at-least-once into *at-most-once*, silently, for exactly the failures you most need to survive.

The fix is to bind claim and work into **one transaction** (3.1's unit of work):

```
await using var tx = await uow.BeginTransactionAsync(ct);
if (!await inbox.TryClaimAsync("stripe", evt.Id, ct))
    { await tx.RollbackAsync(ct); return; } // duplicate → no-op, safely
await ApplyTheEffect(evt, ct); // the business change
await tx.CommitAsync(ct); // claim + effect commit together, or neither
```

Now the claim's durability is *tied* to the work's. Work succeeds → both commit → future redeliveries correctly no-op. Work fails → transaction rolls back → the claim vanishes too → the redelivery re-processes, as it must. The term of art is that this gives ****exactly-once *effect***** (the outcome happens once) on top of at-least-once *delivery* — you can't stop duplicate deliveries, but you can make their *effect* singular. That distinction — once delivery vs once effect — is worth carrying: nobody can promise the former across a network, and this is how you achieve the latter regardless.

5. Red — the tests that pin exactly-once-effect

```
Scenario: first delivery is claimed, duplicate is a no-op
  Given no ledger entry for ("stripe", "evt_1")
  When TryClaimAsync("stripe","evt_1") runs
  Then it returns true
  And a second TryClaimAsync("stripe","evt_1") returns false

Scenario: concurrent claims of the same key – exactly one wins
  When two claims of ("stripe","evt_1") race
  Then exactly one returns true, the other false, and no error is thrown

Scenario: failed work re-processes on redelivery (the correctness rule)
  Given the claim + work run in one transaction and the work throws
  When the transaction rolls back and the event is redelivered
  Then the second delivery is claimed and processed (the first claim did not survive)
```

InboxTests drives these on real Postgres (the `ON CONFLICT` and the concurrent-claim serialization are Postgres semantics — 1.3's decision again). The third scenario is the one that proves the §4 rule and the one a careless implementation fails: it asserts a rolled-back claim does *not* block reprocessing.

6. Run it

```
dotnet test tests/Api.Tests --filter Inbox
```

There's no user-facing surface yet — the inbox is a primitive. Its first real consumer is the billing webhook in 5.2, where you'll see `TryClaimAsync` wrapped exactly as §4 prescribes, deduping Stripe's redeliveries so a subscription event applies once.

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how is inbound at-least-once made safe? (a) Trust the sender to deliver once (it won't); (b) make every inbound handler naturally idempotent by hand; (c) a dedup-ledger gate — claim (source, key) via `ON CONFLICT`, inside the work's transaction, no-op on duplicates.

Chosen: (c) (ADR-007). It works for effects that *can't* be made naturally idempotent (the common webhook case), it's race-free via a database constraint rather than app locks, and binding the claim to the work's transaction delivers exactly-once- effect with the correct failure behavior (rollback → reprocess). It mirrors the outbox, so the two together cover both directions of at-least-once with one mental model.

Rejected: (a) — duplicates are guaranteed, not hypothetical; "assume once" is the bug. (b) — sometimes possible (4.2's email) and preferable *when* it is, but many business effects (apply a payment, activate a plan) have no repetition-safe form, so it can't be the general answer; the inbox is the fallback that always works.

The trade: an extra table and the discipline of remembering to wrap claim+work in a transaction — and getting that wrapping wrong fails silently and dangerously (§4), which is why the rule is documented on the interface and pinned by a test. The ledger also grows unbounded without pruning (a cleanup job — 4.4 — trims old entries, same as the outbox's sent rows).

8. Checkpoint

```
git add -A && git commit -m "feat: inbox — exactly-once-effect dedup for inbound
at-least-once deliveries"
git tag lesson-4.3
```

You should now be able to: explain why inbound deliveries are at-least-once and when that's dangerous; describe the (source, key) ledger and why the key is the sender's id; articulate the claim-inside-the-work's-transaction rule and the at-most-once bug that skipping it causes; and contrast `ON CONFLICT` (dedup) with `SKIP LOCKED` (queue-claim).

Next: Lesson 4.4 — *The scheduler* — the third reliability primitive: recurring in-process work (token cleanup, ledger pruning, the sweeps later lessons rely on).

Lesson 4.4 — The scheduler

Part 4 · Reliability & operations. The third and last of the async primitives. The outbox pushes effects out; the inbox dedups effects coming in; the scheduler does work that no request triggers at all — recurring maintenance on a clock. It's the humblest of the three and the pattern is short, but it closes real loops: the expired tokens from 2.2/2.4, the sent outbox rows and inbox ledger from 4.1/4.3, and the billing sweeps Part 5 will lean on all need *something* to run them on a timer. That something is here, and it follows the same DI-discovered, resilient, testable-core shape as the dispatcher.

Goal: register a recurring job by dropping an `IScheduledJob` into DI — no host edits — and have it run on its own interval, isolated from other jobs' failures, with a testable core separate from the timer loop.

Concepts & principles: in-process scheduled work; DI-discovered jobs (Open/Closed again); failure isolation between jobs; per-run DI scope; the testable-core / background- service split (the outbox shape, reused); why *derived* cleanup, not stored flags.

Maps to: ADR-007 · repo: `src/Core/Abstractions/IScheduledJob.cs`,
`src/Infrastructure/Scheduling/ScheduledJobsHost.cs`, `ScheduledJobsOptions.cs`,
`ExpiredTokenCleanupJob.cs`, tests: `tests/Api.Tests/Scheduling/*`.

Prerequisites: 4.1 (the background-service + testable-core pattern this mirrors), 3.3 (injected clock), 2.9 (the contributor seam's discovery pattern this echoes).

1. Motivate — the work nobody asks for

Some work has no triggering request. Expired login and refresh tokens (2.2, 2.4) pile up in their tables and must be swept. The outbox's sent rows and the inbox's ledger (4.1, 4.3) grow forever without pruning. Later, trials expire and dunning sweeps must run (Part 5). None of this is initiated by a user clicking something — it's *time* that triggers it. You need a component that wakes on a schedule and runs due maintenance.

The temptation is a cron entry or a cloud scheduled-function per task. That works, but it scatters the platform's recurring logic *outside* the platform — into infrastructure config a developer can't see, test, or version alongside the code. For a platform meant to run cheap and self-contained (Postgres-only, one process), an **in-process scheduler** keeps the jobs *in the codebase*, testable and deployable as one unit. That's the choice here, with the honest caveat (§7) that a single sleeping instance has limits.

2. The job — an interface anyone can implement

```
// src/Core/Abstractions/IScheduledJob.cs (the real interface)
/// A recurring background job run by the ScheduledJobsHost on its own Interval (ADR-007).
/// Add a job by registering an IScheduledJob in DI — no host edits. Each run gets a fresh
/// DI scope, so a job may depend on scoped services (a new ApplicationDbContext).
public interface IScheduledJob
{
    string Name { get; } // stable id, tracks last-run; unique across jobs
    TimeSpan Interval { get; } // minimum time between runs
    Task RunAsync(Cancellation token ct = default);
}
```

Three members, and the design intent is entirely in the doc comment: **add a job by registering it in DI — no host edits**. This is the *exact* Open/Closed move you've now seen three times — the data contributors (2.9), the outbox handlers (4.1), and now scheduled jobs. The host doesn't know what jobs exist; it asks DI for all `IScheduledJobs` and runs the due ones. A new sweep is a new class plus one registration; the scheduler is closed to modification, open to extension. When a pattern recurs this often in one codebase, it stops being a trick and becomes the house style — and recognizing it is half of reading the platform fluently.

3. The reference job — and a tenancy-flavored subtlety

```
// src/Infrastructure/Scheduling/ExpiredTokenCleanupJob.cs (the whole job)
/// Reference scheduled job (ADR-007): deletes expired passwordless and refresh tokens so
/// the auth tables don't grow without bound. Only rows whose ExpiresAt is in the past are
/// removed — a revoked-but-unexpired refresh token is kept because its hash still backs
/// rotation/replay detection until it expires. Copy this shape for trial-expiry sweeps,
/// etc.
public sealed class ExpiredTokenCleanupJob(ApplicationDbContext db, TimeProvider clock) :
    IScheduledJob
{
    public string Name => "expired-token-cleanup";
    public TimeSpan Interval => TimeSpan.FromHours(1);

    public async Task RunAsync(Cancellation token ct = default)
    {
        var now = clock.GetUtcNow();
        await db.LoginTokens.Where(t => t.ExpiresAt < now).ExecuteDeleteAsync(ct);
        await db.RefreshTokens.Where(t => t.ExpiresAt < now).ExecuteDeleteAsync(ct);
    }
}
```

Small, but two details are load-bearing:

- ****It deletes only *expired* rows, not *revoked* ones.**** The comment is a direct callback

to 2.2: a refresh token that's been rotated out (revoked) but hasn't yet expired is *deliberately kept*, because its hash is what powers reuse-detection — delete it early and a replayed stolen token would read as merely "unknown" instead of the theft signal 2.2 worked to produce. The cleanup respects the security design; "old" and "safe to delete" are not the same predicate, and getting that wrong would quietly reopen a hole.

- **The clock is injected** (3.3), so "what's expired" is testable by advancing a

`FakeTimeProvider` rather than waiting an hour. And note `ExecuteDeleteAsync` — a set-based delete. Recall from 2.7 that bulk operations bypass the change-tracker (and thus the stamping interceptor); that's *fine here* because this job runs as system maintenance on the system

context (no current tenant), which is exactly the trust level 2.7 reserved for platform code. A feature slice couldn't write this; a platform job can.

4. The host — resilient, scoped, and testable

```
// src/Infrastructure/Scheduling/ScheduledJobsHost.cs (the core)
public async Task RunDueJobsAsync(CancellationToken ct = default)
{
    using var scope = scopeFactory.CreateScope(); // fresh scope per tick (scoped
AppDbContext)
    var now = clock.UtcNow();

    foreach (var job in scope.ServiceProvider.GetServices<IScheduledJob>())
    {
        if (_lastRun.TryGetValue(job.Name, out var last) && now - last < job.Interval)
            continue; // not due yet

        try { await job.RunAsync(ct); }
        catch (OperationCanceledException) when (ct.IsCancellationRequested) { throw; } //
shutdown
        catch (Exception ex)
        {
            logger.LogError(ex, "Scheduled job {Job} failed; will retry next interval",
job.Name);
        }
        finally
        {
            // Advance even on failure so a broken job retries on its interval, not in a hot
loop.
            _lastRun[job.Name] = now;
        }
    }
}
```

Four decisions, each a small reliability lesson:

- **Failure isolation.** One job throwing is caught, logged, and *skipped* — it never stops

the other jobs or the host. A broken trial-sweep must not take down token cleanup. The blast radius of a bad job is that one job.

- **Advance last-run even on failure** (the *finally*). This is the subtle one: if a

failing job's timestamp *weren't* advanced, it would be "due" again on the very next tick and hammer in a hot loop (failing every few seconds, flooding logs, burning the DB). Advancing means a broken job retries on its normal interval — failure degrades to "this maintenance is delayed," not "the host is on fire."

- **A fresh DI scope per tick**, so jobs can depend on scoped services (a new

AppDbContext) — same reason the outbox dispatcher (4.1) scopes per pass. A background service is a singleton; the work it runs is scoped; the scope-per-tick bridges them.

- **The core is split from the loop.** RunDueJobsAsync (pure, testable) is separate from

the ExecuteAsync timer loop (BackgroundService) — the identical split as the outbox processor/dispatcher (4.1), for the identical reason: you can unit-test "did the due jobs run?" by calling the core with a FakeTimeProvider, with no real waiting.

5. Red — deterministic scheduling without waiting

Scenario: a job runs when due, then respects its interval

```
Given a job with Interval = 1 hour, never run
When RunDueJobsAsync runs
Then the job ran once
When RunDueJobsAsync runs again immediately
Then the job did NOT run (not due)
When the clock advances 1 hour and RunDueJobsAsync runs
Then the job ran again
```

Scenario: a failing job is isolated and keeps its schedule

```
Given job A throws and job B succeeds
When RunDueJobsAsync runs
Then B still ran, the host didn't crash, and A is scheduled to retry next interval (not hot-looping)
```

ScheduledJobsHostTests + ExpiredTokenCleanupJobTests drive these with FakeTimeProvider.Advance — an hour passes in a microsecond, and the "isolation" scenario uses a deliberately-throwing fake job. No Task.Delay, no flake. The testable-core split (4.1) is what makes scheduling logic assertable instead of a timing gamble.

6. Run it

```
dotnet test tests/Api.Tests --filter Scheduling
# live: the host ticks on its interval; the LoginTokens/RefreshTokens tables shed expired
# rows within an hour of expiry (or immediately, if you shorten the interval in config to
watch it)
```

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how does recurring maintenance run? (a) OS cron / cloud scheduled functions calling an endpoint; (b) a distributed job framework (Hangfire, Quartz); (c) an in-process BackgroundService running DI-discovered IScheduledJobs.

Chosen: (c) (ADR-007). The jobs live *in the codebase* — versioned, unit-tested, deployed as one unit — with no external scheduler to configure or secure, matching the platform's Postgres-only, run-cheap baseline. DI-discovery keeps it Open/Closed; the core/loop split keeps it testable; failure isolation keeps one bad job cheap.

Rejected: (a) — scatters the platform's recurring logic into infra config that's invisible to the code, untestable in the suite, and drift-prone (the endpoint and the cron entry evolve separately). (b) — Hangfire/Quartz are excellent at *scale* (persistent schedules, multi-node coordination, dashboards) but add a dependency and operational surface the baseline doesn't need; the IScheduledJob seam is the documented swap-in point when multi-node arrives.

The trade — and it's a real one: in-process + single-instance means **if the process isn't running, jobs don't run**. On the free-tier deployment (Part 8) where the instance *sleeps* when idle, a scheduled sweep can be delayed until traffic wakes it — the same caveat as the outbox poller. That's acceptable for cleanup/sweeps (a token pruned an hour late harms nothing) and explicitly *not* acceptable for anything time-critical, which is a recorded deployment decision, not a surprise. And a naive multi-instance run would execute each job on every node — the seam's swap-in (a distributed scheduler) is the answer if that day comes.

8. Checkpoint

```
git add -A && git commit -m "feat: scheduler - DI-discovered recurring jobs,
failure-isolated, testable core"
git tag lesson-4.4
```

You should now be able to: recognize the DI-discovery Open/Closed pattern in its third guise and register a job without touching the host; explain failure isolation and why last-run advances even on failure; connect the token-cleanup predicate back to 2.2's reuse-detection design; and state the single-instance / sleeping-instance limitation and its documented swap-in.

Next: Lesson 4.5 — *Observability* — structured logging, tracing, and health checks: how you see what the platform is doing, and how a deploy knows it's alive.

Lesson 4.5 — Observability

Part 4 · Reliability & operations. You've made the platform *reliable*; now make it *legible* — able to answer "what is it doing, for whom, and is it alive?" Three capabilities, one theme: every signal is enriched with *who* (tenant + user) so a production question ("why is this tenant seeing errors?") has an answer. Structured logs say what happened, traces say where the time went, and health checks tell an orchestrator whether to send traffic. And a discipline threads through all three: the signals carry *identifiers*, never secrets — observability must never become a leak.

Goal: every log line and trace span produced while handling a request carries `tenant_id/user_id`; the app exposes liveness and a database-backed readiness probe; and all of it runs with zero external dependencies by default (exporters are config-gated).

Concepts & principles: structured logging + per-request scope enrichment; distributed tracing (OpenTelemetry); liveness vs readiness; config-gated exporters (no-dependency default); identifiers-not-secrets in telemetry; enrichment that can never fail a request.

Maps to: ADR-008, ADR-C10 · repo:

`src/Api/Observability/RequestLoggingScopeMiddleware.cs`, `TelemetryExtensions.cs`, `DatabaseHealthCheck.cs`, `src/Api/Program.cs` (logging setup + `/health` mapping from 1.2), tests: `tests/Api.Tests/Observability/*`.

Prerequisites: 1.2 (the health endpoint stub, now grown up), 2.6 (`ICurrentTenant`, the enrichment source), 1.4 (config-gating).

1. Motivate — a production question you can't answer

It's 2 a.m. and one tenant reports errors. With `Console.WriteLine`-style logging you have a wall of undifferentiated lines and no way to isolate *that tenant's* requests. With no tracing you can see a request was slow but not *where* the time went. With a naive health check that returns 200 as long as the process is up, your orchestrator keeps routing traffic to an instance whose database connection died. Observability is the difference between "the system is a black box that occasionally misbehaves" and "I can ask it questions and get answers." The three tools below are those three answers — *what*, *where*, *alive* — and the platform's specific contribution is enriching all of them with tenant identity so the questions can be *per-tenant*, which for a SaaS is the question that matters.

2. Structured logging — logs you can query, scoped to who

`Program.cs` (from 1.2, grown up) picks a log format by environment: a readable single-line console in dev, **JSON in production** so a log aggregator can index fields. That's the "structured" in structured logging — a log line isn't a string, it's an *event with fields*. And the platform attaches the fields that matter via a per-request scope:

```
// src/Api/Observability/RequestLoggingScopeMiddleware.cs (core)
/// Opens a per-request logging scope carrying tenant_id and user_id, so every log line
/// emitted while handling the request is filterable by tenant/user (OBS-1). ... Only
/// identifiers go in the scope – never tokens or other secrets.
internal static Dictionary<string, object> BuildScope(HttpContext context, ICurrentTenant
currentTenant)
{
    var scope = new Dictionary<string, object>();
    if (currentTenant.TenantId is { } tenantId) scope["tenant_id"] = tenantId;
    if (context.User.FindFirst(ClaimTypes.NameIdentifier)?.Value is { } userId)
scope["user_id"] = userId;
    return scope;
}
// ... using (logger.BeginScope(scope)) await next(context);
```

BeginScope means *every* log line written anywhere downstream during that request — in a handler, a service, EF — automatically carries `tenant_id` and `user_id`, without any of that code knowing about logging scopes. Now "show me this tenant's errors in the last hour" is a filter, not an archaeology dig. Three details are deliberate:

- **Identifiers only** — the comment is emphatic, and it's the recurring 1.5/4.1 rule:

never a token, never PII, in a log scope. Logs are widely readable and long-retained; a secret in a log is a secret leaked at scale.

- **Anonymous requests add no scope** (empty → skip BeginScope) — no noise on pre-auth traffic.

- ****Registered *after* authentication**** so the claims exist to read. Middleware order is itself a decision; the extension's comment says so.

3. Tracing — where the time actually went

Logs say *what*; traces say *where*. OpenTelemetry (the vendor-neutral standard) produces spans across the request, outbound HTTP, and the database, so a slow request decomposes into "80ms in Postgres, 200ms in that outbound call." The platform's setup, abridged:

```
// src/Api/Observability/TelemetryExtensions.cs (core)
services.AddOpenTelemetry()
    .ConfigureResource(r => r.AddService("Perezosoft.Api"))
    .WithTracing(tracing =>
    {
        tracing
            .AddAspNetCoreInstrumentation(o =>
                o.EnrichWithHttpResponse = (activity, response) => EnrichSpan(activity,
response.HttpContext))
            .AddHttpClientInstrumentation()
            .AddSource("Npgsql"); // stable built-in Npgsql tracing (NOT the beta EF
instrumentation)
        ApplyExporter(tracing, otlpEndpoint, useConsole);
    })
    .WithMetrics(...);
```

Three decisions with real reasoning behind them:

- **Same enrichment, on spans.** `EnrichSpan` tags each request span with

tenant_id/user_id (identifiers only, again), so traces filter per tenant just like logs — one consistent identity story across both signals. And note its defensive guard: context.RequestServices can be null for requests rejected before the pipeline runs (Kestrel bad requests, early aborts), so enrichment null-checks and *never throws* — **observability must not be able to fail the request it observes**. A telemetry bug taking down real traffic is the classic own-goal; the guard forbids it.

- ****AddSource("Npgsql")**, not EF Core instrumentation.** A pointed, documented choice

(ADR-C10): subscribe to Npgsql's *stable, built-in* tracing rather than the perpetually-beta EF Core instrumentation. "Latest stable, never previews" (Golden Rule 6) applied to telemetry — you don't want a beta package deciding whether your DB spans work.

- **Config-gated exporters, no-dependency default.** This is the elegant part: spans are always *produced*, but *exported* only if configured — OTLP when OpenTelemetry:Otlp:Endpoint is set (production, pointing at a collector), the noisy console exporter only if you flip OpenTelemetry:ConsoleExporter=true for local debugging, and **nothing** otherwise. So a fresh checkout runs with a clean console and zero external dependency, while production wires a real backend by setting one config key (1.4's pattern). Instrumentation is always on; where it goes is deployment config. (Both keys must be declared in appsettings.json — the moment TelemetryExtensions reads them, 1.4's doc-sync gate will fail the build until they're documented. The gate doing its job is the reminder; add an OpenTelemetry section with the two keys.)

4. Health — liveness vs readiness

Remember the /health endpoint from 1.2, using the framework's health-check pipeline "so its growth is additive"? This is the growth. There are *two* questions an orchestrator asks, and conflating them is a classic bug:

- **Liveness** — "is the process up?" /health answers this with no dependency probe

(predicate _ => false → no checks run, just "the process answered"). If liveness fails, the orchestrator *restarts* the instance.

- **Readiness** — "can it actually serve?" /health/ready runs the checks *tagged*

ready, including the database probe. If readiness fails, the orchestrator *stops routing traffic* to the instance — but doesn't restart it (the DB might just be briefly unreachable).

```
// src/Api/Program.cs – the two probes (from 1.2, now with a real readiness check)
app.MapHealthChecks("/health", new HealthCheckOptions { Predicate = _ => false });
app.MapHealthChecks("/health/ready", new HealthCheckOptions { Predicate = c =>
    c.Tags.Contains("ready") });
```

```
// src/Api/Observability/DatabaseHealthCheck.cs – the readiness probe
/// Readiness check (OBS-3): reports Healthy only when the database is reachable, so an
/// orchestrator won't route traffic to an instance that can't serve. Status-only – no
/// connection details are surfaced.
public async Task<HealthCheckResult> CheckHealthAsync(HealthCheckContext ctx,
CancellationTokens ct = default)
{
    try
    {
        return await db.Database.CanConnectAsync(ct)
            ? HealthCheckResult.Healthy() : HealthCheckResult.Unhealthy("Database
unreachable");
    }
    catch { return HealthCheckResult.Unhealthy("Database unreachable"); }
}
```

The distinction is operationally load-bearing: a process that's *up* but can't reach its database is *live* (don't restart — restarting won't fix the DB) but *not ready* (don't send it traffic). Get this wrong by making `/health` probe the DB, and a brief DB blip triggers a restart storm across every instance. Note the readiness check is **status-only** — it returns "Database unreachable", never the connection string or error detail (the same no-leak discipline as §2, applied to a *public, unauthenticated* endpoint where it matters most). This `/health/ready` is exactly what Part 8's deploy pipeline polls as its post-deploy smoke gate — the endpoint born trivially in 1.2 becomes the thing that decides whether a release goes live.

5. Red — testing that telemetry enriches and health reports

```
Scenario: request logs carry tenant and user
  Given an authenticated request for tenant T by user U
  When the scope is built
  Then it contains tenant_id = T and user_id = U (and nothing else)

Scenario: an anonymous request adds no scope
  Then BuildScope returns empty (no noise on pre-auth traffic)

Scenario: readiness reflects the database
  Given the database is reachable
  Then /health/ready is Healthy
  Given it is not
  Then /health/ready is Unhealthy – and leaks no connection detail
```

`RequestLoggingScopeMiddlewareTests`, `TelemetryEnrichmentTests`, and `DatabaseHealthCheckTests` drive these — the enrichment logic is extracted into pure static methods (`BuildScope`, `EnrichSpan`) precisely so it's unit-testable without a running server, the same testable-core instinct as 4.1/4.4.

6. Run it

```
dotnet run --project src/Api
curl -k https://localhost:7160/health # Healthy (liveness – process up)
curl -k https://localhost:7160/health/ready # Healthy (readiness – DB reachable); stop
Postgres → Unhealthy
# sign in, make a request, watch the dev console: log lines carry tenant_id/user_id
# set OpenTelemetry:ConsoleExporter=true to see spans, incl. the Npgsql DB span
```

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how is the platform observed? (a) Ad-hoc `Console.WriteLine/string` logs, a single 200-OK health endpoint; (b) a specific vendor's SDK (Datadog/App Insights) wired throughout; (c) structured logging + OpenTelemetry with tenant/user enrichment, config-gated exporters, and split liveness/readiness probes.

Chosen: (c) (ADR-008). Structured + enriched makes signals *queryable per tenant* — the SaaS-critical axis; OpenTelemetry is vendor-neutral, so the backend is a config key (OTLP endpoint), not a code rewrite; config-gated exporters mean zero dependency by default and a real backend in prod without touching code; liveness/readiness split gives orchestrators the two distinct answers they need.

Rejected: (a) — unqueryable logs and a health check that lies (up $\neq$ ready) are exactly the 2 a.m. failure this lesson opened with. (b) — a vendor SDK threaded through the code couples you to that vendor (the anti-seam again); OTLP lets you point at Datadog, Jaeger, Grafana, or nothing by changing config — the vendor choice stays at the edge.

The trade: OpenTelemetry is more setup than `Console.WriteLine`, and enrichment middleware/span-tagging is machinery to maintain. Accepted: it's the difference between a system you can operate and one you can only pray to. Deliberately deferred (Golden Rule 6): the beta EF Core instrumentation, in favor of stable Npgsql tracing.

8. Checkpoint

```
git add -A && git commit -m "feat: observability – enriched structured logs, OTel tracing, liveness/readiness"
git tag lesson-4.5
```

You should now be able to: explain per-request scope enrichment and why `BeginScope` beats threading a logger; distinguish liveness from readiness and the restart-storm bug of conflating them; justify config-gated exporters and vendor-neutral OTLP over a vendor SDK; and state the "identifiers, never secrets" rule and where enrichment must never throw.

Next: Lesson 4.6 — *The append-only audit log* — the difference between *logs* (for operators, ephemeral) and an *audit trail* (for accountability, immutable), and enforcing immutability structurally.

Lesson 4.6 — The append-only audit log

Part 4 · Reliability & operations — finale. One more record-keeping capability, and it is emphatically *not* logging. Lesson 4.5's logs are for operators: ephemeral, secret-scrubbed, thrown away after a retention window. An **audit trail** is for accountability: durable, tenant-owned, queryable, exportable, and — the defining property — *immutable*. "Who changed this member's role, and when?" is a question a customer, a compliance auditor, or a courtroom may ask months later, and the answer must be trustworthy, which means the record cannot have been altered. This lesson builds that trail and enforces its immutability **structurally**, closing Part 4 on the course's signature move.

Goal: security-relevant actions record a tenant-scoped audit event, atomically with the change they describe; the events are append-only — application code can insert but *cannot* update or delete them, enforced by an interceptor, not a policy.

Concepts & principles: audit trail vs operational logs; append-only immutability enforced by a `SaveChanges` interceptor; audit-as-tenant-data (`ITenantScoped`, exportable, dissolvable); staging on the unit of work (atomic with the audited change); the deliberate `ExecuteDelete` bypass for dissolve; identifiers-not-secrets, once more.

Maps to: ADR-008 · repo: `src/Core/Abstractions/IAuditLog.cs`,
`src/Core/Entities/AuditEvent.cs`, `src/Infrastructure/Audit/AuditLog.cs`,
`src/Infrastructure/Persistence/AuditAppendOnlyInterceptor.cs`, tests:
`tests/Api.Tests/AuditLogTests.cs`.

Prerequisites: 4.5 (logs — the thing this is *not*), 2.7 (the interceptor pattern), 4.1 (staging-on-the-unit-of-work), 2.6 (`ITenantScoped`).

1. Motivate — the record that must not be editable

Consider the question "who removed this member?" asked six months after the fact. A log line *might* still exist, buried in an aggregator, scrubbed and rotated. But logs are the wrong tool for three reasons: they're operational (retained briefly, not forever), they're not *tenant data* (a tenant can't query or export their own log lines), and — decisively — they're not *trustworthy as evidence*, because operational logs are mutable by whoever runs the system. An audit trail is a different thing with different requirements:

	Operational logs (4.5)	Audit trail (4.6)
Audience	operators/SRE	the tenant, compliance, legal
Lifetime	short (rotated)	durable (retained, exported)
Ownership	the platform's	the <i>tenant's</i> data
Mutability	doesn't matter	must be immutable
Access	log aggregator	queryable/exportable via the app

The entity's own comment states the line: "*Audit ≠ logs: this is durable, queryable, exportable tenant data, not operational telemetry.*" Once you accept it's *tenant data*, two things follow automatically — it's *ITenantScoped* (a tenant sees only its own trail), and it participates in GDPR export and dissolve (Part 7). And once you accept it's *evidence*, immutability becomes a hard requirement — which §4 enforces in the one way this course trusts: structurally.

2. Recording — semantic events, staged on the caller's transaction

```
// src/Core/Abstractions/IAuditLog.cs (the real interface)
/// Records semantic audit events for the current tenant (OBS-4). The tenant is stamped
/// automatically (the event is ITenantScoped); pass the acting user explicitly. Like the
/// outbox, it STAGES the event on the current unit of work and does NOT save – so the audit
/// commits atomically with the change it records (or rolls back with it). Never put
/// secrets/PII in metadata.
public interface IAuditLog
{
    Task RecordAsync(string action, Guid? actorUserId = null, string? entityType = null,
        string? entityId = null, object? metadata = null, CancellationToken ct = default);
}
```

Two things you already know how to read:

- **It stages, it doesn't save** — exactly the outbox's move (4.1). `AuditLog.RecordAsync`

adds the event to the shared context and stops; it commits with the caller's transaction. This is *essential* for an audit trail: the audit of "role changed" must commit **in the same transaction as the role change itself**. Otherwise you get the audit-specific dual-write — the role change commits but the audit is lost (an unrecorded action), or the audit records something that then rolled back (a phantom event). Atomicity makes the trail *accurate by construction*: an event exists if and only if the thing it describes happened.

- ****action is a stable, greppable verb**** ("member.role_changed"), metadata is

structured JSON, and the rule is stamped on every relevant member: **identifiers only, never secrets or PII beyond actor refs**. An audit trail is *more* exposed than logs (the tenant can read and export it), so the no-secrets discipline (1.5/4.5) is if anything stricter here.

The implementation is unsurprising precisely because the patterns compose — it uses the generic repository (3.1), the injected clock (3.3), and lets the tenant interceptor (2.7) stamp `TenantId`:

```
// src/Infrastructure/Audit/AuditLog.cs (core)
await events.AddAsync(new AuditEvent
{
    Action = action, ActorUserId = actorUserId, EntityType = entityType, EntityId =
entityId,
    Metadata = metadata is null ? null : JsonSerializer.Serialize(metadata),
    CreatedAt = clock.UtcNow(),
}, ct);
// Intentionally no SaveChanges – the audit persists with the caller's unit of work.
```

Callers write `await audit.RecordAsync("member.role_changed", actor, "TenantMembership", id, new { from, to })` right beside the change, inside the same transaction (you'll see RBAC do exactly this in 6.1). One line, and the action is on the

permanent record.

3. Explicit recording, not magic — a recorded decision

A tempting alternative is to auto-audit *everything* by hooking `SaveChanges` and recording every entity change. The platform deliberately doesn't (ADR-008, with an amendment noting the deferral): audit events are **semantic**, not mechanical. "The `Role` column changed from member to admin" is a database fact; "*an owner promoted this member to admin*" is the *meaning*, and only the calling code knows the meaning — the intent, the actor, the business action. An auto-auditor produces a noisy change-feed no compliance officer wants to read; explicit `RecordAsync` calls produce a curated trail of things that *matter*. The cost is that a developer must remember to call it on security-relevant actions — accepted, and exactly the kind of decision the ADR log exists to record so a future maintainer doesn't "helpfully" add the auto-auditor back.

4. Immutability — enforced by the build, not the honor system

Here's the property that makes it an *audit* trail, and the course's signature move one last time in Part 4. It's not enough to *intend* that audit rows are never changed; a bug, a careless `Update`, or a malicious hand could rewrite history. So immutability is enforced by a `SaveChanges` interceptor — the same mechanism as tenant stamping (2.7):

```
// src/Infrastructure/Persistence/AuditAppendOnlyInterceptor.cs (core)
/// Enforces that AuditEvent is append-only (OBS-4): a tracked audit row in Modified or
/// Deleted state at SaveChanges throws, so application code can insert audit events but
/// never alter or remove them.
private static void Guard(DbContext? db)
{
    if (db is null) return;
    foreach (var entry in db.ChangeTracker.Entries<AuditEvent>())
        if (entry.State is EntityState.Modified or EntityState.Deleted)
            throw new InvalidOperationException(
                "AuditEvent is append-only: updates and deletes are not permitted from
                application code.");
}
```

Registered in `OnConfiguring` alongside the tenant stamper (2.7), so — same reasoning — every context enforces it, including ones tests construct directly; the guarantee can't be sidestepped by a code path that forgot to opt in. Try to modify an audit row and `SaveChanges` throws before any SQL runs. Insert is allowed; mutate and delete are not. History is append-only *by construction*, not by hoping no one writes the wrong LINQ.

But notice the deliberate escape hatch, and why it's *safe*:

Set-based deletes (`ExecuteDelete`) don't pass through the change tracker, so tenant dissolve can still purge a tenant's trail.

This is the *same* boundary you learned in 2.7: bulk `ExecuteDelete` bypasses the change tracker and therefore this interceptor. Here that bypass is a *feature*, not a leak — when a tenant is dissolved (2.9) or erased (7.1), its audit trail must go too (it's the tenant's data, and GDPR erasure means *all* of it). So dissolve's contributor uses the set-based delete to purge the trail,

which the interceptor can't and shouldn't block. The rule is precise: *application code* can never mutate an audit row (change-tracked writes are guarded); *platform teardown* can purge a whole tenant's trail (set-based, on the system context). "Append-only" means "no editing history," not "history is undeletable even when the tenant is destroyed" — and knowing exactly where the guarantee ends is, once again, the mark of enforcement done thoughtfully rather than dogmatically.

5. Red — proving append-only holds

```
Scenario: an action is recorded atomically with its change
  Given a role change made inside a transaction with a RecordAsync beside it
  When the transaction commits
  Then exactly one audit event exists for the change
  And if the transaction had rolled back, NO audit event would exist

Scenario: audit rows cannot be modified or deleted by application code
  Given an existing audit event
  When application code sets a field / removes it and calls SaveChanges
  Then SaveChanges throws "append-only" and the row is unchanged

Scenario: dissolve may purge a tenant's trail (set-based)
  When a tenant is dissolved
  Then its audit events are removed via ExecuteDelete (the interceptor doesn't block it)
```

AuditLogTests drives these on real Postgres. The middle scenario is the one that makes it an audit log rather than "a table we promise not to edit" — it asserts the *build-enforced* immutability, the same way 2.6/2.7's tests assert tenancy is structural. The third pins the deliberate boundary so no one later "fixes" the interceptor to block dissolve and thereby breaks erasure.

6. Run it

```
dotnet test tests/Api.Tests --filter Audit
# live (once RBAC's role-change lands in 6.1): change a member's role, then query
# AuditEvents for the tenant → one "member.role_changed" row with the actor + from/to
# metadata
# try to UPDATE that row via EF → InvalidOperationException: append-only
```

7. Architecture Decision

ARCHITECTURE DECISION

The fork: how are security/compliance actions recorded? (a) Rely on operational logs (4.5); (b) an audit table that application convention "doesn't modify"; (c) an `ITenantScoped`, append-only `AuditEvent` written by explicit `RecordAsync` calls, staged on the caller's transaction, immutability enforced by a `SaveChanges` interceptor.

Chosen: (c) (ADR-008). Audit is *tenant data* (queryable, exportable, dissolvable — so `ITenantScoped`), *evidence* (so structurally immutable, not convention-immutable), and *accurate* (staged on the audited change's transaction, so it exists iff the action happened). Explicit recording keeps the trail semantic and curated. Every piece reuses a pattern already built — repository, clock, interceptor, unit of work — so it's a composition, not new machinery.

Rejected: (a) — logs are ephemeral, operator-owned, and mutable: unfit as evidence and invisible to the tenant. (b) — an "audit table by convention" is one careless `Update` from a corrupted trail, and a corrupted audit trail is *worse than none* (it looks authoritative while lying); structural enforcement is the whole point. Also rejected: auto-auditing every `SaveChanges` — it produces mechanical noise, not semantic events (§3).

The trade: developers must *remember* to call `RecordAsync` on security-relevant actions (the price of semantic-over-mechanical), and append-only means corrections are themselves new events (you never edit a mistake, you record a compensating action) — which is exactly how a real ledger works, and the point.

8. Checkpoint

```
git add -A && git commit -m "feat: append-only audit log – tenant-scoped, atomic, immutable by interceptor"
git tag lesson-4.6
```

Part 4 complete. The platform is now *reliable* and *legible*: side effects survive crashes (outbox), inbound duplicates are harmless (inbox), maintenance runs on a clock (scheduler), operators can see per-tenant what's happening (observability), and every consequential action is on an immutable record (audit). Everything that follows — monetization, RBAC, GDPR — stands on these five primitives.

You should now be able to: articulate the audit-vs-logs distinction across all five axes; explain why audit must be staged on the audited change's transaction; describe interceptor- enforced append-only immutability and why `ExecuteDelete` is a deliberate, safe exception; and defend explicit semantic recording over mechanical auto-audit.

Next: *Part 5 — Monetization.* The billing abstraction, Stripe checkout and webhooks (where the inbox and `EnterTenant` finally meet), and plan entitlements + quotas — turning the platform into a product that charges money.

Part 5 — Monetization

Lesson 5.1 — Billing abstraction & entitlements

Part 5 · Monetization. The platform can now do everything except the one thing a *business* must: charge money and gate features by what a tenant paid for. Part 5 adds that — and it opens with two decisions that keep you out of the two worst billing mistakes. First, **never touch card data**: payments happen on the provider's hosted checkout, behind a seam shaped exactly like `ISender`. Second, **the provider is the source of truth for money; your database holds a projection** — and every uncertainty about that projection fails *closed*, to the free plan. Money is the one place where "fail open" isn't a bug, it's a theft.

Goal: a tenant's plan is resolved server-side from a local `Subscription` projection that fails *closed* to `Free`; an endpoint can be gated behind a plan entitlement with one `.RequireEntitlement(...)` call that returns **402**; and the app boots and is testable with no Stripe account at all.

Concepts & principles: the payment-provider seam (never build card forms); projection vs source-of-truth; fail-closed entitlement resolution; the 402-vs-403-vs-401 distinction; plans as code/config not tenant data; the dev fake that *must not* run in production (R1).

Maps to: ADR-006 · R1 · repo: `src/Core/Abstractions/IBillingProvider.cs`, `src/Core/Billing/PlanCatalog.cs`, `src/Core/Entities/Subscription.cs`, `src/Core/Abstractions/IEntitlementService.cs`, `src/Api/Services/EntitlementService.cs`, `src/Api/Endpoints/EntitlementEndpointExtensions.cs`, `src/Infrastructure/Billing/FakeBillingProvider.cs`, the registration guard in `ServiceCollectionExtensions`, tests: `tests/Api.Tests/Billing/EntitlementServiceTests.cs`, `RequireEntitlementFilterTests.cs`, `BillingProviderRegistrationTests.cs`.

Prerequisites: 2.6 (tenant-scoped `Query()`), 2.3 (the seam pattern this mirrors), 1.5 (status-code semantics), 1.4 (config-gating).

1. Motivate — two ways billing goes catastrophically wrong

Mistake one: touching cards. The instant your server *sees* a card number, you're in PCI-DSS scope — audits, compliance, liability, and a breach that ends the company. The entire industry's answer is: don't. The provider (Stripe) hosts the checkout page; the card never touches your servers; you get back a signed *event* saying "this tenant now has Pro." So `IBillingProvider` never has a "charge card" method — it has "create a hosted checkout session" and "parse a webhook." Building a card form is not a feature you're declining to build; it's a liability you're declining to accept.

Mistake two: trusting your own database about money. It's tempting to flip a tenant to Pro when checkout *completes*. But completion and payment aren't the same event, refunds happen, cards expire mid-subscription, disputes reverse charges — and **the provider knows all of that; your database doesn't**. So your `Subscription` row is a *projection* — a cache of the provider's truth, updated by webhooks (5.2) — and whenever it's stale, absent, or

ambiguous, you resolve to **Free**. Granting a paid feature on uncertain state is giving away the product; the safe default when unsure is *less* access, never more.

Both mistakes have the same antidote as the rest of this course: a seam for the first, and fail-closed defaults for the second.

2. The provider seam — money mutations live elsewhere

```
// src/Core/Abstractions/IBillingProvider.cs (shape)
/// Abstraction over the payment provider (ADR-006), mirroring IEmailSender's
Core-abstraction
/// shape so Core stays provider-agnostic and tests run offline. ... Money mutations happen
on
/// the provider's hosted Checkout — we never build card forms or store card data.
public interface IBillingProvider
{
    Task<BillingCheckoutSession> CreateCheckoutSessionAsync(BillingCheckoutRequest r,
CancellationToken ct = default);
    Task<BillingPortalSession> CreatePortalSessionAsync(BillingPortalRequest r,
CancellationToken ct = default);
    BillingWebhookEvent? ParseWebhookEvent(string payload, string? signature); // 5.2
    Task CancelSubscriptionAsync(string stripeSubscriptionId, CancellationToken ct =
default); // dissolve (5.3)
}
```

It's the IEmailSender seam again (2.3), applied to payments: a Core interface speaking the domain's language (checkout, portal, webhook), with the vendor (Stripe) quarantined in Infrastructure. Read one sentence from the checkout method's doc especially: "Access is *NOT granted by completing checkout — only the webhook flips the subscription*." That's Mistake #2 encoded in the API — the checkout call returns a *URL to redirect to*, never a grant. The grant comes later, from a signed webhook, through the projection. Note too the CreateCheckoutSessionAsync carries the TenantId *into* the session, so the eventual webhook can reconcile back to the tenant — the thread that 5.2's EnterTenant will pull.

3. The projection — a cache that fails closed

```
// src/Core/Entities/Subscription.cs (shape)
/// A tenant's billing subscription — a local PROJECTION of the payment provider's state
/// (Stripe is the source of truth for money; ADR-006). A tenant has at most one (unique on
/// TenantId). Absence Free, and any non-active status fails closed to Free —
entitlements
/// are never granted on stale/uncertain state.
public class Subscription : ITenantScoped
{
    public Guid TenantId { get; set; }
    public required string PlanKey { get; set; }
    public required string Status { get; set; } // active | trialing | past_due | canceled
    public string? StripeCustomerId { get; set; }
    public string? StripeSubscriptionId { get; set; }
    public DateTimeOffset? CurrentPeriodEnd { get; set; } // past = lapsed, even if
Status=active
    public DateTimeOffset? LastEventAt { get; set; } // newest applied event — the 5.2
ordering guard
    // ... LapseNotifiedAt (5.3), timestamps
}
```

It's ITenantScoped (2.6 scopes it automatically — a tenant sees only its own subscription), and it carries the provider's ids so the portal and cancel calls can find the customer. Two fields foreshadow later lessons — LastEventAt is the strictly-newer guard 5.2 needs (webhooks arrive out of order), LapseNotifiedAt is 5.3's dunning bookkeeping — but the property that matters *now* is the one repeated everywhere: **absence** □ **Free, non-active** □ **Free, lapsed** □ **Free**. A Subscription is never a reason to grant; it's only ever a *possible* reason, checked live.

Where does a tenant's Free-ness live if absence means Free? Nowhere — and that's the point. There's no "Free subscription" row to create or forget; the *absence* of a paid projection is the Free state. One less thing to keep consistent.

4. Plans are code, not data

```
// src/Core/Billing/PlanCatalog.cs (shape)
public sealed record Plan(string Key, IReadOnlySet<string> Entitlements,
    int? SeatLimit = null, IReadOnlyDictionary<string, int>? UsageLimits = null)
{
    public bool Includes(string entitlementKey) => Entitlements.Contains(entitlementKey);
    // null SeatLimit / absent usage key    unlimited
}

public static class PlanCatalog
{
    {
        private static readonly IReadOnlyDictionary<string, Plan> Plans = new
        Dictionary<string, Plan>
        {
            [PlanKeys.Free] = new(PlanKeys.Free, new HashSet<string>(), SeatLimit: 3, ...),
            [PlanKeys.Pro] = new(PlanKeys.Pro, new HashSet<string> { Entitlements.ProFeature },
            SeatLimit: 10, ...),
        };
        /// Unknown keys fall back to Free (fail-closed).
        public static Plan Get(string planKey) => Plans.TryGetValue(planKey, out var p) ? p :
        Plans[PlanKeys.Free];
    }
}
```

Plans are **code/config, not tenant data** (ADR-006) — the catalog is a static dictionary, because a plan definition ("Pro grants X, caps seats at 10") is a property of *your product*, identical for every tenant, versioned with the code. What's per-tenant is only *which* plan they're on (the projection). And Get fails closed one more time: an unknown plan key resolves to Free, so a typo or a removed plan can never accidentally grant more than Free. (The Entitlements/UsageKeys here are marked EXAMPLE — a real app replaces them with its own gates; the quota fields, SeatLimit/UsageLimits, are 5.3's business.)

5. Entitlements — the server-side gate, failing closed

```
// src/Api/Services/EntitlementService.cs (core)
public async Task<bool> HasAsync(string entitlementKey, CancellationToken ct = default)
{
    var subscription = await subscriptions.Query().FirstOrDefaultAsync(ct); //
tenant-scoped (2.6)
    var planKey = ResolvePlanKey(subscription, clock.UtcNow());
    return PlanCatalog.Get(planKey).Includes(entitlementKey);
}

/// The active plan, or Free if the subscription is absent/inactive/lapsed (fail-closed).
internal static string ResolvePlanKey(Subscription? subscription, DateTimeOffset now)
{
    if (subscription is null) return PlanKeys.Free;
    var active = subscription.Status is SubscriptionStatus.Active or
SubscriptionStatus.Trialing
        && (subscription.CurrentPeriodEnd is null || subscription.CurrentPeriodEnd
> now);
    return active ? subscription.PlanKey : PlanKeys.Free;
}
```

ResolvePlanKey is the whole fail-closed doctrine in five lines, and it's internal static precisely so it can be unit-tested exhaustively without a database: null → Free, past_due → Free, canceled → Free, active-but-CurrentPeriodEnd-in-the-past → Free, only active/trialing-and-unexpired → the real plan. Note it checks the *period end* even when the status says active — because a status can be stale between webhooks, but a lapsed period is lapsed. The clock is injected (3.3), so "is it expired?" is testable. Every path that isn't unambiguously paid-and-current returns Free. **Never trust the client for entitlement** — this reads the server's projection, never a claim or a request field.

6. The 402 gate — a distinct answer for "your plan doesn't include this"

```
// src/Api/Endpoints/EntitlementEndpointExtensions.cs (core)
public static TBuilder RequireEntitlement<TBuilder>(this TBuilder builder, string
entitlementKey)
    where TBuilder : IEndpointConventionBuilder
    => builder.AddEndpointFilter(async (context, next) =>
    {
        var entitlements =
context.HttpContext.RequestServices.GetRequiredService<IEntitlementService>();
        if (await entitlements.HasAsync(entitlementKey,
context.HttpContext.RequestAborted))
            return await next(context);

        var app =
context.HttpContext.RequestServices.GetRequiredService<IApplicationSettings>();
        return Results.Json(new
        {
            error = "payment_required",
            entitlement = entitlementKey,
            message = $"This feature requires a plan that includes '{entitlementKey}'.",
            upgrade_url = $"{app.ClientUrl.TrimEnd('/')}/billing",
        }, statusCode: StatusCodes.Status402PaymentRequired);
    });
```

A feature slice gates a paid capability by adding .RequireEntitlement(Entitlements.Xxx) beside its MapTenantFeatureGroup (3.2) — the same one-liner-guardrail style as

.RequirePermission (6.1). The decision worth dwelling on is the **status code**. Recall 1.5's three "no" answers, and now the set is complete:

- **401 Unauthorized** — "we don't know who you are" (2.1). Fix: sign in.
- **403 Forbidden** — "we know you; your *role* may not do this" (6.1). Fix: get permission.
- **402 Payment Required** — "we know you; your *plan* doesn't include this" (here). Fix:

upgrade — so the response names the missing entitlement and includes an `upgrade_url`.

Three different problems, three different client reactions — and only by distinguishing them can the UI do the right thing (bounce to login / hide the button / show the upgrade page). Returning 403 for an unpaid feature would send the user to a permission dead-end instead of the billing page; returning 500 would look like a bug instead of a sales opportunity. This is 1.5's "the status code is the first bit of meaning" paying off — the E2E journey in 5.3 asserts that a 402 becomes a visible upgrade prompt, the whole point of choosing the right code.

7. The dev fake — and the rule that it must never ship

The app must boot and be testable with no Stripe account, so there's a `FakeBillingProvider`: deterministic fake checkout URLs, recorded requests, and a `ParseWebhookEvent` that trusts a literal "valid" signature so webhook flows can be driven offline. Convenient — and *dangerous if it ever ran in production*, because trusting a literal signature means accepting **forged, unauthenticated, cross-tenant billing writes**. So the registration fails closed at startup (R1):

```
// src/Infrastructure/ServiceCollectionExtensions.cs — the R1 guard
if (!string.IsNullOrEmpty(configuration["Billing:Stripe:SecretKey"]))
    services.AddScoped<IBillingProvider, StripeBillingProvider>();
else if (environment.IsDevelopment())
    services.AddScoped<IBillingProvider, FakeBillingProvider>();
else
    throw new InvalidOperationException(
        "No Billing:Stripe:SecretKey is configured and the environment is not Development.
The " +
        "in-memory FakeBillingProvider trusts a literal webhook signature and must never
run " +
        "outside Development (it would accept forged ... cross-tenant billing writes).
...");
```

Read the `else`: a production deploy with no Stripe key **refuses to boot** rather than silently falling back to the fake. This is the fail-fast-config doctrine (1.4) pointed at the highest-stakes case — better a loud startup crash than a live system accepting forged billing events. It's rule R1, machine-checked (`BillingProviderRegistrationTests`), and it's the sharpest example in the course of "the safe default when unsure is *refuse*, not *proceed*." A convenience that would be catastrophic in production is guarded so it *cannot* reach production.

8. Red — the tests

```
Scenario: entitlement resolution fails closed
  Given no subscription | a past_due one | an active-but-expired one
  Then the tenant resolves to Free and lacks the Pro entitlement
  Given an active, unexpired Pro subscription
  Then the tenant has the Pro entitlement

Scenario: a gated endpoint returns 402 for an unentitled tenant
  Given a Free tenant hitting a .RequireEntitlement("pro-feature") endpoint
  Then the response is 402 with error "payment_required", the entitlement named, and an
  upgrade_url

Scenario: the fake provider must not register outside Development
  Given no Stripe key and a non-Development environment
  Then app startup throws (the fake can never serve production)
```

EntitlementServiceTests hammers ResolvePlanKey's fail-closed matrix (pure, fast, no DB via the internal static split); RequireEntitlementFilterTests asserts the 402 shape; BillingProviderRegistrationTests pins R1. The fail-closed matrix is the one to write exhaustively — every "not clearly paid" input must land on Free.

9. Run it

```
dotnet test tests/Api.Tests --filter Billing
dotnet run --project src/Api # boots with the FakeBillingProvider (no Stripe key needed in
dev)
# hit a .RequireEntitlement endpoint as a Free tenant → 402 { payment_required, upgrade_url
}
```

10. Architecture Decision

ARCHITECTURE DECISION

The fork: how does the platform handle payments and paid gating? (a) Build card capture + charge logic in-house; (b) call the Stripe SDK directly throughout, treat the local DB as the truth; (c) a provider seam (hosted checkout, never card data) + a local `Subscription` *projection* that fails closed to Free + server-side entitlement gates returning 402.

Chosen: (c) (ADR-006). The seam keeps Core provider-agnostic and tests offline (mirroring `ISender`); hosted checkout keeps card data — and PCI scope — entirely out of the platform; projection-not-truth + fail-closed-to-Free means uncertainty costs the *tenant* a feature, never the *business* revenue; 402 gives the client a distinct, actionable answer. The dev fake keeps the app bootable with no Stripe, guarded by R1 so it can't reach production.

Rejected: (a) — building card handling is accepting PCI liability and a breach surface for zero differentiation; nobody should. (b) — a direct-SDK, DB-is-truth design welds Stripe into every call site (anti-seam) *and* trusts local state about money, which drifts from the provider on every refund/dispute/expiry and fails *open* (grants access on stale "active" rows). Both rejected mistakes are the two this lesson opened with.

The trade: access is *eventually* consistent — completing checkout doesn't grant the feature; the webhook does (5.2), so there's a beat of latency and a hard dependency on webhook delivery. Accepted: it's the only design where your grant reflects the provider's actual money state. And the projection must be *reconciled* carefully against out-of-order webhooks — exactly the job of the next lesson.

11. Checkpoint

```
git add -A && git commit -m "feat: billing abstraction + entitlements – provider seam, fail-closed projection, 402 gate"
git tag lesson-5.1
```

You should now be able to: explain why the platform never sees card data and what `IBillingProvider` therefore does and doesn't expose; articulate projection-vs-source-of-truth and why every ambiguity resolves to Free; place 402 among 401/403 and say what client reaction each drives; and justify R1's refuse-to-boot guard as fail-closed applied to money.

Next: Lesson 5.2 — *Stripe: checkout, webhook, portal* — where the projection actually gets updated, and the inbox (4.3) and `EnterTenant` (2.7) finally meet in the webhook handler.

Lesson 5.2 — Stripe: checkout, webhook, portal

Part 5 · Monetization. This is where the projection from 5.1 actually gets written — and it's the lesson the whole reliability toolkit was quietly building toward. The billing webhook handler is the course's most satisfying *composition*: signature verification, the inbox (4.3), `EnterTenant` (2.7), the strictly-newer recency guard, an atomic transaction (3.1), and a dunning notification all cooperate in one method to turn an untrusted, at-least-once, out-of-order HTTP POST into a correct, exactly-once update of a tenant's plan. If Parts 2 and 4 felt like isolated primitives, this is where you see why they were built.

Goal: a tenant can start a hosted checkout and open a management portal; and an inbound Stripe webhook safely reconciles the `Subscription` projection — authentic-only, deduped, tenant-scoped, never regressed by a stale event, and atomic with its notification.

Concepts & principles: hosted checkout/portal (redirect, don't capture); webhook as the *only* source of grants; signature authenticity; inbox dedup on the event id; `EnterTenant` for a JWT-less but authenticated write; the out-of-order recency guard (strictly-newer); one-transaction claim+apply+notify; observable rejection of forged webhooks.

Maps to: ADR-006, ADR-007, ADR-003 · repo: `src/Api/Services/BillingService.cs`, `BillingWebhookHandler.cs`, `src/Api/Controllers/BillingWebhookController.cs`, `BillingController.cs`, `src/Infrastructure/Billing/StripeBillingProvider.cs`, tests: `tests/Api.Tests/Billing/BillingWebhookHandlerTests.cs`, `BillingWebhookControllerTests.cs`, `BillingServiceTests.cs`.

Prerequisites: 5.1 (the seam + projection), 4.3 (inbox), 2.7 (`EnterTenant`), 3.1 (unit of work).

1. Motivate — three interactions, one hard one

Billing has three moving parts, and their difficulty is wildly uneven:

- **Checkout** (start a subscription) and **portal** (manage/cancel) are *easy*: create a

hosted session at the provider, return its URL, redirect the user. The provider does the hard, dangerous work (card capture); you hand off and step back.

- **The webhook** is *hard*, because it's where the money-truth flows back *into* your

system — and it arrives as an untrusted HTTP POST from the public internet, delivered at-least-once, possibly out of order, possibly forged. Everything that makes billing correct or catastrophic lives in how you handle that one request.

So this lesson spends a paragraph each on the easy two and its whole heart on the webhook.

2. Checkout & portal — redirect, grant nothing

```
// src/Api/Services/BillingService.cs (core)
public async Task<CheckoutResult> CreateCheckoutAsync(Guid tenantId, string? planKey,
CancellationTokens ct = default)
{
    if (!IsPurchasablePlan(planKey)) return new(CheckoutOutcome.InvalidPlan, null);
    var baseUrl = app.ClientUrl.TrimEnd('/');
    var session = await billing.CreateCheckoutSessionAsync(
        new BillingCheckoutRequest(tenantId, planKey!, $"{baseUrl}/billing/success",
        $"{baseUrl}/billing/cancel"), ct);
    return new(CheckoutOutcome.Created, session.Url); // ← a URL to redirect to. No
    Subscription written.
}
```

The load-bearing thing is what this method *doesn't* do: it writes **no** Subscription. Its doc comment is blunt — “Checkout writes no Subscription — access is granted only by the BILLING-3 webhook.” The client redirects the user to `session.Url`, the user pays *on Stripe*, and the grant arrives later, through the webhook (§3). Same for the portal: it looks up the tenant's `StripeCustomerId` from the projection (via `tenant-scoped Query()`, 2.6), opens a hosted portal, and any change the user makes there — upgrade, cancel — also comes back as a webhook. There is exactly one write path for subscription state, and it isn't here. `IsPurchasablePlan` fails closed too: only a real, non-Free catalog plan can start a checkout (unknown keys resolve to Free in 5.1, so they're rejected here). Owner-only access is enforced by the controller's `TenantApiControllerBase` (6.1) — checkout is not a thing any member may initiate.

3. The webhook handler — six primitives, one method

Here it is, the composition. Read it once whole, then we'll walk it — every line is a callback to something you built:

```
// src/Api/Services/BillingWebhookHandler.cs (the flow)
public async Task<WebhookResult> HandleAsync(string payload, string? signature,
Cancellation token ct = default)
{
    BillingWebhookEvent? evt;
    try { evt = provider.ParseWebhookEvent(payload, signature); } // 1. AUTHENTICITY
    catch (BillingWebhookSignatureException) { return WebhookResult.InvalidSignature; }
    if (evt is null) return WebhookResult.Ignored; // authentic but irrelevant

    await using var transaction = await unitOfWork.BeginTransactionAsync(ct); // 3.1

    // 2. DEDUP: claim the event id so a redelivery is a no-op (4.3)
    if (!await inbox.TryClaimAsync("stripe", evt.EventId, ct))
    {
        await transaction.CommitAsync(ct);
        return WebhookResult.Duplicate;
    }

    // 3. ENTER TENANT: no JWT here — enter the signature-authenticated tenant so the write
    is
    // stamped + scoped by the normal interceptor/filter, not the escape hatch (2.7)
    bool applied;
    using (tenantContext.EnterTenant(evt.TenantId))
    {
        string? previousStatus;
        (applied, previousStatus) = await UpsertSubscriptionAsync(evt, ct); // 4. RECENCY
    GUARD + upsert
        if (applied)
            await MaybeNotifyDunningAsync(evt, previousStatus, ct); // 5. dunning (5.3)
        await transaction.CommitAsync(ct); // 6. claim + projection + notification commit
    ATOMICALLY
    }

    return applied ? WebhookResult.Applied : WebhookResult.Ignored;
}
```

Walk the six steps as the reunion they are:

1 Authenticity first. A webhook is an anonymous public POST — anyone can send one. So

before anything else, `ParseWebhookEvent` verifies the provider signature; a bad signature returns `InvalidSignature` and applies **nothing**. This is the gate that makes the rest safe to trust: after this line, `evt` is genuinely from Stripe. (`null` means authentic-but-irrelevant — an event type we don't care about — acknowledged, not applied.)

1 Dedup via the inbox (4.3). Stripe delivers at-least-once, so `TryClaimAsync("stripe",`

`evt.EventId)` claims the event id; a redelivery finds it claimed and short-circuits to `Duplicate`. This is the inbox's *first real customer* — and note it's used *exactly* as 4.3's correctness rule demanded: the claim is inside the transaction with the work it guards, so if the apply fails and rolls back, the claim rolls back too and the redelivery re-processes. The billing webhook is why the inbox exists.

1 ****EnterTenant (2.7).** This is the reunion 4.3 and 2.7 were both pointing at. The

webhook has **no JWT** — Stripe doesn't carry your access token — but it *does* carry a trusted tenant id (the metadata you stamped at checkout in 5.1, inside a signature-verified event). So the handler *enters* that tenant, and now the write goes through the **normal** tenant interceptor

and query filter (2.6/2.7): the Subscription is stamped to the entered tenant, scoped reads see only its row — the same structural isolation an authenticated request gets, achieved without a JWT. It does *not* reach for the QueryAllTenants() escape hatch. EnterTenant "scopes, doesn't authorize" (2.7) — the signature did the authorizing; this does the scoping.

1 The recency guard (inside UpsertSubscriptionAsync) — the subtlest correctness

point, and one teams miss:

```
// apply only STRICTLY-NEWER events – webhooks are unordered
if (subscription is not null && subscription.LastEventAt is { } last && evt.OccurredAt <=
last)
    return (false, subscription.Status); // stale/out-of-order → acknowledged, NOT applied
```

Webhooks are not just at-least-once, they're **out of order**. Stripe can deliver a stale canceled event *after* a live active one (retries, network reordering). Without a guard, that stale event would clobber the tenant back to canceled — silently downgrading a paying customer. The guard: the projection remembers LastEventAt (the timestamp of the newest event *applied*), and an incoming event is applied only if it's **strictly newer**. A stale event is acknowledged (its id is claimed, so it won't be retried) but *not applied* — state never regresses. This is the projection-vs-truth discipline from 5.1 enforced against the messiness of real delivery: your cache converges to the provider's latest truth regardless of the order the news arrives in.

5 & 6. **Notify on transition, commit atomically.** If the event was applied *and* the status transitioned into a bad state (past_due/canceled), a dunning notification is staged (5.3) — only on a real transition, never on a stale event or a same-status redelivery. Then the whole thing commits in **one transaction**: the inbox claim, the projection write, and the notification are atomic. Either all three durably happened or none did — so there's never a claimed-but-unapplied event, or an applied change with a lost notification, or a notification for a change that rolled back. Every reliability lesson in Part 4 is load-bearing in this single commit.

4. The endpoint — a system surface that logs its attackers

```
// src/Api/Controllers/BillingWebhookController.cs (core)
[ApiController, AllowAnonymous, Route("api/billing/webhook")]
public class BillingWebhookController(BillingWebhookHandler handler, ILogger<...> logger) :
ControllerBase
{
    [HttpPost] public async Task<IActionResult> Receive(Cancellation token ct)
    {
        var payload = await new StreamReader(Request.Body).ReadToEndAsync(ct);
        var signature = Request.Headers["Stripe-Signature"].ToString();
        var result = await handler.HandleAsync(payload, signature, ct);
        if (result == WebhookResult.InvalidSignature)
        {
            var sourceIp = HttpContext.Connection.RemoteIpAddress?.ToString() ?? "unknown";
            logger.LogWarning("Rejected billing webhook with an invalid signature from
{SourceIp}.", sourceIp);
            return BadRequest(new { error = "invalid_signature" });
        }
        return Ok(); // Applied / Duplicate / Ignored all acknowledge → provider stops
retrying
    }
}
```

Three deliberate choices:

- **It's [AllowAnonymous] and does *not* derive from TenantApiControllerBase.** It's a

system endpoint — the provider calls it with no JWT — so it's the sanctioned exception to "every controller inherits the tenant/admin base" (the R4 arch test's allowlist, from 3.3).

Authenticity is the *signature*, not a token; the handler enters the tenant itself.

- **A rejected signature is logged with the source IP** (never the payload). An anonymous

public endpoint receiving a forged webhook is a *probe* — surfacing it as a warning with the source makes the attack observable instead of a silent 400. (The same "log the source of a rejected inbound" instinct as the SSRF/webhook rules, R19.)

- **Applied / Duplicate / Ignored all return 200.** The provider only cares "did you

receive it?" — retrying is *its* at-least-once mechanism, and a duplicate or an irrelevant event is still "received." Only an *inauthentic* payload gets a 400 (and won't be retried, correctly — it was never a real event).

5. The provider implementation — Stripe behind the seam

StripeBillingProvider (Infrastructure) implements IBillingProvider using the Stripe SDK: CreateCheckoutSessionAsync builds a Checkout session stamping the TenantId into metadata; ParseWebhookEvent calls Stripe's signature verification (with the webhook secret from config) and maps the Stripe event to the provider-agnostic BillingWebhookEvent. The point isn't its detail — it's that *all* the Stripe-specific knowledge (SDK types, event-name strings, price-id ↔ plan-key mapping, signature crypto) lives in this one file, behind the seam, exactly as MailKit lives behind SmtplibEmailSender (2.3). The handler in §3 never sees a Stripe type; it works against BillingWebhookEvent. Swap providers and §3 doesn't change.

6. Red — driving the whole flow with the fake

The FakeBillingProvider (5.1) makes this testable with no Stripe: it treats the literal signature "valid" as authentic and deserializes the payload as a BillingWebhookEvent, so tests drive the handler end-to-end offline.

```
Scenario: an authentic activation applies and grants the plan
  Given a signed "active/pro" event for tenant T
  When the webhook is handled
  Then T's Subscription projection is Pro/active and T now has the Pro entitlement (5.1)
```

```
Scenario: a redelivery is deduped
  When the same event id is delivered twice
  Then the second is Duplicate and applied nothing
```

```
Scenario: a stale out-of-order event does not regress state (the recency guard)
  Given T is active (LastEventAt = t2)
  When a canceled event with OccurredAt = t1 < t2 arrives
  Then it is acknowledged but NOT applied – T stays active
```

```
Scenario: a forged webhook is rejected and logged
  Given a bad signature
  Then the result is InvalidSignature → 400, the source IP is logged, nothing changed
```

The third scenario is the one that separates a correct handler from a plausible-but-wrong one — `BillingWebhookHandlerTests` drives it with `FakeTimeProvider` timestamps to make "older" and "newer" exact. Write that one first; it's the bug most billing integrations ship.

7. Run it

```
dotnet test tests/Api.Tests --filter Billing
# live (fake provider, dev): POST a signed event to /api/billing/webhook →
# the Subscription projection flips to active; a Free-gated endpoint now returns 200 (5.1's
# 402 is gone)
# for real Stripe: `stripe listen --forward-to localhost:7160/api/billing/webhook` replays
# real events
```

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how does subscription state get into the platform? (a) Flip the projection when the checkout *success* redirect returns; (b) poll the provider's API periodically; (c) the provider's *webhook* is the sole write path — verified, deduped via the inbox, applied under `EnterTenant`, guarded against out-of-order delivery, atomic with its notification.

Chosen: (c) (ADR-006 + ADR-007 + ADR-003). The webhook is the provider's own truth pushed to you, so it's the only trustworthy source of a grant; verifying, deduping, and recency-guarding it turns an untrusted, at-least-once, unordered POST into an exactly-once, monotonic projection update; `EnterTenant` gives the JWT-less write the full isolation regime; one transaction keeps claim+apply+notify consistent. It's the reason the inbox and `EnterTenant` were built.

Rejected: (a) — the success redirect is *client-controlled* and fires before payment settles; trusting it grants access to anyone who can reach the success URL, and misses every later change (renewal, failure, cancel). (b) — polling is wasteful, laggy, and still needs all the same reconciliation logic; the webhook is push instead of pull for the same information.

The trade: grants are *eventually* consistent (a webhook round-trip after payment) and the platform depends on webhook delivery — a missed webhook means a stale projection until the next event or a manual re-sync. Accepted: it's the only design that reflects the provider's actual money-truth, and the recency guard + inbox make redelivery (the provider's own recovery mechanism) safe rather than dangerous.

9. Checkpoint

```
git add -A && git commit -m "feat: Stripe checkout/portal + webhook - inbox dedup,
EnterTenant, recency-guarded projection"
git tag lesson-5.2
```

You should now be able to: explain why checkout grants nothing and the webhook is the sole write path; trace the six primitives the handler composes and name the earlier lesson each comes from; articulate the strictly-newer recency guard and the silent-downgrade bug it prevents; and justify the anonymous, signature-authenticated, IP-logging webhook endpoint.

Next: *Lesson 5.3 — Quotas, dunning & dissolve* — the countable limits (seats, metered usage) enforced atomically under concurrency, the lapse sweep, and cleaning up a paid subscription when a tenant is dissolved.

Lesson 5.3 — Quotas, dunning & dissolve

Part 5 · Monetization — finale. Entitlements (5.1) answer a *yes/no* question: does the plan include this feature? Quotas answer a *how many* question: seats used, exports this month — and "how many" is where **concurrency** bites. Two requests consuming the last unit of a quota at the same instant is the classic lost-update race, and getting it wrong lets tenants exceed limits they paid to stay under. This lesson makes quota consumption *atomic under concurrency* (the centerpiece), then closes the billing lifecycle: the lapse sweep that nudges a lapsed tenant, and the dissolve cleanup that stops charging a customer whose tenant no longer exists — both composed from primitives you already built.

Goal: seat and metered-usage limits enforced from the plan, with metered consumption that cannot be raced past the cap; a scheduled sweep that dun-notifies a lapsed tenant exactly once; and billing that cleans itself up on tenant dissolve (cancel the provider subscription, wipe the projection) without blocking a solo owner from leaving.

Concepts & principles: countable quotas vs boolean entitlements; the reserved-seat model; atomic check-and-consume via a conditional UPDATE (the lost-update race and its fix, R30); the first-row-insert race + unique-index retry; the lapse sweep on the scheduler; billing's dissolve contributor + provider-cancel-via-outbox.

Maps to: ADR-006, R30 · repo: `src/Api/Services/QuotaService.cs`,
`src/Core/Entities/UsageCounter.cs`, `src/Api/Services/SubscriptionLapseSweepJob.cs`,
`BillingNotifier.cs`, `BillingDataContributor.cs`,
`src/Infrastructure/Billing/BillingCancelOutboxHandler.cs`, tests:
`tests/Api.Tests/Billing/QuotaServiceTests.cs`, `SubscriptionLapseSweepJobTests.cs`,
`BillingDissolveTests.cs`.

Prerequisites: 5.1 (plan resolution, PlanCatalog quotas), 4.4 (scheduler), 4.1/4.2 (outbox), 2.9 (the contributor seam), 1.3 (real Postgres for the concurrency).

1. Motivate — the last seat, sold twice

A plan caps a tenant at 10 seats. Nine are used. Two admins each invite someone in the same instant. The naive enforcement — *read* the count (9), *check* (9 < 10, OK), *write* (now 10) — runs twice concurrently: both read 9, both pass the check, both write, and the tenant now has 11 seats. That's a **lost update**, and for a quota it means a paying limit silently breached — the tenant got something they didn't pay for, and your revenue model has a hole. The same race applies to metered usage (the 100th export sold twice).

"Check then act" is never atomic across concurrent callers unless the database makes it so. The fix is to fold the check *into* the write — a single conditional statement the database executes atomically — so the limit is enforced by row-locking, not by hope. That's rule R30 ("quota consumption is atomic under concurrency"), and it's the heart of this lesson.

2. Two kinds of quota — seats and metered usage

Quotas come from the plan (5.1's `PlanCatalog`), and null/absent means *unlimited* — so the whole system is inert until a plan sets a real number:

- **Seats** are a *live count*: members + pending invites. `QuotaService.GetSeatUsageAsync`

computes it fresh (never a stored counter that drifts — Golden Rule 4). The subtle part is *pending invites count as seats*:

```
// src/Api/Services/QuotaService.cs
var members = (await tenants.GetMembersAsync(tenantId, ct)).Count;
var pending = (await invitations.GetPendingForTenantAsync(tenantId, ct)).Count;
return new SeatUsage(members + pending, plan.SeatLimit);
```

A pending invite is a **reserved seat** — otherwise a tenant at the cap could send N invitations (each individually "fits" against the current member count) and over-provision when they're all accepted. Counting pending invites reserves the seat at invite time, so the cap holds against the *committed future*, not just the present. (This is why the invite path from 2.9 calls `CanAddSeatsAsync` before creating an invite → a 402 when full.)

- **Metered usage** is a *monthly counter*: a `UsageCounter` row keyed `(TenantId, Key,

Period) where `Period` is "yyyy-MM". Each metered action (an export, say) calls `TryConsumeAsync(key)` which increments the counter if it stays under the cap. This is the one with the race.

3. Atomic consume — the check is the write

Here's the centerpiece. `TryConsumeAsync`, and specifically its increment:

```
// src/Api/Services/QuotaService.cs — the atomic increment
private async Task<bool> TryIncrementAsync(string usageKey, string period, int amount, int
limit, DateTimeOffset now, CancellationToken ct) =>
    await usage.Query()
        .Where(c => c.Key == usageKey && c.Period == period && c.Count + amount <= limit)
// check...
        .ExecuteUpdateAsync(s => s
            .SetProperty(c => c.Count, c => c.Count + amount) // ...and act, in ONE statement
            .SetProperty(c => c.UpdatedAt, now), ct) == 1; // affected 1 row? → consumed
```

Read what makes this correct: the limit check (`c.Count + amount <= limit`) is in the ****WHERE** clause of the `UPDATE` itself**. The database evaluates it while holding a row lock, so two concurrent consumers *serialize* on that row — the first increments 9→10, the second re-evaluates `10 + 1 <= 10` (false), matches no row, and `ExecuteUpdateAsync` returns 0 → not consumed. There is no window between check and write for a second caller to slip through, because there is no *between* — it's one atomic statement. The `== 1` (did we update exactly one row?) is the answer to "did we get to consume?" This is the check-then-act race dissolved by moving the check server-side into the write, and it needs real Postgres row-locking — the exact reason 1.3 banned the in-memory provider, cashed one last time in Part 5.

There's a second, subtler race the full method handles — the *first* consume of a period, when no counter row exists yet:

```
// (abridged) after TryIncrementAsync returns false with no existing row:
var row = new UsageCounter { Key = usageKey, Period = period, Count = amount, UpdatedAt =
now };
try
{
    await usage.AddAsync(row, ct);
    await usage.SaveChangesAsync(ct); // may collide with a concurrent first-insert
    return true;
}
catch (DbUpdateException)
{
    usage.Remove(row); // detach the failed insert
    return await TryIncrementAsync(usageKey, period, amount, limit, now, ct); // retry
    against the now-existing row
}
```

Two requests both seeing "no row yet" would both try to INSERT — and the ****unique index on (TenantId, Key, Period)**** lets exactly one win; the loser catches `DbUpdateException`, detaches its doomed insert, and *retries the conditional increment* against the row that now exists. So even the create-the-first-row case is race-safe, via the same constraint-serialization idea as the inbox (4.3): let the database's unique index be the arbiter, then handle the loser gracefully. Belt (conditional update) and suspenders (unique index + retry).

4. Dunning — nudging a lapsed tenant, once

Payments fail; cards expire. "Dunning" is the polite process of telling a customer their subscription is in trouble. There are two triggers, and you've seen one:

- **Transition-driven** (5.2): when a webhook flips the projection *into* `past_due` or

canceled, `MaybeNotifyDunningAsync` notifies the owner — on the transition only.

- **Time-driven** (this lesson): a subscription can simply *lapse* — its `CurrentPeriodEnd`

passes with no renewal webhook. No event fires, so nothing transition-driven catches it. `SubscriptionLapseSweepJob` is a scheduled job (4.4) that periodically finds subscriptions whose period has ended and nudges the owner **once per lapse**, tracked by `Subscription.LapseNotifiedAt` (5.1): the sweep notifies only if `LapseNotifiedAt` predates the current `CurrentPeriodEnd`, then stamps it — so it fires once, not every run, and re-arms naturally when a new period starts. It composes 4.4 (the schedule) + the `IBillingNotifier` seam — no new machinery, just the reliability primitives pointed at a billing need.

Forward reference — the notifier is a seam, and delivery is deferred. Both dunning paths depend on `IBillingNotifier`, but its *real* implementation — fanning a notification out to the owner's in-app center through the outbox — belongs to the notification epic you'll build in **Part 7**. That's not a problem, and *why* it isn't is the point: `IBillingNotifier` is a Core abstraction (the seam pattern again, 2.3), so you build every *caller* now — `MaybeNotifyDunningAsync`, `SubscriptionLapseSweepJob` — against the interface, and supply a **minimal stub** for the delivery. A faithful stub at this stage records the dunning event on the append-only audit log (4.6): it's tenant-scoped and staged on the caller's unit of work, so it commits atomically with the projection write / lapse stamp — *exactly* the transactional shape the real NOTIFY fan-out will have. When Part 7 lands, you swap only the one implementation class; not a single caller changes. This is the seam earning its keep a second time (after email in 2.3/4.2): it lets a later-taught subsystem be *called* before it's *built*, so the billing lifecycle is complete now and gets richer delivery later without a rewrite. Test the "notify on transition, once" logic against a recording test double, decoupling it from whatever delivers.

The "once per lapse" bookkeeping is the whole trick: a sweep that runs hourly must not email the owner hourly, so the *last-notified* timestamp gates it — the same advance-a-timestamp-to-avoid-repeating instinct as the scheduler's own last-run tracking (4.4).

5. Dissolve — stop charging a tenant that no longer exists

When a tenant is dissolved (2.9), its billing must clean up — or Stripe keeps charging a customer whose tenant is gone. Billing plugs into dissolve through the **contributor seam** (2.9), and its implementation makes three instructive choices:

```
// src/Api/Services/BillingDataContributor.cs (core)
public Task<bool> HasDataAsync(Guid tenantId, CancellationToken ct = default) =>
    Task.FromResult(false); // (1) billing is not "abandonable content"

public async Task WipeAsync(Guid tenantId, CancellationToken ct = default)
{
    var subscription = await subscriptions.QueryAllTenants()
        .FirstOrDefaultAsync(s => s.TenantId == tenantId, ct); // (2) cross-tenant: the
    audited hatch
    if (subscription is null) return;

    if (!string.IsNullOrEmpty(subscription.StripeSubscriptionId))
    {
        await outbox.EnqueueAsync(BillingCancelOutboxHandler.MessageType, // (3) cancel via
    the outbox
        JsonSerializer.Serialize(new
    BillingCancelPayload(subscription.StripeSubscriptionId)), tenantId, ct);
        await subscriptions.SaveChangesAsync(ct);
    }
    await subscriptions.QueryAllTenants().Where(s => s.TenantId ==
    tenantId).ExecuteDeleteAsync(ct);
}

public string ExportKey => "billing";
public async Task<object?> ExportAsync(...) => /* plan/status/period only – never Stripe
ids or card data */;
```

Each choice is a lesson recalled:

1 ****HasDataAsync returns false.**** Recall 2.9: a solo owner leaving is blocked if a

contributor reports abandonable *data*. Billing is *plumbing*, not tenant content — a subscription isn't something the owner would lose regret. So it never blocks the leave; it just cleans itself up. Deciding what counts as "the user's data" versus "our plumbing" is a real modeling judgment, and this is the platform making it explicitly.

1 ****QueryAllTenants()** — and it's allowed here** because this is a **DataContributor.cs*

(3.2's precise rule: the escape hatch is permitted only in contributors, dissolve runs for a tenant other than the current one). Re-constrained to the target tenant, exactly as the rule requires.

1 **The provider cancel goes through the outbox (4.1), not an inline call.** This is the

sharpest choice. Canceling at Stripe is an *external* call — slow, fallible — and it's happening *inside the dissolve transaction*. Making that call inline would either block the teardown on Stripe's latency or, worse, risk a partial dissolve if the call throws mid-transaction. So instead the cancel is **enqueued** on the outbox: staged with the dissolve (atomic — the cancel-intent commits iff the dissolve does), then delivered out-of-band with retry, and the provider cancel is *idempotent* (5.1) so at-least-once redelivery is safe. This is the outbox's whole reason for being (4.1: don't do external effects inline in a transaction) applied to the one external effect in teardown. The *SaveChanges* flushes the staged message before the set-based delete runs, so it survives the dissolve commit.

ExportAsync (GDPR, 7.1) returns plan/status/period — **never the Stripe ids or anything card-related** (5.1's no-card-data rule extends to exports). And its *ExportKey* is "billing", unique per the R13 arch test (3.3).

6. Red — the tests, concurrency first

```
Scenario: metered usage cannot be raced past the cap (R30)
  Given a Free tenant with export cap 3 and 2 already used
  When 5 concurrent TryConsume("export") calls race
  Then exactly 1 succeeds (3rd unit) and 4 are refused — never 4 successes
```

```
Scenario: a pending invite reserves a seat
  Given a tenant at its seat cap via members + pending invites
  Then CanAddSeats(1) is false (a further invite is refused with 402)
```

```
Scenario: the lapse sweep nudges once per lapse
  Given a subscription whose CurrentPeriodEnd has passed and was never lapse-notified
  When the sweep runs twice
  Then the owner is notified exactly once (LapseNotifiedAt gates the second run)
```

```
Scenario: dissolve cancels the provider subscription and wipes the projection
  When a tenant with a live Stripe subscription is dissolved
  Then a provider-cancel is enqueued on the outbox and the Subscription row is gone
  And billing never blocked the (solo-owner) leave
```

The first scenario is the R30 correctness test and *must* run on real Postgres — *QuotaServiceTests* fires genuinely concurrent *TryConsumeAsync* calls and asserts the count never exceeds the cap. A test on a faked database would show four false successes and call it

green — the exact lie 1.3's decision exists to prevent. Write that one first and watch it *fail* against a naive read-check-write to feel why the atomic path matters.

7. Run it

```
dotnet test tests/Api.Tests --filter Billing
# live: consume a metered action past the Free cap → refused; invite past the seat cap → 402;
# dissolve a tenant → OutboxMessages shows a billing-cancel; the Subscription row is gone
```

8. Architecture Decision

ARCHITECTURE DECISION

The fork: how are countable quotas enforced, and how does billing clean up on dissolve? (a) Read-check-write in app code; (b) an in-app lock/mutex around consumption; (c) an atomic conditional UPDATE (check in the WHERE) + unique-index retry for the first insert, plus dissolve via the contributor seam with the provider cancel enqueued on the outbox.

Chosen: (c) (ADR-006, R30). The conditional update makes check-and-consume a single row-locked statement — race-free by the database, not by app coordination; the unique index arbitrates the first-insert race with a graceful retry. Dissolve reuses the contributor seam (2.9) and pushes the one external effect (Stripe cancel) onto the outbox (4.1) so teardown stays fast, atomic, and retry-safe. Pending-invite-as-reserved-seat makes the seat cap hold against committed-but-unaccepted invites.

Rejected: (a) — the lost-update race from §1; it *looks* correct in single-threaded tests and breaks under real concurrency (which is why the R30 test is concurrent and Postgres-backed). (b) — an app-level lock doesn't work across multiple instances (each has its own mutex) and reintroduces the coordination the database already does correctly; pushing concurrency control to the DB is both correct and free. Inline Stripe-cancel in the dissolve transaction — rejected for the outbox, per 4.1's whole thesis.

The trade: the atomic path is less obvious than read-check-write (a conditional UPDATE whose return value is the decision) and leans on Postgres semantics — but it's the only version that's correct under concurrency, and the alternative's simplicity is a mirage that fails exactly when it matters (load). Metered usage is also *monthly* by the `Period` key — a deliberate, simple model; finer windows would need a different scheme.

9. Checkpoint

```
git add -A && git commit -m "feat: quotas + dunning + dissolve – atomic consume (R30), lapse sweep, billing teardown"
git tag lesson-5.3
```

Part 5 complete. The platform is now a *product*: it charges money (hosted checkout), reflects the provider's truth safely (webhook + inbox + recency guard), gates features by plan (402), enforces countable limits without races (atomic quotas), chases lapsed payers (dunning), and cleans up billing when a tenant leaves (dissolve). And every piece composed from Parts 2–4 — the seam, the projection, EnterTenant, the inbox, the outbox, the scheduler, the contributor.

You should now be able to: explain the lost-update race and how a conditional UPDATE dissolves it; justify pending invites as reserved seats; describe once-per-lapse dunning and where its bookkeeping lives; and articulate why the provider cancel goes through the outbox and why billing reports no "abandonable data" to dissolve.

Next: *Part 6 — B2B essentials & security hardening* — RBAC (roles + the permission matrix + 403), file storage, the SSRF seam, and MFA/TOTP step-up.